

Curated Truth: What Drives Selectivity in the Supply of Political Facts?

Yero Samuel Ndiaye*

August 7, 2026

[EARLY DRAFT — DO NOT CIRCULATE — COMMENTS WELCOME]

[Click here for the most recent version](#)

Abstract

The information supply can be distorted not only by falsehoods but also by withheld truths. I study whether, when, and why people selectively share accurate political information. Among 1,200 U.S. adults, people share facts aligned with their political position 9 percentage points more often than unaligned facts. The gap increases when the recipient subsequently makes a real political donation. A second experiment with 1,548 German adults provides evidence on the motives: people want to correct recipients' beliefs but also derive expressive value from sharing facts that support their own political position. Expected persuasion matters only when recipients' choices have real consequences. The pattern extends to 307 professional journalists, who share aligned facts 20 percentage points more often than unaligned facts even without standard market incentives. This *curated truth* **[reduces the accuracy of recipients' beliefs]**. The distortion acts on a margin that current policies such as fact-checking and labeling do not target.

JEL classification: D83, D82, C91, D72, L82

Key words: Information Provision, Selective Disclosure, Political Information, Media Bias

*Ndiaye: University of Cologne and Max Planck Institute for Behavioral Economics; ndiaye@wiso.uni-koeln.de. I am grateful for valuable comments and suggestions from Lara Maria Berger, Markus Dertwinkel-Kalt, Max Rainer Pascal Grossmann, Johannes Münster, Shakked Noy, Axel Ockenfels, Christopher Roth, Nicolas Röver, Carlo Schwarz, Christoph Schottmüller, Sebastian Sieglösch, Till Stange, Johannes Wohlfahrt, Florian Zimmermann, as well as by participants of seminars and conferences at the C-SEB Early Ideas Workshop, IMPRS BeSmart Thesis Workshop, RWI Workshop, DICE Brownbag Workshop, the University of Münster Research Seminar, BCFM PhD Conference, MPI Summer Summit, GfEW Annual Meeting. I thank Jan-Luca Bauss, Frauke Nora Gerster, Linda Plehn, Lena Schmelig, and Moritz Lakenbrink for their excellent research assistance. Support from the Center for Social and Economic Behavior (C-SEB) and the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation) under Germany's Excellence Strategy EXC 2126/1390838866 is gratefully acknowledged.

1 Introduction

People frequently decide which political information to share or withhold. Journalists and editors determine what belongs in an article, and fact-checkers select which claims to scrutinize. Think tanks, advocacy groups, and podcasts covering politics choose which information to disseminate. Much information is supplied outside professional settings (citation), e.g., when people decide what to post on social media and what to raise in conversation.

Distortions in the supply of political information are treated as a threat to democratic decision-making in legislation and platform policy. The EU Digital Services Act requires large platforms to mitigate systemic risks from disinformation, the UK's Online Safety Act prohibits knowingly sending false information intending harm, and Canada's Elections Act forbids knowingly publishing specified false claims about candidates. Alongside these, platform fact-checking programs, content labeling, accuracy nudges, and digital-literacy campaigns aim to detect, correct, or preempt the supply of false content.¹

While false content spreads disproportionately fast and wide (Vosoughi et al. 2018), it is only a small share of what people consume. It is roughly 0.15 percent of Americans' daily media diet and about 1 percent of news consumption (Allen et al. 2020), exposure is concentrated among those with the strongest motivation to seek it out (Budak et al. 2024), and sharing it is rare (Guess et al. 2019). Where its effects on political beliefs have been measured, they appear limited (Allcott and Gentzkow 2017). False information that policy targets is therefore likely a small part of what distorts political information supply and thereby beliefs.

Even if no falsehoods exist, selective supply of true information distorts beliefs unless the receiver can infer what was left out. However, the audience rarely knows what the sender had to choose from, so what is missing cannot be reconstructed from what arrives; and where it could be inferred, receivers appear to read too little into silence (Jin et al. 2021). Under those conditions beliefs polarize even though no signal is false (Bowen et al. 2023), and withholding a fact can mislead as much as asserting a falsehood (Feess et al. 2024).

Selectivity is documented across the information supply chain. People pass on the fact-checks that favor their side (Shin and Thorson 2017), politicians amplify the coverage that favors theirs (Freitag et al. 2021), and fact-checkers are themselves selective about which claims get examined (Louis-Sidois 2025). Importantly, while falsehood is at least verifiable and detectable, withheld truth leaves almost nothing to check.

¹Some instruments also target one-sidedness for licensed broadcasters and market structure: Due-impartiality rules in the United Kingdom, public-service obligations under the European Media Freedom Act, differing-views provisions in Canada's Broadcasting Act, and ownership-pluralism review. Additionally, right-of-reply provisions are triggered by a person being named.

Selectivity of this kind admits several explanations. A sender may share facts that support their own political position because those are the only ones they have encountered, or because they are the only ones they believe to be true. Professionals additionally face market incentives. However, senders may also be *deliberately choosing* among facts they know and believe to be true. The first two explanations concern what a sender *knows*, the third what a sender is extrinsically *paid to do*, and the last what a sender intrinsically *wants*.

Models of strategic disclosure take the sender's preference over the receiver's belief or action as given. I measure which returns senders act on instead. In my framework, the returns divide by whether they depend on the receiver. The *expressive* return does not: it concerns whether a fact is congenial to the sender's own position. The receiver-dependent returns can be *corrective*, concerning whether the receiver ends up better informed; *persuasive*, concerning where provision moves the receiver's position; *signaling*, concerning what provision reveals to the receiver about the sender's stated position; or *reputation*, concerning whether the receiver takes the sender to be informed and credible. I measure the expressive, corrective, and persuasive returns, close the stated-position-signaling channel, and limit reputation concerns by design. Only persuasion is predicted to respond to whether the receiver's beliefs carry political consequences. Which explanation holds determines what could be done about it—an instrument that raises the stakes cannot reach a motive that is indifferent to them—and how well disclosure models describe this setting at all.

I provide causal evidence on the selective provision of true political information in three experiments that hold fixed what senders know and what they are paid to do, and measure the three returns. Senders choose from the same pool of facts, so provision cannot reflect what they happened to encounter. Every fact is true by construction, and I measure perceived credibility, so provision can be examined with it held fixed. One experiment puts the same task to professional journalists, absent the market incentives they would ordinarily face. The facts on offer are either aligned with the sender's position or unaligned with it, and identification comes from comparing the same sender's decisions across facts.

I additionally randomly assign whether the receiver's beliefs may translate into a political action. The treatment shifts none of the elicited returns, so it changes how senders act on their beliefs, not the beliefs. In Experiment 1 the receiver never learns that a disclosure decision was made; in Experiment 2 and 3 the sender's position is disclosed whatever they send, and transmission is probabilistic. In neither can a withheld fact be attributed to the sender, and in neither can provision make the sender's side better identified. That is what separates the expressive return from signaling the sender's stated position. Reputation inference is limited by construction, perceived credibility

is measured and no personally identifying information reaches the receiver.²

In Experiment 1, 1,200 quota-sampled U.S. adults choose which immigration facts to share with receivers who hold inaccurate beliefs. Provision is selective: senders share aligned facts 6.7 percentage points more often than unaligned ones when the receiver takes no political action ($p < 0.01$), a bias that holds when their beliefs about the fact's truth and credibility are held constant. The gap rises to 13.3 percentage points, when the receiver makes a real political donation ($p < 0.05$). This leads to *curated truth*: every shared and received fact is true but the likelihood of sharing and receiving are biased by the sender. However, as alignment here moves with both the expressive and the persuasive reason, the design cannot isolate which motive drives the increase.

Experiment 2 separates them in a setting with 1,548 quota-sampled German adults across two policy domains. I elicit each motive's return directly and make both political positions mutually known. The expressive return predicts provision independent of the receiver acting politically. Absent political action the persuasive return does not predict provision. However, it predicts provision once the receiver can act politically on what they learn. The corrective return predicts more provision without biasing it. Positional alignment does not capture this nuance: the bias is 11.6 percentage points when the receiver does not act and 12.7 percentage points when the receiver does.

Experiment 3 takes that design to the most common suppliers of political information: journalists. In a setting without standard extrinsic incentives such as competition, ownership, and consumer demand, I find that professional German journalists ($n = 307$) provide more information overall. However, their alignment bias is larger than that of the general population (20.4 percentage points; $p < 0.001$): facts that cut against a journalist's own position are passed on only about three-quarters as often as facts that support it. This mostly explained by journalists holding stronger political positions. The results strongly suggest that *some* journalistic selectivity observed outside the laboratory is intrinsic rather than driven by market forces.

Experiment 4 turns to the effect of curated truth receivers and the mechanisms driving it. Receivers are matched to the senders of Experiments 2 and 3 report. **[Selectivity decreases the accuracy of receivers' beliefs and/but ...their political preferences and actions. Receivers ...the selectivity they face and update imperfectly increasing the accuracy loss induced by curated. Importantly, receivers also insufficiently update when receiving information.]**

This paper asks whether, when, and why senders share the true information they encounter. While expression always matters, persuasion matters only where the receiver's beliefs carry political consequences, so the disclosure models fit best exactly where beliefs have impact. The findings also suggest that the bias observed in media

²This is a plausible setting since much political information reaches people with no identifiable individual behind it.

and attributed to extrinsic incentives is actually intrinsic to the journalist. As instruments built against distorted political information assess only whether statements are true, curated truth passes that test by construction. As expressive return predicts selection irrespective of consequences, instruments that change the stakes likely have little purchase on it. [Notably, receivers already expect selectivity but nonetheless do not update accordingly does receiver information campaigns are unlikely to help.]

Related Literature. Sender motives for sharing information have been closely studied. Thaler (2025) shows that paying senders to be rated truthful can lead to a supply of *false* motivated beliefs. I study the withholding of true information without financial incentives linked to the sharing decision, and find that senders’ selectivity is motivated by expressive and persuasive reasons. Guriev et al. (2026) structurally decompose motives for sharing true and false political content into reputation, partisan persuasion, and partisan signaling. I restrict content to facts and control for perceived credibility, thereby limiting the role of reputation concerns. Further, I disclose the senders stance independently of the provision decision, limiting position-signaling, and elicit receiver-specific returns without imposing a structural model.

Related work studies expressive and social-image motives in voicing positions and taking publicly interpretable actions (Bursztyn et al. 2023a, Hillman 2010, Bénabou and Tirole 2006). In every experiment the sender’s position either reaches the receiver regardless of what is sent or never reaches them at all, so provision cannot advertise where the sender stands, and the expressive return measured here is not a return to being seen holding a position. Facts can also carry corrective and persuasive returns beyond their expressive return.

Studies that measure motives from attitude scales or self-report find that partisan polarization is the primary motivation behind sharing from low-credibility sources and predicts real-news sharing in almost the same way (Osmundsen et al. 2021), while moralization and attitude extremity amplify the preference for congruent news, true and false alike (Marie et al. 2023). For politically congruent news, Marie and Petersen (2025) find that reported social motives respond to audience composition while perceived truth generally does not. I elicit the returns of sharing true information for different motives and link them to real provision decisions.

Disclosure theory asks when a sender holding verifiable information can stay silent. Under standard conditions, verifiable information unravels in equilibrium because the receiver interprets silence as bad news (Grossman 1981, Milgrom 1981), and unraveling attenuates when the receiver is unsure the sender is informed (Dye 1985, Jung and Kwon 1988) or when disclosure is costly (Verrecchia 1983).³ Laboratory evidence

³Cheap talk (Crawford and Sobel 1982) and Bayesian persuasion (Kamenica and Gentzkow 2011) vary the message technology and the sender’s commitment power around that same primitive. Shin

shows disclosure falling short of the full-unraveling benchmark across a range of environments (Jin et al. 2021, Degan et al. 2023, Burdea et al. 2023, Benndorf et al. 2025). Farina et al. (2026) separates concealment from selection, varying how much evidence a sender holds and how much of it she may disclose, and finds that receivers do not fully correct for the selection they face. In my paper the evidence is political and the receiver holds a position the sender knows. Neither environment here is predicted to unravel fully: in Experiment 1 the receiver does not know a disclosure decision was made, and in Experiments 2 and 3 a shared fact transmits with probability one half, so silence stays weakly informative rather than decisive.

The media-bias literature explains slant and content choice partly through market incentives. Profit-maximizing outlets slant toward consumers' priors to build a reputation for accuracy, and competition disciplines that slant by raising the chance the true state is later revealed (Gentzkow and Shapiro 2006). Mullainathan and Shleifer (2005) model slant as the selective omission of true evidence and find that competition need not reduce it where readers value congenial coverage, while Zhu and Dukes (2015) model commercial outlets choosing which verifiable facts to report. Consumer demand accounts for more measured slant than owner ideology (Gentzkow and Shapiro 2010), though ownership moves content (Martin and McCrain 2019). Journalists' own ideology travels with them across outlets, and the paper's preferred decomposition attributes part of the between-outlet variation in slant to it (Boxell and Conway 2022).⁴ Multiple extrinsic channels operate simultaneously in these settings. Experiment 3 abstracts from competition, ownership, and consumer demand, so selectivity that persists does not require them.

Individuals disproportionately seek and value congenial information and avoid disconfirming information (Golman et al. 2022, Peterson and Iyengar 2021, Chopra et al. 2022), and the demand for news reflects both accuracy concerns and belief-confirmation motives (Chopra et al. 2024). What arrives is then processed unevenly: receivers underweight belief-threatening signals, while the asymmetry is weaker or absent in domains without an ego stake (Eil and Rao 2011, Möbius et al. 2021, Barron 2021, Zimmermann 2020), and in politics the trust placed in a message tracks its congeniality (Thaler 2024).⁵

(2003) studies a sender who discloses favorable items and conceals unfavorable ones, so that what is withheld enters only as residual uncertainty, and Manili (2024) adds motivated belief formation on the receiver's side. Farina and Leccese (2025) varies the dimensionality of the evidence.

⁴Political control shapes what outlets report (Durante and Knight 2012), and capture by owners, advertisers, and the state constrains media accountability (Besley and Prat 2006, Prat and Strömberg 2013). Exposure and sharing are ideologically asymmetric across the political spectrum (González-Bailón et al. 2023, Ershov and Morales 2024). Ownership transitions sort journalists into compatible outlets (Cagé et al. 2022), and about two-thirds of article-level slant variance lies within outlets rather than between them (Braghieri et al. 2025). Financial attention incentives lead a sample of German journalists similar to mine to choose more emotional and negative headlines (Berger 2026).

⁵Bullock and Lenz (2019) review the evidence that part of the partisan divergence in stated factual beliefs is expressive rather than sincere, and Bursztyn et al. (2023b) find that viewers choose opinion programming to inform incentivized factual estimates even when the reward for accuracy increases.

This literature studies a receiver who chooses what to acquire and how to read it. In this paper, the information the receiver sees is chosen by someone with a (potential) stake in what the receiver believes.

The effect of distorted information supply on beliefs, and the remedies for it, have been studied mostly in the context of falsehoods. Fact-checking reduces misinformation engagement, with the effect concentrated among claims rated outright false and checked promptly (Cagé et al. 2025), and it lowers the rate at which such claims are shared (Henry et al. 2022); accuracy nudges that redirect attention to truth reduce false sharing (Pennycook et al. 2021). Berger et al. (2025) find that media literacy produces broader and more persistent reductions in the perceived credibility of false posts than item-specific fact-checking.⁶ Fact-checking nonetheless leaves intact the persuasive effect of the falsehoods it corrects, even as it restores accurate beliefs (Barrera et al. 2020), and Musolff et al. (2026) find that beliefs keep moving after a false signal has been rendered uninformative, with the residue concentrated in narrative rather than numerical content. Guriev et al. (2026) find that fact-checking reduces the sharing of true posts as well as false ones, while a content-neutral prime curbs false sharing without measurably affecting true. Selectivity among facts is a distortion these measures do not target.

The receiver consequences of selective supply have been modeled theoretically and measured in the laboratory. Bowen et al. (2023) show that when senders pass on only congenial true news and receivers misread the resulting silence, beliefs polarize although no signal is false.⁷ That misperception has a formal counterpart in models of receivers who do not fully account for the tie between what a sender conveys and what she knows (Eyster and Rabin 2005, Fong et al. 2025). For stated positions, Braghieri (2024) separates what a receiver decodes from what the sender encoded.⁸ Experiment 4 measures the belief-accuracy cost of selective provision against ground truth, for political facts, with receivers matched to the provision decisions of senders whose returns to provision were elicited in Experiments 2 and 3.⁹

Aridor et al. (2024) surveys the economics of the platforms that mediate this demand. Using incentivized news-classification tasks, Angelucci and Prat (2024) find high average discernment and substantially larger socioeconomic than partisan differences. Projective paternalism (Ambuehl et al. 2021) concerns intervening in another’s choices by one’s own standards. Related work studies paternalistic intervention more broadly (Bartling et al. 2023) and weighs information against choice restriction (Grossmann 2024); my design sits upstream of both, at the selection of which facts to present.

⁶False content spreads farther and faster than true content on Twitter (Vosoughi et al. 2018); Lazer et al. (2018) review the evidence and the interventions.

⁷However, selection alone cannot generate polarization in their model, and receivers’ misperception of silence is also necessary.

⁸Suvorov et al. (2024) offer a related design joining selective supply to its measured cost, for ego-relevant information among teammates and under a single image motive that is posited rather than estimated. Their receivers’ response to the informative silence is not distinguishable from zero.

⁹The information-provision design this builds on is developed in Kuziemko et al. (2015), Grigorieff et al. (2020) and Haaland et al. (2023), and applied to immigration beliefs in Europe by Alesina et al. (2023).

I structure the rest of the paper as follows. Section 2 presents Experiment 1 on immigration policy in the US. Section 3 develops a conceptual framework for sender behavior. Section 4 presents Experiment 2 on immigration and climate policy in Germany. Section 5 presents Experiment 3 with professional journalists. Section 6 presents Experiment 4 with the receivers. Section 7 discusses the implications, and Section 8 concludes.

2 Experiment 1

In Experiment 1 senders decide whether to share four facts, one at a time, with someone who holds an incorrect belief or reports not knowing. Every sender sees every fact, so provision cannot reflect what they happened to encounter. Every fact is true and shown with its source, so selective provision does not require supplying false content. Senders pass on facts aligned with their own position more often than unaligned facts, and this selectivity increases when the recipient subsequently makes a real political donation.

2.1 Design, Treatment, and Hypotheses

I run a pre-registered online experiment with $N = 1,200$ U.S. adults recruited on Prolific. The policy domain is U.S. immigration policy.

Experimental Timeline The experiment proceeds in three parts (Figure 1). In Part 1, I elicit senders’ prior knowledge about four immigration facts and measure their policy preferences through a donation decision. In Part 2, I elicit senders’ beliefs about the persuasiveness of information. In Part 3, senders make one provision decision per fact, each directed at a specific receiver, and a closing questionnaire records stated motivations, credibility assessments, and demographics.

Facts about Immigration Table 1 presents the four facts, selected from authoritative sources (Bureau of Justice Statistics, Bureau of Labor Statistics, Customs and Border Protection, National Academies of Sciences) and balanced across policy directions. I validated them in a pre-registered survey ($N = 100$) that tested a larger candidate pool against criteria of veracity, perceived credibility, policy relevance, and directional balance (see Appendix A.1).¹⁰ Each is framed as an invertible binary comparison (e.g., “higher” versus “lower”), so beliefs are correct or incorrect.

¹⁰Classifying a fact’s policy direction required agreement on two questions: which side of the immigration debate could use it, and which way learning it would move policy views.




Figure 1: Experimental timeline, Experiment 1.

Table 1: Facts Used in Experiment 1

Direction	Fact
Pro	The crime rate of immigrants is <i>lower</i> than that of U.S.-born Americans.
Pro	The labor force participation of immigrants is <i>higher</i> than that of U.S.-born Americans
Anti	The average number of unauthorized border crossings is <i>higher</i> in the last five years compared to historical averages.
Anti	When comparing what first-generation immigrants use in government services versus what they pay in taxes, the net cost is <i>higher</i> than the same calculation for U.S.-born Americans at the state and local levels.

Notes: Verbatim claims as shown to senders. Direction (Pro) = supports more-permissive immigration policy (Anti) = supports more-restrictive policy. All from statistical offices and research institutes. Each sender faces all four facts in random order. Pre-validation details in Appendix A.1.

Beliefs about Facts I first elicit senders’ prior knowledge by having them complete each comparative statement, in random order, as “higher,” “lower,” or “don’t know.” Responses are incentivized: correct completions receive a bonus, incorrect completions do not, and “don’t know” responses receive the bonus with 25% probability. Senders are then told the correct answer for each. Because senders who already knew a fact may hold it with greater conviction, I control for prior belief accuracy throughout.

Policy Preferences I measure senders’ policy preferences in an incentivized donation task: each sender allocates \$500 between the *American Immigration Council* (AIC), which advocates for increased immigration, and *NumbersUSA*, which advocates for decreased immigration. Senders are shown each organization’s position in their own words before deciding (see Appendix A.2 for details). The split yields a continuous measure of policy preference, and therefore of alignment, that captures both the direction and the intensity of each sender’s position.

Receiver Belief States For each sender, each fact is assigned to exactly one receiver who either completed the comparative statement incorrectly (false belief) or reported not knowing. The sender observes which case applies before deciding. Throughout, senders decide for receivers who are unaware that a disclosure decision was made: provided information appears as part of the survey flow, and withheld information leaves no trace. Non-disclosure therefore carries no inferential content about the sender.

Treatments Table 2 summarizes the four design dimensions. Two vary within sender and two between.

Table 2: Experimental design and identifying variation

Dimension	Variation	Levels
Alignment with statistic	Within-subject	<i>Aligned</i> and <i>Unaligned</i>
Political action (ACTION)	Between-subject	ACTION = 1 and ACTION = 0
Receiver belief state	Within-subject	<i>False belief</i> and <i>Don’t know</i>
Effect of information (INFO)	Between-subject	INFO = 1 and INFO = 0

Experiment 1 ($N = 1,200$ U.S. adults, Prolific, quota-sampled). Within-subject dimensions: each sender evaluates both receiver belief states and all four statistics (two aligned, two unaligned), yielding 4 decisions per sender. Between-subject dimensions: senders are randomly assigned to one of four cells crossing ACTION and INFO.

The primary within-sender dimension is alignment. A fact is *aligned* for sender if the policy side it favors matches her own policy preference, and *unaligned* otherwise. As each side has two facts favoring it, a sender with a directional preference faces two of each, which holds all sender-level heterogeneity fixed. Receiver belief state also varies within sender: each fact is assigned to a receiver who either holds a

false belief or reports not knowing. Between senders, the primary dimension is ACTION, under which receivers make the same incentivized donation decision as senders and otherwise complete the experiment without further decisions.¹¹ INFO supplies the fourth dimension by providing half of senders with the actual results of Haaland and Roth (2020), generating exogenous variation in beliefs about persuasiveness of facts. Crossing the two between-sender treatments gives a 2×2 factorial design: CONTROL (neither ACTION nor INFO), ACTION, INFO, and ACTION $\times$ INFO.

Hypotheses I pre-registered four hypotheses about how senders provide information.

Hypothesis 1.0 (Baseline provision). *Provision rates for both aligned and unaligned information exceed the indifference benchmark of 50% in CONTROL.*

Hypothesis 1.1 (Alignment-based selectivity). *Senders provide aligned information at strictly higher rates than unaligned information in CONTROL.*

Hypothesis 1.2 (ACTION amplifies selectivity). *The alignment gap in provision is larger in ACTION than CONTROL.*

Hypothesis 1.3 (INFO amplifies selectivity). *The alignment gap in provision is larger in INFO than CONTROL.*

2.2 Sample, Incentives, and Protocol

The sender experiment was run on Prolific from July 16 to August 12, 2025, quota-sampled on age, gender, ethnicity, and party identification. (See Table B.4 for details.) This study was approved by the Ethics Council of the Max Planck Society under the Generalized Approval of Experiments Following the Protocol that is Standard in Experimental Economics (approval no. 2018_3 Renewal 2024_28) and was pre-registered on AsPredicted (No. 238343). Of 1,257 starters, 57 were excluded for failing the attention check, switching tabs during the fact questions, or timing out, leaving $N = 1,200$. Senders took a median of 17.26 minutes and earned \$3.32 on average. Receiver data were collected on Prolific between September 1 and September 10, 2025. 96 receivers took a median of 3 minutes and earned \$0.80 on average.

¹¹Through randomization, any difference in alignment-based selectivity is attributable to the downstream political consequence, under the maintained assumption that ACTION operates only through that consequence. ACTION may also shift the salience of the decision and the sender’s construal of the task. I provide suggestive evidence that this is not the case.

2.3 Results

2.3.1 Senders' Beliefs and Policy Preferences

Senders' Beliefs about Facts I find substantial heterogeneity in senders' factual knowledge.¹² A majority correctly identify the facts on immigrants' crime rates (70%), labor force participation (79%), and recent border crossings (61%), but only 30% correctly assess the net fiscal impact of first-generation immigrants. Democrats and immigration supporters are more accurate on the pro-immigration facts, Republicans and immigration opponents on the anti-immigration facts (all differences $p < 0.001$). The partisan sorting motivates controlling for senders' prior beliefs throughout.

Policy Preferences Senders' donation decisions spread across the scale (Figure B.2): 30.2% allocate the entire endowment to the pro-immigration organization, 14.1% to the anti-immigration one, and 7% split evenly. The split correlates with stated policy preferences ($r = 0.56$) and with political affiliation (both $p < 0.001$), each elicited at the end of the experiment. These decisions supply the variation in alignment against which provision is tested.

Senders' Beliefs about the Effect of Information on Political Preferences Senders believe information moves policy views: roughly 90% expect it to raise petition signatures and 72% expect that receivers holding false beliefs update more strongly, though they only predict an increase of 37% against a true effect of 69% (Figure B.1).

Expected Attitude Effects of Information For each sender i and fact j , I define the expected attitude change $\Delta\text{Attitude}_{i,j} \in [-1, 1]$ as the net share of ten hypothetical receivers whom the sender expects to become more, rather than less, supportive of immigration after receiving fact j . These beliefs are elicited unincentivized before the provision decisions. Senders expect each fact to move attitudes in the direction of its lean: 25 percentage points toward pro-immigration for the crime fact and 17 pp for the labor fact, 34 pp toward anti-immigration for the border fact and 32 pp for the fiscal fact (all $p < 0.001$; Appendix Table B.7).

Manipulation Check The INFO treatment moved senders' beliefs about persuasiveness on one side of the debate only. Senders shown the Haaland and Roth (2020) results expect a larger pro-immigration shift from the crime fact (+5.8 pp, $p < 0.01$) and from the labor fact (+3.8 pp, $p < 0.05$). For the two anti-immigration facts the changes

¹²See Table B.5 for details. The fiscal fact, where prior knowledge is lowest, is rated ex post as credible as the other three (Table B.6) and provided at comparable rates. Results are robust to leaving out each fact individually (Appendix Table B.13).




Figure 2: Alignment bias by treatment, Experiment 1 (U.S. adults, $N = 1,200$, Prolific, quota-sampled). Alignment bias is the within-sender difference between the provision rate for aligned statistics and the provision rate for unaligned statistics. Restricted to senders with revealed directional preference ($N = 1,114$, 93%; $AIC_i \neq 250$). Error bars are standard errors.

are small and not distinguishable from zero (Appendix Table B.7; fuller detail in Appendix A.4).

2.3.2 Treatment Effects on Selective Provision of Information

Let $y_{i,j}$ equal 1 if sender i provides fact j , and let $AIC_i \in [0, 500]$ denote i 's donation to the pro-immigration group (AIC). Sender i is pro-immigration if $AIC_i > 250$ and anti-immigration if $AIC_i < 250$. Crime and labor are the pro-immigration facts, border and fiscal the anti-immigration ones, so a pro-immigration sender's two aligned facts are the anti-immigration sender's two unaligned facts.

A sender's *alignment bias* is the share of their aligned facts provided minus the share of their unaligned facts provided. It is positive when provision favors the sender's own policy position and zero when it does not.

Senders share accurate information in all four treatment conditions, including information that cuts against their own position. Unaligned provision rates range from 81.9% (ACTION) to 86.5% (CONTROL), all significantly above 50% ($p < 0.001$, one-sample t -tests).

Result 1.0 (Provision is high, including unaligned information). *In CONTROL, 86.5 percent of unaligned information is shared, significantly above the 50 percent benchmark ($p < 0.001$), confirming Hypothesis 1.0. Unaligned provision exceeds the benchmark in all four treatments ($p < 0.001$).*

Figure 2 shows that provision is nevertheless selective. Senders provide their aligned facts 9.3 percentage points more often than their unaligned facts, and the bias is posi-

tive in all four conditions ($p < 0.01$).¹³ The bias is defined for the 1,114 senders whose donation leaves them on one side of the issue.

I combine the two into a continuous alignment measure:

$$\text{Align}_{i,j} = \begin{cases} \frac{AIC_i - 250}{250} & \text{if } j \in \{\text{crime, labor}\} \\ \frac{AIC_i - 250}{250} \cdot (-1) & \text{if } j \in \{\text{border, fiscal}\} \end{cases}$$

This measure is bounded in $[-1, 1]$:¹⁴ its sign records whether fact j 's direction matches sender i 's position, and its magnitude is the intensity of that position. Only the sign varies within a sender—each sender's four facts carry the same magnitude—so the within-sender comparison is between aligned and unaligned facts and the magnitude enters as a between-sender weight. Thus I impose, rather than test, the assumption that a fact is more aligned, and the provision gap correspondingly larger, the more one-sided the sender's position.¹⁵

Let $\text{Belief}_{i,j}$ denote sender i 's probabilistic prior accuracy about fact j (1 if correct, 0.5 if “don't know,” and 0 if incorrect, as pre-registered), capturing that senders may not update their own beliefs when told the truth. This control adjusts for prior accuracy under the extreme assumption of no updating. I estimate the pre-registered linear probability model:

$$\begin{aligned} y_{i,j} = & \beta_0 + \beta_1 \text{Align}_{i,j} + \beta_2 \text{ACTION}_i + \beta_3 \text{INFO}_i + \beta_4 (\text{Align}_{i,j} \times \text{ACTION}_i) \\ & + \beta_5 (\text{Align}_{i,j} \times \text{INFO}_i) + \beta_6 (\text{ACTION}_i \times \text{INFO}_i) \\ & + \beta_7 (\text{Align}_{i,j} \times \text{ACTION}_i \times \text{INFO}_i) + \gamma_1 \text{Belief}_{i,j} + \gamma_2 \mathbf{S}_j + \gamma_3 \mathbf{X}_i + \eta_{i,j}, \end{aligned} \quad (1)$$

with standard errors clustered at the sender level. Column (1) includes no fixed effects or controls. Column (2) adds fact fixed effects $\mathbf{S}_j$, and column (3) adds demographic and political controls $\mathbf{X}_i$. β_1 captures baseline alignment-based selectivity (Hypothesis 1.1), β_4 whether alignment matters more when receivers act politically (Hypothesis 1.2), and β_5 whether information about persuasiveness amplifies selectivity (Hypothesis 1.3). The control γ_1 estimates whether senders who had more accurate priors

¹³Demeaning provision by fact—the non-parametric counterpart of the fact fixed effects used below—leaves +7.6 pp. The difference reflects the sample's composition rather than selection: the two pro-immigration facts are provided 6.6 pp more often than the two anti-immigration ones, and 62.5 percent of senders hold a pro-immigration position, so the aligned set is on average the more shareable one. The ACTION contrast is unaffected (+4.3 pp raw against +4.5 pp demeaned), since randomization balances that composition across arms.

¹⁴The pre-registration defined alignment on $[0, 1]$. Centering at the indifference point to ease interpretation is a one-to-one transformation that leaves all hypothesis tests invariant, so results under the pre-registered coding are identical.

¹⁵A binary aligned/unaligned indicator, which imposes no proportionality, gives a gap of 5.0 pp ($p = 0.007$) with an ACTION interaction of +6.4 pp ($p = 0.024$). Appendix Table B.11 reports the full specification.

Table 3: Treatment Effects on Information Provision

	Probability of Statistic Provision		
	(1)	(2)	(3)
Align	0.041*** (0.012)	0.036*** (0.012)	0.037*** (0.012)
INFO	−0.013 (0.017)	−0.012 (0.017)	−0.012 (0.017)
ACTION	−0.002 (0.016)	−0.002 (0.016)	−0.003 (0.016)
Belief Accuracy	0.043*** (0.011)	0.034*** (0.012)	0.033*** (0.012)
Align × INFO	0.006 (0.017)	0.004 (0.017)	0.004 (0.017)
Align × ACTION	0.041** (0.018)	0.041** (0.018)	0.041** (0.018)
INFO × ACTION	−0.007 (0.024)	−0.007 (0.024)	−0.008 (0.024)
Align × INFO × ACTION	−0.028 (0.026)	−0.027 (0.026)	−0.027 (0.026)
Clustered SE	Subject	Subject	Subject
Statistics FE	No	Yes	Yes
Controls	No	No	Yes
No. of Senders	1200	1200	1200
Num. obs.	4800	4800	4800
R ²	0.030	0.033	0.040
Adj. R ²	0.028	0.031	0.034

Sample: $N = 1,200$ U.S. adults recruited on Prolific (quota-sampled), 4,800 sender-statistic observations (4 statistics per sender). *Outcome:* binary provision indicator. *Align:* $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, 1]$, where Pref is derived from an incentive-compatible donation split. *Belief Accuracy:* incentivized prior belief correctness, elicited before senders knew they would make a provision decision. *Info:* between-subject treatment sender receives personalized feedback comparing their prediction to the actual results of Haaland and Roth (2020) (see Appendix A.4). *Action:* between-subject treatment receiver makes a downstream political donation. *Controls:* age, gender, education, party affiliation, political knowledge. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

about the fact provide it at higher rates.

Table 3 reports the results.¹⁶ Across all specifications, provision increases significantly in political alignment between the sender and the fact. The alignment coefficient is +0.036 to +0.041 ($p < 0.005$). Because alignment runs from -1 to $+1$, the coefficient is the aligned–unaligned gap for a sender at half the maximum preference intensity, and half that of a sender whose position is entirely one-sided. In CONTROL, the sender-level bias is 6.7 pp ($p < 0.01$), confirming Hypothesis 1.1.

¹⁶Appendix Table B.12 allows asymmetric INFO effects by fact direction, to accommodate the one-sided first stage. Appendix Table B.9 reports the results pooled at the sender level as a pre-registered robustness check. They are essentially the same.

Result 1.1 (Senders favor aligned information). «««< HEAD In CONTROL the alignment bias is 6.7 pp ($p < 0.01$), confirming Hypothesis 1.1. The bias is positive in all four conditions and averages 9.3 pp. The alignment coefficient is +0.036 to +0.041 per unit ($p < 0.005$). ===== In CONTROL the alignment bias is 6.7 pp ($p < 0.01$), confirming Hypothesis 1.1. The bias is positive in all four conditions ($p < 0.005$) and averages 9.3 pp . »»»> dcb760ad569f8f1270335f76418984c28299b9d

Political consequences increase selectivity, confirming Hypothesis 1.2. Comparing CONTROL with ACTION, the sender-level alignment bias rises from 6.7 to 13.3 pp, an increase of 6.5 pp (one-sided $p = 0.011$; Holm-adjusted $p = 0.045$).¹⁷ The decision-level interaction $\text{Align}_{i,j} \times \text{ACTION}_i$ is +0.041 in every specification ($p < 0.05$).¹⁸

Averaging over the independently randomized INFO intervention, the bias rises from 7.1 to 11.4 pp, an increase of 4.3 pp ($p = 0.037$).¹⁹

Result 1.2 (ACTION amplifies selectivity). *Alignment-based selectivity increases when receivers take a downstream political action, confirming Hypothesis 1.2: comparing CONTROL with ACTION the bias rises from 6.7 to 13.3 pp (one-sided $p = 0.011$; Holm-adjusted $p = 0.045$), and averaging over INFO from 7.1 to 11.4 pp.*

Senders with accurate priors about the fact provide it at higher rates (3.3 to 4.3 pp, $p < 0.01$). The effect is smaller in magnitude than the alignment effect and does not attenuate the $\text{Align} \times \text{ACTION}$ interaction.

Hypothesis 1.3 is not confirmed. The pre-registered comparison of CONTROL with INFO gives +0.8 pp (one-sided $p = 0.385$). Neither the main effect of INFO nor its interaction with alignment is distinguishable from zero, and the same holds when INFO is allowed to act asymmetrically by fact direction: Wald tests do not reject that INFO leaves the alignment effect unchanged for the pro-immigration facts ($p = 0.403$) or for the anti-immigration facts ($p = 0.606$), nor that it leaves overall provision unchanged for the pro-immigration facts ($p = 0.390$; Appendix Table B.12). INFO affected beliefs about persuasiveness for two of the four facts and by four to six percentage points. Thus, I do not infer that senders are insensitive to the expected magnitude of attitude change, but I also find no evidence that a shift of this size moves provision.

¹⁷The pre-registration specifies two further comparisons in the same correction family, testing whether combining the two manipulations produces effects beyond either alone: $\text{ACTION} \times \text{INFO}$ against ACTION (−3.7 pp) and against INFO (+2.1 pp). Neither is distinguishable from zero, and the first runs opposite to the direction the pre-registration anticipates. These are pre-specified tests rather than stated hypotheses. Together with Hypotheses 1.2 and 1.3 they complete the family of four over which the Holm correction is applied.

¹⁸Alignment, ACTION and INFO enter fully interacted, so the tabulated interaction is the ACTION effect at INFO = 0: the sender-level test and the interaction are two readings of one comparison rather than independent evidence. Table B.10 shows the results are robust to logistic estimation.

¹⁹Measured from INFO to ACTION $\times$ INFO rather than from CONTROL to ACTION, the increase is 2.1 pp rather than 6.5, but the two increases are not statistically distinguishable (−0.027, standard error 0.026).

Result 1.3 (INFO does not amplify selectivity). *Providing senders with evidence that information shifts political behavior has no detectable effect on the level or the alignment-selectivity of provision. Hypothesis 1.3 is not confirmed.*

The estimates are unchanged under alternative specifications. Sender fixed effects, which identify alignment only from each sender’s own comparisons across facts, leave the interaction at +0.041. A binary alignment coding imposing no proportional scaling in preference intensity, presentation-position fixed effects, and the unrestricted starter sample leave every estimate in place, and no single fact drives the result.²⁰

2.3.3 Heterogeneity and Correlational Evidence

Pre-registered exploratory analyses indicate heterogeneity by political affiliation and preference intensity. Appendix Table B.14 reports the estimates. Neither overall provision nor alignment-based selectivity differs detectably between receivers holding false beliefs and those reporting “don’t know” (Appendix Table B.16), and senders do not expect larger attitude shifts for the former (Appendix Table B.15).

Selective provision could reflect selective belief: senders might withhold unaligned facts because they doubt them. Prior belief accuracy, incentivized and elicited before any provision decision, is controlled in every specification, and the alignment effect does not vary with it: the Align $\times$ Belief Accuracy interaction is +0.009 ($p = 0.58$), and the alignment effect remains +0.031 for facts the sender had answered incorrectly. Perceived credibility, elicited unincentivized after the provision task, is a weaker instrument. Declaring a fact incredible would have been a costless way to rationalize withholding it, yet senders rate 97 percent of facts credible. The remaining 2.8 percent leave too little variation for the corresponding interaction to be informative (-0.021 , $p = 0.65$). Credibility is also slightly higher for aligned facts (98.3 against 96.0 percent), so it is not independent of the decision it would explain. Appendix Table B.17 reports both specifications.

Senders’ own accounts are the only direct evidence on motives in this experiment. A closing question asked what considerations drove the decisions. Two in five senders gave a reason referring to the sharing decision. Those senders are somewhat more issue-extreme and 4.0 pp *less* alignment-biased than the rest ($p = 0.043$). Two large language models classify each response under one frozen coding scheme. I report the first throughout and the second as robustness (Appendix E.1).

Among senders giving a codable reason, correcting the receiver’s beliefs is by far

²⁰Leaving out each fact in turn gives Align $\times$ ACTION between +0.039 and +0.044 (Appendix Table B.13). The binary coding gives +0.064 (Appendix Table B.11). Provision falls about four percentage points after the first fact and is flat thereafter, and position fixed effects leave every estimate unchanged. With four facts there are too few clusters for fact-level inference, so standard errors are clustered on senders throughout.

the most common consideration (85 percent), followed by influencing the receiver’s opinion (8 percent) and expressing one’s own position (3 percent). Relating these to each sender’s alignment bias (Appendix Table B.18), senders citing belief correction are 9.9 pp less alignment-biased ($p < 0.001$), and senders citing persuasion are more biased only when the receiver takes a political action ($p < 0.01$). An expressive reason is seldom stated—twelve senders in all—but where stated it is decisive.²¹

ACTION changes what senders do without changing what they report considering: the share giving a codable reason and the shares citing each consideration are indistinguishable across arms, individually and jointly ($F = 0.44$, $p = 0.72$). This matters twice. The interactions above compare senders citing the same reason across arms rather than different mixes of reasons. Because the question follows the provision task, senders assembling reasons to fit choices already made would have had to assemble them alike in the arm that chose differently. Because the responses are unincentivized, elicited after provision, machine-coded, and sparse in the latter two categories, I treat them as descriptive.

Senders who provide more information subsequently express stronger pro-accuracy norms. Conditional on overall provision, alignment-based selectivity does not predict those norms (Appendix Table B.19).

Importantly, alignment-based selectivity admits two readings: a sender may value passing on facts that favor their own side, or may expect those facts to move the receiver toward it. Here the two point to the same facts. Senders expect net movement toward their own position for 78.0 percent of aligned facts and 17.2 percent of unaligned ones, a gap of 60.8 percentage points ($p < 0.001$).²² Experiment 1 therefore identifies how political consequences change selection on the bundled value of a fact, not which of the two components responds.

3 Conceptual Framework

To distinguish the motives underlying selective disclosure of true information, I develop a framework in which an informed sender decides whether to disclose a true state to a receiver. The framework separates a concern for the receiver’s factual beliefs from three politically directional motives: expression, projective persuasion, and

²¹Twelve of 1,102 senders state an expressive reason, four in CONTROL and eight under ACTION. The coefficient is significant at the one percent level within each arm taken on its own. At that cell size the estimate describes those twelve senders rather than a population magnitude, and the difference between the arms is imprecise. The persuasion-by-ACTION interaction rests on the 23 senders citing persuasion in the ACTION arm. The second classifier gives a smaller but same-signed estimate on the 30 such senders it identifies (Appendix Table B.20).

²²The coincidence is already present at the validation stage, where the two directional criteria—which side of the immigration debate a statistic serves and which way learning it moves policy views—classify the four statistics identically (Appendix A.1).

instrumental persuasion. It characterizes the sender's private disclosure incentives without imposing a complete equilibrium model of receiver inference.²³

3.1 Environment

Nature draws a state of the world $\omega \in \{0, 1\}$, the sender's political position, and an indicator $A \in \{0, 1\}$ for whether the receiver's policy-relevant position may translate into political action. The sender observes ω and chooses whether to disclose it, $d = 1$, or withhold it, $d = 0$. Conditional on disclosure, the state reaches the receiver with probability $\tau \in (0, 1]$; withholding and failed transmission both result in non-receipt. If $A = 1$, the receiver then takes political action.

The sender evaluates disclosure along three dimensions. First, let b^0 and b^1 denote her subjective forecasts of the receiver's posterior probability that $\omega = 1$ under non-receipt and under receipt, with $b^d \in [0, 1]$. The factual belief correction she expects receipt to produce is

$$\Delta b = |\omega - b^0| - |\omega - b^1| \in [-1, 1],$$

so that larger values correspond to larger expected movements of the receiver's belief toward the truth.

Second, let $\alpha \in \mathbb{R}$ denote her assessment of how congenial the state of the world is to her political position. Positive values indicate that she regards it as affirming her position, negative values as conflicting with it, and zero as politically neutral.

Third, let $\Delta\pi \in \mathbb{R}$ denote the movement in the receiver's policy-relevant position toward her own position that she expects receipt to produce, relative to non-receipt. Thus $\Delta\pi > 0$ means she expects receipt to move the receiver toward her, and $\Delta\pi < 0$ that she expects it to move the receiver away. Congeniality and expected persuasion are distinct objects: $\text{sign}(\alpha)$ need not equal $\text{sign}(\Delta\pi)$.

3.2 Sender Utility and the Disclosure Decision

Normalizing the utility of withholding to zero, $U(d = 0) \equiv 0$, the sender's net utility from disclosure is

$$U(d = 1) = \tau\theta_B\Delta b + \theta_E\alpha + \tau c(A) w(\Delta\pi), \quad (2)$$

with $c(A) = \theta_P + \theta_I A$ and $\theta_B, \theta_E, \theta_P, \theta_I \geq 0$. The four weights are a *belief-correction motive* θ_B , the value of correcting the receiver's factual beliefs independent of political content; an *expressive motive* θ_E , the value of disclosing a state of the world that affirms

²³Appendix G collects the proofs.

the sender's own political position, independent of its expected effect on the receiver; a *projective persuasion motive* θ_P , the value of moving the receiver toward that position even absent downstream action; and an *instrumental persuasion motive* θ_I , the additional value of that same movement when the receiver may act politically.

The function $w : \mathbb{R} \rightarrow \mathbb{R}$ maps the expected movement in the receiver's policy-relevant position into utility and is sign-preserving: $w(0) = 0$, $w(\Delta\pi) > 0$ for $\Delta\pi > 0$, and $w(\Delta\pi) < 0$ for $\Delta\pi < 0$. For derivative-based predictions I additionally assume that w is differentiable and strictly increasing on the relevant domain.

The transmission probability multiplies the two receiver-directed terms because factual correction and persuasion both require receipt. It does not multiply the expressive term, which arises from the act of disclosing rather than from the receiver's response: the sender earns no expressive payoff from withholding.

The sender discloses if and only if her utility from disclosure exceeds an idiosyncratic disclosure shock, $d = 1$ if and only if $U(d = 1) > \varepsilon$. The shock is independent of $(\Delta b, \alpha, \Delta\pi, A)$ and drawn from a distribution F with full support on $\mathbb{R}$ and strictly positive density $f = F'$; it captures unmodeled determinants of disclosure such as attention or salience. The probability of disclosure is therefore

$$\Pr(d = 1) = F(\tau\theta_B\Delta b + \theta_E\alpha + \tau c(A)w(\Delta\pi)). \quad (3)$$

3.3 Predictions

The framework yields five predictions: a non-political belief-correction benchmark, two distinct forms of political selectivity, the component of persuasion that operates only when the receiver may act politically, and a magnitude prediction. I state the first four here and the fifth in Appendix G, which also collects the proofs.

Proposition 0 (Belief correction). *If $\theta_B > 0$, disclosure is strictly increasing in the expected factual belief correction:*

$$\frac{\partial \Pr(d = 1)}{\partial \Delta b} > 0.$$

Holding congeniality and expected persuasion fixed, disclosure rises with the correction the sender expects receipt to produce. This benchmark is politically non-directional.

Proposition 1a (Expressive selective disclosure). *If $\theta_E > 0$, disclosure is strictly increasing in congeniality:*

$$\frac{\partial \Pr(d = 1)}{\partial \alpha} > 0.$$

Expression is a property of the act of disclosure itself, not of how the receiver is expected to respond.

Proposition 1b (Persuasive selective disclosure). *If $c(A) > 0$, disclosure is strictly increasing in $w(\Delta\pi)$; if, in addition, w is differentiable and strictly increasing,*

$$\frac{\partial \Pr(d = 1)}{\partial \Delta\pi} > 0.$$

Persuasion depends on the receiver instead: senders favor states of the world they expect to move the receiver toward their own position and disfavor those they expect to move the receiver away from it.

When $\text{sign}(\alpha) = \text{sign}(\Delta\pi)$ —as in Experiment 1—Propositions 1a and 1b make the same prediction. The two channels are separately estimable when the signs diverge.

Proposition 2 (Political action amplification). *If $\theta_I > 0$, allowing the receiver’s policy-relevant position to translate into political action strictly increases the marginal utility of persuasion:*

$$\left. \frac{\partial U(d = 1)}{\partial w(\Delta\pi)} \right|_{A=1} - \left. \frac{\partial U(d = 1)}{\partial w(\Delta\pi)} \right|_{A=0} = \tau\theta_I > 0.$$

Proposition 2 is stated at the utility index, the object the utility function disciplines directly; without further restrictions on the distribution F it does not require a uniformly positive interaction in disclosure probabilities. The instrumental motive is the increment in the value of persuasion when the receiver may act politically, distinct from the always-present projective motive.

The fifth prediction concerns magnitude: when w is strictly increasing, stronger expected persuasion raises disclosure for states of the world that move the receiver toward the sender’s position and lowers it for those that move the receiver away. Appendix G develops it formally as Proposition 3, with the boundary case $w(\Delta\pi) = \text{sign}(\Delta\pi)$ marking the threshold at which the sender responds only to the direction of persuasion and not to its magnitude.

4 Experiment 2

Experiment 1’s result cannot disentangle the channels proposed in the conceptual framework. The baseline alignment gap is consistent with an expressive motive (θ_E), a persuasive one (θ_P), or both, because the facts senders find congenial are also the facts they expect to move the receiver toward them. The increase under political action admits the same ambiguity: it could reflect a distinct instrumental motive (θ_I) that switches on

when the receiver’s beliefs carry consequences, or an expressive motive that strengthens when the stakes are real.

Experiment 2 separates them. Receivers hold known and varying policy positions, so a fact’s congeniality and its expected persuasive effect no longer move together. The three returns to disclosure are elicited directly and with incentives, rather than inferred from choices. Congeniality is measured on facts from a domain the sender was not assigned, which keeps it clear of the provision decisions it is used to explain. Transmission is probabilistic, so a receiver who sees nothing cannot tell withholding from a failed draw, and disclosure does not unravel. The experiment also runs in two policy domains rather than one. Together these changes let me estimate each channel separately and test whether the expressive or persuasive motive responds to political action—the reading Experiment 1 could not rule out.

4.1 Design, Treatment, and Hypotheses

I run a pre-registered online experiment with $N = 1,548$ German adults recruited on Bilendi, quota-sampled on age, gender, and state. Each sender is assigned to one of two policy domains, immigration or climate.

Experimental Timeline The experiment runs in three parts followed by a post-questionnaire (Figure 3). Part 1 collects demographics and both preference measures. Part 2 is the provision task: senders decide whether to share each of five assigned facts, one decision at a time, each directed at a receiver whose policy stance is visible; an attention check follows. Part 3 elicits incentivized beliefs about how each fact corrects the receiver’s factual belief and shifts their policy position. The post-questionnaire asks senders to rate two facts from the domain they were not assigned. Two timing choices differ from Experiment 1. Provision precedes elicitation, so the elicitation cannot prime the provision decision; the quadratic scoring rule addresses the concern that senders report beliefs to justify choices already made. Comprehension checks precede the provision task to account for the more complex design.

Facts I select six facts per domain, balanced for credibility and relevance in a pre-registered pre-validation survey ($N = 300$, AsPredicted #290989; Appendix A.3). Each is presented as a binary, invertible comparison, as in Experiment 1, and sourced from German statistical offices and research institutes. Table 4 lists the twelve statements. All participants are told at the outset that every fact is true and sourced, and receivers are told the same. Each sender is assigned five of the six facts in their domain—two in each political direction plus one drawn at random—so that alignment varies within sender by construction.




Figure 3: Experimental timeline for Experiment 2.

Policy Preferences I measure preferences two ways. Senders place themselves on a five-point policy scale, and they split €250 between the pro-policy and the anti-policy advocacy organization for their domain in a real donation decision, one of which is implemented per session. For immigration the organizations are ProAsyl (advocates open asylum and migration policy) and Republik21 (advocates more restrictive policy); for climate, BUND (supports ambitious climate targets) and Stiftung Marktwirtschaft (emphasizes economic costs of the transition). Senders see each organization’s position in the organization’s own words before deciding.

Receivers Each sender faces five receivers in sequence, one at each position on the five-point scale. Before every decision the sender sees the receiver’s stance stated in full, and knows the receiver has been told the sender’s position in turn. This mutual visibility generates within-sender variation in ideological distance. Two probabilities govern what the receiver experiences, and senders are told both. A sender is matched to a real receiver in at least four of ten sessions. Conditional on a match, a decision to share is implemented with 50% probability: if the draw fails the receiver sees nothing, whatever the sender chose. A receiver who sees nothing therefore cannot tell withholding from a failed draw.

Elicitations Part 3 elicits the two returns that vary with the receiver. For each receiver–fact pair, senders predict the share of receivers holding that stance who would identify

Table 4: Facts Used in Experiments 2 and 3

Direction	Fact
<i>Immigration policy</i>	
Pro	The share of people with a migration background in shortage-occupation fields is <i>higher</i> than their population share.
Pro	The absolute increase in employed foreign nationals between 2013 and 2023 was <i>larger</i> than for German nationals.
Pro	<i>More</i> than half of the persons attending a general integration course for the first time complete it without interruption through the final module.
Anti	Among persons suspected of violent crimes, the share of non-German nationals is <i>higher</i> than their population share.
Anti	The share of persons receiving basic income support (<i>Bürgergeld</i>) is <i>higher</i> among foreign nationals than among German nationals.
Anti	Germany records <i>more</i> persons without a valid residence permit per capita than the EU average.
<i>Climate policy</i>	
Pro	In German electricity generation in 2023 and 2024, the share of renewable energies was <i>higher</i> than the share of fossil fuels.
Pro	The cost of generating one kilowatt-hour of electricity from new, ground-mounted solar installations is <i>lower</i> than from new gas-fired power plants.
Pro	The percentage increase in renewable energy jobs in Germany between 2022 and 2023 was <i>higher</i> than the global average.
Anti	Grid fees for household customers are expected to be <i>more</i> than double their current level by 2045 due to grid expansion.
Anti	There are <i>fewer</i> companies that view the effects of the energy transition positively than companies that view them negatively.
Anti	At higher shares of solar power, <i>more</i> price fluctuations are observed in European electricity markets.

Notes: Verbatim claims as shown to senders in English translation (correct answer in *italics*). Direction (Pro) = supports more-permissive immigration or more-ambitious climate policy; (Anti) = supports more-restrictive or more-cautious policy. Pool of six facts per topic, all verified true from German statistical offices and research institutes. Each sender draws five facts: two pro-policy, two anti-policy, and one random tiebreaker. Topics assigned between subjects; all senders see one topic. Pre-validation details in Appendix A.3.

the fact correctly with and without disclosure. They also predict the position the receiver will take after seeing the fact, and the position the receiver would take without it. Both predictions are scored against matched receivers' realized outcomes under a quadratic scoring rule. The third return is elicited in the post-questionnaire: senders rate two facts from the domain they were *not* assigned—one hypothesized pro-policy and one hypothesized anti-policy—answering whether the information “*shows that your opinion is correct,*” and separately whether they are convinced it is true. The instructions ask senders to judge whether the fact validates their position, not whether their position is correct in the abstract, and not whether the fact would convince anyone else.

Treatments and Hypotheses Senders are assigned between subjects to CONTROL or ACTION. In both arms receivers restate their policy position after possibly receiving a fact.²⁴ In the ACTION arm receivers additionally make a real €250 donation decision between the same two organizations, one of which is implemented per session; senders know this when they decide.

I pre-registered three hypotheses with respect to the treatments, mirroring Experiment 1:

Hypothesis 2.0 (Baseline provision). *Provision rates for unaligned facts exceed the indifference benchmark of 50% in each treatment-topic cell.*

Hypothesis 2.1 (Selective provision). *Senders are more likely to provide a fact aligned with their own political position than one that is not, holding fact and topic fixed.*

Hypothesis 2.2 (Action amplification). *The alignment gap in provision is larger when the receiver makes a real political donation (ACTION) than when the receiver only restates their opinion (CONTROL).*

4.2 Sample, Incentives, and Protocol

Data collection ran from May 27 to June 4, 2026 on Bilendi with quota-sampling for age, gender, and state to approximate the German adult population. (See Appendix Table C.34 for details.) This study was approved by the Ethics Council of the Max Planck Society under the Generalized Approval of Experiments Following the Protocol that is Standard in Experimental Economics (approval no. 2018_3 Renewal 2024_28) and was pre-registered on AsPredicted (No. 293252, an update of No. 292898). Senders took a median completion time of 17 minutes and earned €3 on average. The experiment was programmed in oTree (Chen et al. 2016).

4.3 Policy Preferences and Beliefs

Policy preferences. Figure C.4 reports stated and revealed preference distributions for each topic. For immigration, 66.1% of senders state a policy preference for more restricted immigration and 26.3% donate exclusively to the organization advocating more restrictions. The correlation between the two immigration measures is strong ($r = 0.63$, $p < 0.001$). For climate, a majority states a preference for increased national climate action (52.8%), while the plurality of donations leans toward the organization advocating fewer and more market-based restrictions. The within-person ranking is nonetheless strongly correlated ($r = 0.50$, $p < 0.001$). As pre-registered,

²⁴Both arms include the restatement so that ACTION shifts the political consequentiality of the receiver's beliefs, not the salience of the decision or how senders construe it.

I combine the two measures, each normalized to $[-1, +1]$ with $+1$ the most pro-policy position: revealed preference $\text{Pref}_i^{\text{rev}} = \frac{\text{Donation}_i^{\text{pro}} - 125}{125}$ from the €250 split, stated preference $\text{Pref}_i^{\text{st}} = \frac{3 - \text{Stance}_i}{2}$ from the five-point Likert stance, and their average $\text{Pref}_i = (\text{Pref}_i^{\text{rev}} + \text{Pref}_i^{\text{st}})/2$.²⁵ Each fact carries a pre-validated policy direction, $\text{Dir}_j = +1$ for pro-policy and -1 for anti-policy. Positional alignment is their product, $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, +1]$. Its sign records whether the fact favors the sender’s side, and its magnitude is the sender’s measured preference intensity. Preferences are distributed across the full scale within each topic, allowing for identification of selectivity.

Senders’ fact congeniality rating. The cross-domain rating task asks senders whether two facts from the domain they were *not* assigned “show that your view is correct.” Pooling these ratings gives, for each fact and each stance on that fact’s domain, the share of like-minded senders who regard the fact as validating. I transform this share to $\hat{\alpha}_{ij} = \frac{\text{share_yes}_{ij} - 0.5}{0.5} \in [-1, +1]$, where $+1$ denotes universal validation. Because every rating comes from a sender who encountered the fact outside their own assigned domain, the measure is not built from the provision decisions it is used to explain. For sender i , $\hat{\alpha}_{ij}$ is therefore the average judgment of other participants holding i ’s stated position on that fact’s domain. Using it as i ’s own expressive return assumes that the stance–fact mapping estimated among cross-domain raters transports to the senders deciding on that domain.

Positional alignment and the congeniality index operationalize the same expressive channel, but they impose different structure on it. Alignment crosses a direction fixed in the pre-validation with the sender’s measured policy preference, so one fact carries one direction for every sender. The congeniality index instead aggregates validation judgments collected in the main study from cross-domain raters at the sender’s stated position, so the perceived congeniality of a single fact varies across stance–fact cells. Neither is the focal sender’s own rating of the fact they decide on: both are assigned proxies, but the second captures fact-specific variation among facts carrying the same pre-validated direction, conditional on stated position. Their correlation is 0.61.

The measure takes one value per fact and stance, sixty cells in all, with a realized range of $[-0.43, +0.83]$. Each sender’s five facts draw on average 4.9 distinct values, and the gap between a sender’s highest and lowest is 0.68 on average; that within-sender spread is the variation the expressive channel is estimated from. Congeniality tracks the fact’s political direction: senders at the restrictive pole rate anti-direction facts at $+0.79$ and pro-direction facts at $+0.27$, and the ordering reverses at the opposite pole. Senders at the status quo position record substantial variation as well ($+0.13$ on

²⁵The registration prints the stated component with the opposite sign. Because the stance scale runs from 1 at the most pro-policy position to 5 at the most anti-policy one, the transformation $(3 - \text{Stance}_i)/2$ is what delivers the registered orientation, $+1$ for the most pro-policy sender.

anti-direction and +0.44 on pro-direction facts). Positional alignment, which rests on the composite of the two preference measures, is zero for 52 senders, only 39 of whom place themselves at the status quo.²⁶ Nine of the sixty cells are negative, seven of them belonging to senders who lean rather than commit (stances 2 and 4); senders seldom regard a fact as invalidating their position. Appendix Table C.40 reports the full pattern.

Senders also rate whether they are convinced each fact is true. Trust ratings range from 42% to 93% and follow the same directional gradient ($\beta = +0.130$, $p < 0.001$). Since all twelve facts are true, this alignment-predicted variation in truthfulness judgments is consistent with motivated belief formation.

Senders' beliefs about factual belief correction. Without disclosure, senders expect receivers to answer correctly about half the time (49.3% across fact–receiver pairs), and expect disclosure to raise this to 66.5%, a 17-percentage-point improvement. Panel (A) of Figure 4 shows the distribution of the expected factual belief correction gain $\widehat{\Delta b}_{ijk}$, where k indexes the receiver assigned to sender i 's fact j ; each sender–fact pair faces exactly one receiver, so k is implied by (i, j) and estimation remains at the sender–fact level. The mean is positive for all twelve facts (range 0.12–0.21), climate facts carry higher expected gains than immigration facts (0.19 vs. 0.16), and about 26% of pairs carry a negative expected correction.

This elicitation shows no sign of politically motivated over-rating. Aligned and unaligned facts draw statistically indistinguishable expected corrections (0.173 vs. 0.175; $p = 0.24$, fact FE).²⁷ Expected corrections are largely a property of the sender: adding sender fixed effects raises the adjusted R^2 from 0.01 to 0.41. Within that, neither the fact's alignment with the sender's own position (+0.009, $p = 0.24$), the receiver's stance (joint $p = 0.30$), nor the sender–receiver distance (−0.002, $p = 0.42$) predicts how much correction a sender expects. One regularity remains, directional and small: senders expect a fact to correct more when it points from where the receiver stands toward the sender's side (+0.005 per unit, $p < 0.01$). That is a belief about who is likely to be factually mistaken, not about whose position the sender prefers.

²⁶The other 13 state a position at one pole of the scale and donate against it, so the two measures cancel: the composite records disagreement between them rather than the absence of a side. Conversely, 257 of the 296 senders at the status quo position carry nonzero alignment, at a mean intensity of 0.28 among those 257 (the mean across all 296, counting the zeros, is 0.24). Senders with a zero composite remain in the regressions that use continuous alignment, where they generate no within-sender variation in alignment; they cannot enter contrasts that classify facts as aligned or unaligned. Excluding them moves the alignment coefficient from +0.1278 to +0.1284 (Appendix Table C.29, column 4).

²⁷Splitting by the congeniality rating rather than by positional alignment gives the same null (0.173 vs. 0.169, $p = 0.70$). That split is unbalanced—82 percent of pairs are rated congenial—so it corroborates the result rather than testing it independently.




Figure 4: Elicited returns to disclosure, Experiment 2 (General Population). Distributions of the two incentivized elicitations for positionally aligned facts ($\text{Align}_{ij} > 0$; bars) and unaligned facts ($\text{Align}_{ij} < 0$; black outline), as shares within each group; observations of senders with $\text{Pref}_i = 0$ (260 of 7,740) are omitted from the split. Panel (A): expected factual belief correction $\widehat{\Delta b}_{ijk}$ (0.05-unit bins; group means 0.173 aligned, 0.175 unaligned). Panel (B): expected persuasion $\widehat{\Delta \pi}_{ijk}$ (0.125-unit bins; the bin at $[-0.125, 0)$ contains exact-zero predictions; group means +0.110 aligned, -0.016 unaligned). Shaded regions mark negative values (disclosure expected to move the receiver’s factual belief away from the truth, respectively the receiver away from the sender’s position). Elicitations incentivized via Quadratic Scoring Rule.

Senders’ beliefs about the persuasive effect of disclosure. The persuasion return $\widehat{\Delta \pi}_{ijk} = |\text{Stance}_i - \hat{\pi}_{ijk}^0| - |\text{Stance}_i - \hat{\pi}_{ijk}^1|$ captures expected movement of receiver k ’s stance toward sender i ’s own position from disclosure of fact j ; $\hat{\pi}_{ijk}^0$ and $\hat{\pi}_{ijk}^1$ are the sender’s incentivized expectations of receiver k ’s stance without and with disclosure. Panel (B) of Figure 4 shows its distribution. First, dispersion: the standard deviation is roughly 0.9 against a mean of 0.05, and about 17% of pairs are exact zeros. What characterizes senders’ persuasion beliefs is not their level, which is near zero, but their spread.

Second, these expectations track the design. Senders expect the receiver to move in the fact’s pre-validated direction in 58% of cases—far from mechanical, but far from a coin flip.²⁸ What governs the expected magnitude is the receiver’s position *relative*

²⁸A direction-free component is present but small: the toward-sender drift, which ignores the fact’s direction entirely, is +0.072 ($p < 0.001$) among receivers away from the scale endpoints, against a ten-fold difference in expected movement between facts pointing toward and away from the sender’s side

to the fact rather than the receiver's position as such: expected movement rises with the room the receiver has to move in that direction (+0.056 per scale point, $p < 0.01$, sender and fact fixed effects), while receiver-position fixed effects are jointly indistinguishable from zero ($p = 0.29$). The sender's own position enters through the same geometry—the signed sender–receiver gap interacted with the fact's direction predicts the persuasion return at +0.039 ($p < 0.001$), unchanged when receiver position is absorbed (Appendix Table C.39). Senders do not discount for their own credibility: expected movement rises with ideological distance rather than falling.

Third, and as a consequence, the persuasion return depends on a fact's positional alignment where the belief-correction return does not. Alignment is the fact's direction signed by the sender's own preference, so the two regularities just described imply an aligned/unaligned gap arithmetically. Senders expect aligned facts to move receivers 0.11 toward their own position and unaligned facts slightly away ($\beta = +0.111$, $p < 0.001$, fact FE). That the belief-correction elicitation shows no comparable dependence (0.173 versus 0.175, $p = 0.24$) argues against motivated shading as the source—a sender inclined to rate aligned facts favorably would do so on both. The dependence is modest: the raw gap is 0.14 of a standard deviation, and senders expect movement toward their own position for 46% of aligned facts against 38% of unaligned ones.

Separability of the elicited returns. Whether a fact validates the sender's own position and whether it moves a particular receiver toward that position are distinct questions, and senders' answers come apart: they predict movement counter to a fact's own political direction for 35% of pairs. Empirically the three returns are close to unrelated—pairwise correlations of -0.01 , $+0.10$, and $+0.03$, with within-fact variance inflation factors of 1.01 or below. The design generates this separation. Nor does the ACTION randomization sort senders on any of the three returns or on policy preferences: no balance test rejects, and the distributions match in dispersion as well as in mean.

4.4 Treatment Effects

Senders provide unaligned facts well above the indifference benchmark. Across the four treatment–topic cells, provision of unaligned facts ranges from 55.5% to 61.2%, and every cell rejects the 50% benchmark (one-sample t -tests, all $p < 0.001$), confirming Hypothesis 2.0. Overall provision is 63.9%.

Result 2.0 (Baseline provision). *Senders provide unaligned facts at rates significantly above the 50% indifference benchmark in all four treatment–topic cells ($p < 0.001$).*

Table 5 reports the pre-registered tests. I estimate

(+0.020 versus +0.197).




Figure 5: Alignment bias, Experiments 2 and 3. The bias is the within-sender difference between provision rates for aligned and unaligned facts, averaged across senders. Panel (A): general population ($N = 1,496$ senders; Bilendi panel, Germany, 2026). Panel (B): professional journalists ($N = 297$), who face the ACTION arm only. Sample: senders with a directional policy preference ($\text{Pref}_i \neq 0$), so that alignment is defined; 52 general-population and 10 journalist senders are excluded. Alignment: $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, +1]$. Error bars are standard errors of the sender-level mean. Both panels use the same vertical scale.

$$y_{ij} = \beta_0 + \beta_1 \text{Align}_{i,j} + \beta_2 \text{ACTION}_i + \beta_3 (\text{Align}_{i,j} \times \text{ACTION}_i) + \mathbf{S}_j + \gamma_3 \mathbf{X}_i + \varepsilon_{i,j}, \quad (4)$$

with $\text{ACTION}_i \in \{0, 1\}$ an indicator equal to one if sender i is assigned to the ACTION arm. The unit of observation is a sender–fact pair; standard errors are clustered at the sender level throughout. β_1 identifies baseline alignment-based selectivity (Hypothesis 2.1) and β_3 whether political action amplifies it (Hypothesis 2.2). Column (1) includes no controls, column (2) adds fact fixed effects $\mathbf{S}_j$, and column (3)—the pre-registered preferred specification—adds demographic and political controls $\mathbf{X}_i$. The confirmatory tests use positional alignment Align_{ij} (not the elicited returns); H1 ($\beta_1 > 0$) and H2 ($\beta_3 > 0$) are the pre-registered one-sided directional tests.

Senders share aligned facts at higher rates in both arms. Provision runs 68.2% on aligned facts against 56.6% on unaligned facts in CONTROL, and 71.3% against 58.5% in ACTION. The alignment bias is 11.6 pp in CONTROL and 12.7 pp in ACTION; see Panel (A) of Figure 5. The pre-registered coefficient is +0.128 per unit of alignment ($p < 0.001$, one-sided, column (3)), confirming Hypothesis 2.1. Because alignment runs from -1 to $+1$, flipping a fact from unaligned to aligned moves it by twice the sender’s preference intensity, so under the linear specification the gap is exactly the coefficient itself for a sender at half the maximum intensity. That is close to where the data sit: the gap is 13.3 pp for the median sender, whose intensity is 0.52, and 25.6 pp for a sender whose policy preference is entirely one-sided.²⁹

²⁹This scaling with preference intensity is imposed by the $\text{Pref}_i \times \text{Dir}_j$ construction rather than esti-

Result 2.1 (Selective provision). *The alignment bias is 11.6 pp in CONTROL and 12.7 pp in ACTION; the pre-registered alignment coefficient is +0.128 per unit ($p < 0.001$, one-sided, column (3)).*

The gap is not motivated credibility in disguise. Holding senders' perceived credibility of the facts fixed, it is +11.7 pp and remains significant ($p < 0.001$); I treat this as robustness rather than as the preferred specification, because alignment itself shifts perceived credibility.³⁰

The pre-registered amplification test fails. The Align $\times$ ACTION interaction is +0.014 (one-sided $p = 0.21$, column (3)); I do not confirm Hypothesis 2.2.

Result 2.2 (No action amplification). *The alignment gap is not larger when the receiver takes a downstream political action ($\beta = +0.014$, one-sided $p = 0.21$, column (3)).*

The failure is a property of the proxy rather than of the treatment. Positional alignment measures congeniality: within facts it correlates 0.60 with the elicited expressive return and 0.08 with the elicited persuasion return. If political action operates on the persuasive channel rather than the expressive one, this interaction can transmit almost none of it. I therefore turn to the pre-registered secondary analysis, which estimates the channels directly.

4.5 Separating the Motives

Unlike in Experiment 1, Experiment 2 builds a separate empirical counterpart to each of the model's three belief arguments rather than inferring them from revealed preference: sender-specific forecasts of factual belief correction and of receiver movement, and a congeniality index drawn from cross-domain raters at the sender's own stance. None of the three is assigned, so the coefficients below are within-sender, within-fact associations with provision. The one estimate that rests on randomization is the ACTION interaction. I estimate a linear probability model at the sender-fact level, with standard errors clustered at the sender level throughout:

$$y_{ij} = \beta_0 + \beta_E \alpha_{ij} + \beta_P \Delta\pi_{ijk} + \beta_2 \text{ACTION}_i + \beta_A (\Delta\pi_{ijk} \times \text{ACTION}_i) + \beta_B \Delta b_{ijk} + \mathbf{S}_j + \varepsilon_{ij}, \quad (5)$$

mated; a specification replacing alignment with a binary aligned/unaligned indicator gives an average gap of 12.1 pp.

³⁰The pre-registration specifies a perceived believability control alongside fact fixed effects in column (2). That control was never collected: the cross-domain rating module was shortened after pre-registration to reduce survey length and dropout, a deviation documented in the post-pre-registration decision log. The measure used here is the closest available instrument—senders' cross-domain credibility ratings, aggregated to the fact-by-stance cell—not the individual believability item the pre-registration named.

Table 5: Alignment-Based Selective Provision — Experiment 2 (General Population)

	Probability of Fact Provision		
	(1)	(2)	(3)
Align	0.113*** (0.013)	0.127*** (0.013)	0.128*** (0.013)
Action	0.026** (0.013)	0.026** (0.013)	0.023* (0.012)
Align $\times$ Action	0.011 (0.018)	0.014 (0.018)	0.014 (0.018)
Clustered SE	Sender	Sender	Sender
Fact FE	No	Yes	Yes
Controls	No	No	Yes
Num. obs.	7740	7740	7715
R ²	0.025	0.032	0.045
Adj. R ²	0.025	0.030	0.039

Sample: $N = 1,548$ Experiment 2 senders (Bilendi panel, Germany, 2026), 7,740 sender-fact observations. Each sender makes five provision decisions, one per fact. Column (3) restricts to non-missing demographic controls (7,715 observations). *Outcome:* binary provision indicator ($y_{ij} \in \{0,1\}$). *Alignment:* $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1,1]$ (pre-registered positional alignment measure). *Fact FE:* indicator for each of the 12 facts. *Controls:* age, gender, education, Bundesland, voting intention, left-right self-placement. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$. *Pre-registered tests on column (3):* H1 (Alignment > 0 , one-sided $p < 0.001$, confirmed); H2 (Alignment $\times$ Action > 0 , one-sided $p = 0.21$, not confirmed).

where $y_{ij} \in \{0,1\}$ is sender i 's provision decision for fact j and $\mathbf{S}_j$ are fact fixed effects. The fact fixed effects absorb all time-invariant fact-level characteristics, including political direction, topic-level variation, and fact-specific shareability. Within-sender variation across the five facts assigned to each sender supplies the estimating variation.³¹ Columns (1) and (2) isolate the expressive and persuasion channels in turn; column (3) includes both jointly. Column (4) is the specification implied by the utility function in the conceptual framework: Expression main effect plus the instrumental Persuasion $\times$ ACTION interaction, omitting Expression $\times$ ACTION, which is reported in columns (1) and (3).

The expressive return predicts provision. Moving from a fact that half of like-minded senders regard as validating their position to one that all of them do is associated with a provision rate 14.9 pp higher ($p < 0.001$, column (4)). The receiver learns the sender's position whichever fact arrives. Passing on a congenial fact therefore conveys nothing further about where the sender stands, which rules out reputational position-signaling as a source of this association.

The expressive channel does not respond to political action: the Expression $\times$ AC-

³¹Because identification rests on within-sender variation across facts rather than across experimental conditions, experimenter-demand effects are limited (de Quidt et al. 2018).

Table 6: Action Decomposition: Expression and Persuasion — Experiment 2 (General Population)

	Probability of Fact Provision			
	(1)	(2)	(3)	(4)
Expression	0.137*** (0.028)		0.138*** (0.028)	0.149*** (0.020)
Expression $\times$ Action	0.027 (0.037)		0.020 (0.037)	
Persuasion		0.001 (0.009)	−0.004 (0.009)	−0.004 (0.009)
Persuasion $\times$ Action		0.024* (0.013)	0.024* (0.013)	0.025* (0.013)
Action	0.016 (0.019)	0.025* (0.013)	0.018 (0.019)	0.025** (0.013)
Factual belief correction gain	0.034** (0.017)	0.035** (0.017)	0.034** (0.017)	0.034** (0.017)
Clustered SE	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes
Spec	(1)	(2)	(3)	(4)
Num. obs.	7740	7740	7740	7740
R ²	0.015	0.007	0.016	0.015
Adj. R ²	0.013	0.005	0.013	0.013

Sample: $N = 1,548$ Experiment 2 senders (Bilendi panel, Germany, 2026), 7,740 sender-fact observations; requires non-missing expressive identity, projective persuasion, and factual belief correction gain measures. *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating a_{ij} , elicited post-provision on a different topic. *Persuasion:* incentivized expected preference shift $\Delta\pi_{ijk}$ toward the sender’s position (positional, preference-free reading). *Factual belief correction gain:* incentivized expected correction of the receiver’s factual belief Δb_{ijk} . *Receiver distance (dist):* reported separately in Table C.22 (pre-registered heterogeneity). Column (1): expressive motive and its action interaction. Column (2): the persuasion channel—projective main effect and instrumental action interaction. Column (3): both interactions jointly. Column (4): expressive main effect plus the instrumental (persuasion $\times$ action) interaction. All specifications include fact fixed effects and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

TION interaction is $+0.020$ ($p = 0.59$, column (3)). This bears on Experiment 1. There, the alignment gap increases under political action, and that pattern admits two readings—a distinct instrumental motive that switches on when the receiver acts, or an expressive motive that strengthens when the stakes are real. Experiment 2 tests the second directly and finds no response, which suggests the Experiment 1 amplification operated through the persuasive rather than the expressive component.

The factual belief correction gain Δb_{ijk} predicts provision without tilting it toward congenial facts ($\beta = +0.034$, $p = 0.05$, column (4)), with a positive sign for all twelve facts.³²

Persuasion ($\Delta\pi_{ijk}$) does not predict provision absent consequences, and predicts it once consequences exist. In the CONTROL arm the persuasion return is unrelated to provision ($\beta = -0.004$, $p = 0.65$, column (4)). Under ACTION it becomes a selection

³²Within-arm estimates are directionally consistent ($\beta = +0.053$ in CONTROL, $\beta = +0.018$ in ACTION), and the factual belief correction $\times$ ACTION interaction does not reach significance ($\beta = -0.036$, $p = 0.30$).

criterion: the interaction is $+0.025$ ($p = 0.052$).³³ The treatment changes how senders act on their beliefs, not the beliefs themselves: it shifts neither the mean nor the dispersion of any elicited return, so the activation is not a byproduct of the arm making persuasion more salient, nor of senders reconstructing their beliefs to match choices already made. Appendix Figure C.5 shows the pattern by arm.

This is also why the pre-registered amplification test failed. If political action operates on the persuasion channel, positional alignment carries only $r = 0.084$ of that channel within facts. At the estimated interaction, the implied $\text{Align} \times \text{ACTION}$ coefficient is $+0.003$, against a minimum detectable effect of 0.049 for that test, or 6.5 percent of it. Bringing the implied coefficient to the detection threshold would require a correlation of 1.29 between the two measures. The pre-registered test could not have detected the instrumental motive at the magnitude the data support. This is why the returns are elicited.

Substituting the pre-registered positional measure for the elicited expressive return changes nothing (Appendix Table C.27): positional alignment enters at $+0.133$, the belief-correction coefficient is unmoved at $+0.034$, and the $\text{Persuasion} \times \text{ACTION}$ interaction stands at $+0.025$. The two expressive measures carry the same channel, and neither carries persuasion—each correlates with the persuasion return at 0.08 and 0.10 within facts. The congeniality rating correlates marginally more strongly of the two, which does not contaminate the channel estimates, since both enter jointly here and their shared variance is absorbed; it matters only for the single-proxy interaction test of the previous subsection, where nothing absorbs it.

Result 2.3 (Motives). *Expression predicts provision ($\beta = +0.149$, $p < 0.001$) and does not respond to political action ($\text{Expression} \times \text{ACTION} = +0.020$, $p = 0.59$). The expected factual belief correction predicts provision ($\beta = +0.034$, $p = 0.05$) without tilting it toward congenial facts. The persuasion return does not predict provision in the CONTROL arm ($\beta = -0.004$, $p = 0.65$) and does under ACTION ($\beta = +0.025$, $p = 0.052$). Column (4).*

The elicited returns cohere with what senders say about their own decisions: senders who independently state a persuasion motive in open text select more steeply on the persuasion return ($+0.107$, $p < 0.001$, Appendix Table B.21), on the subsample giving a codable reason.³⁴

³³Inference on this interaction is mixed and I report it in full. Logit and probit give $p = 0.049$ and $p = 0.051$. Because ACTION is randomized at the sender level, the design-based test matching the assignment mechanism is randomization inference over sender-level reassignment: one-sided $p = 0.024$ (5,000 permutations). An exact wild-cluster bootstrap clustering instead on the twelve facts gives $p = 0.108$. With only twelve fact clusters, this more conservative inferential result tests robustness to fact-level dependence rather than corroborating the sender-level inference. Across 84 non-redundant specifications the interaction is positive in all 84, though only 1 percent reach conventional significance (median $p = 0.129$).

³⁴27 percent of senders give a reason that refers to the sharing decision, and those who do are somewhat more issue-extreme; the comparison conditions on giving a codable reason.

The estimates are robust to how the sample is drawn and to the order facts appear in. Every estimate keeps its sign under alternative sample definitions (Appendix Table C.35); excluding attention-check failers, who complete the study in twelve minutes against a median of seventeen, strengthens the expressive and belief-correction channels and leaves the action-conditional interaction at +0.026. Fact order is randomized, and although provision declines mildly over the five positions (−0.8 pp per position, $p = 0.037$), adding presentation-position fixed effects leaves every estimate in the table unchanged to the third decimal.

A direction-only analog $\delta_{ijk}^{\text{dir}} \equiv -\Delta\text{stance}_{ijk} \cdot \text{sign}(\text{pref}_i)$ scales expected receiver movement by the sign rather than the magnitude of the sender’s preference. Appendix C.5 compares the two proxies and their binary codings (Appendix Tables C.37 and C.38): the instrumental interaction survives only under codings that carry magnitude, and not under binary ones.

4.6 Heterogeneity

I pre-registered three dimensions of heterogeneity as exploratory: sender–receiver ideological distance, sender political extremity, and policy domain. Appendix Table C.22 covers the treatment-effects contrast, Appendix Tables C.23 and C.24 the motives.³⁵

Provision falls with the ideological distance between sender and receiver. Distance is $\text{dist}_{ij} \equiv |\text{Stance}_i - \text{Stance}_{ij}^R|$, the absolute gap on the five-point scale between sender i ’s self-placement and the revealed position $\text{Stance}_{ij}^R \in \{1, \dots, 5\}$ of the receiver paired with that decision. Each Likert unit of distance reduces provision by 9.4 pp ($p < 0.001$). Political action offsets part of the penalty ($\text{dist} \times \text{ACTION}$, +0.019, $p = 0.041$). Neither selection channel varies with distance on its own: the alignment slope is flat (+0.001), as are the expressive and persuasion slopes.³⁶ When a distant receiver will act politically, senders select more strongly on the persuasion return ($\text{Persuasion} \times \text{ACTION} \times \text{dist}$, +0.017, $p = 0.048$), consistent with strategic targeting. The pre-registered proxy shows the same triple in the same direction (+0.020, $p = 0.097$).

Political extremity shifts the level of provision but not its slope. More polarized senders provide slightly more overall (+0.010, $p < 0.10$); relative to center voters, supporters of extreme parties (AfD, Linke, BSW) provide 4.5 pp less and non-affiliates 9.3 pp less ($p = 0.021$ and $p < 0.001$). No group displays a different alignment slope,

³⁵The specifications estimate twenty interaction terms; Benjamini–Hochberg q -values within that family appear in the table notes. Nine of the twenty fall below $p = 0.10$ against two expected under the null, and none survives false-discovery control (smallest $q = 0.13$). I read the patterns in this subsection as suggestive.

³⁶Senders with extreme preferences face more high-distance receiver cells by construction, potentially conflating behavioral effects with a mechanical alignment–distance correlation. Within-sender variation in distance across the five receivers is the primary identifying variation and is independent of sender extremity.

and polarization does not interact with alignment (+0.007). One channel does vary: the action-conditional persuasion interaction is concentrated among center voters (+0.050, $p = 0.006$) and is not detected among supporters of extreme parties.

The action-conditional channel is confined to one domain. On the elicited persuasion return, the ACTION interaction is +0.038 in immigration against +0.011 in climate. The pre-registered alignment proxy places it in the same domain (+0.050 versus -0.042), so two measures built from different data agree on where it sits. Neither survives false-discovery control, and I take the agreement as suggestive rather than as confirmation of Hypothesis 2.2 in one domain. The expressive channel behaves differently: it is present in both (+0.143 immigration, +0.158 climate). The belief-correction channel is detected in immigration only (+0.053, $p = 0.037$).

Result 2.4 (Heterogeneity). *Sender–receiver ideological distance and sender political extremity shift the level of provision rather than its selectivity: provision declines by 9.4 pp per Likert unit of distance ($p < 0.001$), and no group displays a different alignment slope. The action-conditional persuasion channel is the exception—concentrated among center voters (+0.050, $p = 0.006$) and in the immigration domain (+0.038).*

5 Experiment 3 (Journalists)

Journalists disproportionately decide which facts about contested policy reach the public. Models of media bias explain those decisions through incentives outside the disclosure decision itself—advertising revenue, audience demand for confirmation, and owner preferences (Gentzkow and Shapiro 2006, Mullainathan and Shleifer 2005, Besley and Prat 2006). Bias that persists once these incentives are absent is therefore evidence of intrinsic motives and preferences. Experiment 3 removes them: journalists choose which facts to pass to a single receiver, with no audience, no editor, and no commercial stake in the outcome. One extrinsic incentive survives the design, and it points the other way—German journalists work under the codified obligations of veracity and balance of the *Pressekodex*, enforced by the Presserat (Deutscher Presserat 2025, Steindl et al. 2017).

Bias persists. Experiment 3 applies the ACTION arm of Experiment 2 to a different supplier pool—the same facts, the same receivers, the same alignment measure—so I test the first two hypotheses of intrinsic bias unchanged and compare the two populations.

5.1 Design and Sample

Three features of the Experiment 2 protocol change. Every journalist is matched to a real receiver, where Experiment 2 matched senders in at least four sessions in ten. The

attention check is omitted.³⁷ The post-questionnaire replaces the general-population items with measures of medium type, reach, and editorial autonomy. The experiment was pre-registered separately (AsPredicted No. 292902). Because journalists completed no cross-domain rating task, I take the congeniality index $\hat{\alpha}_{ij}$ from the general population’s ratings, matched on the journalist’s own stance and on the individual fact. Section 4.3 states the assumption this rests on. For journalists it spans a further margin, since the raters differ in occupation as well as in the domain they were assigned.

$N = 307$ journalists made 1,535 provision decisions, above the pre-registered target of 300. They were recruited through the Deutscher Journalisten-Verband and the Deutsche Journalistinnen- und Journalisten-Union mailing lists.³⁸ They took a median of 22 minutes and earned €15 on average.

Against the stratified benchmark of Loosen et al. (2023), the sample matches on the share working for news agencies (3.6% against 3.5%), and differs on who these journalists are and where they publish—younger, more often female, more educated, and concentrated in online outlets (45.5% against 3.1% native-online). A majority hold an editorial role (57.5%); the benchmark reports no comparable breakdown by position, so that share stands without one. Appendix Table D.41 reports the comparison in full.³⁹

The sample leans left, as is documented for German journalists (Steindl et al. 2017, Loosen et al. 2023). On immigration it sits on the other side of the median from the public: 58.6% of journalists favor eased policy, while 66.1% of the general population favors more restriction; on climate, 83.4% favor expanded policy compared with 52.8%. Grüne and Linke are the closest party for 58.3% of the sample, and mean left–right self-placement is 3.2, compared with 5.0. Mean preference intensity $|\text{Pref}_i|$ is 0.69 against 0.54: the same response to alignment produces a wider realized bias in a pool with more one-sided preferences.

5.2 Selective Provision

Provision of unaligned facts is well above indifference in both domains: 65.3% in climate and 67.7% in immigration, both rejecting the 50% benchmark (one-sample t -tests, $p < 0.001$) and confirming Hypothesis 2.0. Overall provision is 76.7%, against 63.9% in the general population.

Provision is selective. Journalists provide 87.0% of aligned facts against 66.6% of

³⁷The professional journalist population is small enough that exclusions would be costly and uninformative.

³⁸One journalist completed the study seven times, identified by the payout details they submitted; only the first participation is retained. This study was approved by the Ethics Council of the Max Planck Society under the same Generalized Approval as Experiment 2.

³⁹This is a convenience sample of German journalists recruited online, not a probability sample. Loosen et al. required eighteen months of multi-source population tracing to assemble their sampling frame.

Table 7: Pre-Registered Treatment Effects: Journalists

	Probability of Fact Provision		
	(1)	(2)	(3)
Alignment	0.141*** (0.015)	0.128*** (0.020)	0.129*** (0.020)
Clustered SE	Journalist	Journalist	Journalist
Fact FE	No	Yes	Yes
Controls	No	No	Yes
Num. obs.	1535	1535	1535
R ²	0.062	0.076	0.085
Adj. R ²	0.062	0.069	0.075

Sample: $N = 307$ Experiment 3 journalists, 1,535 journalist-fact observations. All journalists in the ACTION arm; no between-subject Action treatment variation. *Alignment:* $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, 1]$ (pre-registered expressive proxy). *Controls (col. 3):* age. Standard errors clustered at journalist level. *Pre-registered H1 test, col. (3):* Alignment > 0 , one-sided p reported in text. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

unaligned ones, an alignment bias of 20.4 pp (Figure 5); the corresponding bias is 11.6 and 12.7 pp in the general population's two arms. The pre-registered coefficient is +0.129 per unit of alignment ($p < 0.001$, one-sided, column (3) of Table 7), confirming Hypothesis 2.1. As in the general population, and again under the linear specification, the gap is exactly that coefficient for a journalist at half the maximum preference intensity; it is 19.3 pp for the median journalist, whose intensity is 0.75, and 25.8 pp for one whose policy preference is entirely one-sided.⁴⁰

Result 3.0 (Selective provision among journalists). *Journalists provide 87.0% of aligned facts against 66.6% of unaligned ones, an alignment bias of 20.4 pp; the pre-registered alignment coefficient is +0.129 per unit ($p < 0.001$, one-sided, column (3)).*

The larger alignment bias reflects journalists' preferences rather than a steeper response to congeniality. Their per-unit response is +0.128 against +0.135 (fact fixed effects), and comparing senders at matched preference intensity accounts for much of the difference. This is the same selection rule, applied by a supplier pool with more one-sided preferences.⁴¹

⁴⁰The effect is not an artifact of the linear probability model. Across the 297 journalists who see at least one aligned and one unaligned fact, the mean within-journalist gap between the two provision shares is +20.4 pp; a one-sided Wilcoxon signed-rank test gives $p < 0.001$, and an exact sign test on the 182 journalists whose two shares differ gives 145 positive ($p < 0.001$). Demeaning provision by fact—the non-parametric counterpart of the fact fixed effects—leaves +10.3 pp ($p < 0.001$; within-journalist permutation test, 10,000 draws). The difference reflects journalists' ideological homogeneity: the facts they find congenial are also the facts they collectively provide most, and the fact fixed effects absorb that.

⁴¹Within quartiles of $|\text{Pref}_i|$ the aligned–unaligned gap runs 2.9, 9.4, 17.1, and 31.0 pp in the general

Table 8: Provision Decisions: General Population and Journalists Pooled

	Probability of Fact Provision		
	(1)	(2)	(3)
Align	0.126*** (0.013)	0.123*** (0.013)	0.124*** (0.013)
Action	0.026** (0.013)	0.023* (0.012)	0.023* (0.012)
Align $\times$ Action	0.005 (0.016)	0.008 (0.016)	0.017 (0.017)
Journalist	0.118*** (0.016)	0.097*** (0.019)	0.097*** (0.019)
Align $\times$ Journalist			−0.029 (0.021)
Clustered SE	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes
Controls	No	Yes	Yes
Sample	Pooled	Pooled	Pooled
Num. obs.	9275	9245	9245
R ²	0.046	0.099	0.099
Adj. R ²	0.044	0.094	0.094

Sample: pooled Experiment 2 general-population senders ($N = 1,548$; both arms; 7,740 observations) and Experiment 3 journalists ($N = 307$; ACTION arm only; 1,535 observations). *Outcome:* binary provision indicator. *Align:* positional alignment $\text{Pref}_i \times \text{Dir}_j \in [-1, +1]$. *Journalist:* indicator for Experiment 3 participants (all in the ACTION arm). *Controls* (columns (2)–(3)): sender-receiver ideological distance, age, gender, education, Bundesland, left-right self-placement, and closest-party affiliation—the demographic variables measured in both samples. All specifications include fact fixed effects. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

5.3 Motives and Comparison with the General Population

The expressive channel is present among journalists. I estimate the Experiment 2 channel specification on the journalist sample (Appendix Table D.44): Expression predicts provision at +0.134 ($p = 0.003$), against +0.158 in the general population’s ACTION arm, and the factual belief correction gain enters at +0.076 ($p = 0.055$) against +0.018 ($p = 0.46$). The persuasion return is +0.020 ($p = 0.17$)—the same point estimate as the general population’s ACTION arm (+0.021, $p = 0.02$), estimated on 1,535 decisions rather than 3,965. Experiment 3 is uninformative about the persuasion channel. With no CONTROL arm, θ_I is not separately identified; I recover $\theta_P + \theta_I$.

I now pool the two samples. Journalists differ in the level of provision and on no motive-channel slope. In the pre-registered specification (Table 8) they provide 11.8 pp population; no journalist quartile departs from its counterpart, and a joint test of the four interactions does not reject ($p = 0.27$). In the pooled specification of Table 8 the two alignment slopes differ by −0.029 (standard error 0.021), so the data rule out a journalist slope more than roughly a tenth steeper than the general population’s.

Table 9: Motive Decomposition: General Population and Journalists Pooled

	Probability of Fact Provision			
	(1)	(2)	(3)	(4)
Expression	0.148*** (0.017)	0.149*** (0.017)	0.149*** (0.017)	0.153*** (0.020)
Factual belief correction gain	0.039** (0.016)	0.045*** (0.016)	0.045*** (0.016)	0.039** (0.017)
Persuasion	0.011* (0.006)	0.012** (0.006)	−0.002 (0.009)	−0.002 (0.009)
Persuasion × Action			0.025** (0.012)	0.023* (0.013)
Action			0.023* (0.013)	0.023* (0.013)
Journalist	0.132*** (0.015)	0.103*** (0.019)	0.092*** (0.020)	0.085*** (0.027)
Expression × Journalist				−0.027 (0.043)
Factual belief correction gain × Journalist				0.049 (0.042)
Persuasion × Action × Journalist				0.010 (0.017)
Clustered SE	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes
Controls	No	Yes	Yes	Yes
Sample	Pooled	Pooled	Pooled	Pooled
Num. obs.	9275	9245	9245	9245
R ²	0.027	0.038	0.039	0.040
Adj. R ²	0.025	0.033	0.034	0.033

Sample: pooled Experiment 2 general-population senders ($N = 1,548$; both arms; 7,740 observations) and Experiment 3 journalists ($N = 307$; ACTION arm only; 1,535 observations). *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating $a_{ij} \in [-1, +1]$ from the general-population cross-domain rating cell lookup (merged for journalists under the stance-equivalence assumption). *Factual belief correction gain:* incentivized expected correction of the receiver’s factual belief $\Delta b_{ijk} \in [0, 1]$. *Persuasion:* incentivized expected movement of the receiver toward the sender’s position $\Delta \pi_{ijk}$ (positional reading). *Journalist:* indicator for Experiment 3 participants. Column (3) adds the ACTION main effect and the Persuasion × ACTION interaction, identified from the general population (who vary across arms). Column (4) adds journalist heterogeneity: the Expression and Factual belief correction interactions with the journalist indicator, and the persuasion channel’s heterogeneity as the Persuasion × ACTION × Journalist term; because journalists are observed only in the ACTION arm, the Persuasion × Journalist double is collinear with this triple and omitted. *Controls* (columns (2)–(4)): age, gender, education, Bundesland, left-right self-placement, and closest-party affiliation—the demographic variables measured in both samples. Sender-receiver ideological distance is a control in Table 8 but not here, because it bounds the persuasion return mechanically. All specifications include fact fixed effects. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

more than the general population’s ACTION arm, and 9.7 pp more once demographic and political characteristics are held constant (both $p < 0.001$). The alignment slope does not differ (Align × Journalist = -0.029 , $p = 0.16$). The same holds for the elicited returns (Table 9): the journalist indicator is $+0.085$ ($p = 0.002$), and none of the three motive-channel interactions is statistically distinguishable from zero.⁴²

Result 3.1 (Level, not structure). *Journalists provide 9.7 pp more than the general population’s ACTION arm ($p < 0.001$) and differ on no motive-channel slope. Expression predicts provision ($+0.134$, $p = 0.003$); the persuasion coefficient is too imprecise to speak to that channel.*

Journalists hold different expectations about how information moves receivers, and

⁴²Expression -0.027 ($p = 0.53$); factual belief correction $+0.049$ ($p = 0.25$); the action-conditional persuasion term $+0.010$ ($p = 0.54$).

the difference is one of resolution rather than of kind. The same geometry governs both samples—a fact’s direction interacted with the signed sender–receiver gap—but journalists read it far more sharply: they expect movement in the fact’s pre-validated direction in 77% of cases against 58% in the general population, and the geometry coefficient is roughly three times as large (Appendix Table C.39). What follows from that reading is the entanglement. They expect an aligned fact to move the receiver toward them in 51.3% of cases against 24.2% for an unaligned one, a gap of 27.1 pp where in the general population it is 7.7 pp, and positional alignment correlates with the elicited persuasion return at 0.20 within facts against 0.09. Alignment and expected persuasion are far more entangled in this sample, so the larger alignment bias cannot be read as evidence of a stronger expressive motive alone.

Journalists’ beliefs about factual correction differ in level, not in structure. They expect a fact to raise the share of receivers answering correctly from 44.7 to 76.8 percent, a 32-percentage-point improvement against the general population’s 17, and only 7.9% of their fact–receiver pairs carry a negative expected correction against 25.7%. What predicts those expectations is the same in both samples: the receiver’s position relative to the fact (+0.012 per unit, $p < 0.01$), and neither the fact’s alignment with the journalist’s own position (+0.019, $p = 0.26$) nor the receiver’s stance (joint $p = 0.77$). One difference: journalists expect more correction the further the receiver sits from them (+0.011 per Likert unit, $p = 0.04$), where the general-population slope is flat. They treat ideological distance as a marker of factual misinformation.

Journalists’ own accounts of their decisions line up with the pattern in their choices. A closing question asked for their reasons, and 74 percent gave one that refers to the sharing decision. Classifying those reasons by primary motive, journalists who cite correcting the receiver’s factual beliefs show an alignment bias 11.2 pp lower ($p < 0.01$) and those who cite influencing the receiver’s opinion one 24.0 pp higher ($p < 0.001$); these patterns mirror the results in Experiment 1 (Appendix Table D.43), in the two categories whose coefficients are insensitive to the coding scheme.⁴³ Eleven journalists cite an expressive reason; the coefficient is +22.6 pp and is not distinguishable from zero, so that category cannot be read either way. These reasons are recorded after the provision task and are unincentivized: they document coherence between what journalists did and what they say about it, not the motives themselves.

Agreement with norms of information sharing, elicited after the provision task, predicts how much information journalists and the general population provide. For journalists only, it also predicts selectivity: agreement that one should pass on true infor-

⁴³The two experiments’ open-text responses are classified under different schemes, so I checked which categories that choice moves. Re-coding the Experiment 1 corpus under seven schemes leaves the belief-correction coefficient within -0.074 and -0.110 and the persuasion coefficient within $+0.169$ and $+0.321$, while the expressive coefficient ranges from $+0.243$ to $+0.630$. The comparison drawn here is restricted to the two stable categories; the scheme-sensitive one is the category not read below.

mation contradicting one’s own position is associated with a weaker alignment slope (-0.091 per scale point, $p < 0.001$), and the slope halves from journalists who agree to those who strongly agree without reaching zero ($+0.091$, $p < 0.001$). Appendix Table F.50 reports the battery in full, and Appendix Table D.42 reports the pre-registered exploratory heterogeneity battery.

6 Experiment 4

[PLACEHOLDER — Section opener, two paragraphs: what Experiments 1–3 leave unmeasured on the receiver side, and what this design changes.]

6.1 Design

Every receiver is matched to one sender from Experiment 2 or 3 and inherits a decision that sender actually made. The sender saw one fact and chose whether to pass it to a receiver holding the matched receiver’s policy stance, and that choice stands as it was made: the receiver in this experiment is the receiver that sender was deciding about.

Experimental Timeline The experiment runs in four parts followed by a post-questionnaire (Figure 6). Demographics and the baseline policy position come first, and the match to a sender is claimed at that point. Part 1 collects the receiver’s prior belief about the matched fact, before senders are mentioned. Part 2 covers the instructions and the comprehension checks, and closes with the reception screen; the policy position is asked again immediately after. Part 3 elicits the posterior belief on the same scale as the prior. Part 4 elicits what the receiver believes about how the senders in their cell disclosed. Two placements carry analytical weight. The transmission draw is made at the end of the instructions, so it is fixed before the comprehension checks and before the receiver sees anything. For receivers matched to ACTION-arm senders, the donation sits between the policy position and the posterior belief, placing a real decision between reception and the posterior for those receivers and not for the others.

Reception Disclosure reaches the receiver with probability one half, the same transmission rate the senders were promised. The draw is made for every matched receiver, including those whose sender withheld, so the data distinguish three states: the fact was sent and arrived, the fact was sent and the draw failed, or the fact was withheld. Receivers see only whether something arrived. Each receiver is told the sender’s policy stance and that the sender knew theirs. Each is also told that a coin decides whether any decision is implemented, so the instructions name both causes of silence, and the




Figure 6: Experimental timeline for Experiment 4.

two are observationally identical.⁴⁴ The sender’s disclosure decision is not randomized; the transmission draw is, which makes receipt against failed transmission the one comparison in this experiment that rests on randomization alone.

Belief accuracy Receivers report their belief about the fact twice, once before any mention of senders and once after reception. Each fact completes a sentence with one of two words, one of which is correct, and receivers report on a single slider the probability they attach to each completion; the second elicitation uses the same scale, with the earlier answer neither shown nor used to position the slider.⁴⁵ Both elicitations are

⁴⁴Two comprehension questions on exactly this point precede reception, and each permits two bonus-eligible attempts.

⁴⁵The two words appear on the same side of the slider throughout for a given participant, and which side is randomized across participants. The stimulus set is lopsided—the correct answer is the increase

incentivized by a quadratic scoring rule.

I score a belief by its squared distance from the realized state, $L = (b - \omega)^2$, a strictly proper rule that is uniquely minimized in expectation by the true posterior and therefore rewards calibration.⁴⁶ A received message identifies the state, so the loss borne after receipt reflects the receiver's failure to adopt what they were shown. Silence leaves the state unresolved; loss after silence carries both what the environment withheld and how the receiver read it.

Beliefs about senders After all other outcomes, receivers estimate how many of a set of real senders passed the fact on. The set is drawn from the same cell of sender stance, receiver stance, topic, and fact that produced their own match, so the quantity they estimate is one the sender data measures directly. They estimate it twice, once for each possible correct answer, because the object of interest is not how often senders share but whether they share differently depending on what the fact says.⁴⁷ The estimate is incentivized against one real decision drawn from that cell.

Policy position and donation The matched-topic policy item is asked again after reception, on the scale used at intake, without reference to the earlier answer. Receivers matched to senders in the ACTION arm additionally split €250 between two organizations, pooled with the sender donations so that one decision is drawn and implemented.

Outcomes I measure four quantities. Belief accuracy is the squared distance defined above, taken after reception. The policy position is the matched-topic item, signed by the fact's direction so that a positive value denotes movement toward the position the fact supports. Perceived selectivity is the difference between a receiver's two estimates of sender behavior. Some comparisons are at the level of environments, which I index by a cell g collecting topic, fact, sender stance, and receiver stance, with σ_g the realized disclosure rate in that cell among Experiment 2 and 3 senders and λ_g its share among matched receivers. The pre-registration states the specifications and the exclusions; it registers no directional prediction, and the sign of each quantity is an empirical question.

word for ten of twelve facts—so under any fixed layout display position would proxy the truth. The stored quantity is the probability on a named word and does not depend on the layout.

⁴⁶An absolute-distance loss is minimized by the median, so a calibrated probability does not minimize it; I therefore use squared distance.

⁴⁷The question asks for counts: it concerns a specific set of real people, and respondents handle counts more reliably than shares. Receivers who received the fact are asked about a different fact of the same topic at the same stances, since being shown the answer would collapse one of the two branches.

6.2 Sample, Incentives, and Protocol

[PLACEHOLDER — bracketed values are filled after fielding.] Data collection ran from [dates] on Bilendi with quota-sampling for age, gender, and state to approximate the German adult population, yielding $N = [N]$ receivers. The design commits to matching every journalist from Experiment 3, as promised to them; the realized shares are [PLACEHOLDER — realized journalist matching share] for journalists and [share] for general-population senders. This study was approved by the Ethics Council of the Max Planck Society under the Generalized Approval of Experiments Following the Protocol that is Standard in Experimental Economics (approval no. 2018_3 – Renewal 2024_28) and was pre-registered separately on AsPredicted (No. [number]). Receivers took a median completion time of [time] and earned [amount] on average.

The matching queue also determines what happens when a receiver drops out. A sender is consumed at the transmission draw, so a receiver who leaves before the draw returns that sender to the queue, while one who leaves after it does not. Receivers removed in post-hoc cleaning do not release their sender, since the reception event occurred. [PLACEHOLDER — progression from the draw through reception, position, donation, posterior and completion by transmission outcome, and the balance of the draw on baseline characteristics.]

6.3 The Effect of Information vs. Silence on Beliefs

Table ?? reports what receiving the fact does to belief accuracy. Among the decisions in which the sender chose to disclose, the transmission draw alone decides whether the fact arrives, so receipt against failed transmission is randomized. Receipt moves posterior epistemic loss by [PENDING DATA] (column 1), conditioning on the prior loss and fact fixed effects, with a randomization-inference p -value of [PENDING DATA] from re-drawing the coin under its own assignment mechanism.

Result 4.0 (The price of a fact). *Receiving one fact moves posterior epistemic loss by [PENDING DATA] relative to an otherwise identical receiver whose transmission failed.*

The loss remaining after receipt is [PENDING DATA]. A receiver who adopts a true and sourced fact they have just been shown bears none, so this level is itself the departure, and nothing about the disclosure rule enters it.

I measure adoption as the share of the available stated-belief gap that a receiver closes. The headline specification allows that share to vary linearly with a continuous fact-by-stance congeniality index taken from the Experiment 2 cross-domain rating cells; the index is assigned from prior raters rather than elicited from the receiver. Its use extends one maintained stance-equivalence restriction across three margins: the cross-domain rating must transport to the provision topic in Experiment 2, from the

general population to journalists in Experiment 3, and from senders to the receiver role in Experiment 4. A separately registered companion replaces congeniality with positional alignment, constructed only from the receiver’s pre-treatment stated stance and the fact direction. The two proxies are never entered together or selected by their results. Both are centred at zero, but their zeroes differ: zero congeniality is a fact rated as non-validating at that stance, whereas zero alignment is a midpoint-stance receiver. Full adoption is the joint restriction that adoption at zero equals one and its moderator slope equals zero; the estimates are [PENDING DATA].

Table ?? puts the three reception states side by side. Receipt, transmission failure, and withholding account for [PENDING DATA] of matched receivers respectively, with mean losses of [PENDING DATA]. Loss relative to a reference in which every receiver receives splits as $\bar{\mathcal{L}} - \mu_R = p_F(\mu_F - \mu_R) + p_W(\mu_W - \mu_R)$, which is [PENDING DATA]. Only the first of those two terms rests on the coin. The withholding row carries the sender’s selection, so I report it as accounting against a universal-receipt reference and draw no causal reading from it.

The two silence states differ by [PENDING DATA] within cell. Receivers arrive in queue order and take whichever sender is next, so which sender a receiver draws within a cell is not chosen. Whether that sender withheld is the sender’s own decision, and a difference between the two silence states therefore measures selection into silence; it does not measure how receivers read it.

Receipt moves the matched-topic policy position, signed by the fact’s direction, by [PENDING DATA] (column 3). The design is powered to detect a movement of roughly a fifth of a standard deviation, which is a large amount to ask of a single fact, so I report the estimate with its interval. An interval covering zero bounds what one fact moves; it does not show that a fact moves nothing.

6.4 The Cost of Selectivity

Senders did not withhold at random, so the volume withheld is not the only channel through which disclosure bears on belief accuracy. Let g index environments—a fact, the state it is in, and the receiver facing it—with frequencies λ_g summing to one, disclosure rate σ_g , and mean loss μ_g^R under receipt and $\mu_g^\emptyset$ under non-receipt. Mean realized loss is $\sum_g \lambda_g [\tau \sigma_g \mu_g^R + (1 - \tau \sigma_g) \mu_g^\emptyset]$. Replacing every σ_g by the volume-matched $\bar{\sigma} = \sum_g \lambda_g \sigma_g$, and holding τ , the frequencies, and the conditional losses fixed,

$$\mathcal{L}(\sigma) - \mathcal{L}(\bar{\sigma}) = \tau \text{Cov}_\lambda \left(\sigma_g, \mu_g^R - \mu_g^\emptyset \right), \quad (6)$$

so selection damages accuracy when disclosure concentrates where receipt improves it least. Realized disclosure rates and realized conditional losses are its only inputs.⁴⁸ I match volume within topic, since senders were assigned a topic and never chose between topics, so a difference in disclosure between them is not selection. The equation is written with a single $\bar{\sigma}$ for readability; I report the global variant alongside the within-topic one, and essentially all of the dispersion in σ_g lies within topic rather than between.

The accuracy cost of selection is [PENDING DATA] (Table ??, column 2), which is [PENDING DATA] of mean loss, with a randomization-inference p -value of [PENDING DATA]. Within the decisions senders chose to disclose, the coin alone decides receipt, so the within-cell gap the estimate uses is the one between receipt and a failed draw, and receivers whose sender withheld do not enter it. The estimate asks whether that gap covaries with how often the cell's senders disclosed.

Result 4.1 (The accuracy cost of selection). *Relative to a counterfactual that discloses the same volume without selecting on it, the realized disclosure profile changes belief accuracy by [PENDING DATA], or [PENDING DATA] of mean loss.*

The sign of the covariance says where senders concentrated disclosure. [PENDING DATA] Because σ_g is measured from Experiment 2 and 3 decisions, the interaction coefficient is attenuated by the sampling error in those cell rates; the product reported here is not, because the same error enters the variance it is multiplied by.⁴⁹

The estimate is the realized cost to belief accuracy, and it is not the Bayesian information effect of selection—the change in how much silence itself conveys. That object turns on disclosure in the state that did not occur, which realized behavior does not pin down; Appendix G derives it and the body does not report it.

6.5 Mechanisms

Receivers could in principle undo some of this themselves. Doing so would require knowing how senders disclose and reading silence accordingly, so I report what receivers believe about the disclosure rule and how they respond to silence.

Environment. Nature draws a binary state of the world $\omega \in \{0, 1\}$. The receiver does not observe the state and holds prior $p = \Pr(\omega = 1)$, with $p \in (0, 1)$. The sender observes ω and chooses whether to disclose it, $d \in \{0, 1\}$; disclosure is truthful, so she

⁴⁸It is not assumption-free. Holding μ_g^R and $\mu_g^\emptyset$ fixed while the rule changes restricts how the receiver reads silence, since a receiver who tracked the rule would read it differently under each. Appendix G derives the equation and states the restriction.

⁴⁹Cell disclosure rates rest on a median of 29 sender decisions, so a share of the dispersion in σ_g is sampling noise. Appendix G reports the noise-corrected dispersion alongside the raw one.

can transmit the true state or stay silent but cannot report a state that did not occur. Let

$$\sigma_1 = \Pr(d = 1 \mid \omega = 1), \quad \sigma_0 = \Pr(d = 1 \mid \omega = 0)$$

denote the sender's true state-contingent disclosure probabilities. A disclosed state reaches the receiver with probability $\tau \in (0, 1]$; a withheld state, or a disclosed state that fails to transmit, leaves no message. The receiver observes $m \in \{\omega, \emptyset\}$: the true state after successful transmission, and nothing otherwise. She never observes d , so non-receipt ($m = \emptyset$) pools a withheld state with a disclosed state that failed to transmit.

Receiver beliefs. I separate the information a message carries from the receiver's response to it, and let a Bayesian receiver who knows the disclosure rule (σ_1, σ_0) and the transmission rate τ fix the reference. A received message identifies the state, so Bayes posterior beliefs b_m are $b_1^{\text{Bayes}} = 1$ and $b_0^{\text{Bayes}} = 0$. Non-receipt does not identify the state: it arises in both states, through withholding or failed transmission, and Bayesian belief is

$$b_{\emptyset}^{\text{Bayes}} = \Pr(\omega = 1 \mid m = \emptyset) = \frac{p(1 - \tau\sigma_1)}{p(1 - \tau\sigma_1) + (1 - p)(1 - \tau\sigma_0)}. \quad (7)$$

Silence carries information only when the two disclosure rates differ: if $\sigma_1 = \sigma_0$, non-receipt leaves belief at the prior. A smaller τ pulls $b_{\emptyset}^{\text{Bayes}}$ toward p , because failed transmission then explains non-receipt more readily than withholding does. This is a property of the reference and not an assumption about behavior. Let $\tilde{b}_1, \tilde{b}_0, \tilde{b}_{\emptyset} \in [0, 1]$ denote the receiver's actual beliefs after each message. I place no restriction on how they relate to the reference: the receiver may update as Bayes prescribes, under- or over-react, ignore silence, or read silence as confirmation of her own position, moving opposite the Bayesian direction. Her beliefs may depart from the reference either because she misreads silence given the disclosure rule or because she holds mistaken beliefs about the rule itself.

Silence leaves even a Bayesian receiver in error, so belief error after non-receipt reflects both its ambiguity and the receiver's reading of it. The framework restricts neither the direction nor the magnitude of that reading, nor how it varies with τ .

Receivers' beliefs about the disclosure rule are [PENDING DATA]. Asked how many senders in their own cell passed the fact on, they estimate [PENDING DATA] against a realized σ_g averaging 0.66 across cells. Each receiver gives two estimates, one for each possible correct answer, and the difference between them is what they believe about how far senders select. Regressing that difference on the true difference in the receiver's own cell gives [PENDING DATA], the share of senders' selectivity that

receivers perceive.⁵⁰

Whether receivers respond to silence in the direction the benchmark prescribes is estimated by regressing belief change on the gap the benchmark implies, holding the fact’s continuous cell-assigned congeniality fixed; the slope is [PENDING DATA]. A separately registered companion replaces congeniality with the five-level stated-alignment control; the two are never entered jointly or selected by their results.⁵¹ Any departure from the benchmark is a different question, and this design bounds it without estimating it. A transmission rate of one half was promised to the senders, so silence is weakly diagnostic here whatever receivers do. The slope on the benchmark-implied gap is [PENDING DATA] against the values of one and zero.⁵²

Result 4.2 (Beliefs about senders and responses to silence). *Receivers perceive [PENDING DATA] of the selectivity senders exercise, and the slope of belief change on the gap the benchmark implies is [PENDING DATA]; the design bounds any departure from the benchmark without estimating it.*

The three quantities reported above differ in what they require. The loss borne after receipt requires least: its benchmark is zero. The gap between loss under non-receipt and loss under receipt also depends on the mix of environments the disclosure rule itself generates. The covariance in (6) requires in addition that the conditional losses hold fixed as the rule changes. Under a known disclosure rule, systematic selection places information in the omission itself. What this experiment measures is the belief accuracy receivers are left with, not the ambiguity a Bayesian would still face; Appendix G separates the two.

7 Discussion

In three experiments—quota-sampled U.S. and German general-population samples and 307 professional German journalists—senders decide which true, sourced facts reach a receiver whose political position they know, and provision favors facts supporting the sender’s own side. The realized alignment bias is 11.6 and 12.7 percentage

⁵⁰The level of $\hat{\sigma}$ is reported as a descriptive alongside it. The level alone cannot answer this question: the true gap flips sign with sender stance and averages to nearly zero across receivers, so a receiver who perceives a third of the selectivity still returns a level that reads as well calibrated.

⁵¹The raw contrast between receivers whose benchmark gap points one way and the other is not used. Congeniality predicts both belief change and the sign of that gap, so the contrast carries an offset that has nothing to do with updating. The congeniality index is assigned from prior Experiment 2 raters under the maintained stance-equivalence assumption and is not the receiver’s own elicited belief.

⁵²I report the slope against three benchmarks for the disclosure rule: the measured mirror rate, which is the registered primary, a flat one half, and receivers’ own elicited estimates, which give [PENDING DATA]. The own-estimate coefficient is defined on receivers’ stated counts. Count granularity, clamping and the nonlinear Bayesian transformation can change its sampling behavior in either direction, so no generic attenuation or lower-bound interpretation is imposed; the registered HC1 estimator is read against its stated-scale validation.

points across Experiment 2's two arms and 20.4 among journalists, and in Experiment 1 it rises from 6.7 to 13.3 pp, when the receiver's beliefs can translate into a political donation. Nothing false is transmitted; the distortion lies in which facts are more likely to reach the receiver.

Three elicited returns predict provision, and they behave differently. The expressive return does not respond to whether the receiver's beliefs acquire consequences ($+0.020$, $p = 0.59$). By contrast, the persuasion return predicts nothing absent political action and predicts provision once it is present; the corrective return predicts more provision without tilting it. Models of selective disclosure take the sender's objective as a primitive and derive the disclosure rule from it. Here, behavior and elicited returns show that expressive considerations shape selection even when receiver beliefs carry no consequences. Persuasive considerations enter when those beliefs can affect political action, so the canonical objective fits the consequential setting more closely.

Senders do not appear to favor congenial facts simply because they find them more credible. The closest available credibility controls are Experiment 1's sender-level ex-post rating and, in Experiments 2 and 3, an assigned fact-by-stance proxy constructed from other raters; the alignment result survives both. The experiments also remove the public and attributable reputation mechanisms central to models of audience signaling (Bénabou and Tirole 2006, Guriev et al. 2026). Senders are anonymous, interact once, and in Experiments 2 and 3 non-arrival cannot be separated from failed transmission. These features do not eliminate every social-image concern, but they remove the audience-facing reputational return generated by an identifiable disclosure record. These results have different evidentiary status. Randomization identifies how political consequences change the weight senders place on elicited returns; the returns themselves are not assigned, so their channel coefficients are within-sender, within-fact associations. Experiment 3 has no CONTROL arm, and its motive estimates remain associational.

The experiments also differ in what non-arrival reveals. In Experiment 1 the receiver does not observe that a disclosure opportunity occurred; in Experiments 2 and 3 non-arrival is observed but pools withholding with failed transmission, making silence weakly informative without producing full unraveling. They differ too in what positional alignment detects, and the source of that difference is how far a fact's alignment coincides with the direction senders expect it to move the receiver: $+60.8$ pp among U.S. senders facing a documented false belief, $+27.1$ among journalists, $+7.6$ in the German general population's ACTION arm. Experiment 1's measure is directional and unincentivized, whereas Experiments 2 and 3 elicit a positional return under a scoring rule, so the ordering is suggestive. Where a control arm exists to test it, the ordering tracks the alignment interaction: political action increases the gap in Experiment 1, in which alignment bundles the expressive and persuasive value of a fact, and leaves it

unmoved in Experiment 2, in which alignment predominantly loads on the expressive channel. The expressive return does not strengthen under ACTION, which suggests that the Experiment 1 amplification operated through the persuasive component.

Journalists select by the same rule. Their per-unit alignment slope is statistically indistinguishable from the general population's; they hold stronger political positions, which widens the realized gap under the linear specification without implying a steeper response. Because the task removes competition, ownership, and consumer demand, what persists is a component of selection those mechanisms do not account for. Professional norms covary with selection among journalists: agreeing that one should pass on true information contradicting one's own position predicts a weaker alignment slope, 9.1 pp less per scale point, and moving from agreement to strong agreement halves the slope without eliminating it. The general population shows no such attenuation, and its one surviving interaction points the other way. All items were elicited after the provision task, so this is heterogeneity consistent with norms rather than evidence that norms can be moved. Journalists also provide 9.7 pp more under common controls while differing on no motive slope; professional norms, task engagement, demand, and recruitment selection are all consistent with that level difference, and the design does not distinguish them.

Provision decisions were observed, raising the concern that selectivity reflects what senders believed the study wanted. A uniform demand effect cannot generate alignment-contingent selection within sender, although an alignment-contingent demand effect could. Among journalists, professional self-presentation points toward broader provision and a smaller alignment gap, since veracity and balance are codified obligations. This is a directional argument rather than a test; the observed gap nevertheless remains 20.4 pp.

Each sender addresses a single receiver of known position, whereas journalists in the field supply many at once and work under editors, markets, deadlines, legal constraints, and reputational stakes that the task removes. The general-population samples are quota-sampled, not representative, and the journalist sample is a convenience sample recruited online, overrepresenting online outlets and left-leaning positions. Each experiment covers one country and at most two policy domains, and I make no cross-national claim. The facts are true, sourced, assigned, and costless to transmit, which makes the design a clean test of selection conditional on possessing known facts but leaves it silent on the level and pattern of supply when facts must first be found, verified, produced, or prioritized.

Existing accuracy-based instruments do not reach this margin: fact-checking and labels evaluate what is communicated, not which facts are omitted. The corrective return is associated with how much senders share rather than which side they favor, so accuracy-oriented interventions may raise volume without redressing selection. Mak-

ing beliefs consequential increases selection along the expected persuasive-effect margin rather than dampening it. Expressive selection responds to neither, so reaching it requires interventions that alter its own payoff; norms and audience composition are candidate margins rather than interventions tested here. Rules aimed at completeness are institutionally possible but hard to define and enforce, because accuracy can be checked against ground truth while completeness is relative to an open-ended universe of relevant facts. Regulation has bound the balance of supplied content before: the Fairness Doctrine obliged U.S. broadcasters to present contrasting viewpoints on controversial issues until the Commission abrogated it in 1987, concluding that it no longer served the public interest, was not statutorily mandated, and was inconsistent with First Amendment values ([Federal Communications Commission 1987, 2011](#))—a bound on what has been attempted, not a proposal these experiments imply. Changing who selects information is another candidate margin, but the experiments do not test organizational composition. What curated truth costs the receivers who face it is [\[PENDING DATA\]](#).

8 Concluding Remarks

The political information environment can be distorted without a single false statement. I study why senders choose which true political facts to pass on and which to withhold. Across three experiments—quota-sampled U.S. and German general-population samples and professional German journalists—provision favors the sender’s own position, and in Experiment 1 that bias increases when the receiver’s beliefs carry political consequences. The channels are separately estimable: sharing aligned facts is expressive and operates regardless of consequences; the wish to correct factual beliefs predicts more sharing without tilting what senders share; and selection on a fact’s expected persuasive effect appears only where the receiver acts politically, an activation positive in every specification but statistically fragile in some. Journalists, deciding without the extrinsic pressures media-bias models invoke, select by the same rule at a higher level: some of the selectivity those models attribute to markets survives their removal.

A receiver can scrutinize what arrives. What was never sent is a different object: in Experiment 1 they do not know a decision was made, and in Experiments 2 and 3 they can infer that something may have been withheld but cannot separate withholding from a transmission that failed. Whether supply-side selectivity translates into distorted beliefs therefore turns on how receivers read that inference—one who reads it correctly undoes part of the curation, one who underreads it inherits the sender’s selection. The receiver experiment closes this loop, and its answer to what curated truth costs its audience, against an epistemic accuracy benchmark rather than a total-welfare one, is [\[PENDING DATA\]](#). What is already clear is the shape of the problem: the sup-

ply of political information is curated by the motives of those who hold it, the curation survives professional norms and visible positions, and no amount of verifying what is said will reveal what is not.

References

- Abramitzky, R., Ager, P., Boustan, L., Cohen, E., and Hansen, C. W. (2024). The effect of immigration on the economy: Lessons from the 1920s border closure. *American Economic Review: Insights*.
- Alesina, A., Miano, A., and Stantcheva, S. (2023). Immigration and redistribution. *The Review of Economic Studies*, 90(1):1–39.
- Allcott, H. and Gentzkow, M. (2017). Social media and fake news in the 2016 election. *Journal of Economic Perspectives*, 31(2):211–236.
- Allen, J., Howland, B., Mobius, M., Rothschild, D., and Watts, D. J. (2020). Evaluating the fake news problem at the scale of the information ecosystem. *Science Advances*, 6(14):eaay3539.
- Ambuehl, S., Bernheim, B. D., and Ockenfels, A. (2021). What motivates paternalism? an experimental study. *American economic review*, 111(3):787–830.
- Angelucci, C. and Prat, A. (2024). Is journalistic truth dead? measuring how informed voters are about political news. *American Economic Review*, 114(4):887–925.
- Aridor, G., Jiménez-Durán, R., Levy, R., and Song, L. (2024). The economics of social media. *Journal of Economic Literature*, 62(4):1422–1474.
- Barrera, O., Guriev, S., Henry, E., and Zhuravskaya, E. (2020). Facts, alternative facts, and fact checking in times of post-truth politics. *Journal of public economics*, 182:104123.
- Barron, K. (2021). Belief updating: Does the ‘good-news, bad-news’ asymmetry extend to purely financial domains? *Experimental Economics*, 24(1):31–58.
- Bartling, B., Cappelen, A. W., Hermes, H., Tungodden, B., and Skivenes, M. (2023). Free to fail? paternalistic preferences in the united states. *NHH Dept. of Economics Discussion Paper*, (09).
- Bénabou, R. and Tirole, J. (2006). Incentives and prosocial behavior. *American economic review*, 96(5):1652–1678.
- Benndorf, V., Kübler, D., and Normann, H.-T. (2025). Information unraveling and limited depth of reasoning. *Games and Economic Behavior*.
- Berger, L. M. (2026). How digital media markets amplify news sentiment. Working paper, University of Cologne.
- Berger, L. M., Kerkhof, A., Mindl, F., and Münster, J. (2025). Debunking “fake news” on social media: Immediate and short-term effects of fact-checking and media literacy interventions. *Journal of Public Economics*, 245:105345. Open access, CC BY. Available online 5 April 2025. Pre-registered at AEA Registry AEARCTR-0008199.

- Besley, T. and Prat, A. (2006). Handcuffs for the grabbing hand? media capture and government accountability. *American Economic Review*, 96(3):720–736.
- Bowen, T. R., Dmitriev, D., and Galperti, S. (2023). Learning from shared news: When abundant information leads to belief polarization. *The Quarterly Journal of Economics*, 138(2):955–1000.
- Boxell, L. and Conway, J. (2022). Journalist ideology and the production of news: Evidence from movers. Working paper.
- Braghieri, L. (2024). Political correctness, social image, and information transmission. *American Economic Review*, 114(12):3877–3904.
- Braghieri, L., Eichmeyer, S., Levy, R., Mobius, M. M., Steinhardt, J., and Zhong, R. (2025). Level slant and polarization of news consumption on social media. *SSRN Working Paper*.
- Budak, C., Nyhan, B., Rothschild, D. M., Thorson, E., and Watts, D. J. (2024). Misunderstanding the harms of online misinformation. *Nature*, 630:45–53.
- Bullock, J. G. and Lenz, G. (2019). Partisan bias in surveys. *Annual Review of Political Science*, 22(1):325–342.
- Burdea, V., Montero, M., and Sefton, M. (2023). Communication with partially verifiable information: An experiment. *Games and Economic Behavior*, 142:113–149.
- Bursztyn, L., Egorov, G., Haaland, I., Rao, A., and Roth, C. (2023a). Justifying dissent. *The Quarterly Journal of Economics*, 138(3):1403–1451.
- Bursztyn, L., Rao, A., Roth, C., and Yanagizawa-Drott, D. (2023b). Opinions as facts. *The Review of Economic Studies*, 90(4):1832–1864.
- Cagé, J., Gallo, N., Hengel, M., Henry, E., and Huang, Y. (2025). Fact-checking and misinformation: Evidence from the market leader. CESifo Working Paper 12319, CESifo.
- Cagé, J., Hengel, E., Hervé, N., and Urvoy, C. (2022). Hosting media bias: Evidence from the universe of French broadcasts 2002–2020. HAL preprint hal-03878119.
- Chen, D. L., Schonger, M., and Wickens, C. (2016). otree—an open-source platform for laboratory, online, and field experiments. *Journal of Behavioral and Experimental Finance*, 9:88–97.
- Chopra, F., Haaland, I., and Roth, C. (2022). Do people demand fact-checked news? evidence from us democrats. *Journal of Public Economics*, 205:104549.
- Chopra, F., Haaland, I., and Roth, C. (2024). The demand for news: Accuracy concerns versus belief confirmation motives. *The Economic Journal*, 134(661):1806–1834.
- Crawford, V. P. and Sobel, J. (1982). Strategic information transmission. *Econometrica*, 50(6):1431–1451.
- de Quidt, J., Haushofer, J., and Roth, C. (2018). Measuring and bounding experimenter demand. *American Economic Review*, 108(11):3266–3302.
- Degan, A., Li, M., and Xie, H. (2023). An experimental investigation of persuasion through selective disclosure of evidence. *Canadian Journal of Economics/Revue canadienne d’économie*, 56(4):1490–1516.

- Deutscher Presserat (2025). Publizistische Grundsätze (Pressekodex): Richtlinien für die publizistische Arbeit nach den Empfehlungen des Deutschen Presserats. Fassung vom 19. März 2025; first adopted 1973.
- Durante, R. and Knight, B. (2012). Partisan control, media bias, and viewer responses: Evidence from berlusconis italy. *Journal of the European Economic Association*, 10(3):451–481.
- Dye, R. A. (1985). Disclosure of nonproprietary information. *Journal of accounting research*, pages 123–145.
- Eil, D. and Rao, J. M. (2011). The good news-bad news effect: Asymmetric processing of objective information about yourself. *American Economic Journal: Microeconomics*, 3(2):114–138.
- Ershov, D. and Morales, J. S. (2024). Sharing news left and right: Frictions and misinformation on twitter. *The Economic Journal*, 134:2391–2417.
- Eyster, E. and Rabin, M. (2005). Cursed equilibrium. *Econometrica*, 73(5):1623–1672.
- Farina, A., Fréchette, G. R., Ispano, A., Lizzeri, A., and Perego, J. (2026). The selective disclosure of evidence: An experiment. *The Review of Economic Studies*. Forthcoming.
- Farina, A. and Leccese, M. (2025). Hiding a flaw? experimental evidence on multi-dimensional information disclosure. Working paper, September 2025. Princeton University and Boston University.
- Federal Communications Commission (1987). Syracuse Peace Council, Memorandum Opinion and Order. 2 FCC Rcd 5043; recon. denied, 3 FCC Rcd 2035 (1988); aff’d sub nom. Syracuse Peace Council v. FCC, 867 F.2d 654 (D.C. Cir. 1989), cert. denied, 493 U.S. 1019 (1990).
- Federal Communications Commission (2011). Order (DA 11-1432): Deletion of obsolete broadcast and cable rules referencing the fairness doctrine. Adopted and released August 24, 2011.
- Feess, E., Jost, P.-J., and Ressi, A. (2024). Fake news and the asymmetry between type I and type II errors: Dismissing true information is more prevalent than following false information. Working paper. AEA RCT Registry AEARCTR-0010984.
- Fong, M.-J., Lin, P.-H., and Palfrey, T. R. (2025). Cursed sequential equilibrium. *American Economic Review*, 115(8):2616–2658.
- Freitag, J., Kerkhof, A., and Münster, J. (2021). Selective sharing of news items and the political position of news outlets. *Information Economics and Policy*, 56:100926.
- Gentzkow, M. and Shapiro, J. M. (2006). Media bias and reputation. *Journal of Political Economy*, 114(2):280–316.
- Gentzkow, M. and Shapiro, J. M. (2010). What drives media slant? evidence from us daily newspapers. *Econometrica*, 78(1):35–71.
- Golman, R., Loewenstein, G., Molnar, A., and Saccardo, S. (2022). The demand for, and avoidance of, information. *Management Science*, 68(9):6454–6476.
- González-Bailón, S., Lazer, D., Barberá, P., Zhang, M., Allcott, H., Brown, T., Crespo-Tenorio, A., Freelon, D., Gentzkow, M., Guess, A. M., et al. (2023). Asymmetric ideological segregation in exposure to political news on facebook. *Science*, 381(6656):392–398.

- Grigorieff, A., Roth, C., and Ubfal, D. (2020). Does information change attitudes toward immigrants? *Demography*, 57:1117–1143.
- Grossman, S. J. (1981). The informational role of warranties and private disclosure about product quality. *The Journal of law and Economics*, 24(3):461–483.
- Grossmann, M. R. P. (2024). Knowledge and freedom: Evidence on the relationship between information and paternalism.
- Guess, A., Nagler, J., and Tucker, J. (2019). Less than you think: Prevalence and predictors of fake news dissemination on facebook. *Science advances*, 5(1):eaau4586.
- Guriev, S., Henry, E., Marquis, T., and Zhuravskaya, E. (2026). Curtailing false news, amplifying truth. *Econometrica*. Forthcoming.
- Haaland, I. and Roth, C. (2020). Labor market concerns and support for immigration. *Journal of Public Economics*, 191:104256.
- Haaland, I., Roth, C., and Wohlfart, J. (2023). Designing information provision experiments. *Journal of Economic Literature*, 61(1):3–40.
- Henry, E., Zhuravskaya, E., and Guriev, S. (2022). Checking and sharing alt-facts. *American Economic Journal: Economic Policy*, 14(3):55–86.
- Hillman, A. L. (2010). Expressive behavior in economics and politics. *European Journal of political economy*, 26(4):403–418.
- Jin, G. Z., Luca, M., and Martin, D. (2021). Is no news (perceived as) bad news? an experimental investigation of information disclosure. *American Economic Journal: Microeconomics*, 13(2):141–173.
- Jung, W.-O. and Kwon, Y. K. (1988). Disclosure when the market is unsure of information endowment of managers. *Journal of Accounting research*, pages 146–153.
- Kamenica, E. and Gentzkow, M. (2011). Bayesian persuasion. *American Economic Review*, 101(6):2590–2615.
- Kuziemko, I., Norton, M. I., Saez, E., and Stantcheva, S. (2015). How elastic are preferences for redistribution? evidence from randomized survey experiments. *American Economic Review*, 105(4):1478–1508.
- Lazer, D. M., Baum, M. A., Benkler, Y., Berinsky, A. J., Greenhill, K. M., Menczer, F., Metzger, M. J., Nyhan, B., Pennycook, G., Rothschild, D., et al. (2018). The science of fake news. *Science*, 359(6380):1094–1096.
- Light, M. T., He, J., and Robey, J. P. (2020). Comparing crime rates between undocumented immigrants, legal immigrants, and native-born us citizens in texas. *Proceedings of the National Academy of Sciences*, 117(51):32340–32347.
- Loosen, W., von Garmissen, A., Bartelt, E., and van Olphen, T. (2023). Journalismus in Deutschland 2023: Aktuelle Befunde zu Situation und Wandel. Arbeitspapiere des Hans-Bredow-Instituts 68, Hans-Bredow-Institut, Hamburg.
- Louis-Sidois, C. (2025). Both judge and party? investigating the political unbiasedness of fact-checkers. *Journal of the European Economic Association*, 23(6):2137–2164.

- Manili, J. (2024). Disclosure to an heterogeneous audience: The role of motivated beliefs. Working paper, Northwestern University. SSRN 4696951.
- Marie, A., Altay, S., and Strickland, B. (2023). Moralization and extremism robustly amplify myside sharing. *PNAS Nexus*, 2(4):pgad078.
- Marie, A. and Petersen, M. B. (2025). Motivations to connect with like-minded audiences increase partisan sharing on social media. *PNAS nexus*, 4(7):pgaf197.
- Martin, G. J. and McCrain, J. (2019). Local news and national politics. *American Political Science Review*, 113(2):372–384.
- Milgrom, P. R. (1981). Good news and bad news: Representation theorems and applications. *The Bell Journal of Economics*, pages 380–391.
- Möbius, M. M., Niederle, M., Niehaus, P., and Rosenblat, T. S. (2021). Managing self-confidence: Theory and experimental evidence. Working paper, October 2021.
- Mullainathan, S. and Shleifer, A. (2005). The market for news. *American Economic Review*, 95(4):1031–1053.
- Musolf, R., Roth, C., and Zimmermann, F. (2026). Learning from false stories. ECONtribute Discussion Paper No. 412.
- National Academies of Sciences, Engineering, and Medicine (2017). *The Economic and Fiscal Consequences of Immigration*. The National Academies Press, Washington, DC.
- Osmundsen, M., Bor, A., Vahlstrup, P. B., Bechmann, A., and Petersen, M. B. (2021). Partisan polarization is the primary psychological motivation behind political fake news sharing on twitter. *American political science review*, 115(3):999–1015.
- Pennycook, G., Epstein, Z., Mosleh, M., Arechar, A. A., Eckles, D., and Rand, D. G. (2021). Shifting attention to accuracy can reduce misinformation online. *Nature*, 592:590–595.
- Peterson, E. and Iyengar, S. (2021). Partisan gaps in political information and information-seeking behavior: Motivated reasoning or cheerleading? *American Journal of Political Science*, 65(1):133–147.
- Prat, A. and Strömberg, D. (2013). The political economy of mass media. *Advances in economics and econometrics*, 2:135.
- Shin, H. S. (2003). Disclosures and asset returns. *Econometrica*, 71(1):105–133.
- Shin, J. and Thorson, K. (2017). Partisan selective sharing: The biased diffusion of fact-checking messages on social media. *Journal of Communication*, 67(2):233–255.
- Steindl, N., Lauerer, C., and Hanitzsch, T. (2017). Journalismus in Deutschland: Aktuelle Befunde zu Kontinuität und Wandel im deutschen Journalismus. *Publizistik*, 62(4):401–423.
- Suvorov, A., van de Ven, J., and Villeval, M. C. (2024). Selective information sharing and group delusion. Working paper, New Economic School, University of Amsterdam, and CNRS-GATE Lyon.
- Thaler, M. (2024). The fake news effect: Experimentally identifying motivated reasoning using trust in news. *American Economic Journal: Microeconomics*, 16(2):1–38.
- Thaler, M. (2025). The supply of motivated beliefs. Working paper.

- Verrecchia, R. E. (1983). Discretionary disclosure. *Journal of Accounting and Economics*, 5:179–194.
- Vosoughi, S., Roy, D., and Aral, S. (2018). The spread of true and false news online. *Science*, 359(6380):1146–1151.
- Zhu, Y. and Dukes, A. (2015). Selective reporting of factual content by commercial media. *Journal of Marketing Research*, 52(1):56–76.
- Zimmermann, F. (2020). The dynamics of motivated beliefs. *American Economic Review*, 110(2):337–361.

A Experimental Design Details

A.1 Statistic Selection and Validation Survey

I selected statistics for the main experiment from a larger pool of candidates using a pre-registered validation survey ($N = 100$, AsPredicted No. 232064) conducted on Prolific between June 10 and June 21, 2025. The survey evaluated eight candidate statistics—four hypothesized as pro-immigration and four as anti-immigration—against three survey-measured criteria (perceived credibility, policy relevance, and perceived political direction) and one researcher-verified criterion (veracity). The selection rule was pre-specified: within each direction, select the two statistics with the highest policy relevance among those with adequate perceived credibility, subject to the constraint that the final set contains two pro-immigration and two anti-immigration statistics.

Candidate Statistics and Selection Criteria. Table A.1 reports the eight candidate statistics with their survey-measured scores. The four criteria were operationalized as follows:

Veracity was established by the researcher through source verification. Each selected statistic is supported by multiple independent, authoritative sources: peer-reviewed academic publications, federal statistical agencies, or nonpartisan research institutions. Source documentation is provided below.

Policy relevance was measured by asking respondents: “How relevant do others consider the statistic when forming immigration policy positions?” (1 = completely irrelevant, 5 = extremely relevant).

Perceived credibility was measured by asking respondents: “How credible is the statistic?” (1 = not at all, 5 = entirely).

Perceived political direction was measured through two items: *usability* (“For whom would the statistic be most useful in a policy debate?”; 1 = exclusively supporting expansive immigration, 5 = exclusively supporting restrictive immigration) and *opinionability* (“How does learning about the statistic change someone’s immigration policy views?”; 1 = convince to support expansion, 5 = convince to support limitation). These directional measures serve the classification of statistics as pro- or anti-immigration and the balance criterion.

Table A.1: Candidate Statistics: Survey-Measured Indicators

Statistic	Policy	Perceived	Perceived Direction	
	Relevance	Credibility	Usability	Opinionability
<i>Pro-immigration</i>				
Crime Rate [†]	3.74 (1.18)	3.50 (0.97)	2.23 (1.06)	2.36 (1.11)
Labor [†]	3.64 (1.04)	3.81 (1.02)	2.40 (1.09)	2.57 (1.23)
Entrepreneurship	3.12 (1.08)	3.54 (1.00)	2.09 (1.05)	2.18 (1.15)
Healthcare	3.28 (1.03)	3.59 (1.07)	2.03 (1.05)	2.19 (1.15)
<i>Anti-immigration</i>				
Border Crossings [†]	3.85 (1.14)	3.97 (0.90)	3.94 (1.14)	4.08 (1.11)
Fiscal Impact [†]	3.79 (1.01)	3.55 (0.93)	3.68 (1.15)	3.73 (1.01)
Poverty Rate	3.41 (1.01)	3.93 (0.93)	3.40 (1.13)	3.52 (1.11)
Wage Inequality	3.35 (0.95)	3.53 (0.90)	3.41 (1.06)	3.43 (0.99)

Notes: $N = 100$. Cells report means with standard deviations in parentheses. Selected statistics with [†]. Policy relevance: 1 = completely irrelevant, 5 = extremely relevant. Perceived credibility: 1 = not at all credible, 5 = entirely credible. Usability and opinionability: 1 = supports expansive immigration, 5 = supports restrictive immigration. One-sample t -tests against the neutral midpoint of 3 reject the null for all eight statistics on both directional measures ($p < 0.001$ for all selected statistics; see text).

Selection. The selection rule applied within each direction independently. Among the four pro-immigration candidates, Crime Rate (3.74) and Labor Force Participation (3.64) had the highest policy relevance scores, both with adequate perceived credibility (≥ 3.50). Among the four anti-immigration candidates, Border Crossings (3.85) and Fiscal Impact (3.79) had the highest policy relevance, again with adequate credibility. This rule uniquely determines the four selected statistics reported in the main text (Section 2.1). The one rejected statistic with notably high credibility—Poverty Rate (3.93)—scored lower on policy relevance (3.41) than both selected anti-immigration statistics, confirming that the selection was not driven by credibility alone.

Perceived directional scores did not vary significantly by respondents' political affiliation (Democrat, Republican, Independent), indicating that the directional classification reflects consensus perception rather than partisan projection.

Direction Validation. I classify Crime Rate and Labor Force Participation as pro-immigration and Border Crossings and Fiscal Impact as anti-immigration. This classification is validated by the survey's directional measures. For usability and opinionability respectively, the pro-immigration statistics score below the neutral midpoint of 3 (Crime: 2.23, 2.36; Labor: 2.40, 2.57) and the anti-immigration statistics score above it (Border: 3.94, 4.08; Fiscal: 3.68, 3.73). One-sample t -tests against the midpoint of 3

reject the null of no directional content for all four selected statistics on both measures: Crime ($t = -7.25$ and $t = -5.74$), Labor ($t = -5.50$ and $t = -3.49$), Border Crossings ($t = 8.21$ and $t = 9.75$), and Fiscal Impact ($t = 5.89$ and $t = 7.20$), all $p < 0.001$.

This classification is independently confirmed by the main experiment. In the expected attitude shift elicitation (Section 2.3.1, Table B.7), senders in the main experiment ($N = 1,200$) expect Crime and Labor to shift receiver preferences toward pro-immigration positions (25 and 17 pp, both $p < 0.001$) and Border and Fiscal to shift preferences toward anti-immigration positions (-34 and -32 pp, both $p < 0.001$). The convergence of two independent samples using two independent methods—survey-based directional perception in the validation study and unincentivized attitude-shift predictions in the main experiment—makes the directional classification robust. The two directional criteria coincide fact by fact: usability (which side the statistic serves) and opinionability (which way learning it moves views) yield the same classification for all four selected facts (Table A.1). A statistic’s congeniality to a sender’s side and its expected persuasive direction are thus inseparable at the design stage.

Source Documentation. Veracity was established through source verification rather than survey measurement. Each selected statistic is supported by multiple independent, authoritative sources. I report the specific values shown to participants, the primary source for each, and corroborating evidence. All sources were available at the time of the experiment (July–August 2025).

Crime Rate (immigrants have lower crime rates than U.S.-born Americans). Participants were shown violent crime *arrest* rates per 100,000 residents in Texas (2012–2018): 213.0 for U.S.-born citizens, 185.3 for legal immigrants (13% lower), and 96.2 for undocumented immigrants (55% lower). These figures are from Light et al. (2020), who analyze the universe of felony arrests recorded in the Texas Computerized Criminal History database cross-referenced with Department of Homeland Security immigration status checks.⁵³ Other studies use Texas Department of Public Safety conviction data over a longer period (2013–2022) and reports homicide conviction rates of 3.0 per 100,000 for native-born Americans, 2.2 for unauthorized immigrants (26% lower), and 1.2 for legal immigrants (61% lower).⁵⁴ At the national level, Abramitzky et al. (2024) analyze census-linked incarceration records from 1870–2020 and find that immigrants have been consistently less likely to be incarcerated than the U.S.-born population throughout this period.⁵⁵ Texas is the only U.S. state that systematically records the immigration status of all arrestees, making it the sole source of population-level crime-

⁵³Published in *Proceedings of the National Academy of Sciences*; data and replication materials available at <https://www.openicpsr.org/openicpsr/project/124923>.

⁵⁴Cato Institute Policy Analysis no. 977; available at <https://www.cato.org/policy-analysis/illegal-immigrant-murderers-texas-2013-2022>.

⁵⁵Published in *American Economic Review: Insights*.

rate comparisons by immigration status. The national-level incarceration evidence in [Abramitzky et al. \(2024\)](#) suggests the Texas pattern is not anomalous.

Labor Force Participation (immigrants have higher labor force participation than U.S.-born Americans). Participants were shown 2023 annual average labor force participation rates: 66.6% for the foreign-born population and 61.8% for the native-born population (4.8 percentage points higher). These figures are from Table 1 of the Bureau of Labor Statistics annual report *Foreign-Born Workers: Labor Force Characteristics—2023*.⁵⁶ The same data are available through the Federal Reserve Bank of St. Louis FRED database, Current Population Survey Table A-7.⁵⁷ The foreign-born participation rate has exceeded the native-born rate in every year of the BLS series.

Border Crossings (unauthorized border crossings have been higher in the past five years relative to historical averages). Participants were shown two figures computed from U.S. Customs and Border Protection (CBP) enforcement statistics:⁵⁸ a historical average of 890,435 apprehensions per year over fiscal years 1970–2019, and a recent five-year average of 1,584,573 encounters per year over fiscal years 2020–2024 (78% higher). The comparison requires combining two CBP reporting series: “apprehensions” (reported through FY2019) and “encounters” (reported from FY2020 onward, which combine apprehensions and expulsions under Title 42 and Title 8 processing).⁵⁹ The statistic reflects data available at experiment time (July–August 2025). Subsequent enforcement changes under the second Trump administration reduced encounters to record lows in FY2025, as documented by Pew Research Center analysis of CBP data.⁶⁰ The FY2025 decline does not affect the validity of the statistic as presented to participants, which concerned the five-year average through FY2024.

Fiscal Impact (net fiscal costs of first-generation immigrants are higher than U.S.-born Americans at the state and local level). Participants were shown that first-generation immigrant households generate \$57.4 billion more in annual state and local government costs than they contribute in taxes, based on 2011–2013 data. This figure is from the National Academies of Sciences, Engineering, and Medicine comprehensive report on the economic and fiscal consequences of immigration, which remains the most authoritative and comprehensive fiscal analysis of immigration in the United States.⁶¹

⁵⁶The annual report is available at <https://www.bls.gov/news.release/pdf/forbrn.pdf>. At the time of writing, the 2024 edition (released May 2025) reports 66.5% versus 61.7%, confirming the same pattern.

⁵⁷<https://fred.stlouisfed.org/release/tables?rid=50&eid=1736>

⁵⁸CBP historical data available at <https://www.cbp.gov/newsroom/stats/cbp-enforcement-statistics>.

⁵⁹Beginning in March 2020, CBP shifted from reporting apprehensions to reporting “encounters,” a broader category. The historical series (1970–2019) uses the apprehension metric; the recent series (2020–2024) uses the encounter metric. Since encounters $\geq$ apprehensions by construction, the 78% comparison is an upper bound—i.e., it may slightly overstate the increase by comparing a narrower historical measure to a broader recent measure.

⁶⁰<https://www.pewresearch.org/short-reads/2026/02/02/migrant-encounters-at-the-us-mexico-border-ar>

⁶¹National Academies of Sciences, Engineering, and Medicine (2017); the state and local fiscal impact

The report notes that this fiscal gap reverses at the federal level and across generations: the children and grandchildren of immigrants are among the strongest net fiscal contributors. I present only the state-and-local first-generation finding because it represents the anti-immigration direction in my balanced design. The Cato Institute provides complementary analysis of immigrants' fiscal effects over a longer period (1994–2023), confirming that the fiscal picture varies substantially by level of government and generational horizon.⁶²

A.2 Belief Elicitation and Donation Task (Experiment 1)

This section documents two implementation details of the Experiment 1 sender protocol referenced in Section 2.1: the incentivization of the prior-knowledge elicitation and the revealed-preference donation task.

Prior-Knowledge Elicitation. Senders complete each comparative statement by choosing “higher,” “lower,” or “don’t know.” Correct completions receive a bonus, incorrect completions do not, and “don’t know” responses receive the bonus with 25% probability. The 25% payoff for “don’t know” is strictly dominated under expected utility maximization with binary alternatives, where uninformed guessing yields a 50% expected bonus. The option may nonetheless be selected by highly ambiguity-averse participants who prefer an objective 25% lottery over an ambiguous 50% guess. This implies that reported “don’t know” rates likely understate genuine ignorance among expected-utility maximizers.

Donation Task. Senders allocate \$500 between the *American Immigration Council* (AIC), which advocates for increased immigration, and *NumbersUSA*, which advocates for decreased immigration—both well-established organizations representing opposing positions in the immigration debate. I briefly describe the stated goals of each organization and present their missions in their own words, with the display order randomized. Participants are informed that one donation decision from the experiment will be randomly selected and implemented, and they receive a verification link to confirm the donation was placed.

The framing text and both mission statements are reproduced verbatim below. Senders saw the two organizations side by side as bulleted lists, with their left–right order randomized; the version shown here is the one presented to senders who saw the AIC on the left, and the alternative reverses both the blocks and the corresponding clauses of

analysis is in Chapter 8. Available at <https://www.nationalacademies.org/read/23550/chapter/14>.

⁶²<https://www.cato.org/white-paper/immigrants-recent-effects-government-budgets-1994-2023>

the framing text without altering any wording. The ellipsis in the AIC block appears in the stimulus. Quotation marks are set in L^AT_EX conventions; the wording is unaltered.

Framing text. American immigration organizations range from those advocating for more open policies to those pushing for greater restrictions. The *American Immigration Council* champions seeing immigrants as valuable contributors who deserve stronger protections and family-friendly policies. In contrast, *NumbersUSA* advocates for reduced immigration levels to prevent overcrowding and promotes stricter enforcement of limitations.

Below, you will find a brief description of each organization in their own words.

American Immigration Council (AIC)

- The American Immigration Council envisions a nation where immigrants are embraced, communities are enriched, and justice prevails for all.
- We strive to create a society that values immigrants as vital contributors to our nation [...].
- Our immigration system and the way it treats immigrants often fail to live up to our highest values and ideals.
- Not only have we failed, in many cases, to afford basic due process rights, but in countless ways, our outdated immigration laws and policies do not respect the family unit, discourage innovation, and turn away those in need of protection.

NumbersUSA

- NumbersUSA advocates for sensible immigration policies that ensure Americans' safety and promote economic fairness while protecting the environment and quality of life.
- The current system creates endless "chain migration," of distant relatives, with backlogs, frustration, and a level of immigration that is incompatible with conservation and economic justice.
- The current system displaces qualified Americans with exploitable foreign workers.
- The current system raffles away green cards and puts tourists who give birth in the U.S. on a path to citizenship.

A.3 Fact Selection and Pre-Validation Survey — Experiments 2 and 3

I selected facts for Experiments 2 and 3 from a larger candidate pool using a pre-registered pre-validation survey (AsPredicted #290989, quota target $N = 330$, final analysis sample $N \approx 300$ after speeder and consent exclusions) conducted on the Bilendi panel in Germany in May 2026. The survey evaluated 36 candidate facts—six hypothesized as pro-policy and six as anti-policy in each of three domains (immigration, climate, redistribution)—against four criteria, and additionally validated the domain-specific advocacy organizations used in the ACTION treatment.

Criteria and Selection Rule. Four criteria were applied to each candidate fact. *Prior knowledge* was elicited via an incentivized fill-in-the-blank completion of the fact statement with a confidence slider [50, 100]; the probabilistic prior knowledge score is $PK = \text{confidence}/100$ for correct completions and $1 - \text{confidence}/100$ for incorrect completions. *Directional consistency* was measured by a 5-point rating (1 = strongly pro-policy, 5 = strongly anti-policy); a fact passes if the share of respondents rating it in its assigned direction exceeds the share rating it in the opposite direction. *Policy relevance* and *perceived credibility* were each measured on a 1–5 scale.

The pre-specified selection rule selected the top three facts by policy relevance within each topic $\times$ direction cell among those passing directional consistency, subject to balance constraints across the pro and anti sets: $|\bar{c}_{\text{pro}} - \bar{c}_{\text{anti}}| \leq 1.0$ for credibility, $|\overline{PK}_{\text{pro}} - \overline{PK}_{\text{anti}}| \leq 0.25$, and $|\bar{r}_{\text{pro}} - \bar{r}_{\text{anti}}| \leq 1.0$ for policy relevance.

The ACTION treatment organizations were additionally validated against two criteria: (i) at least 50% of respondents correctly recognize each organization as representing the pro-policy or anti-policy position respectively, and (ii) a positive Spearman correlation between topic-specific attitude and hypothetical donation to the pro-organization, restricted to respondents with directional attitudes.

Domain Selection. The redistribution domain failed on both grounds and is excluded. All six redistribution anti-direction candidate facts failed the directional consistency criterion, and the redistribution anti-organization had a recognition rate of 33.3%, well below the 50% threshold. Immigration passed all fact and organization criteria (anti-organization recognition: 93.6%; Spearman $r = 0.57$). Climate’s facts all passed directionality and the organization correlation criterion ($r = 0.57$); the climate anti-organization fell marginally below the recognition threshold (42.6% vs. 50%). Since climate’s failure was marginal and isolated to one metric while redistribution’s failures were compounding, climate was retained and this deviation is noted here.

Deviations from Algorithmic Fact Selection. *Immigration.* The algorithm selected *Bildungsabschluss* ($PK = 0.81$), which produced a pro–anti PK gap of 0.29, exceeding

the 0.25 balance threshold. It was replaced by *Illegale Aufenthalte* (PK = 0.51), closing the gap to 0.15 while satisfying all other balance constraints. The label of one anti-fact was updated from *Arbeitslosigkeit* to *Bürgergeld* to reflect the 2026 German policy context; the underlying statistic is unchanged.

Climate. The pro set was revised to *Erneuerbare Energien*, *Stromgestehungskosten*, and *Grüne Beschäftigung* for higher directional consistency. *Netzausbaukosten* (anti) was reformulated between the pre-validation survey and the main experiment; directional ratings in Table A.2 reflect the pre-reformulation version, which failed the consistency criterion (17.5% consistent vs. 42.1% inconsistent). The reformulated version is what was fielded in the main experiment.

Selected Facts and Pre-Validation Ratings. Table A.2 reports the 12 selected facts with their pre-validation survey scores. Table A.3 reports the balance check.

Table A.2: Selected Facts: Pre-Validation Ratings (Experiments 2 and 3)

Fact	Dir.	Relevance	Credibility	PK	% Consistent
<i>Immigration</i>					
Fachkräftemangel	Pro	3.79	3.65	0.52	45.8
Beschäftigungswachstum	Pro	3.43	3.45	0.55	38.8
Integrationskurs	Pro	3.61	3.29	0.37	32.1
Kriminalität	Anti	4.04	3.76	0.70	90.0
Bürgergeld [†]	Anti	3.98	3.74	0.67	78.7
Illegale Aufenthalte	Anti	3.69	3.45	0.51	76.5
<i>Climate policy</i>					
Erneuerbare Energien	Pro	3.67	3.61	0.52	51.0
Stromgestehungskosten	Pro	3.46	3.51	0.56	61.4
Grüne Beschäftigung	Pro	3.28	3.46	0.57	47.8
Netzausbaukosten [†]	Anti	3.84	3.53	0.72	17.5
Energiekosten Industrie	Anti	4.09	3.60	0.81	35.8
Preisvolatilität	Anti	3.73	3.36	0.66	33.9

Notes: $N \approx 300$. Policy relevance and perceived credibility: 1 = lowest, 5 = highest. PK is the confidence-weighted prior knowledge probability (see text). % Consistent is the share of respondents whose directional rating matches the assigned fact direction. [†] Reformulated after pre-validation; ratings are pre-reformulation and failed the consistency criterion. [‡] Relabeled from *Arbeitslosigkeit*; underlying statistic unchanged.

Table A.3: Balance Check: Selected Facts by Domain (Experiments 2 and 3)

Domain	Δ Relevance	Δ Credibility	Δ PK	Passes
Immigration	0.29	0.19	0.15	Yes
Climate	0.42	0.03	0.18	Yes

Notes: Δ = absolute difference between mean pro-direction and mean anti-direction facts. Pre-specified thresholds: Δ Relevance ≤ 1.0 , Δ Credibility ≤ 1.0 , Δ PK ≤ 0.25 . Climate Δ PK uses pre-reformulation statistics for *Netzausbaukosten*.

A.4 Information Condition Implementation

Part 2 of the experiment proceeds in three stages: study presentation and elicitation (all senders), information feedback (INFO = 1 only), and expected attitude change elicitation (all senders). I describe each stage and reproduce the text shown to participants.

Stage 1: Study Presentation and Elicitation (All Senders). All senders read the following description of the [Haaland and Roth \(2020\)](#) study:

A recent study published in the prestigious academic journal, *Journal of Public Economics*, examined how people form political opinions. The researchers surveyed more than 3,000 Americans that are representative of the US population.

People were told about a 1980 event when 55,000 Cuban immigrants unexpectedly arrived in Miami, increasing the city's workforce by 8% almost overnight. They were asked how this affected the wages and employment of U.S.-born workers.

Then, half of the participants learned that research found this wave of immigration had virtually no effect on the wages or unemployment for U.S.-born workers, because the new immigrants increased demand for goods and services, which created more jobs.

The other half of the participants received no information.

Senders then make two incentivized predictions:

1. *Behavioral effect (continuous)*: "By what percentage did learning about the research findings change people's petition signing rates?" Senders enter a percentage (positive or negative).
2. *Differential updating (categorical)*: "After learning about the research findings, who do you think changed their attitudes toward immigration more?" Senders choose among: (a) people who initially believed immigration harmed U.S.-born workers,

- (b) people who initially believed immigration did not harm U.S.-born workers, or
- (c) both groups changed about the same amount.

Each question is equally likely to be selected for the bonus payment. For the continuous estimate, the probability of receiving the bonus is 100% if the estimate equals the true value (69%) and decreases linearly to 0% at a deviation of 20 percentage points or more.⁶³ For the categorical estimate, the probability is 100% if correct and 0% otherwise. Senders are instructed not to search for the answers and tab-switching is monitored to enforce compliance; senders who switched tabs during this stage were excluded as pre-registered.

Stage 2: Information Feedback (INFO = 1 Only). Senders in the INFO condition receive personalized feedback comparing their predictions to the actual results. Senders in the no-INFO condition proceed directly to Stage 3. The feedback screen displays the following:

Petition Signatures

The percentage of people who signed the White House petition actually increased by **69%**.

You estimated that the percentage of people who signed the White House petition changed by *[participant's estimate]*%.

That means *you* *[underestimated/overestimated]* the change by *[participant-specific gap]* percentage points.

Effect on Attitudes toward Immigration

Actual Result: People who initially believed immigration harmed U.S.-born workers changed their attitudes toward immigration **significantly more** than those who already held correct beliefs.

Your Answer: *[participant's answer]*.

The feedback is accompanied by a bar chart that displays the sender's estimate and the true value (69%) side by side, with bar heights scaled proportionally. The sender's categorical prediction is marked as correct or incorrect with a color-coded indicator. A summary reiterates the two key findings: (i) receiving information strongly changed behavior and (ii) the effect was strongest for participants who initially held false beliefs.⁶⁴

⁶³This implements a piecewise-linear proper scoring rule that is incentive-compatible under risk neutrality and approximately so under moderate risk aversion.

⁶⁴The gap between the participant's estimate and the true value is computed individually and displayed dynamically in the experimental interface. Senders who estimated the effect exactly see a confirmation message rather than a deviation.

Stage 3: Expected Attitude Change Elicitation (All Senders). All senders then predict how receiving each of the four immigration statistics would shift the attitudes of hypothetical receivers. For each statistic j , senders allocate 10 hypothetical participants—described as holding the receiver belief state realized for that statistic (false belief or “don’t know”)—among three categories: those who become *more supportive* of immigration, those who become *less supportive* of immigration, and those whose opinions *do not change*. The allocations must sum to 10. This procedure generates the expected attitude change measure $\Delta\text{Attitude}_{i,j}$ defined in Section 2.3.1, which provides within-sender variation in perceived persuasiveness across statistics, and between-sender variation across receiver belief states.

You will estimate how ten participants who hold **false beliefs** or **do not know the answer** would change their attitudes about immigration if they **are provided with research findings**.

In particular, you will estimate how many of them...

- become less supportive of immigration.
- don’t change their opinion.
- become more supportive of immigration.

The interface displays 10 person icons to make the allocation task concrete. Senders complete this task four times: once for each of the four statistics, with the hypothetical receivers described by the belief state of the receiver group matched to that statistic. The order of statistics is randomized across senders.

A.5 Experimental Instructions and Interface

Available upon request.

B Experiment 1 Supplementary Tables and Figures

B.1 Sample Characteristics

Table B.4: Sample characteristics and treatment balance

	Population (1)	Sample (2)	CONTROL (3)	ACTION (4)	INFO (5)	ACTION×INFO (6)
<i>Gender</i>						
Female	0.509	0.502	0.507	0.490	0.503	0.510
Male	0.491	0.478	0.470	0.477	0.477	0.487
<i>Ethnicity</i>						
White	0.578	0.587	0.577	0.590	0.587	0.593
Black	0.121	0.129	0.137	0.120	0.140	0.120
Hispanic	0.187	0.182	0.180	0.183	0.177	0.187
Asian	0.059	0.058	0.063	0.063	0.057	0.050
Other	0.055	0.044	0.043	0.043	0.040	0.050
<i>Political Affiliation</i>						
Democrat	0.386	0.388	0.377	0.393	0.393	0.387
Independent	0.310	0.308	0.323	0.303	0.303	0.300
Republican	0.304	0.298	0.297	0.297	0.300	0.300

Sample: $N = 1,200$ U.S. adults recruited on Prolific (quota-sampled), Experiment 1. Column (1): 2020 U.S. Census population proportions. Column (2): full sample proportions. Columns (3)–(6): proportions by treatment cell (CONTROL = no-Action, no-Info; ACTION = Action only; INFO = Info only; ACTION×INFO = Action and Info).

B.2 Additional Summary Statistics

Table B.5: Prior knowledge of the four immigration statistics

	Share of Responses		
	Correct	False	Don't know
Crime	0.700	0.223	0.076
Labor	0.794	0.141	0.065
Border	0.612	0.317	0.071
Fiscal	0.304	0.553	0.143

Sample: $N = 1,200$ senders, Experiment 1. Elicited before senders were informed they would make provision decisions. *Correct:* stated belief matches the verified fact direction; *False:* stated belief contradicts the fact. *Don't know:* explicit non-answer (senders who select “Don't know” receive a 25% probability of being scored as correct, incentivizing truthful reporting of ignorance).




Figure B.2: Distribution of senders donation allocations to the advocacy groups ($N = 1,200$, Experiment 1). Horizontal axis: allocation to the pro-immigration group collapsed into five categories (“All pro = \$500, “Pro-leaning = \$251–499, “Equal-split = \$250, “Anti-leaning = \$1–249, “All anti = \$0) from a total \$500 stake. Vertical axis: proportion of senders.




Figure B.1: Senders’ beliefs about the persuasive effect of information (pre-treatment).

Notes: Distribution of senders’ estimates of the effect of providing information on the receiver’s likelihood of signing a petition, elicited before the provision task. $N = 1,200$ senders, Experiment 1; bins are 5 percentage points wide. The dashed line marks the sample mean (37.5); the solid green line marks the true effect of 69% from [Haaland and Roth \(2020\)](#).

B.3 Robustness of Main Results

Logistic Regression. Table B.10 re-estimates specification (1) using logistic regression, which is the correctly specified model if the disclosure shock ε of Section 3.2 is logistically distributed. The framework assumes only that F has full support and a strictly positive density, so the logit is a parametric special case rather than an implication of the framework. All results replicate: the $\text{Align} \times \text{ACTION}$ interaction is positive and significant across all specifications ($p < 0.05$), and the $\text{Align} \times \text{INFO}$ interaction remains null. I report the linear probability model in the main text for interpretability of coefficients as percentage-point effects.

Table B.6: Perceived Credibility of Statistics by Political Affiliation

	Percentage Perceiving Statistic as Credible			
	All	Democrats	Independents	Republicans
Crime	0.962	0.996	0.949	0.933
Labor	0.978	0.998	0.965	0.964
Border	0.978	0.985	0.976	0.972
Fiscal	0.968	0.981	0.962	0.958

Sample: $N = 1,200$ senders, Experiment 1. Credibility assessed post-provision (potentially endogenous to the provision decision). Values shown are the share of senders rating each statistic as credible (1 = credible, 0 = not credible).

Table B.7: Expected effects of information on political preferences

Statistic	Expect Effect of Information on Political Preferences			
	Crime (1)	Labor (2)	Border (3)	Fiscal (4)
Constant	0.248*** (0.016)	0.168*** (0.016)	-0.338*** (0.018)	-0.322*** (0.020)
INFO	0.058*** (0.019)	0.038** (0.019)	-0.033 (0.021)	0.031 (0.023)
Estimated Effect Size	0.000 (0.000)	0.001*** (0.000)	0.001*** (0.000)	0.002*** (0.000)
Clustered SE	No	No	No	No
Controls	No	No	No	No
R ²	0.009	0.018	0.017	0.022
Adj. R ²	0.007	0.016	0.015	0.020
Num. obs.	1200	1200	1200	1200

Sample: $N = 1,200$ senders, Experiment 1. *Outcome:* expected effect of each statistic on receivers' immigration policy preferences, $[-1, 1]$ (positive = pro-immigration shift). *Info:* binary indicator for the between-subject feedback treatment. *Estimated Effect Size:* sender's incentivized estimate of the treatment effect magnitude in Haaland and Roth (2020), included as a covariate. Standard errors are not clustered. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Binary Alignment. Table B.11 re-estimates specification (1) with a binary aligned indicator in place of the continuous alignment measure. Because $\text{Align}_{i,j}$ is the product of the sender's preference and the statistic's direction, its magnitude is the sender's preference intensity and is constant across that sender's four statistics; the continuous specification therefore imposes that the provision gap is proportional to intensity rather than testing it. The binary coding drops that restriction, and the coefficient becomes the aligned-unaligned contrast itself, identical for every sender. Every conclusion is unchanged: alignment and its ACTION interaction are significant in all three columns (+5.0 pp, $p < 0.01$, and +6.4 pp, $p < 0.05$, in column 3), and every term

Table B.8: Expected effects of information on political preferences with controls

Statistic	Expect Effect of Information on Political Preferences			
	Crime (1)	Labor (2)	Border (3)	Fiscal (4)
<i>Panel (A): INFO</i>				
Constant	0.236*** (0.067)	0.033 (0.071)	−0.336*** (0.073)	−0.259*** (0.080)
INFO	0.055*** (0.019)	0.035* (0.019)	−0.034 (0.021)	0.028 (0.023)
Estimated Effect Size	0.000 (0.000)	0.001*** (0.000)	0.001*** (0.000)	0.001*** (0.000)
Belief Accuracy	0.116*** (0.023)	0.133*** (0.027)	0.003 (0.024)	0.129*** (0.026)
<i>Panel (B): Receiver Belief State</i>				
Constant	0.249*** (0.068)	0.043 (0.071)	−0.347*** (0.074)	−0.255*** (0.080)
Receiver Belief False	0.014 (0.019)	0.006 (0.019)	−0.000 (0.021)	0.018 (0.023)
Estimated Effect Size	0.000 (0.000)	0.001*** (0.000)	0.001*** (0.000)	0.001*** (0.000)
Belief Accuracy	0.116*** (0.023)	0.134*** (0.027)	0.001 (0.024)	0.129*** (0.026)
Clustered SE	No	No	No	No
Controls	Yes	Yes	Yes	Yes
<i>Panel (A): R²</i>	0.049	0.054	0.047	0.076
<i>Panel (B): R²</i>	0.042	0.051	0.044	0.075
<i>Panel (A) Adj. R²</i>	0.031	0.036	0.029	0.058
<i>Panel (B) Adj. R²</i>	0.025	0.033	0.027	0.058
Num. obs.	1200	1200	1200	1200

Sample: $N = 1,200$ senders, Experiment 1. *Outcome:* expected effect of each statistic on receivers' immigration policy preferences, $[-1, 1]$ (positive = pro-immigration shift). *Panel (A):* regressor is the INFO indicator. *Panel (B):* regressor is Receiver Belief False (= 1 if the matched receiver holds an incorrect prior; = 0 if the receiver reports not knowing). *Estimated Effect Size:* sender's incentivized estimate of the Haaland & Roth (2020) treatment effect. *Belief Accuracy:* incentivized prior knowledge indicator. Controls: demographic covariates. Standard errors are not clustered. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

involving INFO remains null. The interaction is larger than its continuous counterpart because the two are different objects—a contrast rather than a per-unit effect.

Alternative Information Provision Specification. Table B.12 allows asymmetric INFO effects by fact direction to accommodate the asymmetric belief-shifting effect documented in Table B.7. The additional terms $\text{INFO}_i \times \text{PROSTAT}_j$ and $\text{Align}_{i,j} \times \text{INFO}_i \times \text{PROSTAT}_j$ are not statistically significant in the specifications with fact fixed effects.

Table B.9: Treatment Effects on Alignment Bias

	Alignment Bias	
	(1)	(2)
Intercept	0.067*** (0.021)	−0.026 (0.071)
<i>Treatment</i> : ACTION	0.065** (0.029)	0.066** (0.029)
<i>Treatment</i> : INFO	0.008 (0.029)	−0.000 (0.029)
<i>Treatment</i> : ACTION × INFO	0.029 (0.029)	0.024 (0.029)
Controls	No	Yes
R ²	0.005	0.048
Adj. R ²	0.003	0.029
Num. obs.	1114	1114

Sample: $N = 1,114$ senders with revealed directional policy preference ($AIC_i \neq 250$), Experiment 1. *Outcome*: sender-level alignment bias (mean aligned minus mean unaligned provision rate). Reference: CONTROL (no-Action, no-Info). *Treatment* ACTION: Action only. *Treatment* INFO: Info only. *Treatment* ACTION × INFO: Action and Info. Column (2) adds age, gender, education, ethnicity, and party identification. Standard errors are not clustered. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Wald tests fail to reject that INFO changes the alignment gradient for pro-immigration facts ($p = 0.403$) or for anti-immigration facts ($p = 0.606$), and fail to reject that INFO changes overall provision for pro-immigration facts ($p = 0.390$). The core result—Align × ACTION—is unchanged.

Leave-One-Statistic-Out Robustness. Table B.13 re-estimates the preferred specification (Table 3, column 3) four times, each time dropping one of the four statistics. Both coefficients are stable in magnitude across all four runs—Align stays within 0.031–0.043 (full sample: 0.037) and remains significant at $p < 0.05$ in every run; Align × ACTION stays within 0.039–0.044 (full sample: 0.041) and is significant at $p < 0.05$ in three of four runs and at $p < 0.10$ in the fourth. No single statistic drives either result.

B.4 Heterogeneous Treatment Effects

B.5 Mechanism-Consistent Evidence

Table B.10: Treatment Effects on Information Provision Logistic

	Probability of Statistic Provision		
	(1)	(2)	(3)
Align	0.460*** (0.133)	0.416*** (0.134)	0.430*** (0.137)
INFO	-0.126 (0.183)	-0.126 (0.184)	-0.126 (0.188)
ACTION	0.173 (0.195)	0.172 (0.196)	0.167 (0.200)
Belief Accuracy	0.399*** (0.103)	0.308*** (0.114)	0.317*** (0.113)
Align $\times$ INFO	0.005 (0.177)	-0.009 (0.177)	-0.010 (0.183)
Align $\times$ ACTION	0.535** (0.210)	0.538** (0.209)	0.553*** (0.214)
INFO $\times$ ACTION	-0.234 (0.259)	-0.233 (0.260)	-0.241 (0.268)
Align $\times$ INFO $\times$ ACTION	-0.434 (0.270)	-0.425 (0.270)	-0.433 (0.277)
Clustered SE	Subject	Subject	Subject
Statistics FE	No	Yes	Yes
Controls	No	No	Yes
No. of Senders	1200	1200	1200
AIC	3258.054	3247.312	3248.924
BIC	3316.341	3325.029	3449.691
Log Likelihood	-1620.027	-1611.656	-1593.462
Deviance	3240.054	3223.312	3186.924
Num. obs.	4800	4800	4800

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. Logistic regression; coefficients are log-odds (not percentage-point effects). Column structure matches Table 3: column (1) no fixed effects or controls; column (2) adds statistic fixed effects; column (3) adds statistic fixed effects and demographic controls. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.11: Treatment Effects on Information Provision (Binary Alignment)

	Probability of Statistic Provision		
	(1)	(2)	(3)
Aligned	0.056*** (0.019)	0.050*** (0.019)	0.050*** (0.019)
INFO	-0.017 (0.026)	-0.014 (0.026)	-0.014 (0.025)
ACTION	-0.046* (0.027)	-0.045* (0.027)	-0.048* (0.027)
Belief Accuracy	0.047*** (0.012)	0.036*** (0.013)	0.035*** (0.013)
Aligned $\times$ INFO	0.009 (0.027)	0.004 (0.027)	0.004 (0.027)
Aligned $\times$ ACTION	0.065** (0.028)	0.064** (0.028)	0.064** (0.028)
INFO $\times$ ACTION	0.028 (0.038)	0.025 (0.038)	0.024 (0.038)
Aligned $\times$ INFO $\times$ ACTION	-0.047 (0.041)	-0.042 (0.041)	-0.041 (0.041)
Clustered SE	Subject	Subject	Subject
Statistics FE	No	Yes	Yes
Controls	No	No	Yes
No. of Senders	1114	1114	1114
Num. obs.	4456	4456	4456
R ²	0.028	0.032	0.041
Adj. R ²	0.026	0.030	0.034

Sample: 1,114 senders with a directional policy preference, 4,456 sender-statistic observations, Experiment 1. Linear probability model. *Aligned* is an indicator equal to one if the statistic's policy direction matches the sender's own position, replacing the continuous alignment measure of Table 3. The binary coding imposes no proportionality between the provision gap and preference intensity, so the coefficient is the aligned-unaligned contrast itself rather than a per-unit effect. Senders who split the donation evenly have no aligned statistic and are excluded; they enter Table 3 but supply no alignment variation there. Column structure matches Table 3: column (1) no fixed effects or controls; column (2) adds statistic fixed effects; column (3) adds statistic fixed effects and demographic controls. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.12: Treatment Effects on Information Provision

	Probability of Statistic Provision		
	(1)	(2)	(3)
Align	0.041*** (0.012)	0.037*** (0.012)	0.037*** (0.013)
INFO	-0.034* (0.019)	-0.018 (0.020)	-0.018 (0.020)
ACTION	-0.002 (0.016)	-0.002 (0.016)	-0.003 (0.016)
Belief Accuracy	0.038*** (0.012)	0.035*** (0.012)	0.034*** (0.012)
Align $\times$ INFO	-0.015 (0.021)	-0.009 (0.021)	-0.011 (0.022)
Align $\times$ ACTION	0.041** (0.018)	0.041** (0.018)	0.041** (0.018)
INFO $\times$ ACTION	-0.007 (0.024)	-0.007 (0.024)	-0.008 (0.024)
Align $\times$ INFO $\times$ ACTION	-0.027 (0.026)	-0.027 (0.026)	-0.027 (0.026)
INFO $\times$ PROSTAT	0.038*** (0.015)	0.005 (0.020)	0.005 (0.020)
Align $\times$ INFO $\times$ PROSTAT	0.025 (0.023)	0.025 (0.023)	0.029 (0.025)
Clustered SE	Subject	Subject	Subject
Statistics FE	No	Yes	Yes
Controls	No	No	Yes
No. of Senders	1200	1200	1200
Num. obs.	4800	4800	4800
R ²	0.032	0.034	0.041
Adj. R ²	0.030	0.031	0.034

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. Extends Table 3 with asymmetric INFO effects by statistic direction. *ProStat*: binary indicator for pro-immigration statistics (Crime and Labor). INFO $\times$ ProStat tests whether the INFO treatment shifts provision differentially by statistic direction; Align $\times$ INFO $\times$ ProStat tests whether the alignment gradient shifts differentially. Column structure matches Table 3. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.13: Leave-One-Statistic-Out Robustness of the Alignment and Action-Amplification Effects — Experiment 1

Coefficient	Sign-positive	Min β	Max β	Max p
Align	4/4	0.031	0.043	0.028
Align $\times$ Action	4/4	0.039	0.044	0.068
Statistics FE	Yes			
Controls	Yes			
Clustered SE	Sender			

Sample: Experiment 1 senders. Each coefficient is re-estimated four times in the preferred specification (Table 3, column 3; statistic fixed effects, full demographic and political controls, sender-clustered standard errors), each time dropping one of the four statistics (crime, labor, border, fiscal). Columns report the number of leave-one-out estimates with a positive coefficient, the minimum and maximum coefficient across the four runs, and the largest p -value observed. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.




Figure B.3: Provision rates by expected effect on receiver's immigration preferences, Experiment 1 ($N = 1,200$). Horizontal axis: *AlignPersuasion* = expected preference shift toward receiver's policy position, signed by sender direction, $\in [-1, 1]$. Vertical axis: share of statistics disclosed. Red dots: sender expects a preference-unaligned shift ($\delta < 0$); gray squares: no expected shift ($\delta = 0$); blue dots: sender expects a preference-aligned shift ($\delta > 0$). Dashed horizontal lines indicate the within-category mean provision rate (red for unaligned, blue for aligned category). Error bars are 95% confidence intervals.

Table B.14: Heterogeneity in Selective Provision (Binary Alignment)

	Probability of Statistic Provision		
	(1)	(2)	(3)
Aligned	−0.008 (0.024)	0.083*** (0.024)	0.004 (0.033)
Aligned × ACTION	0.075** (0.035)	0.021 (0.033)	−0.011 (0.048)
Aligned × Partisanship	0.123*** (0.038)		
Aligned × ACTION × Partisanship	−0.056 (0.059)		
Aligned × Independent		−0.071** (0.033)	
Aligned × Republican		−0.006 (0.035)	
Aligned × ACTION × Independent		0.096* (0.050)	
Aligned × ACTION × Republican		−0.031 (0.050)	
Aligned × Preference intensity			0.082* (0.042)
Aligned × ACTION × Preference intensity			0.069 (0.061)
Moderator	Partisanship	Political party	Pref. intensity
Clustered SE	Subject	Subject	Subject
Statistics FE	Yes	Yes	Yes
Controls	Yes	Yes	Yes
No. of Senders	1114	1114	1114
Num. obs.	4456	4456	4456
R ²	0.041	0.040	0.044
Adj. R ²	0.036	0.033	0.038

Sample: senders with a directional policy preference, Experiment 1. Pre-registered exploratory heterogeneity. Each column interacts a binary aligned indicator, the ACTION treatment, and one moderator, with statistic fixed effects and demographic controls. Reference category for political party is Democrat. Two features of the specification matter for interpretation. Alignment enters as a binary indicator rather than the continuous measure: because $|\text{Align}_{i,j}|$ is the sender's preference intensity, an interaction between the continuous measure and any intensity-correlated moderator partly re-encodes the proportionality the construction imposes rather than heterogeneity in the provision gap. And statistic fixed effects are required because alignment maps the sender's side onto a fixed pair of statistics, and the two pro-immigration statistics are provided 6.6 percentage points more often than the two anti-immigration ones; a sender-level specification cannot absorb that difference. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.15: Receiver belief state: perceived effect of providing information

Statistic	Expect Effect of Information on Political Preferences			
	Crime (1)	Labor (2)	Border (3)	Fiscal (4)
<i>Panel (A): Receiver Belief State</i>				
Constant	0.265*** (0.016)	0.185*** (0.017)	−0.352*** (0.018)	−0.313*** (0.020)
Receiver Belief False	0.020 (0.019)	0.004 (0.019)	−0.003 (0.021)	0.011 (0.023)
Estimated Effect Size	0.000 (0.000)	0.001*** (0.000)	0.001*** (0.000)	0.002*** (0.000)
R ²	0.002	0.015	0.015	0.021
Adj. R ²	0.000	0.013	0.013	0.019
Num. obs.	1200	1200	1200	1200

*** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$. Panel (A): Dependent variable is the sender's expected effect of each statistic on receivers' immigration preferences (range $[-1, 1]$). Receiver Belief False equals one if the matched receiver holds a false belief, zero if the receiver reports not knowing. The null effect confirms that the Δb variation (receiver belief state) does not contaminate senders' perceived δ : the two channels are empirically separable.

Table B.16: Heterogeneous Treatment Effects of Receiver Beliefs

	Probability of Statistic Provision	
	(1)	(2)
Intercept	0.863*** (0.018)	0.894*** (0.045)
Align	0.026* (0.014)	0.025* (0.014)
INFO	-0.034 (0.021)	-0.034 (0.022)
ACTION	-0.021 (0.020)	-0.022 (0.021)
Receiver Belief False	-0.025 (0.018)	-0.024 (0.018)
Accuracy Prior Belief	0.035*** (0.012)	0.034*** (0.012)
Align $\times$ INFO	0.025 (0.023)	0.025 (0.023)
Align $\times$ ACTION	0.053** (0.024)	0.054** (0.024)
INFO $\times$ ACTION	0.014 (0.030)	0.014 (0.031)
Align $\times$ Receiver Belief False	0.019 (0.022)	0.020 (0.022)
INFO $\times$ Receiver Belief False	0.043 (0.027)	0.043 (0.027)
ACTION $\times$ Receiver Belief False	0.034 (0.026)	0.034 (0.026)
Align $\times$ INFO $\times$ ACTION	-0.050 (0.035)	-0.051 (0.035)
Align $\times$ INFO $\times$ Receiver Belief False	-0.042 (0.032)	-0.042 (0.032)
Align $\times$ ACTION $\times$ Receiver Belief False	-0.020 (0.033)	-0.022 (0.033)
INFO $\times$ ACTION $\times$ Receiver Belief False	-0.041 (0.038)	-0.043 (0.038)
Align $\times$ INFO $\times$ ACTION $\times$ Receiver Belief False	0.047 (0.045)	0.048 (0.045)
Clustered SE	Subject	Subject
Statistics FE	Yes	Yes
Controls	No	Yes
No. of Senders	1200	1200
R ²	0.034	0.041
Adj. R ²	0.030	0.033
Num. obs.	4800	4800

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. Outcome: binary provision indicator. Full four-way interaction: Align, ACTION, INFO, and Receiver Belief False (= 1 if the matched receiver holds a factually incorrect prior belief; = 0 if the receiver reports not knowing). Column (2) adds demographic controls. Summary statistics from this specification are reported in Appendix Table B.15, Panel (B). Statistic fixed effects and sender-clustered standard errors throughout. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.17: Does the Alignment Effect Vary with the Sender's Beliefs?

	Probability of Statistic Provision		
	(1)	(2)	(3)
Align	0.031* (0.016)	0.052 (0.048)	0.044 (0.049)
Belief Accuracy	0.035*** (0.012)	0.024** (0.012)	0.026** (0.012)
Credibility		0.253*** (0.063)	0.253*** (0.063)
Align $\times$ Belief Accuracy	0.009 (0.016)		0.015 (0.015)
Align $\times$ Credibility		-0.021 (0.047)	-0.025 (0.047)
Align $\times$ ACTION	0.041** (0.018)	0.044** (0.018)	0.044** (0.018)
Align $\times$ INFO	0.004 (0.017)	0.007 (0.017)	0.007 (0.017)
Align $\times$ INFO $\times$ ACTION	-0.027 (0.026)	-0.032 (0.025)	-0.032 (0.025)
Clustered SE	Subject	Subject	Subject
Statistics FE	Yes	Yes	Yes
Controls	Yes	Yes	Yes
No. of Senders	1200	1200	1200
Num. obs.	4800	4800	4800
R ²	0.040	0.058	0.058
Adj. R ²	0.034	0.052	0.052

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. Each column is specification (1) with statistic fixed effects and demographic controls, augmented by an interaction between alignment and one of the two belief measures. *Belief Accuracy* is the sender's incentivized prior about the statistic, elicited before any provision decision (1 correct, 0.5 "don't know," 0 incorrect). *Credibility* is the sender's unincentivized rating of the statistic, elicited after the provision task. The two interactions carry different weight: belief accuracy varies across the full sample, whereas 97.2% of statistics are rated credible, so the credibility interaction is identified from 136 decisions and its confidence interval is several times the alignment effect. The main effect of alignment in each column is the alignment effect at zero on the interacted measure. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.18: Self-Reported Motivations and Alignment Bias (Experiment 1)

	Alignment Bias			
	(1)	(2)	(3)	(4)
Intercept	0.109*** (0.014)	0.091*** (0.017)	0.098*** (0.016)	0.097*** (0.016)
<i>Belief correction</i>	−0.099*** (0.017)	−0.099*** (0.017)	−0.096*** (0.017)	−0.096*** (0.017)
<i>Expression</i>	0.636*** (0.105)	0.630*** (0.105)	0.605*** (0.106)	0.749*** (0.129)
<i>Persuasion</i>	0.276*** (0.084)	0.276*** (0.083)	0.070 (0.115)	0.056 (0.118)
<i>Other stated reason</i>	−0.000 (0.129)	0.001 (0.131)	0.001 (0.130)	0.001 (0.130)
ACTION		0.036* (0.020)	0.021 (0.019)	0.022 (0.019)
<i>Persuasion</i> × ACTION			0.392*** (0.146)	0.419*** (0.155)
<i>Expression</i> × ACTION				−0.219 (0.190)
Receiver Action	No	Yes	Yes	Yes
Num. obs.	1102	1102	1102	1102
R ²	0.099	0.102	0.114	0.115
Adj. R ²	0.096	0.098	0.109	0.109

Sample: $N = 1,102$ senders with a revealed directional policy preference ($AIC_i \neq 250$), Experiment 1. *Outcome:* sender-level alignment bias (mean aligned minus mean unaligned provision rate). *Regressors:* indicators for each motive the sender’s open-ended decision reason was coded as citing, by `claude-sonnet-5` under frozen coding scheme v4.4 (Appendix E.1); the omitted reference is a sender whose response was coded as citing none of the listed motives. *Register:* exploratory. The reasons are unincentivized and were recorded after the provision task, so the estimates are associations between what senders said and what they did, not measurements of motives. Categories are not mutually exclusive, so the indicators are not a partition. *Belief correction* marks direction-blind truth-telling or the receiver’s right to accurate information; *Expression* marks reasons citing the fact’s fit with the sender’s own position; *Persuasion* marks receiver-directed reasons. ACTION is the receiver political-action treatment. Columns (3)–(4) add the interaction of each political motive with ACTION. The *Expression* indicator is affirmed for 12 senders (4 in CONTROL, 8 in ACTION) and the *Persuasion* indicator for 43 (20 and 23). Heteroskedasticity-robust standard errors in parentheses. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.19: Stated Accuracy Norms and Provision Behavior

	Correct false beliefs	Facts matter for opinions	Platforms should fact-check	Disagreement is about facts
	(1)	(2)	(3)	(4)
Provision rate	0.960*** (0.159)	0.791*** (0.138)	0.640*** (0.186)	0.490*** (0.162)
Alignment bias	0.028 (0.075)	0.061 (0.073)	−0.091 (0.105)	0.033 (0.098)
Controls	Yes	Yes	Yes	Yes
No. of Senders	1114	1114	1114	1114
R ²	0.086	0.083	0.100	0.030
Adj. R ²	0.075	0.072	0.089	0.019
Num. obs.	1114	1114	1114	1114

Sample: senders with a directional policy preference, Experiment 1. The dependent variable is agreement with one post-task statement about accurate information, on a five-point scale; column (4) is reverse-coded so that higher values indicate the pro-accuracy pole on every item. Behavior is the regressor: *Provision rate* is the share of statistics the sender provided and *Alignment bias* is the share of aligned statistics provided minus the share of unaligned statistics provided. The items were elicited after the provision task, so the direction of this specification follows [Ambuehl et al. \(2021\)](#), who relate stated policy positions to behavior rather than the reverse; the estimates are coherence evidence and identify no causal channel. Because provision averages roughly 89%, a high-provision sender has a mechanically compressed feasible alignment gap, so the coefficient on alignment bias conditional on the provision rate should not be read as an absence of association between norms and selectivity. All columns include demographic and political controls and treatment indicators. Heteroskedasticity-robust standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.20: Self-Reported Motivations and Alignment Bias, Secondary Classifier (Experiment 1)

	Alignment Bias			
	(1)	(2)	(3)	(4)
Intercept	0.112*** (0.014)	0.093*** (0.017)	0.100*** (0.017)	0.099*** (0.017)
<i>Belief correction</i>	-0.099*** (0.017)	-0.100*** (0.017)	-0.096*** (0.017)	-0.096*** (0.017)
<i>Expression</i>	0.628*** (0.092)	0.620*** (0.093)	0.545*** (0.113)	0.757*** (0.112)
<i>Persuasion</i>	0.187*** (0.070)	0.187*** (0.069)	0.034 (0.086)	0.027 (0.085)
<i>Other stated reason</i>	-0.045 (0.053)	-0.048 (0.053)	-0.046 (0.053)	-0.046 (0.053)
ACTION		0.037* (0.020)	0.022 (0.020)	0.023 (0.020)
<i>Persuasion</i> × ACTION			0.308** (0.126)	0.336** (0.132)
<i>Expression</i> × ACTION				-0.320* (0.188)
Receiver Action	No	Yes	Yes	Yes
Num. obs.	1102	1102	1102	1102
R ²	0.078	0.081	0.091	0.092
Adj. R ²	0.075	0.077	0.086	0.086

Sample: $N = 1,102$ senders with a revealed directional policy preference ($AIC_i \neq 250$), Experiment 1. *Outcome:* sender-level alignment bias (mean aligned minus mean unaligned provision rate). *Regressors:* indicators for each motive the sender's open-ended decision reason was coded as citing, from the *secondary* classifier (`gpt-5.6-terra`) applied to the same frozen prompt and the same 1,200 responses as the primary classifier of Table B.18 (Appendix E.1); the omitted reference is a sender whose response was coded as citing none of the listed motives. Categories are not mutually exclusive. ACTION is the receiver political-action treatment. Heteroskedasticity-robust standard errors in parentheses. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.21: Stated Motives and Revealed Selection Returns (Experiments 2 and 3)

	Provision	
	(1)	(2)
Expression (α)	0.151*** (0.020)	0.124*** (0.042)
Stated Expression	−0.019 (0.038)	−0.253*** (0.078)
Persuasion ($\Delta\pi$)	0.004 (0.007)	0.004 (0.015)
Stated Persuasion	−0.071*** (0.023)	−0.091*** (0.031)
Belief correction (Δb)	0.028 (0.018)	0.023 (0.052)
Stated Belief correction	0.236*** (0.033)	0.228*** (0.032)
$\alpha \times$ Stated Expression	−0.086 (0.077)	0.336** (0.165)
$\Delta\pi \times$ Stated Persuasion	0.107*** (0.032)	0.090*** (0.029)
$\Delta b \times$ Stated Belief correction	0.082 (0.068)	0.004 (0.066)
Sample	General population (Exp. 2) Journalists (Exp. 3)	
Fact FE	Yes	Yes
Num. obs.	7740	1530
R ²	0.029	0.166
Adj. R ²	0.027	0.155

Sample: decision-level provision, Experiment 2 general-population senders (column 1) and Experiment 3 journalists (column 2). *Specification:* each revealed selection return—the expressive rating α_{ij} (Expression), the elicited persuasion return $\Delta\pi_{ijk}$ (Persuasion), and the elicited factual belief correction Δb_{ijk} (Belief correction)—is interacted with an indicator for whether the sender’s open-ended decision reason states the corresponding motive (Appendix E.1). A positive interaction means senders who *state* a motive select more strongly on the matching *revealed* return. Fact fixed effects included; standard errors clustered by sender. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

C Experiment 2 Supplementary Tables and Figures

C.1 Political Preferences of Senders




Figure C.4: Distribution of policy preferences among Experiment 2 senders ($N = 1,548$). *Top row (Panels A–B):* stated preferences from the 5-point Likert item; tick labels complete the stance statement shown in the panel title (e.g., *Immigration to Germany should be ... greatly eased*). *Bottom row (Panels C–D):* revealed preferences from the real donation task, grouped into five buckets from All pro (€250 to the pro-policy organization) to All anti (€0), displayed pro-to-anti left to right. *Left column:* immigration senders ($N = 809$). *Right column:* climate senders ($N = 739$).

C.2 Heterogeneity, Controls, and Provision Figures

I omit demographic controls from the main-text provision tables because the alignment effect is identified within sender and the ACTION effect is randomized. Adding the full set of sender demographics to the pre-registered tests and to the action decomposition leaves the coefficients stable (see Tables C.25 and C.26).

Positional alignment combines two preference measures, and Table C.29 rebuilds it from each alone. Both single-measure coefficients are smaller than the composite

Table C.22: Heterogeneity by Sender Profile: Crude Alignment and Action — Experiment 2 (General Population)

Moderator Topic	Probability of Fact Provision				
	(1) Distance Pooled	(2) polarization Pooled	(3) Pol. group Pooled	(4) Immigration	(5) Climate
Alignment	0.119*** (0.019)	0.115*** (0.019)	0.124*** (0.021)	0.119*** (0.017)	0.143*** (0.020)
Action	-0.006 (0.017)	0.040** (0.018)	0.004 (0.018)	0.026 (0.018)	0.026 (0.017)
Alignment × Action	-0.018 (0.027)	0.002 (0.026)	0.008 (0.028)	0.050** (0.022)	-0.042 (0.029)
Receiver distance	-0.094*** (0.007)				
Alignment × Distance	0.001 (0.008)				
Action × Distance	0.019** (0.009)				
Alignment × Action × Distance	0.020* (0.012)				
Sender polarization		0.010* (0.006)			
Alignment × Polarization		0.007 (0.009)			
Action × Polarization		-0.010 (0.009)			
Alignment × Action × Polarization		0.007 (0.012)			
Extremes (ref: center)			-0.045** (0.020)		
Others (ref: center)			-0.093*** (0.026)		
Alignment × Extremes			0.001 (0.028)		
Alignment × Others			0.024 (0.040)		
Action × Extremes			0.024 (0.027)		
Action × Others			0.067* (0.037)		
Alignment × Action × Extremes			0.016 (0.038)		
Alignment × Action × Others			-0.007 (0.052)		
Clustered SE	Sender	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes	Yes
Moderator Topic	Distance Pooled	Polarization Pooled	Pol. group Pooled	—	—
Spec	(1)	(2)	(3)	Immigration (4)	Climate (5)
Num. obs.	7740	7740	7715	4045	3695
R ²	0.079	0.033	0.035	0.035	0.030
Adj. R ²	0.077	0.030	0.032	0.033	0.028

Sample: $N = 1,548$ Experiment 2 general-population senders (Bilendi panel, Germany, 2026). Columns (1)–(2): 7,740 sender-fact observations. Column (3): 7,715 (non-missing political group). Columns (4)–(5): topic subsamples (immigration: $N = 809$ senders; climate: $N = 739$ senders). Outcome: binary provision indicator. Alignment: $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, 1]$ (pre-registered confirmatory measure). Receiver distance: $|\text{sender_likert} - \text{rcv_stance}| \in \{0, 1, 2, 3, 4\}$. Sender polarization: absolute left-right distance from center. Political group: Extremes = AfD, Linke, BSW; Others = other/no vote; Center = CDU/CSU, SPD, Grüne, FDP (reference). Columns (1)–(3) include the moderator and all interactions; columns (4)–(5) are subsample estimates with no moderator. All specifications include fact fixed effects and sender-clustered standard errors. Multiple testing: the exploratory heterogeneity family spans twenty interaction terms across this table and Table C.23; Benjamini–Hochberg-adjusted q -values within that family are approximately 0.20 for every interaction significant at raw $p < 0.10$. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.23: Heterogeneity by Sender Profile: Motive Decomposition — Experiment 2 (General Population)

	Probability of Fact Provision		
	(1) Distance	(2) Polarization	(3) Political group
Expression	0.165*** (0.028)	0.121*** (0.028)	0.140*** (0.029)
Persuasion	0.014 (0.014)	−0.020 (0.013)	−0.020 (0.013)
Action	−0.006 (0.017)	0.039** (0.018)	0.002 (0.019)
Persuasion × Action	−0.015 (0.020)	0.027 (0.018)	0.050*** (0.018)
Receiver distance	−0.101*** (0.009)		
Expression × Distance	0.004 (0.015)		
Persuasion × Distance	−0.001 (0.006)		
Action × Distance	0.018* (0.010)		
Persuasion × Action × Distance	0.017** (0.009)		
Sender polarization		0.004 (0.008)	
Expression × Polarization		0.018 (0.013)	
Persuasion × Polarization		0.011* (0.006)	
Action × Polarization		−0.009 (0.009)	
Persuasion × Action × Polarization		−0.002 (0.008)	
Extremes (ref: center)			−0.070*** (0.027)
Others (ref: center)			−0.096*** (0.033)
Expression × Extremes			0.038 (0.043)
Expression × Others			−0.003 (0.053)
Persuasion × Extremes			0.018 (0.019)
Persuasion × Others			0.046* (0.026)
Action × Extremes			0.026 (0.027)
Action × Others			0.068* (0.037)
Persuasion × Action × Extremes			−0.032 (0.028)
Persuasion × Action × Others			−0.070* (0.037)
Clustered SE	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes
Moderator	Distance	Polarization	Political group
Spec	(1)	(2)	(3)
Num. obs.	7740	7740	7715
R ²	0.069	0.017	0.019
Adj. R ²	0.066	0.014	0.016

Sample: $N = 1,548$ Experiment 2 general-population senders (Bilendi panel, Germany, 2026); 7,740 sender-fact observations in columns (1)–(2), 7,715 in column (3) (non-missing political group). Outcome: binary provision indicator. Expression: cross-domain behavioral rating a_{ij} , elicited post-provision on a different topic. Persuasion: incentivized expected preference shift $\Delta\pi_{ijk}$ toward the sender's position. Receiver distance: $|\text{sender_likert} - \text{recv_stance}| \in \{0, 1, 2, 3, 4\}$. Sender polarization: absolute left-right distance from center. Political group: Extremes = AfD, Linke, BSW; Others = other/no vote; Center = CDU/CSU, SPD, Grüne, FDP (reference). All specifications include fact fixed effects and sender-clustered standard errors. Multiple testing: the exploratory heterogeneity family spans twenty interaction terms across this table and Table C.22; Benjamini–Hochberg-adjusted q -values within that family are approximately 0.20 for every interaction significant at raw $p < 0.10$. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.24: Motive Decomposition by Topic — Experiment 2 (General Population)

	Probability of Fact Provision	
	(1)	(2)
Expression	0.143*** (0.025)	0.158*** (0.033)
Factual belief correction	0.053** (0.025)	0.016 (0.024)
Persuasion	−0.005 (0.013)	−0.003 (0.013)
Action	0.024 (0.018)	0.026 (0.018)
Persuasion × Action	0.038** (0.017)	0.011 (0.019)
Clustered SE	Sender	Sender
Fact FE	Yes	Yes
Topic	Immigration	Climate
Spec	(1)	(2)
Num. obs.	4045	3695
R ²	0.016	0.016
Adj. R ²	0.013	0.013

Sample: $N = 1,548$ Experiment 2 general-population senders (Bilendi panel, Germany, 2026). Column (1): immigration subsample ($N = 809$ senders). Column (2): climate subsample ($N = 739$ senders). *Outcome:* binary provision indicator. Each column applies the preferred motive specification (col. (4) of Table 6) within topic subsample: Expression (α_{ij} , cross-domain behavioral), Factual belief correction (Δb_{ijk} , incentivized), Persuasion ($\Delta \pi_{ijk}$, incentivized positional), and the Persuasion × ACTION interaction. All specifications include fact fixed effects and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

(+0.100 and +0.102 against +0.128). That is consistent with averaging reducing measurement noise, if the two components are imperfect readings of the same underlying preference, though the comparison cannot separate measurement error from a difference in what each measure captures. The sender-level biases do not order the same way: 11.3 percentage points for the revealed measure alone (on 1,404 senders), 13.8 for the stated measure alone (on 1,252), and 12.2 for the composite (on 1,496), which falls between them. The three constructions differ in who receives a directional classification at all. Across every construction the alignment coefficient is positive and its interaction with ACTION remains indistinguishable from zero.

Table C.28 replaces continuous alignment with a binary aligned/unaligned indicator in all three experiments. The continuous measure imposes a gap proportional to preference intensity; the binary one imposes a constant gap instead, so the two bracket the functional form rather than nesting. Observations with zero alignment are dropped from both columns of each pair, so the specifications are estimated on

Table C.25: Provision Decisions with Demographic Controls — Experiment 2 (General Population)

	Probability of Fact Provision	
	(1)	(2)
Align	0.110*** (0.013)	0.128*** (0.013)
Action	0.021* (0.012)	0.023* (0.012)
Align $\times$ Action	0.014 (0.017)	0.014 (0.018)
Clustered SE	Sender	Sender
Fact FE	No	Yes
Controls	Yes	Yes
Num. obs.	7715	7715
R ²	0.085	0.045
Adj. R ²	0.080	0.039

Sample: $N = 1,548$ Experiment 2 general-population senders, both arms; 7,715 observations with non-missing controls. *Outcome:* binary provision indicator. *Align:* positional alignment $\text{Pref}_i \times \text{Dir}_j \in [-1, +1]$. *Controls:* sender-receiver ideological distance, age, gender, education, Bundesland, closest-party affiliation, and left-right self-placement. Column (1) omits fact fixed effects; column (2) includes them and reproduces the pre-registered preferred specification. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

identical samples. Signs, significance and the treatment interactions are unchanged throughout: Experiment 1's amplification under ACTION survives (+0.064, $p = 0.023$, against +0.041 continuous) and Experiment 2's remains indistinguishable from zero under either measure. Estimating the gap separately within intensity bins does not reject the proportional form in any experiment, or in either arm of Experiments 1 and 2 taken separately.⁶⁵

C.3 Robustness of the Motive-Channel Estimates

The action decomposition of Table 6 is robust to alternative inference, ordering, functional-form, and sample-restriction choices; all checks use the canonical positional persuasion measure $\Delta\pi_{ijk}$. Table C.30 re-estimates the preferred specification under two-way clustering by sender and fact and under an exact wild-cluster bootstrap over all $2^{12} = 4,096$ fact-level sign patterns. Table C.31 adds the within-sender presentation order and a

⁶⁵Bin-wise gaps rise with preference intensity in all three samples. A joint test of the bin-specific deviations from each specification's linear form does not reject it (Experiment 1 $p = 0.47$, Experiment 2 $p = 0.77$, Experiment 3 $p = 0.81$; by arm, $p = 0.85$ and $p = 0.92$ in Experiment 1 and $p = 0.92$ and $p = 0.37$ in Experiment 2). Because senders with zero alignment are excluded, these tests bear on the slope over the observed range of intensity and not on the behaviour at zero, where the proportional form is untestable here.

Table C.26: Action Decomposition with Demographic Controls — Experiment 2 (General Population)

	Probability of Fact Provision	
	(1)	(2)
Expression	0.149*** (0.020)	0.151*** (0.020)
Persuasion	−0.004 (0.009)	−0.003 (0.009)
Persuasion × Action	0.025* (0.013)	0.023* (0.013)
Action	0.025** (0.013)	0.022* (0.013)
Factual belief correction gain	0.034** (0.017)	0.042** (0.017)
Clustered SE	Sender	Sender
Fact FE	Yes	Yes
Controls	No	Yes
Num. obs.	7740	7715
R ²	0.015	0.029
Adj. R ²	0.013	0.022

Sample: $N = 1,548$ Experiment 2 general-population senders, both arms; column (1) 7,740 observations, column (2) 7,715 with non-missing controls. *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating $\alpha_{ij} \in [-1, +1]$. *Persuasion:* incentivized expected movement of the receiver toward the sender’s position $\Delta\pi_{ijk}$ (positional reading). *Factual belief correction gain:* incentivized expected correction of the receiver’s factual belief $\Delta b_{ijk} \in [0, 1]$. Column (1) reproduces the preferred action-decomposition specification; column (2) adds demographic controls (age, gender, education, Bundesland, closest-party affiliation, left-right self-placement). Sender-receiver ideological distance is excluded as a heterogeneity variable. All specifications include fact fixed effects. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

lagged provision indicator. Table C.32 probes the shape of $w(\Delta\pi)$ by replacing $\Delta\pi_{ijk}$ with its sign and with a version winsorized at the ninetieth percentile of $|\Delta\pi|$: the sign coding carries neither the Persuasion main effect ($\beta = +0.010$, $p = 0.23$) nor its ACTION interaction ($\beta = -0.001$), whereas winsorizing, which preserves magnitude, leaves the interaction near its canonical value ($\beta = +0.021$, $p = 0.18$). Table C.33 drops low- $|\Delta\pi|$ senders, perfect sharers, and low- α facts in turn; the instrumental interaction stays positive throughout, though it is not separately significant in these restricted subsamples under the positional measure. Table C.36 reports leave-one-fact-out estimates for each motive; the within-fact variance inflation factors are near unity ($VIF \approx 1.00$) for all three channels.

Table C.35 varies the sample definition. Column (2) relaxes the pre-registered speeder exclusion (completion time below $0.25\times$ the realised-sample median). Column (3) adds the 106 participants who started before the pre-registration update of 27 May 2026, which dropped a non-essential elicitation module after the survey’s length caused

Table C.27: Motive Channels with the Pre-Registered Positional Measure Substituted for the Elicited Expressive Return — Experiment 2

	Probability of Fact Provision			
	(1)	(2)	(3)	(4)
Expression	0.138*** (0.028)	0.149*** (0.020)		
Expression $\times$ Action	0.020 (0.037)			
Alignment			0.127*** (0.013)	0.133*** (0.009)
Alignment $\times$ Action			0.011 (0.017)	
Factual belief correction gain	0.034** (0.017)	0.034** (0.017)	0.034* (0.017)	0.033* (0.017)
Persuasion	-0.004 (0.009)	-0.004 (0.009)	-0.006 (0.009)	-0.006 (0.009)
Persuasion $\times$ Action	0.024* (0.013)	0.025* (0.013)	0.024* (0.013)	0.025* (0.013)
Action	0.018 (0.019)	0.025** (0.013)	0.025** (0.013)	0.025** (0.013)
Clustered SE	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes
Expressive measure	Elicited	Elicited	Positional	Positional
Spec	(1)	(2)	(3)	(4)
Num. obs.	7740	7740	7740	7740
R ²	0.016	0.015	0.033	0.033
Adj. R ²	0.013	0.013	0.031	0.031

Sample: $N = 1,548$ Experiment 2 senders (Bilendi panel, Germany, 2026), 7,740 sender–fact observations. *Outcome:* binary provision indicator. The specification is the motive specification of Table 6, estimated twice: once with the elicited expressive return α_{ij} (the cross-domain congeniality rating), and once with positional alignment $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, +1]$, the pre-registered confirmatory proxy, in its place. Columns (1) and (3) include the expressive measure interacted with ACTION; columns (2) and (4) omit it. All specifications include fact fixed effects and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

elevated dropouts; they completed the longer instrument (median 29 vs. 17 minutes) and are excluded from the canonical sample for instrument comparability. Column (4) adds both groups. Column (5) implements the pre-specified attention-check robustness: failers are retained in the primary analysis per the pre-registration. Excluding them leaves the Persuasion $\times$ ACTION point estimate unchanged (+0.026, still significant at the ten-percent level) and strengthens the factual belief correction slope (+0.053 vs. +0.034). Failers complete the study faster (median 12 vs. 18 minutes) and report flatter belief elicitation (mean Δb_{ijk} of 0.08 vs. 0.20 for passers), but their provision rate is nearly identical (0.641 vs. 0.638). The sign of every estimate is unchanged in every variant; the instrumental interaction attenuates and loses precision as the excluded groups

Table C.28: Alignment Measured Continuously and as a Binary Indicator

	Experiment 1		Experiment 2		Experiment 3	
	Continuous (per unit) (1)	Binary (gap) (2)	Continuous (per unit) (3)	Binary (gap) (4)	Continuous (per unit) (5)	Binary (gap) (6)
Alignment	+0.037*** (0.012)	+0.050*** (0.019)	+0.128*** (0.013)	+0.122*** (0.017)	+0.130*** (0.020)	+0.168*** (0.028)
Alignment $\times$ ACTION	+0.041** (0.018)	+0.064** (0.028)	+0.014 (0.018)	+0.025 (0.023)	—	—
Observations	4,456	4,456	7,455	7,455	1,485	1,485
Fact/topic FE	Yes	Yes	Yes	Yes	Yes	Yes
Controls	Yes	Yes	Yes	Yes	Yes	Yes

Notes. Linear probability models; the dependent variable is provision. Continuous alignment is $\text{Pref}_i \times \text{Dir}_j \in [-1, +1]$; the binary indicator is $1(\text{Align}_{ij} > 0)$. Observations with $\text{Align}_{ij} = 0$ are excluded from both columns of each pair rather than classified as unaligned, so the two specifications are estimated on identical samples. **The two entries in each pair are on different scales and are not competing estimates of one quantity:** a continuous coefficient is a per-unit slope, so the aligned–unaligned gap it implies is twice the coefficient times the sender’s preference intensity, while a binary coefficient is that gap directly, averaged over the sample. Experiment 1 retains the INFO treatment structure of Table 3. Experiment 3 is a single-condition journalist sample with no randomised ACTION treatment, so no interaction is estimable. Standard errors clustered at the sender level in parentheses. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

Table C.29: Provision and positional alignment under alternative preference measures

	(1) Composite (registered)	(2) Revealed only	(3) Stated only	(4) Excluding zero index	(5) Excluding offsetting
Positional alignment	+0.128*** (0.013)	+0.100*** (0.012)	+0.102*** (0.012)	+0.128*** (0.013)	+0.129*** (0.013)
Positional alignment $\times$ ACTION	+0.014 (0.018)	+0.009 (0.017)	+0.009 (0.016)	+0.014 (0.018)	+0.014 (0.018)
Observations	7,715	7,715	7,715	7,455	7,650
Senders	1,548	1,548	1,548	1,496	1,535
Fact fixed effects	Yes	Yes	Yes	Yes	Yes
Controls	Yes	Yes	Yes	Yes	Yes

Notes. Linear probability models of provision on positional alignment, its interaction with ACTION, and the pre-specified controls (age, gender, education, state, vote intention, left–right placement), with fact fixed effects and standard errors clustered at the sender level. Column (1) is the pre-registered specification, in which the sender’s policy preference is the average of the revealed measure from the €250 donation split and the stated measure from the five-point stance item. Columns (2) and (3) rebuild alignment from each component alone. Column (4) drops the 52 senders whose composite preference is exactly zero; column (5) drops only the 13 of those who state a position at one pole of the scale and donate against it. Stars denote two-sided significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

are added back (+0.022, +0.020, and +0.017 across columns (2)–(4), significant only in column (2)).

C.4 Distributional Detail of the Elicited Returns

Senders expect partial updating from disclosure: 11.9% of fact–receiver pairs receive a post-disclosure prediction of 100%, and 23.7% remain below 50% even after disclosure. A quarter of fact–receiver pairs (26%) carry negative expected factual belief correction, i.e., the sender expects disclosure to leave the receiver slightly further from the truth.

Table C.30: Inference Robustness: Clustering and Wild Bootstrap — Experiment 2 (General Population)

	(1) Sender SE (preferred)	(2) Two-way SE (Sender + Fact)
Expression	0.149*** (0.020)	0.149*** (0.016)
Factual belief correction gain	0.034** (0.017)	0.034 (0.021)
Persuasion	−0.004 (0.009)	−0.004 (0.009)
Persuasion × Action	0.025* (0.013)	0.025 (0.015)
Action	0.025** (0.013)	0.025* (0.012)
Clustered SE	Sender	Sender + Fact
Fact FE	Yes	Yes
Wild bootstrap p (fact, exact 2^{12})	—	0.108
Num. obs.	7,740	7,740

Sample: Experiment 2 general-population senders; 7,740 sender-fact observations. Both columns estimate the preferred action-decomposition specification (*Expression*, *Factual belief correction gain*, *Persuasion*, *Persuasion* × *Action*, *Action*) with fact fixed effects. Column (1): standard errors clustered at the sender level (preferred). Column (2): two-way clustering by sender and fact (Cameron–Gelbach–Miller; $K = 12$ fact clusters, asymptotic standard errors indicative only). *Wild bootstrap:* exact enumeration of all $2^{12} = 4,096$ Rademacher sign patterns at the fact level (impose-null variant); p -value for *Persuasion* × *Action*. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.31: Order and Serial-Autocorrelation Robustness — Experiment 2 (General Population)

	(1) + Order	(2) + Lag	(3) + Both
Expression	0.149*** (0.020)	0.164*** (0.022)	0.164*** (0.022)
Factual belief correction gain	0.033* (0.017)	0.038** (0.019)	0.038** (0.019)
Persuasion	−0.004 (0.009)	0.002 (0.011)	0.002 (0.011)
Persuasion × Action	0.025* (0.013)	0.030** (0.015)	0.030** (0.015)
Action	0.025** (0.013)	0.025* (0.013)	0.025* (0.013)
Presentation order	−0.007** (0.004)	—	0.005 (0.005)
Lag(provision)	—	0.076*** (0.014)	0.076*** (0.014)
Clustered SE	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes
Num. obs.	7,740	6,192	6,192

Sample: Experiment 2 general-population senders. Column (1): 7,740 observations (all five slots). Columns (2)–(3): slots 2–5 only (slot 1 has no lag). All columns add the listed control(s) to the preferred action-decomposition specification with fact fixed effects and sender-clustered standard errors. *Presentation order:* linear slot index $\in \{1, \dots, 5\}$. *Lag(provision):* the sender’s provision decision for the immediately preceding fact. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.




Figure C.5: Mean provision rates by expected preference-aligned belief shift and arm, Experiment 2 (General Population). Dark green: facts for which the sender expects the receiver’s policy position to move toward the sender’s own position ($\Delta\pi_{ijk} > 0$). Dark red: facts for which no such shift is expected ($\Delta\pi_{ijk} \leq 0$). x -axis: experimental arm (CONTROL or ACTION). Error bars are 95% confidence intervals (sender-level means). Sample: Experiment 2 senders with non-missing $\Delta\pi_{ijk}$ elicitation; $\Delta\pi_{ijk}$ is incentivized via Quadratic Scoring Rule.

Table C.32: Functional Form of $w(\delta)$: Direction versus Magnitude — Experiment 2 (General Population)

	(1) $\text{sign}(\delta)$	(2) Winsorized δ
Expression	0.148*** (0.020)	0.148*** (0.020)
Factual belief correction gain	0.034* (0.017)	0.033* (0.017)
Persuasion (transform)	0.010 (0.009)	0.001 (0.011)
Persuasion (transform) $\times$ Action	−0.001 (0.012)	0.021 (0.016)
Action	0.026** (0.013)	0.026** (0.013)
Persuasion transform	$\text{sign}(\delta)$	Winsorized δ
Fact FE	Yes	Yes
Num. obs.	7,740	7,740

Sample: Experiment 2 general-population senders; 7,740 observations. Both columns replace *Persuasion* ($\Delta\pi_{ijk}$, positional reading) in the preferred action-decomposition specification with a transformed version, keeping fact fixed effects and sender-clustered standard errors. Column (1): $\text{sign}(\delta) \in \{-1, 0, +1\}$ (direction only). Column (2): δ winsorized at the 90th percentile of $|\delta|$ (magnitude, extremes capped; N unchanged). A significant direction main effect with a null direction $\times$ Action interaction, alongside a surviving winsorized magnitude $\times$ Action interaction, indicates w is strictly increasing rather than a step function. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.33: Sample-Exclusion Robustness for the Instrumental Motive — Experiment 2 (General Population)

	(1) Drop low- $ \delta $	(2) Drop perfect sharers	(3) Drop low- α facts
Expression	0.153*** (0.023)	0.198*** (0.022)	0.151*** (0.022)
Factual belief correction gain	0.024 (0.019)	0.042** (0.018)	0.016 (0.021)
Persuasion	-0.004 (0.009)	-0.002 (0.010)	-0.002 (0.011)
Persuasion $\times$ Action	0.024* (0.013)	0.021 (0.014)	0.025 (0.015)
Action	0.031** (0.014)	0.017 (0.011)	0.032** (0.015)
Fact FE	Yes	Yes	Yes
Num. obs.	5,805	6,355	5,171

Sample: Experiment 2 general-population senders; preferred action-decomposition specification, fact fixed effects, sender-clustered standard errors. Column (1) drops senders in the bottom quartile of mean $|\delta|$ (expect near-zero persuasion across facts). Column (2) drops perfect sharers (provision rate = 1.0, all five facts shared). Column (3) drops the bottom tertile of facts by mean expressive magnitude $|\alpha|$. The instrumental motive (*Persuasion* $\times$ *Action*) is positive in every restriction. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

This rate is flat across receiver stances (23.4%–26.9%), receiver–fact alignment (counter-attitudinal vs. same-side), and sender–receiver distance (all regression coefficients $p > 0.35$). The mean expected factual belief correction is likewise unrelated to receiver position and flat in sender–receiver distance (pooled slope $\beta = -0.002$, $p = 0.42$, fact fixed effects, sender-clustered standard errors; Figure C.6, Panel (A)).

In 59% of fact–receiver pairs, senders expect disclosure to bring the receiver’s stance weakly closer to their own relative to non-disclosure ($\widehat{\Delta\pi}_{ijk} \geq 0$; 41.7% strictly closer, 17.3% unchanged). The complementary 41% counter-directional rate falls as the receiver sits further from the sender (-0.028 per Likert unit, $p < 0.001$, fact FE), so senders expect fewer backfires from distant receivers; across the five receiver stances it runs from 35.0% to 44.4%. The average persuasion return at distance zero is essentially zero (-0.067). From distance 2 onward, senders expect disclosure to close a roughly constant share of the sender–receiver gap, about six percent (5.6%, 5.8%, and 5.9% at distances 2, 3, and 4; equality Wald $p = 0.98$; Figure C.6, Panel (B)). The aligned profile is the steeper of the two: the distance $\times$ alignment interaction is $+0.073$ ($p < 0.001$, fact FE), and $+0.047$ ($p = 0.01$) with sender fixed effects added. Part of the raw distance gradient reflects the bounded five-point stance scale, since predictions for receivers at the scale ends can move only inward. Among receivers at interior stance positions, the gradient persists ($+0.081$, $+0.215$, and $+0.214$ at distances 1, 2, and 3 relative to distance zero, sender and fact fixed effects, all $p < 0.01$).

Table C.34: Sample characteristics and quota targets, Experiment 2

	Population (1)	Sample (2)	CONTROL (3)	ACTION (4)
<i>Gender</i>				
Female	0.499	0.481	0.501	0.462
Male	0.501	0.517	0.497	0.536
Diverse	—	0.003	0.003	0.003
<i>Age band</i>				
18–29	0.179	0.197	0.199	0.195
30–39	0.178	0.176	0.177	0.175
40–49	0.170	0.160	0.150	0.169
50–59	0.193	0.200	0.188	0.212
60–75	0.279	0.267	0.286	0.248
<i>Bundesland</i>				
Baden-Württemberg	0.134	0.129	0.135	0.122
Bayern	0.159	0.156	0.144	0.166
Berlin	0.045	0.042	0.040	0.044
Brandenburg	0.030	0.032	0.036	0.029
Bremen	0.008	0.010	0.011	0.010
Hamburg	0.023	0.023	0.021	0.024
Hessen	0.075	0.072	0.068	0.076
Mecklenburg-Vorpommern	0.019	0.023	0.023	0.024
Niedersachsen	0.095	0.099	0.102	0.096
Nordrhein-Westfalen	0.216	0.215	0.204	0.226
Rheinland-Pfalz	0.050	0.050	0.046	0.054
Saarland	0.012	0.010	0.016	0.005
Sachsen	0.047	0.048	0.050	0.047
Sachsen-Anhalt	0.025	0.032	0.029	0.034
Schleswig-Holstein	0.035	0.031	0.041	0.021
Thüringen	0.025	0.028	0.034	0.021
<i>School-leaving qualification</i>				
Hauptschulabschluss or none	0.297	0.265	0.277	0.253
Mittlerer Abschluss (Realschule)	0.314	0.293	0.270	0.315
Fachhochschul-/Hochschulreife or higher	0.390	0.442	0.453	0.431

Sample: $N = 1,548$ German adults recruited on Bilendi (quota-sampled), Experiment 2. Column (1): German population proportions. Gender, age band, and Bundesland: population aged 18–75 as of 31 December 2024 (Bevölkerungsfortschreibung on the Zensus-2022 basis; Statistisches Bundesamt, Statistischer Bericht 12411, tables 12411-06 and 12411-09). No registry benchmark exists for “Diverse.” School-leaving qualification: Mikrozensus 2023 (Statistisches Bundesamt, Bildungsstand der Bevölkerung), persons aged 15 and over excluding those still in school—a broader base than the sample’s 18–75. Column (2): full-sample proportions. Columns (3)–(4): proportions by experimental arm.

C.5 Directional vs. Positional Persuasion Proxy

The canonical persuasion proxy $\Delta\pi_{ijk}$ measures reduction in Likert distance between the receiver’s predicted post-disclosure stance and the sender’s own position—a positional measure. A direction-only analog, $\delta_{ijk}^{\text{dir}}$ (defined in Section 4.5), strips preference magnitude from the signal: it is positive whenever the receiver moves toward the sender’s preferred side, regardless of how far the receiver is from the sender’s exact position or how strong the sender’s preference is. Because $\delta_{ijk}^{\text{dir}}$ is undefined for neutral senders, 260 observations from 52 neutral senders (3.4 % of observations) are excluded

	Canonical	+ Speeders	+ Early starts	+ Both	Attention passers
Alignment (H1)	0.128*** (0.013)	0.131*** (0.013)	0.123*** (0.013)	0.125*** (0.012)	0.135*** (0.015)
Alignment $\times$ ACTION (H2)	0.014 (0.018)	0.012 (0.017)	0.015 (0.017)	0.013 (0.017)	0.007 (0.020)
Expression (α_{ij})	0.148*** (0.020)	0.147*** (0.019)	0.142*** (0.019)	0.141*** (0.019)	0.167*** (0.023)
Factual belief correction gain (Δb_{ijk})	0.034* (0.017)	0.031* (0.017)	0.034** (0.017)	0.031* (0.017)	0.053*** (0.020)
Persuasion ($\Delta\pi_{ijk}$)	0.008 (0.007)	0.009 (0.006)	0.007 (0.006)	0.008 (0.006)	0.010 (0.008)
Persuasion $\times$ ACTION	0.025* (0.013)	0.022* (0.013)	0.020 (0.012)	0.017 (0.012)	0.026* (0.015)
Senders	1548	1588	1654	1694	1156
Observations (confirmatory)	7715	7910	8245	8440	5760
Observations (motive specifications)	7740	7940	8270	8470	5780

Dependent variable: binary provision decision ($y_{ij} \in \{0, 1\}$), Experiment 2 general-population senders, under alternative sample definitions. Canonical: the main analysis sample (no speeders, no pre-launch starts). + Speeders: adds completers faster than $0.25 \times$ the median completion time. + Early starts: adds participants who started before the launch of the main collection on 27 May 2026, 10:00 CET. + Both: adds both groups. Attention passers: restricts the canonical sample to senders answering the attention-check item correctly (the attention check was administered in Experiment 2 only). Rows (1)–(2) are from the pre-registered confirmatory specification: provision on positional alignment, ACTION, their interaction, demographic controls (age, gender, education, Bundesland, voting intention, left–right), and fact fixed effects. Rows (3)–(5) are from the three-channel motive specification (Expression, factual belief correction gain, Persuasion, fact fixed effects); row (6) adds the Persuasion $\times$ ACTION interaction to that specification. The two observation rows differ because the confirmatory specification drops observations with missing demographic controls (25 observations from five senders in the canonical column). Standard errors clustered at the sender level (in parentheses). *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.35: Sample-inclusion robustness of the Experiment 2 estimates.

Table C.36: Leave-One-Fact-Out Robustness of the Motive Decomposition — Experiment 2 (General Population)

Motive	Sign-positive	Min β	Max β	Max p
Expression	12/12	0.143	0.156	0.000
Factual belief correction gain	12/12	0.022	0.041	0.232
Persuasion	12/12	0.006	0.011	0.374
Fact FE	Yes			
Clustered SE	Sender			

Sample: Experiment 2 general-population senders. Each motive is re-estimated twelve times in the three-motive specification (Expression, Factual belief correction gain, Persuasion), each time dropping one of the twelve facts; fact fixed effects, sender-clustered standard errors. Columns report the number of leave-one-out estimates with a positive coefficient, the minimum and maximum coefficient across the twelve runs, and the largest p -value observed.

from the directional columns.

Table C.37 re-estimates the specification replacing $\Delta\pi_{ijk}$ with $\delta_{ijk}^{\text{dir}}$, pooled (columns 1–3) and with the ACTION interaction (columns 4–6). Standalone, the directional proxy is marginally significant while the positional one is not (positional $\beta = +0.008$, $p = 0.20$, all senders; directional $\beta = +0.011$, $p = 0.06$, non-neutral senders; columns 1–2). When they enter jointly, the directional proxy absorbs the shared variance ($\beta_{\text{dir}} = +0.016$, $p = 0.054$; $\beta_{\text{pos}} = -0.008$, not statistically distinguishable from zero; column 3). The two measures correlate at $r = 0.74$ after fact fixed effects, so the horse




Figure C.6: Elicited returns to disclosure by sender-receiver distance, Experiment 2 (General Population). One marker per fact $\times$ sender-receiver distance $\times$ alignment (cell mean over ≥ 20 elicitations), positionally aligned facts ($\text{Align}_{ij} > 0$) as filled black circles, unaligned ($\text{Align}_{ij} < 0$) as open gray squares, horizontally offset for legibility; facts of senders with $\text{Pref}_i = 0$ are omitted from the split. Solid (dashed) lines connect the observation-weighted means of aligned (unaligned) elicitations at each distance. Panel (a): expected factual belief correction $\widehat{\Delta b}_{ijk}$ —flat in distance (pooled slope $\beta = -0.003$, $p = 0.38$, fact fixed effects, sender-clustered standard errors). Panel (b): expected persuasion $\widehat{\Delta \pi}_{ijk}$ —rising in distance ($\beta = +0.072$, $p < 0.001$); the aligned and unaligned mean profiles are statistically indistinguishable (distance $\times$ alignment interaction $p = 0.12$). Dashed lines mark zero. Elicitations incentivized via Quadratic Scoring Rule.

race partitions shared variance; it is descriptive and does not adjudicate between the proxies. A common-sample check rules out sample composition as the driver of the joint-specification pattern: positional estimates on the non-neutral subsample remain close to their full-sample values (pooled $\beta = +0.005$, $p = 0.43$, against $+0.008$, $p = 0.20$ on the full sample; $\times \text{ACTION}$ $\beta = +0.022$, $p = 0.083$, against $+0.025$, $p = 0.052$). The θ_I interaction is directionally consistent under both codings ($\beta = +0.025$, $p = 0.052$ positional; $\beta = +0.019$, $p = 0.092$ directional; columns 4–5).

Table C.38 isolates the functional form: it repeats the ACTION-interaction specification under binary codings of both proxies, replacing each continuous measure with an indicator for a strictly positive value. The ACTION amplification—the instrumental motive θ_I —requires *magnitude*: replacing either continuous proxy with its binary version eliminates the interaction entirely (positional binary $\times \text{ACTION}$ $\beta = +0.004$, $p = 0.86$; directional binary $\times \text{ACTION}$ $\beta = +0.013$, $p = 0.59$), while both continuous codings retain it (positional $p = 0.052$; directional $p = 0.092$). Absent political action, the directional binary is positive but not distinguishable from zero ($\beta = +0.024$, SE

	Pooled			With Action Interaction		
	Positional	Directional	Horse race	Positional	Directional	Horse race
Expression α_{ij}	0.148*** (0.020)	0.147*** (0.020)	0.148*** (0.020)	0.149*** (0.020)	0.148*** (0.020)	0.148*** (0.020)
Factual belief correction Δb_{ijk}	0.034* (0.017)	0.032* (0.018)	0.033* (0.018)	0.034** (0.017)	0.032* (0.017)	0.033* (0.017)
Persuasion-positional $\Delta \pi_{ijk}^{\text{pos}}$	0.008 (0.007)		-0.008 (0.010)	-0.004 (0.009)		-0.015 (0.013)
Persuasion-directional $\Delta \pi_{ijk}^{\text{dir}}$		0.011* (0.006)	0.016* (0.008)		0.001 (0.008)	0.011 (0.012)
ACTION				0.025** (0.013)	0.025* (0.013)	0.025* (0.013)
Persuasion-pos. $\times$ ACTION				0.025* (0.013)		0.014 (0.019)
Persuasion-dir. $\times$ ACTION					0.019* (0.011)	0.010 (0.017)
Clustered SE	Sender	Sender	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes	Yes	Yes
Action level + interaction	No	No	No	Yes	Yes	Yes
Sample	All senders	Non-neutral	Non-neutral	All senders	Non-neutral	Non-neutral
Num. obs.	7740	7480	7480	7740	7480	7480
R ²	0.014	0.014	0.015	0.015	0.016	0.016
Adj. R ²	0.012	0.013	0.013	0.013	0.013	0.013

Dependent variable: binary provision decision ($y_{ij} \in \{0, 1\}$). Sample: Experiment 2, $N = 1,548$ senders. Directional columns exclude 260 observations from senders with $\text{pref}_i = 0$ (neutral donors; $N = 7,480$); positional columns retain the full sample ($N = 7,740$). Persuasion-positional ($\Delta \pi_{ij}^{\text{pos}}$) is `delta_gain_2`: reduction in Likert distance to the sender's own position. Persuasion-directional ($\Delta \pi_{ij}^{\text{dir}}$) is $\delta_{ij}^{\text{dir}} \equiv -\Delta \text{distance}_{ij} \cdot \text{sign}(\text{pref}_i)$: whether the receiver moves toward the sender's preferred side, preference-magnitude-free. The two proxies correlate at $r = 0.74$ after fact fixed effects; horse-race columns partition shared variance and are descriptive, not an adjudication between the proxies. Standard errors clustered at the sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.37: Positional vs. directional persuasion proxies, Experiment 2.

0.017), so these data do not establish that baseline senders respond to direction alone.

Table C.39 reports the design-geometry specifications summarized in Sections 4.3 and 5.2: the elicited persuasion return regressed on directional opportunity—the fact's pre-validated direction interacted with the signed sender–receiver stance gap—with sender, fact, and distance fixed effects, for the general population and for journalists, with and without receiver-position fixed effects. No elicited quantity appears on the right-hand side.

	Positional		Directional	
	Continuous	Binary	Continuous	Binary
Expression α_{ij}	0.149*** (0.020)	0.149*** (0.020)	0.148*** (0.020)	0.146*** (0.020)
Factual belief correction Δb_{ijk}	0.034** (0.017)	0.034* (0.017)	0.032* (0.017)	0.033* (0.017)
Persuasion-pos. (continuous)	-0.004 (0.009)			
Persuasion-pos. (binary)		0.011 (0.016)		
Persuasion-dir. (continuous)			0.001 (0.008)	
Persuasion-dir. (binary)				0.024 (0.017)
ACTION	0.025** (0.013)	0.025 (0.016)	0.025* (0.013)	0.021 (0.017)
Persuasion-pos. (cont.) $\times$ ACTION	0.025* (0.013)			
Persuasion-pos. (bin.) $\times$ ACTION		0.004 (0.023)		
Persuasion-dir. (cont.) $\times$ ACTION			0.019* (0.011)	
Persuasion-dir. (bin.) $\times$ ACTION				0.013 (0.023)
Clustered SE	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes
Sample	All senders	All senders	Non-neutral	Non-neutral
Num. obs.	7740	7740	7480	7480
R ²	0.015	0.015	0.016	0.016
Adj. R ²	0.013	0.013	0.013	0.014

Dependent variable: binary provision decision ($y_{ij} \in \{0, 1\}$). Sample: $N = 1,548$ Experiment 2 senders (7,740 observations); directional columns exclude 260 observations from senders with $\text{pref}_i = 0$ (neutral donors; $N = 7,480$). Each column re-estimates the ACTION-interaction specification with one coding of the persuasion proxy. Continuous codings: positional $\Delta\pi_{ijk}^{\text{pos}}$ (`delta_gain_2`) and directional $\Delta\pi_{ijk}^{\text{dir}}$. Binary codings: an indicator equal to one when the corresponding continuous measure is strictly positive. All specifications include Expression, Factual belief correction, fact fixed effects, and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.38: Persuasion proxy functional form: continuous vs. binary codings, Experiment 2.

Table C.39: Design Geometry and the Persuasion Return

	General population		Journalists	
	(1)	(2)	(3)	(4)
Directional opportunity (signed gap $\times$ direction)	0.039*** (0.006)	0.039*** (0.006)	0.114*** (0.012)	0.112*** (0.012)
Distance FE	Yes	Yes	Yes	Yes
Receiver-position FE	No	Yes	No	Yes
Sender FE	Yes	Yes	Yes	Yes
Fact FE	Yes	Yes	Yes	Yes
Clustered SE	Sender	Sender	Sender	Sender
Num. obs.	7740	7740	1535	1535
R ²	0.323	0.323	0.368	0.372
Adj. R ²	0.152	0.152	0.200	0.203

Dependent variable: the elicited persuasion return $\Delta\pi_{ijk}$. Directional opportunity is the fact's pre-validated direction interacted with the signed sender-receiver stance gap, $-\text{dir}_j \times (\text{Stance}_i - \text{Stance}_k)$; it is determined entirely by the design and uses no elicited quantity. Distance fixed effects absorb $|\text{Stance}_i - \text{Stance}_k|$. Standard errors clustered at the sender.

Companion statistics not shown as rows. Senders expect movement in the fact's pre-validated direction in 58% of cases in the general population and 77% among journalists (receivers away from the scale endpoints). Expected movement rises with the room the receiver has to move in the fact's direction in the general population (+0.056 per scale point, $p < 0.01$) but not detectably among journalists (+0.030, $p = 0.43$, which does not rule out the general-population value). A joint test of the receiver-position fixed effects in columns (2) and (4) does not reject for the general population ($p = 0.29$) but does for journalists ($p = 0.013$): journalists respond to where the receiver stands, not only to the relative gap. The journalist sample is smaller (307 senders against 1,548), so differences that fail to reach significance there are bounds rather than nulls.

Table C.40: Expression Measure (α_{ij}) by Sender Stance and Individual Fact — Experiment 2

Panel A: Immigration facts

Sender stance	Anti-policy direction (–)			Pro-policy direction (+)		
	Violent crime	Bürgergeld	Res. permits	Shortage fields	Employment	Integration
1 (most pro)	+0.00	+0.29	+0.00	+0.43	+0.00	+0.67
2	–0.33	–0.16	–0.43	+0.44	+0.37	+0.77
3 (neutral)	+0.29	–0.07	+0.00	+0.64	+0.33	+0.59
4	+0.66	+0.83	+0.54	+0.00	+0.28	–0.16
5 (most anti)	+0.82	+0.82	+0.79	+0.38	+0.34	+0.24

Panel B: Climate facts

Sender stance	Anti-policy direction (–)			Pro-policy direction (+)		
	Grid fees	Business surv.	Price fluct.	Renewable share	Solar cost	Renew. jobs
1 (most pro)	+0.23	+0.45	+0.38	+0.49	+0.71	+0.55
2	–0.15	+0.23	+0.26	+0.46	+0.59	+0.46
3 (neutral)	+0.14	+0.10	+0.30	+0.31	+0.50	+0.30
4	+0.64	+0.58	+0.61	+0.00	–0.40	–0.33
5 (most anti)	+0.78	+0.79	+0.54	+0.26	+0.20	–0.10

Notes: Expression (α_{ij}) = $2(\bar{g}_{ij} - 0.5)$, where $g = 1$ if the sender rates the cross-domain fact as “supporting my opinion.” Per-cell sizes are approximately 30–100 observations (individual fact $\times$ stance); values in small cells (fewer than 40 observations) should be interpreted with caution. Fact identifiers in brackets. Sender stance = stated Likert position on the provision topic (1 = most pro-policy; 5 = most anti-policy). Cross-domain ratings are elicited on the *other* topic; all senders in the cell have the same stated stance.

D Experiment 3 Supplementary Tables and Figures

D.1 Sample Composition

Table D.41 compares the demographic and occupational characteristics of the Experiment 3 journalist sample to Loosen et al. (2023), the most recent representative survey of the German journalist population. The post-questionnaire collected gender, age, education, medium type, conditional broadcaster affiliation, geographic reach, role (multiple selection), and editorial autonomy. Fields not elicited in Experiment 3—years of experience, full-time versus freelance status, and topics covered—cannot be compared to the benchmark.

Figure D.7 reports the stated and revealed policy-preference distributions of the journalist sample, by topic, on the axes and bucket definitions of Figure C.4 for the Experiment 2 general-population sample.

D.2 Heterogeneity

Table D.42 reports the pre-registered exploratory heterogeneity of the alignment gap among journalists—by sender–receiver distance, polarization, issue extremity, and the journalist background variables (medium type, broadcaster type, reach, role, editorial

Table D.41: Experiment 3 Sample Characteristics: Comparison to Population Benchmark

	Experiment 3 (<i>N</i> = 307)	Loosen et al. (2023) (<i>N</i> = 1,221)
<i>Gender</i>		
Female	53.4	44.0
Male	44.7	55.9
<i>Age</i>		
Under 50	77.1	62.4
<i>Education</i>		
University degree (B.A.+)	84.8	67.1
<i>Medium type</i>		
Print	28.6	57.0
Online	45.5	3.1
Television	7.1	17.3
Radio/podcast	10.7	17.2
News agency	3.6	3.5
<i>Broadcaster type</i> (TV/radio respondents only)		
Public broadcaster	80.0	61.4
<i>Role</i> (multiple selection; shares exceed 100 %)		
Editor	57.5	—
<i>Editorial autonomy</i>		
High ($\geq 4/5$)	76.0	62.0
<i>Political affiliation</i> (no population benchmark)		
Grüne or Linke (closest party)	58.5	—
Mean left-right self-placement (0–10)	3.2	—

Sample: *N* = 307 professional German journalists recruited via DJV and DJU mailing lists, Experiment 3 (convenience sample, not a probability sample). *Benchmark:* Loosen et al. (2023), stratified random sample of active German journalists (*N* = 1,221, September 2022–February 2023, HBI Hamburg; Loosen et al.). All figures are percentages. *Medium type:* Experiment 3 uses a single online category for all digital-native outlets; the Loosen et al. benchmark reports native-online journalists (3.1%) as a distinct population segment; print aggregates newspapers and magazines (57.0%). *Broadcaster type* is conditional on TV or radio respondents only. *Role* allows multiple selection; Loosen et al. report no breakdown by editorial position, so this row appears without a benchmark. *Editorial autonomy* is not a common item: Experiment 3 asks how often the respondent can decide which topics to work on (1 = never to 5 = very often), while Loosen et al. report the share describing themselves as very or completely free in topic selection. *Age:* Loosen et al. report a population mean of 45.3 years; Experiment 3 elicited age as a categorical variable. *Political affiliation* and left-right self-placement were not elicited in Loosen et al. and appear without a benchmark.

autonomy)—together with the corresponding interactions for the three motives.

D.3 Self-Reported Decision Motives

Table D.43 relates each journalist’s alignment bias to the motives their free-text account of their decisions was classified as citing, in the specification of Table B.18 for Experiment 1. Because Experiment 3 assigns every journalist to the ACTION arm, the treat-




Figure D.7: Distribution of policy preferences among Experiment 3 journalists ($N = 308$). *Top row (Panels A–B):* stated preferences from the 5-point Likert item; tick labels complete the stance statement shown in the panel title (e.g., *Immigration to Germany should be ... greatly eased*). *Bottom row (Panels C–D):* revealed preferences from the real donation task, grouped into five buckets from All pro (€250 to the pro-policy organization) to All anti (€0), displayed pro-to-anti left to right. *Left column:* immigration senders ($N = 163$; $N = 162$ for the stated item, one non-response). *Right column:* climate senders ($N = 145$). Axes and bucket definitions are identical to Figure C.4, which reports the same distributions for the Experiment 2 general-population sample.

ment interactions of that specification are not estimable here; columns (2) to (4) report robustness to preference intensity, topic fixed effects, and restriction to journalists who gave an account referring to the sharing decision. The two tables share a specification but not a coding version: Experiment 3 is coded under the earlier version of the scheme and Experiment 1 under the later one (Appendix E.1), so the signs are comparable and the magnitudes are not.

Moderator	Interaction of moderator with:				Obs.
	Align _{ij}	α_{ij}	Δb_{ijk}	$\Delta \pi_{ijk}$	
Distance Stance _i – Stance _k	0.026*** (0.009)	0.030 (0.031)	0.010 (0.029)	0.006 (0.012)	1535
Polarization	0.014 (0.010)	0.028 (0.026)	–0.032 (0.028)	0.006 (0.011)	1530
Issue extremity	–0.039 (0.030)	0.037 (0.066)	0.042 (0.057)	0.012 (0.023)	1535
Online medium	–0.036 (0.029)	–0.035 (0.076)	–0.001 (0.081)	–0.004 (0.028)	1530
Public broadcaster	–0.140* (0.074)	–0.318* (0.187)	0.148 (0.225)	–0.048 (0.067)	275
National/international reach	–0.056* (0.030)	–0.102 (0.078)	0.051 (0.081)	0.029 (0.030)	1535
Editor role	–0.021 (0.030)	0.030 (0.079)	–0.068 (0.079)	0.013 (0.028)	1535
Editorial autonomy	–0.016 (0.017)	–0.050 (0.044)	0.101** (0.041)	0.019 (0.011)	1530

Dependent variable: binary provision decision ($y_{ij} \in \{0, 1\}$), Experiment 3 journalists. Each row reports interaction coefficients between one moderator and the selection measures: column (1) is from a regression of provision on positional alignment, the moderator, and their interaction; columns (2)–(4) are from a single regression interacting the moderator with Expression (α_{ij}), the factual belief correction gain (Δb_{ijk}), and Persuasion ($\Delta \pi_{ijk}$) jointly. All specifications include fact fixed effects and journalist-clustered standard errors (in parentheses); moderators are centered at the estimation-sample mean. Moderators: sender–receiver distance, polarization (|left-right – 5|), issue extremity (|stance – 3|), and the pre-registered journalist background variables, coarsened as online medium (vs. all other medium types), public broadcaster (defined only for television and radio senders), national or international reach (vs. local or regional), any editor role in the multi-select role item, and editorial autonomy (Likert 1–5). Benjamini–Hochberg correction is applied within two families: the pre-registered alignment battery (column 1; 8 tests, minimum $q = 0.11$) and the channel battery (columns 2–4; 24 tests), in which only Expression $\times$ Online medium survives at the 10 % level ($q = 0.074$). Stars mark unadjusted p -values: *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table D.42: Heterogeneity of alignment- and channel-based selection among journalists, Experiment 3.

	(1)	(2)	(3)	(4)
<i>Belief correction</i>	−0.112*** (0.041)	−0.136*** (0.042)	−0.116*** (0.042)	−0.125** (0.053)
<i>Expression</i>	0.226 (0.139)	0.201 (0.139)	0.219 (0.139)	0.213 (0.144)
<i>Persuasion</i>	0.240*** (0.063)	0.213*** (0.063)	0.240*** (0.063)	0.228*** (0.067)
<i>Misc.</i>	−0.183 (0.163)	−0.107 (0.178)	−0.197 (0.165)	−0.196 (0.167)
$ \text{Pref}_i $		0.233*** (0.080)		
Constant	0.183*** (0.038)	0.033 (0.064)		0.196*** (0.051)
Topic fixed effects	No	No	Yes	No
Sample	All	All	All	Referrers
Observations	296	296	296	220
Adjusted R^2	0.106	0.128	0.104	0.163

Dependent variable: the journalist-level alignment bias, the within-journalist difference between the provision share for aligned facts and the provision share for unaligned facts, defined on the sample of journalists holding a directional policy preference and at least one fact of each type. Regressors are indicators for the motives a journalist’s free-text account of their decisions was classified as citing; a journalist may cite more than one, and the omitted category is giving no account that refers to the sharing decision. Classification is by large language model under the coding scheme of Appendix E.1. Column (1) is the specification of Table B.18 for Experiment 1; column (2) adds preference intensity $|\text{Pref}_i|$; column (3) adds topic fixed effects; column (4) restricts to journalists who gave an account referring to the sharing decision. Experiment 3 assigns every journalist to the ACTION arm, so the treatment interactions of the Experiment 1 specification are not estimable here. Heteroskedasticity-robust standard errors in parentheses. The expressive category is cited by eleven journalists. These items are recorded after the provision task and are not incentivized; the analysis is exploratory. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table D.43: Self-reported decision motives and the alignment bias among journalists, Experiment 3.

Table D.44: Motive Decomposition: General Population versus Journalists — Experiments 2 and 3

	Probability of Fact Provision		
	(1)	(2)	(3)
Expression	0.148*** (0.020)	0.158*** (0.027)	0.134*** (0.045)
Factual belief correction gain	0.034* (0.017)	0.018 (0.024)	0.076* (0.039)
Persuasion	0.008 (0.007)	0.021** (0.009)	0.020 (0.014)
Clustered SE	Sender	Sender	Journalist
Fact FE	Yes	Yes	Yes
Arm	Both	Action only	Action only
Population	General population	General population	Journalists
Num. obs.	7740	3965	1535
R ²	0.014	0.016	0.061
Adj. R ²	0.012	0.013	0.052

Sample: Experiment 2 general-population senders (columns 1–2; $N = 1,548$, Bilendi panel, Germany, 2026) and Experiment 3 journalists (column 3; $N = 307$ in the regression sample, recruited via DJV/DJU mailing lists). *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating α_{ij} ; for journalists, merged from the general-population cross-domain rating cell lookup under the stance-equivalence assumption. *Factual belief correction gain:* incentivized expected correction of the receiver’s factual belief Δb_{ijk} . *Persuasion:* incentivized expected preference shift $\Delta \pi_{ijk}$ toward the sender’s position (positional reading, `delta_gain_2`). All columns include fact fixed effects; standard errors clustered at the sender (journalist) level. All journalists are in the ACTION arm. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

E LLM Classification Methodology

E.1 Classification Procedure

I classify the open-ended decision-reason responses from all three experiments—Experiment 1 (1,200 U.S. general-population senders, in English), Experiment 2 (1,548 German general-population senders), and Experiment 3 (306 German journalists)—into theory-derived motive categories using large language models. Each distinct response is classified in its own API call returning one structured JSON object, with the model pinned and no temperature or sampling settings changed; a response that is an exact duplicate of another takes that response’s classification rather than a second call. The coding is a reclassification under a new category scheme, not a relabeling of an earlier one: the categories, the gating step, and the boundary rules below all differ from the scheme used before.

Two versions of the scheme are in play, and any comparison between Experiment 1 and Experiment 3 is therefore a comparison across versions. Experiments 2 and 3 are classified once, by Anthropic’s Claude Sonnet 5 (`claude-sonnet-5`), under the earlier version. Experiment 1 is classified under a later and stricter version, and twice: Claude Sonnet 5 and OpenAI’s GPT-5.6 (`gpt-5.6-terra`) each code all 1,200 responses under one prompt frozen before either run, neither run returned a classification failure, and I report Claude Sonnet 5 as the primary classification throughout with GPT-5.6 as robustness. What the versions share is the gating logic and the three categories that carry the theory; what they differ in is the boundary rules and the auxiliary categories (Table E.45).

The scheme is applied in two stages. A first-step gate asks whether the response refers to the sharing decision, the receiver, or the information provision—explicitly or elliptically. Responses that describe only the sender’s own epistemic process, that state a non-directional opinion with no link to the decision, or that describe only a different task in the experiment fail the gate and receive no further coding. A response may also refer to the decision without giving any reason for it. Responses that pass are coded into the categories in Table E.45. Categories are not mutually exclusive; a label is assigned only where the text gives distinct support for it, and no motive is inferred beyond the text.

The two versions record a passing response differently. The earlier version assigns each response exactly one primary motive, the most central justification, alongside the full set of categories it mentions, and records a think/do subtype for persuasion that the analysis does not use. The Experiment 1 version assigns no primary motive. It instead records each category in one of three states, *absent*, *affirmed*, or *mentioned but rejected*, and requires the model to return the span of text supporting any state other

than absent. All regressions use the affirmed state. The third state carries information the earlier version discarded, and it is not rare here: a sender who names their own political position and then says they disclosed regardless of it is coded as mentioning and rejecting expression rather than as absent, which the primary classifier records for 23 responses against 12 affirmations.

Three categories carry the theory in both versions. Expression marks the fact's fit with the sender's own position as the stated decision factor, with no receiver-directed goal. Persuasion marks a receiver-directed goal—moving what the receiver thinks or does. Belief correction marks truth-telling or the receiver's right to accurate information and is direction-blind. These three are written to correspond to the channels in the main text: expression to the expressive channel, persuasion to the persuasion channel, and belief correction to factual belief correction. The correspondence is a choice of category definitions and nothing more; these self-reports carry none of the standing of the elicited returns, and they are not used as proxies for them. The Experiment 1 version draws two of these boundaries more tightly than the earlier one, and affirms expression for twelve of its 1,200 responses. Expression is affirmed only where the disclosure act and a directional political selection criterion are both explicit, so a bare statement of political position or identity fails the gate rather than counting as expression, and truth, accuracy, or universal disclosure stated as overriding a political preference makes expression mentioned but rejected. Persuasion, in turn, admits an elliptical receiver-impact consideration offered as the answer, without requiring the sender to restate that it governed the decision.

The complete classification prompt, including all decision rules, boundary conditions, and worked examples, is available upon request. The Experiment 1 prompt and the response manifest were hashed before either classification run, and both runs record the same hashes, so the two classifications differ only in the model.

Figure E.8 reports the resulting distribution of coded motives for Experiment 1 under each of the two classifiers, and Figure E.9 the distribution for Experiments 2 and 3 under the earlier version of the scheme. Because the two versions differ in their category sets, the two figures are read side by side but not compared bar for bar.

E.2 Validation

Experiment 1's coding is checked in two ways: against the second classifier, and against my own reading of the cases the successive versions of the scheme had disagreed about, which I resolved by hand before either run. The two models agree on 91 percent of gate decisions ($\kappa = 0.81$), on 93 percent of the no-stated-reason decisions ($\kappa = 0.86$), on 94 percent of belief-correction codings ($\kappa = 0.88$), and on 98 percent of persuasion codings ($\kappa = 0.73$). On the resolved hard cases, the primary classifier matches my

Table E.45: Coding Scheme and Decision Rules

Category	Definition	Key boundary rule	Applies to
Belief correction	Truth-telling, honesty, or the receiver's right to accurate information; direction-blind	Does not depend on the fact's political direction	All
Expression	The fact's fit with the sender's own position is the stated decision factor, with no receiver goal	No receiver-directed goal. Exp. 1: the disclosure act and a directional selection criterion must both be explicit, so a bare statement of position fails the gate	All
Persuasion	A receiver-directed goal: moving what the receiver thinks or does	Any reference to the receiver changing is coded here. Exp. 1: an elliptical receiver-impact consideration offered as the answer qualifies	All
Balance	Deliberately sharing counter-attitudinal or both-sides content	The receiver is meant to see the other side	All
Role neutrality (journalists only)	Professional role integrity: it is not the journalist's place to filter what the reader sees	Distinct from <i>balance</i> ; concerns the sender's role, not both-sidedness	Exps. 1, 3
Receiver tailoring	Decision conditioned on the receiver's beliefs or characteristics, without reference to attitude change	Any reference to the receiver changing → <i>persuasion</i>	Exps. 2–3
Disbelief	Credibility doubts about the fact or source as the reason for withholding	Politically motivated skepticism → <i>expression</i> or <i>persuasion</i> . In Exp. 1 a flag, not a category	Exps. 2–3
Relevance	The fact is judged irrelevant or unimportant for the receiver, without a directional stake	Misleadingness with a directional stake → <i>persuasion</i>	Exps. 2–3
Futility	"Facts do not change minds" offered as the stated reason	A general claim about receivers, not about one fact	Exps. 2–3
Other stated reason (Exps. 2–3: miscellaneous)	Refers to the decision and gives a reason outside the listed categories, or none that is interpretable	Residual category; enters the regression only as a control	All
Gate outcomes	Describes only the sender's own epistemic process or a different task.	Receives no category coding	All




Figure E.8: Distribution of Experiment 1 senders' stated decision motives, among the senders whose open-ended decision reason refers to the disclosure decision (39.1 percent under the primary classifier, 36.3 percent under the secondary). Bars give the share of those senders whose response was coded as citing each motive; categories are not mutually exclusive, so the bars do not sum to one. Both classifiers apply the same frozen prompt to the same 1,200 responses (Appendix E.1).

resolution for 66 of 78 and the secondary for 58 of 78.

Expression is the weak point, and the relevant number is not the one κ reports. Raw agreement is 99 percent and κ is 0.54, which at a base rate of twelve and ten affirmations out of 1,200 says little either way. The informative figure is the overlap: the two models agree on six of those responses. The set of senders behind the expression coefficient is therefore itself model-dependent, which is the reason to read that row of Table B.18 as a description of a handful of senders.

Two limits on what this establishes. Agreement between two models applying the same prompt is cross-model reliability, not intercoder reliability in the usual sense, and it cannot show that either coding is correct; blinded coding by research assistants will supply an independent human benchmark and is not yet available. And the two classifications do not deliver identical estimates. Every coefficient in Table B.18 keeps its sign under the second classifier and the persuasion-by-ACTION result holds at the five percent level under both, but the two persuasion interactions fall from the one to the five percent level, and the expression interaction moves from insignificant to the ten percent level with its negative sign unchanged (Appendix Table B.20). The arm balance of the stated-motive distribution holds under both.

Experiments 2 and 3 are classified once, so no agreement statistic is available for them. The only indicator is the model's self-reported confidence, which is calibrated against nothing and is better read as a rough signal than as validation: the median is 0.85, and the model reports confidence at or above 0.7 for 86 percent of Experiment 2 responses and 81 percent for Experiment 3. The three categories central to the analysis—belief correction, expression, and persuasion—carry explicit boundary rules in the prompt.




Figure E.9: Distribution of senders’ stated decision motives in Experiments 2 and 3, among senders whose open-ended decision reason refers to the disclosure decision (referral rates: Experiment 2 27.0%, Experiment 3 73.9%). Panel (a): share for whom each motive is the *primary* stated reason. Panel (b): share mentioning each motive at all (multiple motives may be coded per response). Motives coded from open-ended text by a large language model (claude-sonnet-5; Appendix E.1). Experiment 1 is coded under a later scheme with a different category set and is reported separately in Figure E.8.

E.3 Coding Examples

Table E.46 presents representative Experiment 1 responses classified into each motive category, illustrating where the boundaries fall. Tables E.47 and E.48 present the corresponding examples for the Experiment 2 general-population sample and the Experiment 3 journalist sample, translated from German.

Table E.46: Coding Examples (Experiment 1)

Category	Status	Example
Belief correction	Affirmed	<i>"People should know the truth." "I wanted the people who got the questions wrong to know the correct answer. People should be more knowledgeable about things."</i>
Expression	Affirmed	<i>"When I thought that the fact was unfavorable towards immigration, I chose to not share the fact." "I wanted to show a more positive view of immigrants."</i>
	Mentioned but	<i>"Everyone should be informed about facts even if it goes against my beliefs or hiding facts to sway decisions. Nothing but the facts."</i>
	rejected	<i>"I think everyone deserves to know the truth and be informed, even if the facts don't align with my personal preferences."</i>
Persuasion	Affirmed	<i>"If the question had a direct impact on the way people view immigrants." "People who are ill informed, when they are enlightened, may choose to reallocate more funds towards helping immigrants if they know the truth."</i>
	Mentioned but	<i>"Everyone should know the truth regardless of how it influences their decisions on immigration."</i>
	rejected	<i>"Wanting to be honest instead of manipulative."</i>
Other stated reason	Affirmed	<i>"Basically only on whether or not I personally would want/care to know. For a couple of them, I WOULD want to know. And, for the other two, not as much."</i>
Referred, no stated reason	Gate outcome	<i>"Basic knowledge and provided information." "The information that was provided at the end."</i>
Different task only	Gate outcome	<i>"I made decisions based on what I truly thought was the correct answer and was surprised to find I was mistaken on a few."</i>
Did not refer to disclosure	Gate outcome	<i>"My own beliefs and also taking into consideration the information from research." "I think that immigration is a valuable thing for the US. I think immigrants help the economy and do important jobs US citizens do not want to do."</i>

Open-ended responses from Experiment 1 senders, illustrating the boundaries of the coding scheme. Every example shown is coded identically by both classifiers. *Affirmed* means the response gives distinct support for the category; *mentioned but rejected* means the response raises the consideration and states that it did not govern the decision—the two expression examples are senders who name their own position and say they disclosed regardless of it. The last three rows are gate outcomes and receive no category coding: a response may refer to the disclosure decision without giving a reason for it, may describe only a different task in the experiment, or may state a political opinion with no link to the decision. Responses are lightly excerpted for length; full text available upon request.

Table E.47: Coding Examples — Experiment 2 (General Population)

Primary motive	Example
Belief correction	<i>"Since the information is always truthful, I wanted every person to receive it—regardless of whether it matched their view."</i> <i>"I shared all the information. Everyone has the right to be informed."</i>
Expression	<i>"I checked whether it fit my own attitude."</i> <i>"As a rule, I sent things that support my opinion."</i>
Persuasion	<i>"I wanted to influence the person so that they would come around to my opinion."</i> <i>"I tried to reinforce views that favour climate protection and to convince the others."</i>
Receiver tailoring	<i>"If the person had thought positively, send the message; if they thought negatively, don't send it."</i> <i>"I looked at whether the information matched the person's opinion and then decided whether to send it."</i>
Balance	<i>"I looked at what might broaden their opinion, so that they would also see the other side once."</i> <i>"I tried to be impartial, so also to send information favourable to migration, as long as they were correct facts."</i>

Representative open-ended responses from Experiment 2 senders (German general-population sample), illustrating the classification boundaries between primary motives. Responses are translated from German and lightly excerpted for length; original text available upon request. Primary motive classification by Claude Sonnet 5.

Table E.48: Coding Examples — Experiment 3 (Journalists)

Primary motive	Example
Belief correction	<i>“All readers should have all the information so that they can form their own opinion. That is why I sent everything.”</i> <i>“Everyone has a right to the facts, whatever their opinion. So I always sent them.”</i>
Expression	<i>“I checked whether the information supported my opinion or not, and sent it accordingly.”</i> <i>“Weighing the reader’s existing opinion against the slant of the information—admittedly in favour of my own view.”</i>
Persuasion	<i>“You try to nudge people a little towards your own opinion.”</i> <i>“Starting from my own opinion, I tried to gauge in which case the person would move closer to my view.”</i>
Role neutrality	<i>“I always send information. It is my job.”</i> <i>“As a journalist I send facts neutrally—independent of my personal opinion. The readers’ opinion was really secondary.”</i>
Balance	<i>“Balance was important to me.”</i> <i>“I tried to create a counterweight to the constant political propaganda.”</i>

Representative open-ended responses from Experiment 3 senders (German journalists), illustrating the classification boundaries between primary motives. “Role neutrality” captures the professional-norm rationale of transmitting facts irrespective of the sender’s or receiver’s stance. Responses are translated from German and lightly excerpted for length; original text available upon request. Primary motive classification by Claude Sonnet 5.

F Self-Reported Motives, Norms, and Beliefs

F.1 Stated Norms and Perceived Design Balance

Both experiments closed with the question whether respondents perceived a political imbalance in the study. A large majority perceived none (85.5 % of the general population, 82.4 % of journalists), and the shares are flat across treatment arms and topics. Among the minority perceiving an imbalance, a left-favoring reading outnumbers a right-favoring one roughly two to one in both samples (9.5 % vs. 5.0 % in the general population; 11.7 % vs. 5.9 % among journalists), against a fact menu balanced by construction (two pro and two anti facts, plus one randomly drawn). Table F.49, Panel A, reports the shares.

Table F.49, Panel B, reports the distributions of the three normative items, and Table F.50 the item $\times$ selection interaction battery discussed in Section 5.3.

F.2 Stated Receiver-Model Beliefs

Beyond their sharing decisions, senders explained what they based their receiver-belief predictions on; a large language model codes these prediction rationales into stated beliefs about whether receivers update their position (persuadable, unpersuadable, and related claims) and whether they already know the facts (Appendix E.1). Experiment 1 elicited no such predictions and is omitted. Figure F.10 reports the distribution of these stated beliefs. Table F.51 relates them to the incentivized elicited returns, and Table F.52

	General population (Exp. 2)			Journalists (Exp. 3)		
	Mean	SD	% Agree	Mean	SD	% Agree
<i>Panel A: Perceived political imbalance (% of valid answers)</i>						
No imbalance perceived			85.5			82.4
Imbalance favoring the left			9.5			11.7
Imbalance favoring the right			5.0			5.9
<i>Panel B: Normative and policy items (1–5 agreement)</i>						
Share opposing information	4.13	0.90	80.0	4.60	0.68	95.4
Social-media transparency	4.05	0.96	75.7	4.61	0.65	95.1
Viewpoint-diversity obligation	4.17	0.90	80.4	4.22	0.91	82.7

Post-questionnaire self-reports, one per sender. Panel A: answers to “Did you perceive a political imbalance in the study?” (options: no imbalance, imbalance favoring the left, imbalance favoring the right), in percent of valid answers; Experiment 2 $N = 1,548$ (no nonresponse), Experiment 3 $N = 307$ (one nonresponse). Panel B: agreement on a 1–5 scale (1 = fully disagree, 5 = fully agree) with three items: whoever passes on political information should also pass on factually correct information that contradicts their own position (worded for journalists in Experiment 3); social media should be more transparent about why users see some political information and not others; platforms and media should be obligated to surface relevant information from different perspectives. % Agree is the share answering 4 or 5. Experiment 2 $N = 1,548$; Experiment 3 $N = 306$ (two nonresponses).

Table F.49: Perceived design balance and normative items, Experiments 2 and 3.

Item	Interaction of item with:			Obs.
	Align _{ij}	α_{ij}	$\Delta\pi_{ijk}$	
<i>Panel A: General population (Experiment 2)</i>				
Share opposing information	0.006 (0.010)	0.027 (0.021)	0.006 (0.007)	7740
Social-media transparency	0.008 (0.009)	0.003 (0.019)	0.005 (0.006)	7740
Viewpoint-diversity obligation	0.017* (0.009)	0.054*** (0.020)	0.010 (0.007)	7740
<i>Panel B: Journalists (Experiment 3)</i>				
Share opposing information	−0.091*** (0.021)	−0.109* (0.058)	−0.041 (0.026)	1530
Social-media transparency	−0.058** (0.023)	−0.013 (0.047)	−0.051** (0.025)	1530
Viewpoint-diversity obligation	0.001 (0.022)	0.023 (0.042)	−0.009 (0.016)	1530

Dependent variable: binary provision decision ($y_{ij} \in \{0, 1\}$). Each cell reports the interaction coefficient between the row item (centered at 3, the scale midpoint) and one selection measure, from a separate regression of provision on the measure, the item, and their interaction, with fact fixed effects and sender-clustered standard errors (in parentheses). Selection measures: positional alignment (Align_{ij}), Expression (α_{ij}), and Persuasion ($\Delta\pi_{ijk}$). Items as in Table F.49, Panel B. Benjamini–Hochberg correction is applied across the 18-test family (3 items $\times$ 3 measures $\times$ 2 samples). Three interactions survive at the 10% level: in Experiment 3, the opposing-information item with positional alignment ($q < 0.001$) and the transparency item with positional alignment ($q = 0.079$); in Experiment 2, the viewpoint-diversity item with Expression ($q = 0.078$). The next smallest is $q = 0.185$. Stars mark unadjusted p -values: *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table F.50: Normative items and selection: channel $\times$ item interactions, Experiments 2 and 3.

to the independently coded stated motives of Appendix E.1.

Table F.51: Convergent Validity of the Incentivized Belief Elicitations (Experiments 2 and 3)

	Persuasiveness $ \Delta\pi $		Belief correction Δb	
	(1)	(2)	(3)	(4)
Intercept	0.649*** (0.013)	0.656*** (0.032)	0.170*** (0.006)	0.303*** (0.013)
Stated <i>persuadable</i>	0.191*** (0.058)	−0.005 (0.059)		
Stated <i>unpersuadable</i>	−0.138** (0.055)	−0.267*** (0.049)		
Stated <i>not-know</i>			0.073* (0.040)	0.159*** (0.041)
Stated <i>know</i>			0.111** (0.054)	0.078** (0.036)
Sample	General population	Journalists	General population	Journalists
Num. obs.	1548	306	1548	306
R ²	0.008	0.060	0.004	0.054

Sample: sender-level, Experiment 2 general-population senders ($N = 1,548$) and Experiment 3 journalists ($N = 306$). *Specification:* OLS of each sender's mean incentivized elicited belief on indicators for the receiver-updating claims stated in the sender's open-ended prediction rationale (Appendix E.1). Columns 1–2: dependent variable is the mean predicted movement magnitude $|\Delta\pi| = |\text{delta_shift}|$. Columns 3–4: dependent variable is the mean elicited factual belief correction Δb (*acc_gain*). Stating that receivers are *unpersuadable* predicts smaller elicited movement; stating that receivers do *not know* the facts predicts larger elicited belief correction. The knowledge axis is corroborated by the predicted baseline-knowledge share (*delta_pct_none*), on which *not-know* is negative (−6.9pp in column 3, −9.9pp in column 4) and *know* is null. Heteroskedasticity-robust standard errors in parentheses. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table F.52: Cross-Source Coherence: Stated Motive Predicts Stated Receiver-Model (Experiments 2 and 3)

	Stated <i>persuadable</i>		Stated <i>not-know</i>	
	(1)	(2)	(3)	(4)
Intercept	0.033*** (0.005)	0.176*** (0.025)	0.010*** (0.003)	0.052*** (0.016)
Stated persuasion motive	0.214*** (0.048)	0.162** (0.063)		
Stated belief-correction motive			0.039 (0.028)	0.087 (0.036)
Sample	General population	Journalists	General population	Journalists
Num. obs.	1548	306	1548	306
R ²	0.054	0.027	0.005	0.023

Sample: sender-level, Experiment 2 general-population senders ($N = 1,548$) and Experiment 3 journalists ($N = 306$). *Specification:* linear-probability regressions of an indicator for a stated receiver-updating *belief* (from the *delta_reason* coding of the prediction rationale) on an indicator for the corresponding stated *motive* (from the independent *decision_reason* coding of the sharing rationale). The two open-text questions are coded in separate classification runs. The intercept is the base rate of stating the belief among senders who do not state the motive; the slope is the enrichment among senders who state a motive. Senders who state a persuasion motive are markedly more likely to state that receivers are *persuadable*, and senders who state a belief-correction motive are more likely to state that receivers do not know the facts. Because the stated-belief cells are rare, association is best assessed by Fisher's exact test: $p < 0.001$ (column 1), $p = 0.007$ (column 2), $p = 0.031$ (column 3), $p = 0.011$ (column 4); all four coherences are significant. Heteroskedasticity-robust standard errors in parentheses. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.




Figure F.10: Distribution of senders' stated beliefs about how receivers update, among senders whose open-ended prediction rationale yields a codable claim (codable rates: Experiment 2 11.2%, Experiment 3 54.2%). Panel (a): beliefs about whether receivers' policy position moves. Panel (b): beliefs about whether receivers already know the facts. Shares sum to more than one because a rationale may state several claims. Beliefs coded from open-ended text by a large language model (claude-sonnet-5, delta_reason scheme; Appendix E.1). Experiment 1 elicited no receiver-updating predictions and is omitted.

G Proofs

This appendix collects the proofs of the propositions in Section 3.3 and develops the magnitude prediction in full.

Proof of Proposition 0 (Belief correction). From equation (3),

$$\frac{\partial \Pr(d = 1)}{\partial \Delta b} = f(U(d = 1)) \tau \theta_B.$$

Since $f > 0$ everywhere, $\theta_B > 0$, and $\tau > 0$ the derivative is strictly positive. $\square$

Proof of Proposition 1a (Expressive selective disclosure). From equation (3),

$$\frac{\partial \Pr(d = 1)}{\partial \alpha} = f(U(d = 1)) \theta_E.$$

Since $f > 0$ everywhere and $\theta_E > 0$, the derivative is strictly positive. $\square$

Proof of Proposition 1b (Persuasive selective disclosure). From equation (3),

$$\frac{\partial \Pr(d = 1)}{\partial w(\Delta\pi)} = f(U(d = 1)) \tau c(A).$$

Since $f > 0$ everywhere, $c(A) > 0$, and $\tau > 0$, this derivative is strictly positive.

If w is differentiable and strictly increasing, then $w'(\Delta\pi) > 0$ on the relevant domain. Therefore,

$$\frac{\partial \Pr(d = 1)}{\partial \Delta\pi} = f(U(d = 1)) \tau c(A) w'(\Delta\pi) > 0.$$

$\square$

Proof of Proposition 2 (Political action amplification). Differentiating equation (2) with respect to $w(\Delta\pi)$ gives

$$\frac{\partial U(d = 1)}{\partial w(\Delta\pi)} = \tau(\theta_P + \theta_I A).$$

The difference in this marginal utility between the action and no-action states is

$$\left[\frac{\partial U(d = 1)}{\partial w(\Delta\pi)} \mid A = 1 \right] - \left[\frac{\partial U(d = 1)}{\partial w(\Delta\pi)} \mid A = 0 \right] = \tau(\theta_P + \theta_I) - \tau\theta_P = \tau\theta_I.$$

Since $\tau > 0$ and $\theta_I > 0$, political action strictly increases the marginal utility of persuasion; this increment is the instrumental motive. $\square$

Persuasiveness amplification. The framework also yields a magnitude prediction, stated here in full.

Proposition 3 (Persuasiveness amplification). *Fix $A \in \{0, 1\}$ and suppose $c(A) > 0$. If w is differentiable and strictly increasing, then stronger perceived persuasion increases disclosure when persuasion moves the receiver toward the sender's own position and decreases disclosure when persuasion moves the receiver away from it. Specifically, for $\Delta\pi > 0$,*

$$\frac{\partial \Pr(d = 1)}{\partial \Delta\pi} > 0,$$

while for $\Delta\pi < 0$, disclosure is strictly decreasing in $|\Delta\pi|$:

$$\frac{\partial \Pr(d = 1)}{\partial |\Delta\pi|} < 0.$$

Proof. For $\Delta\pi > 0$, equation (3) implies

$$\frac{\partial \Pr(d = 1)}{\partial \Delta\pi} = f(U(d = 1)) \tau c(A) w'(\Delta\pi).$$

Since $f > 0$, $\tau > 0$, $c(A) > 0$, and $w'(\Delta\pi) > 0$, this derivative is strictly positive.

For $\Delta\pi < 0$, write $\Delta\pi = -|\Delta\pi|$, so that $d\Delta\pi/d|\Delta\pi| = -1$. By the chain rule,

$$\frac{\partial \Pr(d = 1)}{\partial |\Delta\pi|} = \frac{\partial \Pr(d = 1)}{\partial \Delta\pi} \cdot \frac{d\Delta\pi}{d|\Delta\pi|} = -f(U(d = 1)) \tau c(A) w'(\Delta\pi).$$

Since $f > 0$, $\tau > 0$, $c(A) > 0$, and $w'(\Delta\pi) > 0$, this derivative is strictly negative. $\square$

Proposition 3 adds a magnitude prediction to the directional prediction in Proposition 1b. When the state of the world is perceived to move the receiver toward the sender's position, disclosure increases with the perceived strength of that movement. When it is perceived to move the receiver away, disclosure decreases with the perceived strength of that movement. The boundary case is a function such as $w(\Delta\pi) = \text{sign}(\Delta\pi)$, where the sender responds only to the direction of persuasion and not to its magnitude.

Expected loss and the ambiguity of silence. Taking expectations over the state and over transmission, the epistemic loss generated by the environment is

$$\mathcal{L} = p\tau\sigma_1(1 - \tilde{b}_1)^2 + (1 - p)\tau\sigma_0\tilde{b}_0^2 + p(1 - \tau\sigma_1)(1 - \tilde{b}_\emptyset)^2 + (1 - p)(1 - \tau\sigma_0)\tilde{b}_\emptyset^2. \quad (8)$$

Evaluated with the receiver's beliefs, $\mathcal{L}$ is the belief-accuracy loss relative to the realized state and requires no assumption on how she updates. Writing $q_1 = p(1 - \tau\sigma_1)$ and $q_0 = (1 - p)(1 - \tau\sigma_0)$ for the silence weights, with $b_\emptyset^{\text{Bayes}} = q_1/(q_1 + q_0)$, the quadratic metric decomposes it as

$$\begin{aligned} \mathcal{L} = & \underbrace{q_1(1 - b_\emptyset^{\text{Bayes}})^2 + q_0(b_\emptyset^{\text{Bayes}})^2}_{\text{ambiguity of non-receipt}} \\ & + \underbrace{p\tau\sigma_1(1 - \tilde{b}_1)^2 + (1 - p)\tau\sigma_0\tilde{b}_0^2 + (q_1 + q_0)(\tilde{b}_\emptyset - b_\emptyset^{\text{Bayes}})^2}_{\text{departure from Bayes}}. \end{aligned}$$

The first term is the loss a Bayesian still bears because silence overlaps across states—a function of $(p, \sigma_1, \sigma_0, \tau)$ alone, and zero only when non-receipt separates the states.

The remaining terms are the receiver's departures from Bayes, each a squared deviation $(\tilde{b}_m - b_m^{\text{Bayes}})^2$ that vanishes under Bayesian updating and is strictly positive otherwise. Their nonnegativity holds given the benchmark $b_\emptyset^{\text{Bayes}}$, the single point at which a receiver-side rational-updating restriction enters, deliberately and minimally.

Selective disclosure is not intrinsically information-destroying. The ambiguity term equals $q_1 q_0 / (q_1 + q_0)$, a product of the two silence weights whose sum $q_1 + q_0 = 1 - \tau \bar{\sigma}$ is fixed by the unconditional disclosure rate $\bar{\sigma} = p\sigma_1 + (1 - p)\sigma_0$. A product with a fixed sum is largest when its two factors are equal, so selection lowers the term whenever it moves the weights apart, and whether it does so depends on the prior. Imperfect transmission bounds how far this goes, since for $\tau < 1$ both weights stay strictly positive and non-receipt never fully reveals the state.

I compare an environment with its volume-matched non-selective counterfactual $\sigma_1 = \sigma_0 = \bar{\sigma}$, holding τ fixed, which isolates the contribution of selection in two senses that rest on different restrictions. The first is what selection does to the ambiguity term, which I call the *Bayesian information effect of selection*. Its sign is ambiguous: at a symmetric prior $p = 1/2$ the volume-matched counterfactual sits at the maximum of the ambiguity term, so selection weakly sharpens what silence conveys, while away from $p = 1/2$ it can take either sign. Appendix G derives the effect and the threshold where it turns. It depends on $\sigma_1 - \sigma_0$, and disclosure is observed only at the state that occurred: what a sender discloses knowing $\omega = 1$ restricts σ_0 only through an assumption linking her behavior across states, since the counterfactual is her disclosure of the same state had the truth run the other way.

The Bayesian information effect of selection. This paragraph and the next derive the two claims about selection stated in Section 6.5. Throughout, $\tau \bar{\sigma} < 1$, so non-receipt has positive probability. Write $\Delta\sigma = \sigma_1 - \sigma_0$ and $K = 1 - \tau \bar{\sigma} > 0$, so that $\sigma_1 = \bar{\sigma} + (1 - p)\Delta\sigma$ and $\sigma_0 = \bar{\sigma} - p\Delta\sigma$. The silence weights are then $q_1 = pK - h$ and $q_0 = (1 - p)K + h$ with $h = p(1 - p)\tau\Delta\sigma$, and $q_1 + q_0 = K$. The ambiguity term is

$$\mathcal{A}(\tau, \sigma_1, \sigma_0) = \frac{q_1 q_0}{q_1 + q_0} = p(1 - p)K + (2p - 1)h - \frac{h^2}{K}.$$

The volume-matched non-selective rule sets $\Delta\sigma = 0$ and hence $h = 0$, giving $\mathcal{A} = p(1 - p)K$. The Bayesian information effect of selection is the difference,

$$\mathcal{A}(\tau, \sigma_1, \sigma_0) - p(1 - p)(1 - \tau \bar{\sigma}) = p(1 - p)\tau(2p - 1)\Delta\sigma - \frac{p^2(1 - p)^2\tau^2(\Delta\sigma)^2}{1 - \tau \bar{\sigma}},$$

which is positive if and only if $\Delta\sigma$ and $2p - 1$ share a sign and $|\Delta\sigma| < |2p - 1|(1 - \tau \bar{\sigma}) / (p(1 - p)\tau)$. At $p = 1/2$ the linear term vanishes and the effect reduces to $-\tau^2(\Delta\sigma)^2 / (16(1 - \tau \bar{\sigma})) \leq 0$, so selection weakly reduces the ambiguity of non-receipt

at a symmetric prior.

The realized accuracy cost of selection. Let g index environments with frequencies λ_g summing to one, disclosure rates σ_g , and conditional mean losses μ_g^R under receipt and $\mu_g^\emptyset$ under non-receipt. Receipt occurs with probability $\tau\sigma_g$, so

$$\mathcal{L}(\sigma) = \sum_g \lambda_g \left[\tau\sigma_g \mu_g^R + (1 - \tau\sigma_g) \mu_g^\emptyset \right].$$

Replacing every σ_g by $\bar{\sigma} = \sum_g \lambda_g \sigma_g$ and holding τ , the frequencies, and the conditional losses fixed,

$$\mathcal{L}(\sigma) - \mathcal{L}(\bar{\sigma}) = \tau \sum_g \lambda_g (\sigma_g - \bar{\sigma}) \left(\mu_g^R - \mu_g^\emptyset \right) = \tau \text{Cov}_\lambda \left(\sigma_g, \mu_g^R - \mu_g^\emptyset \right),$$

which is (6). Holding the conditional losses fixed is a restriction on receiver behavior rather than an accounting identity: it requires that the receiver's reading of receipt and of silence in environment g not respond to the disclosure rule that generated them.