

Who Should Know What and Why? Curated Supply of Political Information

Yero Samuel Ndiaye*

July 10, 2026

[EARLY DRAFT — COMMENTS WELCOME]

[Click here for the most recent version](#)

Abstract

The information environment is distorted not only by false content but by the selective supply of true content. I study why senders choose which accurate political facts to share, and decompose the motives behind it. Among 1,200 U.S. adults, senders are 8 percentage points more likely to share aligned than unaligned facts; the gap roughly doubles when the receiver acts politically, and operates primarily through withholding unaligned facts. Incentivized belief elicitation across 1,548 German adults reveal the mechanism: an expressive motive drives baseline selectivity; a persuasion motive—shifting the receiver’s beliefs toward one’s own position—activates specifically when the receiver acts. Professional journalists show the same motive pattern, and their selectivity persists even absent the extrinsic incentives the media-bias literature invokes. The *curated truth* distortion operates entirely on true content, a margin that current policies such as fact-checking and labeling do not reach.

JEL classification: D83, D82, C91, D72, L82

Key words: Information Provision, Selective Disclosure, Political Information, Media Bias

*Ndiaye: University of Cologne and Max Planck Institute for Behavioral Economics; ndiaye@wiso.uni-koeln.de. I am grateful for valuable comments and suggestions from Arno Appfelstaedt, Lara Maria Berger, Khanh Hoang, Christoph Feldhaus, Markus Dertwinkel-Kalt, Max Rainer Pascal Grossmann, Melisa Kurtis, Felix Mindl, Johannes Münster, Shakked Noy, Axel Ockenfels, Eva Maria Oess, Malte Rattenborg, Christopher Roth, Nicolas Röver, Christoph Schottmüller, Sebastian Siegloch, Till Stange, Johannes Wohlfahrt, and Tom Zimmermann, as well as by participants of seminars and conferences at the C-SEB Early Ideas Workshop and IMPRS BeSmart Thesis Workshop. I thank Jan-Luca Bauss, Frauke Nora Gerster, Linda Plehn, Lena Schmelig, Anselm Bührlein, and Moritz Lakenbrink for their excellent research assistance. Support from the Center for Social and Economic Behavior (C-SEB) and the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation) under Germany’s Excellence Strategy EXC 2126/1390838866 is gratefully acknowledged.

1 Introduction

People decide which political information to share everywhere and all the time. For example, everybody decides what to share on social media and what to mention in conversations, journalists and editors decide what to include in their articles, and fact-checkers decide which statements to verify.¹ What is shared matters: information shapes beliefs, beliefs shape political preferences, and in democracies these preferences translate into collective decisions.

Distortions in the supply of political information are widely seen as a threat to democratic decision-making, as evidenced by the policy and research focus on the topic. Both have mostly concentrated on fake news as the culprit. Policy targets it directly: the EU Digital Services Act mandates content moderation, while platform fact-checking programs, content labeling, accuracy nudges, and digital-literacy interventions aim to detect, correct, or pre-empt false content. Research has moved in parallel, studying how false content shapes political beliefs ([Allcott and Gentzkow 2017](#), [Barrera et al. 2020](#)), why senders supply it ([Thaler 2021](#), [Marie et al. 2023](#), [Guriev et al. 2026](#)), and which remedies curb its spread and effects ([Cagé et al. 2025](#), [Aridor et al. 2024](#), [Berger et al. 2025](#)). Yet false content is a small share of what people consume. It is roughly 0.15 percent of Americans' daily media diet and about 1 percent of news consumption ([Allen et al. 2020](#)),² exposure is concentrated among a narrow fringe with the strongest motivation to seek it out ([Budak et al. 2024](#)), and sharing it is rare—in one study linking survey and behavioral data, over 90 percent of users shared no articles from fake-news domains at all ([Guess et al. 2019](#)). Where its effects on political beliefs have been measured, they appear limited and concentrated among the already committed ([Allcott and Gentzkow 2017](#), [Budak et al. 2024](#)). The supply of false information that policy targets and research focuses on is therefore only part of what distorts

¹Ordinary individuals are now routine suppliers of political information: about a third of U.S. adults use social media to voice political views ([Pew Research Center 2025](#)), and on Facebook and Instagram more than 70 percent of those who get news on the platform get it from friends, family, and acquaintances ([Pew Research Center 2024](#)).

²[Allen et al. \(2021\)](#) show that Facebook's public-sharing dataset inflates the fake-news share by roughly 4× relative to a representative panel due to URL censoring; the corrected figure is approximately 2.5 percent of Facebook news clicks.

political beliefs.

Selective provision of *true* information is the other part, and it is well-documented observationally. Citizens pass on the fact-checks that favor their side (Shin and Thorson 2017) and politicians amplify the coverage that favors theirs (Freitag et al. 2021); fact-checkers, too, selectively choose which claims to scrutinize (Louis-Sidois 2025), and the pattern holds for true and false content alike (Marie et al. 2023). Selecting among true facts can mislead as much as fabrication: when senders share only the congenial findings and receivers read too little into the silence, beliefs polarize (Bowen et al. 2023), and a true message left unheeded can do as much harm as a false one believed (Feess et al. 2024). Worse, falsehood is at least verifiable and detectable, while withheld truth is *invisible*.

I provide the first causal evidence on the selective provision of true political information and decompose the motives behind it. I cast provision as the outcome of four separable motives. A sender may value correcting the receiver’s factual beliefs; she may gain expressively from passing on facts that fit her own political position; she may want to move the receiver’s beliefs toward that position (projective persuasion); and this persuasive motive gains weight when the receiver’s beliefs may translate into political action (the instrumental motive). Identification rests on two sources of variation: within-sender variation across facts—each sender evaluates aligned and unaligned facts about the same topic, holding sender-level heterogeneity constant—and between-subject assignment to a political-action condition that shifts the return to persuasion. I proceed in three steps: I document selective provision and show that political action amplifies it (Experiment 1), decompose the motives that drive it (Experiments 2 and 3), and target its downstream cost with a receiver experiment (Experiment 4).

Experiment 1 establishes the phenomenon. Among 1,200 quota-sampled U.S. adults, each holding a true immigration fact, provision exceeds 80 percent in every condition, but it is selective: senders are about 8 percentage points more likely to share aligned than unaligned facts ($p < 0.01$), controlling for prior beliefs. The gap roughly dou-

bles, to 16 percentage points, when the receiver makes a downstream political donation ($p < 0.05$), and it widens primarily through the withholding of unaligned facts (-4.6 pp) rather than additional sharing of aligned facts ($+1.9$ pp). Because the receiver never learns that a disclosure decision was made, the pattern is not driven by reputation or signaling. Experiment 1 cannot, however, say *why* senders select: alignment bundles the expressive value of sharing a congenial fact with its persuasive value, and the two cannot be separated within this design.

Experiment 2 separates them. With 1,548 quota-sampled German adults across two policy domains, I add incentivized belief elicitation that distinguish a fact's expected factual belief correction from its expected persuasive effect, and receivers of known political position that make audience-conditional provision observable. I find that expressive identity drives selective provision absent political action ($p < 0.01$), whereas instrumental persuasion activates specifically when the receiver acts politically ($p < 0.10$)—the motive that positional alignment alone cannot isolate. Belief-correction raises provision without making it selective ($p < 0.05$).

Experiment 3 takes the same design to 307 professional German journalists. I find their provision is driven by belief correction and persuasion (both $p < 0.01$; the expressive component is positive but not significant), with larger factual belief correction and persuasion coefficients than the citizen action arm, and it persists once the extrinsic incentives that media-bias models invoke are held constant. *Some* journalistic selectivity is therefore intrinsic. Experiment 4 turns to the receiver: when audiences under-read the silence that selective sharing leaves behind, a diet of entirely true facts can still move beliefs apart, and the experiment is designed to quantify this downstream cost ...

I provide causal evidence on a distortion carried entirely by true facts, and separate the motives behind it: belief correction, expressive value, projective persuasion, and instrumental persuasion. The decomposition points to two policy implications, and existing tools reach neither. Accuracy-based content policy (fact-checking and content-labeling) misses *curated truth* because the distortion rides on the truth of individual

facts rather than on their falsehood. Information quality depends on completeness as well as accuracy, and completeness is the dimension current tools leave untouched; this is a distortion that fact-checking and labeling do not reach. Media-bias remedies premised on extrinsic incentives, such as competition, ownership, and consumer demand, face a separate limit: they cannot reach the intrinsic component. Experiment 3 isolates, the selectivity that survives once those incentives are held constant. The decomposition sharpens the concern, because the instrumental motive is the component policy can least detect, and it grows strongest where beliefs translate into collective decisions—the stakes at which the gap in policy is most costly.

1.1 Related Literature

I connect to five strands of literature.

Individual selective supply. Citizens selectively pass on the fact-checks that favor their side (Shin and Thorson 2017) and politicians amplify coverage that favors theirs (Freitag et al. 2021); fact-checkers, too, are selective in which claims they scrutinize, with the bias residing in target selection rather than in verdict accuracy (Louis-Sidois 2025, Andreottola and Louis-Sidois 2025). The pattern extends to true and false content alike (Marie et al. 2023), is driven by partisanship (Osmundsen et al. 2021), and intensifies before like-minded audiences (Marie and Petersen 2025). This literature documents that selective provision of true information exists but cannot separate its underlying motives: naturalistic sharing data conflate the expressive benefit of sharing politically congenial facts with the persuasive benefit of shifting receivers' beliefs, and both are confounded with belief-correction concerns in real-world settings. A parallel literature shows that senders supply opinions and stated beliefs strategically for expressive and social-image reasons (Braghieri 2024, Bursztyn et al. 2023b,a, Hillman 2010); I identify the corresponding selective-disclosure mechanism when the content is factual and verifiable. I provide the first motive decomposition, using controlled within-sender variation and incentivized belief elicitation to separate what observational work could only document in aggregate.

Fake news: effects and decomposition. A body of work establishes that false content is a small share of consumption whose measured effects on political beliefs are limited and concentrated among the already committed (Allen et al. 2020, Allcott and Gentzkow 2017, Budak et al. 2024, Guess et al. 2019, Feess et al. 2024), that it nonetheless spreads faster and farther than true content (Vosoughi et al. 2018, Lazer et al. 2018), and that voters vary widely in how well they tell true news from false (Angelucci and Prat 2024). The remedies policy deploys reach only adjudicable falsehoods: fact-checks correct false claims imperfectly and incompletely, media literacy interventions show limited durable effects (Barrera et al. 2020, Berger et al. 2025, Cagé et al. 2025), and accuracy nudges that redirect attention to truth reduce false sharing (Pennycook et al. 2021)—none reaching the selective omission of true facts. Guriev et al. (2026) structurally decompose the motives for sharing political content—reputation, partisan persuasion, and partisan signaling—in a setting where the receiver observes that a sharing decision has been made and where content may be true or false; related survey-experimental work studies the checking and sharing of alternative facts (Henry et al. 2022). The two designs partition the information environment: their setting covers visible sharing of true and false content, with motives recovered via a structural model; mine covers verified-true content with unattributable withholding, identifying motives from elicited beliefs without a structural model, and isolating the instrumental persuasion motive through an exogenous political-action treatment rather than recovering it structurally. Thaler (2021) shows that incentivizing senders to be rated truthful can backfire into the supply of *false* motivated beliefs; my object is the mirror image, the selective withholding of verified-true content.

Institutional media bias. Outlet-level selectivity arises from extrinsic incentives. Gentzkow and Shapiro (2006) show that profit-maximizing outlets slant toward consumers' prior beliefs to build a reputation for accuracy, with competition disciplining slant by raising the probability of ex-post verification of the true state. Gentzkow and Shapiro (2010) establish, across 433 U.S. daily newspapers, that consumer demand is the dominant driver of slant—accounting for roughly 20 percent of cross-newspaper variance—

while owner ideology contributes negligibly and political capture through advertising is a null (Besley and Prat 2006, Prat and Strömberg 2013, Martin and McCrain 2019). Political control nonetheless shapes what outlets report (Durante and Knight 2012), and exposure is segregated asymmetrically across the ideological spectrum, with platform architecture amplifying partisan disparities in news sharing (González-Bailón et al. 2023, Ershov and Morales 2024). Evidence from journalist movers establishes that individual ideology leaves a fingerprint even within these settings: journalist fixed effects explain around 16 percent of between-outlet slant variation and travel with journalists across outlets (Boxell and Conway 2022); ownership transitions amplify this through ideological sorting of journalists into compatible outlets (Cagé et al. 2022). Yet in naturalistic settings all extrinsic channels operate simultaneously, so a controlled setting is needed to isolate the intrinsic component. Most slant, moreover, lives within outlets rather than across them (Braghieri et al. 2025), so outlet-level accounts—whether from journalist ideology or, as theory shows, from competitive incentives that reduce factual provision in equilibrium (Zhu and Dukes 2015)—miss the margin at which individuals select. Experiment 3 finds that selectivity persists once these extrinsic incentives are held constant—including the financial-incentive channel Berger (2026) documents for a similar population of German online journalists. *Some* journalist selectivity is therefore intrinsic: it survives absent the incentives these models invoke.

Information demand and belief updating. Individuals disproportionately seek and value congenial information and avoid disconfirming information (Golman et al. 2022, Peterson and Iyengar 2021, Chopra et al. 2022, Grigorieff et al. 2020), partisans diverge even in their stated factual beliefs (Bullock and Lenz 2019), and the platforms that mediate this demand have their own economics (Aridor et al. 2024); provision experiments show that factual information nonetheless shifts beliefs and behavior (Haaland et al. 2023, Alesina et al. 2023, Barrera et al. 2020, Kuziemko et al. 2015). These papers study the demand side. Chopra et al. (2024) decompose news demand into accuracy and belief-confirmation motives—the demand analog of my supply-side decomposition. Bowen et al. (2023) supplies the downstream theory: when senders pass along

only congenial true news and receivers misread the resulting silence, beliefs polarize even though no signal is false. His “dogmatic” sharer is the behavior my expressive motive micro-founds, and the setting his result needs—selective sharing whose omissions go unnoticed—is the one my design creates. How receivers read that silence is itself patterned: people underweight belief-threatening signals, an asymmetry that recedes once the ego stake is removed (Eil and Rao 2011, Möbius et al. 2021, Barron 2021, Zimmermann 2020), in politics the trust placed in a message tracks its congeniality (Thaler 2024), and receivers are only partly cursed about the tie between a sender’s silence and her private information (Eyster and Rabin 2005, Fong et al. 2025). The receiver experiment is designed to measure this updating from non-disclosure directly, matched to the disclosure rates the sender experiments estimate.

Voluntary disclosure. The disclosure literature shows that informed senders reveal verifiable evidence in equilibrium because receivers can infer bad news from silence—the unraveling principle (Grossman 1981, Milgrom 1981). Unraveling attenuates when receivers are uncertain whether the sender is informed (Dye 1985, Jung and Kwon 1988, Verrecchia 1983) or engage in motivated belief formation (Manili 2024), and laboratory evidence confirms that non-disclosure is common even when unraveling should hold (Jin et al. 2021, Degan et al. 2023, Burdea et al. 2023, Farina and Leccese 2025). Experiments 2 and 3 implement the Dye probabilistic-informed-sender structure directly: a 50% transmission probability means the receiver cannot attribute a missing fact to the sender’s decision—silence is weakly informative but not perfectly so, because non-receipt is equally consistent with a failed transmission. Experiment 1 adds a stronger suspension: the receiver does not know a disclosure decision was made at all, so silence carries no inferential content. Unlike cheap talk (Crawford and Sobel 1982), both designs use verifiable content; unlike full-disclosure settings, silence never perfectly reveals selective omission in either.

Farina et al. (2026) shows that selection (choosing which favorable signals to reveal from a larger pool) is the dominant distortion, with receivers failing to fully discount it. My evidence is true political facts rather than abstract urn draws, enabling direct elic-

itation of politically motivated sender preferences; non-disclosure is unattributable to the sender rather than revealed by an empty message, eliminating full disclosure as the equilibrium benchmark; the receiver holds a known political position rather than being anonymous, making audience-conditional provision directly observable; and I decompose four sender motives from incentivized belief elicitation rather than recovering behavioral types *ex post*. My framework also relates to Bayesian persuasion (Kamenica and Gentzkow 2011)—I impose no commitment and study motives for selection rather than optimal signal design—and the selective-omission model of Mullainathan and Shleifer (2005) is an outlet-level mechanism related to the individual-level mechanism I identify experimentally. Projective paternalism (Ambuehl et al. 2021), with related work on substituting information for choice restriction (Grossmann 2024, Bartling et al. 2023), is conceptually related to the projective persuasion motive, which my design isolates at the stage of selecting which facts to present rather than restricting choices.

I structure the rest of the paper as follows. Section 2 presents Experiment 1 as motivating evidence. Section 3 develops a conceptual framework for sender behavior. Section 4 presents Experiment 2. Section 5 extends the design to professional journalists. Section 6 develops the receiver-updating framework and the receiver experiment. Section 7 discusses the implications, and Section 8 concludes.

2 Experiment 1

Experiment 1 documents a basic phenomenon: people who know a true fact share it selectively, and the selectivity intensifies when the recipient makes a downstream political decision.

2.1 Design, Treatment, and Hypotheses

I run a pre-registered online experiment with $N = 1,200$ U.S. adults recruited on Prolific using quota-sampling on age, gender, ethnicity, and party identification. I choose policy domain U.S. immigration as (i) opinions are polarized and (ii) misconceptions

about factual information are common.




Figure 1: Experimental timeline, Experiment 1.

Experimental Timeline The experiment proceeds in three parts (Figure 1). In Part 1, I elicit senders' prior knowledge about four immigration facts and measure their policy preferences through a donation decision. In Part 2, I elicit senders' beliefs about the persuasive effect of each fact. In Part 3, senders make one provision decision per fact, each directed at a specific receiver, and a closing questionnaire records motivations, credibility assessments, and demographics. Identification rests on two dimensions varied within sender (alignment between sender and fact; receiver belief state) and two varied between sender (whether the receiver takes a political action; whether senders

learn about the effect of information). The primary within-sender dimension is alignment: each sender faces two aligned and two unaligned facts, so the design holds all sender-level heterogeneity fixed. A fact is *aligned* for sender i if its directional implication favors i 's political position—that is, if the fact affirms what the sender prefers to be true—and *unaligned* otherwise. Alignment is simultaneously an expressive property (sharing a politically congenial fact) and a persuasive one (an aligned fact also tends to move a receiver toward the sender's position). The primary between-sender dimension is ACTION: receivers make the same incentivized donation decision as senders in the ACTION condition and complete the experiment without further decisions otherwise. Receivers are unaware that a disclosure decision was made: provided information appears as part of the survey flow, withheld information leaves no trace, and non-disclosure therefore carries no inferential content. The INFO treatment provides exogenous variation in senders' beliefs about how persuasive each fact is; the INFO manipulation is analyzed in Appendix A.4.

Facts about Immigration I select four verified-true facts from authoritative sources (Bureau of Justice Statistics, Bureau of Labor Statistics, Customs and Border Protection, National Academies of Sciences), balanced across policy directions, and validate them in a pre-registered survey ($N = 100$) that tested a larger candidate pool against criteria of veracity, perceived credibility, policy relevance, and directional balance (see Appendix A.1). Each is framed as an invertible binary comparison (e.g., “higher” versus “lower”), so beliefs are correct or incorrect rather than varying in degree. I classify crime and labor as pro-immigration and border and fiscal as anti-immigration, and verify the classification using senders' expected attitude shifts (Section 2.3.1). The four statements are:

1. *Crime*: The crime rate of immigrants is *lower* than that of U.S.-born Americans.
2. *Labor*: The labor force participation of immigrants is *higher* than that of U.S.-born Americans.
3. *Border*: The average number of unauthorized border crossings is *higher* in the last

five years compared to historical averages.

4. *Fiscal*: When comparing what first-generation immigrants use in government services versus what they pay in taxes, the net cost is *higher* than the same calculation for U.S.-born Americans at the state and local levels.

Beliefs about Facts I first elicit senders' prior knowledge by having them complete each comparative statement, in random order, as "higher," "lower," or "don't know." Responses are incentivized: correct completions receive a bonus, incorrect completions do not, and "don't know" responses receive the bonus with 25% probability (Appendix A.2). Senders are then told the correct answer for each. Because senders who already knew a fact may hold it with greater conviction, I control for prior belief accuracy throughout.

Policy Preferences I measure senders' policy preferences through an incentive-compatible donation task: each sender allocates \$500 between the *American Immigration Council* (AIC), which advocates for increased immigration, and *NumbersUSA*, which advocates for decreased immigration. Senders are shown each organization's position in their own words before deciding. The split yields a continuous measure of policy preference, and therefore of alignment, that captures both the direction and the intensity of each sender's position. Appendix A.2 reports the implementation details.

Receiver Belief States Each fact is assigned to one receiver who either completed the comparative statement incorrectly (false belief) or reported not knowing; the sender observes which case applies before deciding. Both types benefit from disclosure, but correcting a false belief arguably yields a larger expected accuracy gain than updating from ignorance.

Treatments Table 1 summarizes the four design dimensions.

In the ACTION condition, receivers make the same incentive-compatible \$500 donation decision as senders; in the no-ACTION condition, they complete the experiment

Table 1: Experimental design and identifying variation

Dimension	Variation	Levels
Receiver belief state	Within-subject	<i>False belief</i> and <i>Don't know</i>
Alignment with statistic	Within-subject	<i>Aligned</i> and <i>Unaligned</i>
Political action (ACTION)	Between-subject	ACTION = 1 and ACTION = 0
Effect of information (INFO)	Between-subject	INFO = 1 and INFO = 0

Experiment 1 ($N = 1,200$ U.S. adults, Prolific, quota-sampled). Within-subject dimensions: each sender evaluates both receiver belief states and all four statistics (two aligned, two unaligned), yielding 4 decisions per sender. Between-subject dimensions: senders are randomly assigned to one of four cells ($n = 300$ each) crossing ACTION and INFO. The Identification column indicates the utility-function proposition each design dimension is intended to separately identify.

without it.³ The INFO treatment provides half of senders with the actual results of [Haaland and Roth \(2020\)](#), generating exogenous variation in beliefs about persuasiveness; the manipulation is analyzed in Appendix A.4. Crossing the two between-sender treatments gives a 2×2 factorial design with $n = 300$ per cell: CONTROL (neither), A (ACTION only), I (INFO only), and A $\times$ I (both).

Hypotheses I pre-registered four hypotheses about how senders provide information.

Hypothesis 0 (Baseline provision). *Provision rates for both aligned and unaligned information exceed the indifference benchmark of 50% in all treatments.*

Hypothesis 1 (Alignment-based selectivity). *Senders provide aligned information at strictly higher rates than unaligned information in treatment CONTROL.*

Hypothesis 2 (Political action amplifies selectivity). *The alignment gap in provision increases when receivers take political actions.*

Hypothesis 3 (Magnitude of preference shift amplifies selectivity). *The alignment gap in provision is increasing in senders' beliefs about the magnitude of the attitude shift that disclosure induces.*

I test Hypothesis 3 two ways: through the between-sender INFO treatment, which

³Through randomization any difference in alignment-based selectivity is attributable to the downstream political consequence, under the maintained assumption that ACTION operates only through that consequence. ACTION may also shift the salience of the decision and the sender's construal of the task.

shifts beliefs about persuasiveness, and through within-sender variation in expected attitude change, which provides a correlational test.

2.2 Sample, Incentives, and Protocol

The sender experiment was run on Prolific between July 16 and August 12, 2025, quota-sampled on age, gender, ethnicity, and party identification to approximate the U.S. adult population. (See Table B.4 for details.) This study was approved by the Ethics Council of the Max Planck Society under the Generalized Approval of Experiments Following the Protocol that is Standard in Experimental Economics (approval no. 2018_3 Renewal 2024_28) and was pre-registered on AsPredicted (No. 238343). Of 1,257 starters, I stopped collection at the pre-registered target of 300 per treatment passing the attention check, not switching tabs during the fact questions, and not timing out. Senders took a median of 17.26 minutes and earned \$3.32 on average. Receiver data were collected on Prolific between September 1 and September 10, 2025; 96 receivers took a median of 3 minutes and earned \$0.80 on average.

2.3 Results

2.3.1 Senders' Beliefs and Policy Preferences

Before turning to treatment effects, I test the prerequisite for my hypotheses: senders have policy preferences. Additionally, I show that senders believe information can shift receivers' policy preferences and that they hold heterogeneous prior beliefs about the facts with partisan sorting.

Senders' Beliefs about Facts I find substantial heterogeneity in senders' factual knowledge.⁴ A majority correctly identify facts about immigrants' crime rates (70%), labor force participation (79%), and recent border crossings (61%), but only 30% correctly assess the net fiscal impact of first-generation immigrants. Democrats and immigration

⁴See Table B.5 for details. The fiscal fact, where prior knowledge is lowest, is rated as credible ex post as the other three (Table B.6) and provided at comparable rates; results are robust to leaving out each fact individually (Appendix Table B.12).

supporters are more accurate on pro-immigration facts (crime and labor) and Republicans and immigration opponents on anti-immigration facts (border and fiscal; all differences $p < 0.001$). This partisan sorting validates my classification of facts as pro- or anti-immigration and motivates controlling for senders' prior beliefs throughout.

Policy Preferences The distribution of senders' donation decisions is markedly bimodal (Figure B.2): 30.2% allocate their entire endowment to the pro-immigration group, 14.1% give everything to the anti-immigration group, and only 7% split equally. Donations correlate strongly with stated policy preferences ($r = 0.56$, $p < 0.001$) and political affiliation ($p < 0.001$), providing the variation in alignment needed to test for selective provision.

Senders' Beliefs about the Effect of Information on Political Preferences Senders already believe that information strongly affects policy preferences (Figure B.1): roughly 90% expect information to increase petition signatures, and 73% correctly identify that individuals with initially false beliefs update more strongly. They do, however, underestimate the magnitude—predicting a 38% increase on average against a true effect of 69%—leaving room for the INFO condition to shift beliefs upward.

Expected Attitude Effects of Information Alignment-based selectivity requires that senders expect each fact to shift receivers' preferences in a predictable direction. For each sender i and fact j , I define the expected attitude change $\Delta\text{Attitude}_{i,j} \in [-1, 1]$ as the net share of ten hypothetical receivers whom the sender expects to become more (rather than less) supportive of immigration after receiving fact j (unincentivized, elicited before the provision decisions), with 1 (−1) implying all receivers become more (less) supportive of immigration when seeing fact j .

Senders expect information to shift attitudes in the predicted directions. Providing the crime fact is expected to shift net attitudes 25 percentage points toward pro-immigration and labor 17 pp; the border fact is expected to shift net attitudes 34 pp toward anti-immigration and fiscal 32 pp (all $p < 0.001$; Appendix Table B.7). These ex-




Figure 2: Mean provision rates by treatment and statistic alignment, Experiment 1 (U.S. adults, $N = 1,200$, Prolific, quota-sampled). Panel (A): mean sender-level provision rates for aligned (dark green) and unaligned (dark red) statistics, by treatment cell. Panel (B): alignment bias (aligned minus unaligned provision rate) by treatment cell. Restricted to senders with revealed directional preference ($N = 1,114$, 93%; $AIC_i > 250$ = pro-immigration, $AIC_i < 250$ = anti-immigration, $AIC_i = 250$ excluded). Error bars are 95% confidence intervals.

pected attitude shifts establish that senders believe disclosure moves receivers' policy attitudes in the hypothesized directions and corroborate the classification. The classification holds pair by pair, not only on average: in 76% of sender–fact pairs the predicted net movement runs in the fact's own direction. For a sender with a directional position, a congenial fact and a fact expected to move receivers toward that position are therefore almost always the same fact. Senders expect net movement toward their own position for 78% of aligned facts but only 17% of unaligned ones (a gap of 61 pp, $p < 0.001$). The INFO treatment, however, generates only modest and one-sided variation in expected effect on attitudes, raising it for the pro-immigration facts alone (Appendix Table B.7; fuller detail in Appendix A.4).

2.3.2 Treatment Effects on Selective Provision of Information

Let $y_{i,j}$ equal 1 if sender i provides fact j , and let $AIC_i \in [0, 500]$ denote i 's donation to the pro-immigration group (AIC). Sender i is pro-immigration if $AIC_i > 250$ and anti-immigration if $AIC_i < 250$; the few with $AIC_i = 250$ reveal no directional preference and are excluded from the sender-level analysis. I classify $\mathcal{J}^+ = \{\text{crime, labor}\}$ as pro-

immigration and $\mathcal{J}^- = \{\text{border, fiscal}\}$ as anti-immigration. Sender i 's aligned set is $\mathcal{A}_i = \mathcal{J}^+$ if pro-immigration and $\mathcal{J}^-$ otherwise, with the unaligned set $\mathcal{U}_i$ defined analogously.

I define the sender's alignment bias as:

$$\text{Alignment Bias}_i = \frac{1}{|\mathcal{A}_i|} \sum_{j \in \mathcal{A}_i} y_{i,j} - \frac{1}{|\mathcal{U}_i|} \sum_{j \in \mathcal{U}_i} y_{i,j}. \quad (1)$$

A positive alignment bias indicates selective provision favoring the sender's own policy position. A zero alignment bias indicates no selectivity, which includes both sharing all information and sharing no information.⁵

I now test whether senders selectively provide accurate information and whether political action amplifies this selectivity. Figure 2 summarizes senders' provision decisions, restricting the sender-level analysis to the 1,114 senders (93%) who revealed a directional policy preference.

Panel (A) shows that senders overwhelmingly share accurate information across all four treatment conditions. Unaligned provision rates range from 81.9% (A) to 86.5% (CONTROL), all significantly above 50% ($p < 0.001$, one-sample t -tests).

Result 0 (Information provision is high). *Senders share accurate information at high rates, including information unaligned with their policy preferences ($p < 0.001$ against the 50% benchmark in all treatments).*

Panel (B) reveals that provision is selective. The alignment bias is significantly positive in all four conditions ($p < 0.01$), ranging from 6.7 to 13.3 pp. In both pairwise comparisons holding the INFO condition constant, the alignment bias is larger under ACTION (Appendix Table B.9).

⁵At the sender level, I test Hypothesis 0 using a one-sample t -test of the unaligned provision rate against 0.5; Hypothesis 1 using a one-sample t -test of $\text{Alignment Bias}_i > 0$ in CONTROL; and Hypotheses 2 and 3 using directed two-sample t -tests comparing Alignment Bias_i across the relevant treatment conditions. I Holm-correct tests for Hypotheses 2 and 3.

Decision-Level Analysis I capture within-sender variation in preference intensity with a continuous alignment measure:

$$\text{Align}_{i,j} = \begin{cases} \frac{AIC_i - 250}{250} & \text{if } j \in \mathcal{J}^+ \\ \frac{AIC_i - 250}{250} \cdot (-1) & \text{if } j \in \mathcal{J}^- \end{cases}$$

This measure is bounded in $[-1, 1]$: positive values indicate alignment with sender i 's preference (relative to the midpoint $AIC_i = 250$), negative values misalignment, and $|\text{Align}_{i,j}|$ captures preference intensity. Because $\text{Align}_{i,j} = 0$ for indifferent senders, it incorporates all senders. Equivalently, $\text{Align}_{i,j} = \text{Pref}_i \times \text{Dir}_j$, where $\text{Pref}_i = (AIC_i - 250)/250$ is revealed preference intensity and $\text{Dir}_j = +1$ if $j \in \mathcal{J}^+$ and -1 if $j \in \mathcal{J}^-$.⁶

Let $\text{Belief}_{i,j}$ denote sender i 's prior accuracy about fact j (1 if correct, 0.5 if "don't know," and 0 if incorrect as pre-registered), capturing the stronger accuracy norm faced by senders who already knew a fact was true; without this control, the uniformly higher provision rates for pre-known facts would conflate the belief-correction motive with the alignment effect. I estimate the pre-registered linear probability model:

$$\begin{aligned} y_{i,j} = & \beta_0 + \beta_1 \text{Align}_{i,j} + \beta_2 \text{ACTION}_i + \beta_3 \text{INFO}_i + \beta_4 (\text{Align}_{i,j} \times \text{ACTION}_i) \\ & + \beta_5 (\text{Align}_{i,j} \times \text{INFO}_i) + \beta_6 (\text{ACTION}_i \times \text{INFO}_i) \\ & + \beta_7 (\text{Align}_{i,j} \times \text{ACTION}_i \times \text{INFO}_i) + \gamma_1 \text{Belief}_{i,j} + \gamma_2 \mathbf{S}_j + \gamma_3 \mathbf{X}_i + \eta_{i,j}, \end{aligned} \quad (2)$$

with standard errors clustered at the sender level, fact fixed effects $\mathbf{S}_j$, and demographic and political controls $\mathbf{X}_i$. Column (1) includes no fixed effects or controls; column (2) adds $\mathbf{S}_j$; column (3) additionally adds demographic and political controls $\mathbf{X}_i$ and is the preferred specification. β_1 captures baseline alignment-based selectivity (Hypothesis 1), β_4 whether alignment matters more when receivers act politically (Hypothesis 2), and β_5 whether exogenously shifting beliefs about persuasiveness amplifies selectivity (Hypothesis 3). The control γ_1 estimates whether senders who knew

⁶The pre-registration defined alignment on $[0, 1]$. Centering at the indifference point to ease interpretation is a one-to-one transformation that leaves all hypothesis tests invariant, so results under the pre-registered coding are identical.

Table 2: Treatment Effects on Information Provision

	Probability of Statistic Provision		
	(1)	(2)	(3)
Align	0.041*** (0.012)	0.036*** (0.012)	0.037*** (0.012)
INFO	-0.013 (0.017)	-0.012 (0.017)	-0.012 (0.017)
ACTION	-0.002 (0.016)	-0.002 (0.016)	-0.003 (0.016)
Belief Accuracy	0.043*** (0.011)	0.034*** (0.012)	0.033*** (0.012)
Align $\times$ INFO	0.006 (0.017)	0.004 (0.017)	0.004 (0.017)
Align $\times$ ACTION	0.041** (0.018)	0.041** (0.018)	0.041** (0.018)
INFO $\times$ ACTION	-0.007 (0.024)	-0.007 (0.024)	-0.008 (0.024)
Align $\times$ INFO $\times$ ACTION	-0.028 (0.026)	-0.027 (0.026)	-0.027 (0.026)
Clustered SE	Subject	Subject	Subject
Statistics FE	No	Yes	Yes
Controls	No	No	Yes
No. of Senders	1200	1200	1200
Num. obs.	4800	4800	4800
R ²	0.030	0.033	0.040
Adj. R ²	0.028	0.031	0.034

Sample: $N = 1,200$ U.S. adults recruited on Prolific (quota-sampled), 4,800 sender-statistic observations (4 statistics per sender). *Outcome:* binary provision indicator. *Align:* $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, 1]$, where Pref is derived from an incentive-compatible donation split. *Belief Accuracy:* incentivized prior belief correctness, elicited before senders knew they would make a provision decision. *Info:* between-subject treatment sender receives personalized feedback comparing their prediction to the actual results of Haaland and Roth (2020) (see Appendix A.4). *Action:* between-subject treatment receiver makes a downstream political donation. *Controls:* age, gender, education, party affiliation, political knowledge. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

a fact provide it at higher rates.

Table 2 reports the results.⁷ Across all specifications, provision increases significantly in political alignment. A shift from maximally unaligned to maximally aligned increases provision by 7.2 to 8.2 pp ($p < 0.003$).

⁷Appendix Table B.11 reports specifications allowing asymmetric INFO effects by fact direction to accommodate the asymmetric first-stage effect. The additional terms $\text{INFO}_i \times \text{PROSTAT}_j$ and $\text{Align}_{i,j} \times \text{INFO}_i \times \text{PROSTAT}_j$ are not statistically significant in the specifications with fact fixed effects. Wald tests fail to reject that INFO changes the alignment gradient for pro-immigration facts ($p = 0.403$) or anti-immigration facts ($p = 0.606$), and fail to reject that INFO changes overall provision for pro-immigration facts ($p = 0.390$).

Result 1 (Provision is higher for aligned than unaligned information). *Provision is selectively biased: senders are 7 to 8 pp more likely to share aligned information ($p < 0.003$).*

The interaction $\text{Align}_{i,j} \times \text{ACTION}_i$ is positive and statistically significant across all specifications ($p < 0.05$). Political action widens the alignment gap by an additional 8.2 pp, approximately doubling the total effect from 7.2–8.2 pp to 15.4–16.4 pp.⁸

The action effect operates primarily by reducing provision of *unaligned* information rather than by increasing provision of aligned information. Moving from CONTROL to A reduces unaligned provision by 4.6 pp (from 86.5% to 81.9%) while increasing aligned provision by only 1.9 pp (from 93.3% to 95.2%). This asymmetry reflects, in part, a mechanical ceiling: aligned provision already exceeds 93%, leaving limited margin for further increase.

Result 2 (Selectivity is higher when receivers take a downstream political action). *Alignment-based selectivity approximately doubles when receivers make a downstream political action ($p < 0.05$). The action effect operates primarily through reduced provision of unaligned information rather than increased provision of aligned information.*

Senders who directly know a fact to be true provide it at higher rates (3.3 to 4.3 pp, $p < 0.01$). This effect is smaller in magnitude than the alignment effect and does not attenuate the $\text{Align} \times \text{ACTION}$ interaction.

Neither the main effect of INFO nor its interactions with alignment are statistically distinguishable from zero, including in specifications that allow asymmetric effects for the pro-immigration facts where INFO successfully shifted beliefs about persuasiveness. This null has two informative interpretations. First, the first-stage results in Section 2.3.1 show that INFO generated only modest and asymmetric shifts in expected persuasiveness, so the manipulation had little to amplify. Second, conditional on the first stage, the pattern indicates that senders respond to the *direction* of expected persuasion, not its magnitude: provision is unchanged when the direction is fixed and only the expected degree of attitude change moves.

⁸Table B.10 shows that the results are robust to logistic regression estimation.

Result 3 (Beliefs about the effect of information do not affect selectivity). *Information about the effect of information does not affect the likelihood of overall information provision nor the selectivity of information provision.*

2.3.3 Heterogeneous Treatment Effects

Having established that alignment-based selectivity exists at baseline and approximately doubles under political action, I now ask who drives each component, examining heterogeneity by political polarization as pre-registered. Two distinct groups drive the two effects. The baseline alignment bias is concentrated among strong partisans: moving from non-partisanship to full partisanship raises it by 14 percentage points ($p = 0.017$), while Independents show no baseline bias. The action-induced amplification, by contrast, is driven by Independents: political action raises their alignment bias by 14.6 pp ($p < 0.05$) and insignificantly lowers it for strong partisans (-10.2 pp, $p = 0.214$). Appendix Table B.13 reports the full specification by three measures of polarization.⁹

Receiver belief state does not drive provision decisions. Each fact is matched to a receiver who either holds a false belief or reports not knowing; a belief-correction motive would predict higher provision toward false-belief receivers, whose accuracy gap is larger.¹⁰ Table B.16 shows this does not obtain: senders are not more likely to share with false-belief receivers (-2.5 pp, $p = 0.17$), and the alignment gap is comparable across both receiver types (5.2 pp for don't-know receivers, 9.0 pp for false-belief receivers; difference $p = 0.39$). Senders respond to the *direction* of the fact, not the epistemic state of the recipient.

⁹The sender-level specification is $\text{Alignment Bias}_i = \pi_0 + \pi_1 A_i + \pi_2 Z_i + \pi_3 A_i Z_i + \gamma' X_i + \eta_i$, with moderator Z_i and controls X_i , estimated on CONTROL and treatment A. Appendix Table B.14 pools across the INFO condition and Appendix Table B.17 reports the decision-level specification; conclusions are identical in direction and significance.

¹⁰Senders also do not expect larger attitude shifts for false-belief than for don't-know receivers (Appendix Table B.15).

2.3.4 Mechanism-Consistent Correlational Evidence

The observed selective provision could reflect veracity concerns: senders may share what they find credible and withhold what they doubt, and they may find aligned facts more credible. Belief-correction measures predict the *level* of provision but not its alignment-based selectivity. Facts that are deemed credible ex-post are 25.6 to 26.9 pp more likely to be shared ($p < 0.01$), yet adding ex-post credibility and prior belief accuracy and their interactions with ACTION leaves the alignment effect (6 to 6.2 pp, $p < 0.05$) and the Align $\times$ ACTION interaction (9 to 9.4 pp, $p < 0.05$) intact, and neither veracity concern measure interacts with ACTION (both null; Appendix Table B.18). When receivers make political decisions, senders do not raise their epistemic standards; they become more selective along the alignment dimension. The selectivity is thus not an epistemic-standards story (Appendix Table B.18).

I find the same amplification in senders' own words. A closing question asked what considerations drove their decisions; about half of senders gave a reason that refers to the sharing decision, and these senders are somewhat more partisan and issue-extreme than those who do not, so I treat their answers as correlational.¹¹ A large language model classifies each response into a primary motive (Appendix E.1). Among these senders, belief-correction is the most common reason (66 percent), followed by persuasion (13 percent) and expression (7 percent).

Table 3 relates these stated motives to each sender's alignment bias. Senders who cite belief-correction are 9.1 percentage points less alignment-biased ($p < 0.01$). Senders who cite an expressive reason are more biased, and equally so whether or not the receiver acts (21.8 pp, $p < 0.05$; the interaction with ACTION is 3.9 pp, not distinguishable from zero). Senders who cite persuasion are more biased only when the receiver takes a political action: their bias is not distinguishable from zero absent the action and rises by 35.5 pp when it is present ($p < 0.01$).¹² The self-reports locate the action-driven half of the selectivity in persuasion-motivated senders, mirroring the behavioral finding

¹¹Referring senders average higher partisanship and issue extremity than non-referrers ($p < 0.01$); the analysis conditions on giving a codable reason.

¹²The persuasion-by-ACTION interaction rests on the 41 senders who cite a persuasion motive in the ACTION arm; I read it as suggestive.

Table 3: Self-Reported Motivations and Alignment Bias (Experiment 1)

	Alignment Bias			
	(1)	(2)	(3)	(4)
Intercept	0.096*** (0.014)	0.078*** (0.017)	0.090*** (0.016)	0.091*** (0.016)
<i>Belief correction</i>	-0.091*** (0.016)	-0.092*** (0.017)	-0.091*** (0.017)	-0.091*** (0.017)
<i>Expression</i>	0.243*** (0.063)	0.241*** (0.063)	0.240*** (0.063)	0.218** (0.102)
<i>Persuasion</i>	0.246*** (0.063)	0.243*** (0.062)	0.047 (0.087)	0.045 (0.087)
<i>Misc.</i>	-0.084 (0.077)	-0.089 (0.078)	-0.085 (0.078)	-0.084 (0.078)
ACTION		0.036* (0.020)	0.013 (0.020)	0.011 (0.019)
<i>Persuasion</i> × ACTION			0.353*** (0.118)	0.355*** (0.119)
<i>Expression</i> × ACTION				0.039 (0.130)
Receiver Action	No	Yes	Yes	Yes
Num. obs.	1102	1102	1102	1102
R ²	0.077	0.080	0.096	0.096
Adj. R ²	0.074	0.076	0.091	0.090

Sample: $N = 1,102$ senders with a revealed directional policy preference ($AIC_i \neq 250$), Experiment 1. *Outcome:* sender-level alignment bias (mean aligned minus mean unaligned provision rate). *Regressors:* indicators for the primary self-reported motivation from an LLM classification of senders' open-ended decision reasons (Appendix E.1); reference category is *Did Not Refer to Disclosure*. *Expression* marks reasons citing the fact's fit with the sender's own position; *Persuasion* marks receiver-directed reasons. ACTION is the receiver political-action treatment. Columns (3)–(4) add the interaction of each political motive with ACTION. Heteroskedasticity-robust standard errors in parentheses. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

that ACTION widens the alignment gap.

Senders' agreement with pro-accuracy norms shapes how much they share and how selectively. After the provision task, senders rated their agreement with four statements endorsing accurate information: that false factual beliefs should be corrected, that facts matter for forming opinions, that social media should fact-check, and that disagreement stems from facts rather than values. Senders who endorse these norms more strongly provide more overall (three of the four statements, $p < 0.01$), and norm endorsement narrows the alignment gap ($\text{Align} \times \text{Agreement}$, -1.8 to -3.6 pp for those three statements, $p < 0.05$; Appendix Table B.19).

3 Conceptual Framework

I develop a conceptual framework of selective disclosure of true information. A sender who knows a true fact decides whether to disclose it to a receiver. She may value improving the receiver's accuracy, derive expressive utility from sharing a politically congenial state of the world, value moving the receiver's beliefs or preferences toward her own position, or—when the receiver may act politically—care about the political consequences of that movement. I nest these four motives in a single sender-side utility function and derive behavioral predictions. The framework characterizes the sender's private utility rather than imposing a full receiver-inference structure in the main text, departing in purpose from classical disclosure models (Grossman 1981, Milgrom 1981) that study equilibrium inference.¹³

3.1 Environment

Consider a sender, a receiver, and a state of the world. First, nature determines the true state of the world $\omega \in \{0,1\}$, the sender's political position, and whether the receiver's policy preference may translate into political action, $A \in \{0,1\}$. Without loss of generality, I normalize the true state of the world to $\omega = 1$.

Second, the sender observes the true state of the world and chooses whether to disclose it, $d = 1$, or withhold it, $d = 0$. Finally, if $A = 1$, the receiver's resulting policy-relevant position may translate into political action.

The sender evaluates disclosure using four objects. First, let b^0 denote the belief that the sender expects the receiver to hold if she withholds the state of the world, and let b^1 denote the belief that she expects the receiver to hold if she discloses it. Both are the sender's subjective forecasts of the receiver's posterior probability that $\omega = 1$, with $b^d \in [0,1]$. Let $\Delta b \in [-1,1]$ denote the factual belief correction the sender expects

¹³Appendix G embeds this utility representation in a receiver-inference environment and characterizes equilibrium behavior.

disclosure to produce:

$$\Delta b = |\omega - b^0| - |\omega - b^1|.$$

As $\omega = 1$, larger values of Δb correspond to larger expected movements of the receiver's belief toward the truth.

Second, let $\alpha \in \mathbb{R}$ denote the sender's assessment of the state of the world's congeniality to her political position. Positive values of α indicate that the sender regards the state of the world as affirming her political position. Negative values indicate that she regards the state of the world as not affirming her political position. A value of zero denotes a state of the world she regards as politically neutral. Expressive utility arises from the act of disclosing the state of the world, not from privately knowing it; the sender receives no expressive payoff from withholding.

Third, let $\Delta\pi \in \mathbb{R}$ denote the movement in the receiver's policy-relevant position toward the sender's own position that the sender expects disclosure to produce. Thus, $\Delta\pi > 0$ means that the sender expects disclosure to move the receiver toward her own position, while $\Delta\pi < 0$ means that she expects disclosure to move the receiver away from it. The distinction between α and $\Delta\pi$ is important. The former captures how the sender regards the state of the world's political congeniality to her own position; the latter captures what the sender expects disclosure to do to the receiver.

Finally, $A \in \{0, 1\}$ indicates whether the receiver's policy-relevant position may translate into political action. When $A = 0$, the sender may still value moving the receiver's beliefs or preferences toward her own position. When $A = 1$, such movement also has potential instrumental value because it may affect political outcomes.

3.2 Sender Utility

The sender's utility from disclosure combines four motives. First, a *belief-correction motive*, $\theta_B \geq 0$: the sender values correcting the receiver's factual beliefs, independent of political content. Second, an *expressive motive*, $\theta_E \geq 0$: the sender derives utility from

disclosing the state of the world when it affirms her own political position, independent of its expected effect on the receiver. Third, a *projective persuasion motive*, $\theta_P \geq 0$: the sender values moving the receiver's policy-relevant beliefs or preferences toward her own position, even absent downstream action. Fourth, an *instrumental persuasion motive*, $\theta_I \geq 0$: when the receiver's beliefs or preferences may translate into political action, the sender places additional value on moving the receiver toward her own position.

I normalize the utility of withholding to zero, $U(d = 0) \equiv 0$, so that all terms represent gains from disclosure relative to withholding. Combining these motives, the sender's net utility from disclosure is

$$U(d = 1) = \theta_B \Delta b + \theta_E \alpha + (\theta_P + \theta_I A) w(\Delta \pi), \quad (3)$$

where $w : \mathbb{R} \rightarrow \mathbb{R}$ maps the sender's perceived movement in the receiver's policy-relevant position into utility. I assume that w is sign-preserving: $w(0) = 0$, $w(\Delta \pi) > 0$ for $\Delta \pi > 0$, and $w(\Delta \pi) < 0$ for $\Delta \pi < 0$. For derivative-based predictions, I additionally assume that w is differentiable and strictly increasing on the relevant domain. Define $c(A) = \theta_P + \theta_I A$ as the composite persuasion weight on $w(\Delta \pi)$, combining the projective sub-motive (θ_P , present whether or not the receiver may act) and the instrumental sub-motive ($\theta_I A$, present only when $A = 1$).

The additive specification separates the four motives conceptually. The belief-correction term depends on the perceived factual belief correction from disclosure. The expressive term depends on whether the state of the world itself affirms the sender's political position. The projective term depends on the perceived movement of the receiver toward or away from the sender's position. The instrumental term places additional value on that same movement when the receiver's beliefs or preferences may translate into political action; it is a distinct motive operating on the same persuasion channel, not an amplification of the projective persuasion motive.¹⁴

¹⁴The additive structure follows [Bénabou and Tirole \(2006\)](#) in nesting distinct motives in a single utility representation. The model nests pure belief-correction, pure expressive disclosure, pure projective disclosure, and pure instrumental disclosure as special cases. The expressive motive draws on the social-

3.3 Disclosure Decision

The sender discloses if and only if the utility from disclosure exceeds an idiosyncratic disclosure shock:

$$d = 1 \quad \text{iff} \quad U(d = 1) > \varepsilon.$$

The shock ε is drawn from a distribution F with full support on $\mathbb{R}$ and strictly positive density $f = F'$. It captures unmodeled determinants of disclosure, such as attention, salience, idiosyncratic features of the sender-receiver interaction, or idiosyncratic inclinations toward disclosure.

The probability of disclosure is therefore

$$\Pr(d = 1) = F(\theta_B \Delta b + \theta_E \alpha + (\theta_P + \theta_I A)w(\Delta \pi)). \quad (4)$$

3.4 Predictions

The framework yields four behavioral predictions. The first is a non-political belief-correction benchmark. The next two distinguish two forms of political selectivity: expressive selectivity and projective selectivity. The fourth identifies the instrumental motive, the component of persuasion that operates only when the receiver may act politically. All proofs are in Appendix G.

Proposition 0. *If $\theta_B > 0$, disclosure is strictly increasing in the perceived factual belief correction from disclosure:*

$$\frac{\partial \Pr(d = 1)}{\partial \Delta b} > 0.$$

Proposition 0 is the belief-correction benchmark. Holding political content and expected persuasion fixed, senders with a belief-correction motive are more likely to disclose the state of the world when they perceive disclosure to correct the receiver's image account of communication, in which senders gain utility from transmitting position-affirming content (Braghieri 2024).

factual beliefs more.

Proposition 1a (Expressive selective disclosure). *If $\theta_E > 0$, disclosure is strictly increasing in congenial alignment:*

$$\frac{\partial \Pr(d = 1)}{\partial \alpha} > 0.$$

Proposition 1a captures expressive disclosure. Holding factual belief correction and expected receiver movement fixed, senders with an expressive motive are more likely to disclose the state of the world when it affirms their own political position.

Proposition 1b (Persuasive selective disclosure). *If $c(A) > 0$, disclosure is strictly increasing in $w(\Delta\pi)$:*

$$\frac{\partial \Pr(d = 1)}{\partial w(\Delta\pi)} > 0.$$

If, in addition, w is differentiable and strictly increasing, disclosure is strictly increasing in $\Delta\pi$:

$$\frac{\partial \Pr(d = 1)}{\partial \Delta\pi} > 0.$$

Proposition 1b captures projective selective disclosure. Holding factual belief correction and expressive alignment fixed, senders are more likely to disclose the state of the world when they perceive it as moving the receiver toward their own position, and less likely to disclose it when they perceive it as moving the receiver away.

When $\text{sign}(\alpha) = \text{sign}(\Delta\pi)$ —the typical case for a politically aligned fact—Propositions 1a and 1b predict the same direction. When the signs differ, they diverge: the expressive motive favors disclosure of facts the sender finds politically congenial ($\alpha > 0$) regardless of persuasive direction, while the projective persuasion motive favors disclosure only when the fact is expected to move the receiver toward the sender's position ($\Delta\pi > 0$).

Proposition 2 (Political action amplification). *If $\theta_I > 0$, political action strictly increases the marginal utility of persuasion:*

$$\frac{\partial^2 U(d=1)}{\partial w(\Delta\pi) \partial A} = \theta_I > 0.$$

Proposition 2 is stated at the utility-index level. This is the object directly disciplined by the utility function. Without additional restrictions on the distribution F , the model does not require a uniformly positive interaction effect in disclosure probabilities for all parameter values. The behavioral prediction is that the latent value of moving the receiver toward the sender’s position is larger when the receiver’s beliefs or preferences may translate into political action; this increment in the value of persuasion is the instrumental motive, a sub-motive distinct from the always-present projective persuasion motive.

The framework also yields a magnitude prediction: when w is strictly increasing, stronger expected persuasion raises disclosure for facts that move the receiver toward the sender’s position and lowers it for facts that move the receiver away. I develop this prediction formally in Appendix G, where the boundary case $w(\Delta\pi) = \text{sign}(\Delta\pi)$ marks the threshold at which the sender responds only to the direction of persuasion and not to its magnitude.

4 Experiment 2

Experiment 1 establishes three findings: provision is high, there is a baseline alignment gap consistent with an expressive motive (θ_E) or a projective persuasion motive (θ_P). Additionally, political action approximately doubles the gap consistent with an instrumental persuasion motive (θ_I) or expression increasing under action. Experiment 1’s design cannot decompose these and therefore, three questions remain open: (i) which component of the baseline gap is expressive and which is projective persuasion, (ii) is the increase in selectivity explained by expression or instrumental persuasion, and (iii) what the H3 null implies about $w(\Delta\pi)$. We adjust the experimental design in Experi-

ment 2 to answer these questions.

4.1 Design, Treatment, and Hypotheses

I run a pre-registered online experiment with $N = 1,548$ German adults recruited on Bilendi using quota-sampling on age, gender, and state (Bundesland).¹⁵ I study two policy domains assigned between subjects: immigration policy and climate policy, extending the topic space beyond Experiment 1.

Experimental Timeline The experiment proceeds in three parts, followed by a post-questionnaire (Figure 3). Part 1 elicits senders' demographics, stated policy preferences and revealed policy preferences. Part 2 is the information provision task: senders decide whether to share each of five assigned facts, one per decision, directed at receivers whose policy stance is visible. An attention check follows Part 2. Part 3 elicits incentivized beliefs about how each fact corrects receivers' factual beliefs and shifts their policy views. In the post-questionnaire subjects rate two facts from the other domain in terms of their congeniality and trustworthiness. Two timing design choices differ from Experiment 1. First, provision decisions precede belief elicitation: placing provision before elicitation ensures that elicitation does not influence provision behavior, and the Part 3 elicitation provides a timing-independent check on the behavioral results. Second, comprehension checks gate the Part 2 bonus: senders must answer all three questions correctly on the first attempt; the questions loop until all answers are correct.

Facts I select six verified-true facts per topic, balanced for credibility and relevance in a pre-registered pre-validation survey ($N = 300$, AsPredicted #290989), see Appendix A.3. Each fact is presented as a binary, invertible comparison with a documented correct answer as in Experiment 1, sourced from German statistical offices and

¹⁵Two cumulative exports of the final panel batch overlapped; eight completers appearing in both exports were deduplicated. The speeder cutoff (one quarter of the median completion time), recomputed on the deduplicated sample, excluded one additional participant.




Figure 3: Experimental timeline for Experiment 2.

research institutes. Table 4 lists the twelve statements. All participants are informed at the outset that every fact is verified true and sourced, and receivers know the same.

Policy Preferences I measure senders' policy preferences in two ways. First, senders indicate their policy position on a five-point Likert scale. Second, senders split €250 between the pro-policy and anti-policy advocacy organization for their topic in a real donation decision: one decision per session is implemented with actual money. For immigration, the organizations are ProAsyl (advocates open migration policy) and Republik21 (advocates more restrictive policy); for climate, BUND (supports ambitious climate targets) and Stiftung Marktwirtschaft (emphasizes economic costs of the transition). Senders are shown each organizations position in their own words before deciding.

Receiver Policy Stance Each sender faces five receivers sequentially, one per position on the five-point Likert scale policy stance. Before each provision decision, the sender observes the receiver's stance as a full sentence and knows the receiver in turn is told

Table 4: Facts Used in Experiments 2 and 3

Direction	Fact
<i>Immigration policy</i>	
Pro	The share of people with a migration background in shortage-occupation fields is <i>higher</i> than their population share.
Pro	The absolute increase in employed foreign nationals between 2013 and 2023 was <i>larger</i> than for German nationals.
Pro	<i>More</i> than half of the persons attending a general integration course for the first time complete it without interruption through the final module.
Anti	Among persons suspected of violent crimes, the share of non-German nationals is <i>higher</i> than their population share.
Anti	The share of persons receiving basic income support (<i>Bürgergeld</i>) is <i>higher</i> among foreign nationals than among German nationals.
Anti	Germany records <i>more</i> persons without a valid residence permit per capita than the EU average.
<i>Climate policy</i>	
Pro	In German electricity generation in 2023 and 2024, the share of renewable energies was <i>higher</i> than the share of fossil fuels.
Pro	The cost of generating one kilowatt-hour of electricity from new, ground-mounted solar installations is <i>lower</i> than from new gas-fired power plants.
Pro	The percentage increase in renewable energy jobs in Germany between 2022 and 2023 was <i>higher</i> than the global average.
Anti	Grid fees for household customers are expected to be <i>more</i> than double their current level by 2045 due to grid expansion.
Anti	There are <i>fewer</i> companies that view the effects of the energy transition positively than companies that view them negatively.
Anti	At higher shares of solar power, <i>more</i> price fluctuations are observed in European electricity markets.

Notes: Verbatim claims as shown to senders in English translation (correct answer in *italics*). Direction (Pro) = supports more-permissive immigration or more-ambitious climate policy; (Anti) = supports more-restrictive or more-cautious policy. Pool of six facts per topic, all verified true from German statistical offices and research institutes. Each sender draws five facts: two pro-policy, two anti-policy, and one random tiebreaker. Topics assigned between subjects; all senders see one topic. Climate policy extends the topic space beyond Experiment 1. Pre-validation details in Appendix A.3.

their policy position. This mutually visible stance structure generates within-sender variation in sender-receiver ideological distance. Senders know that receivers are told a fact may or may not have been shared with them — with 50% probability the match occurs and the receiver receives the fact (if sent) — so the receiver cannot attribute a missing fact to the sender’s deliberate withholding; silence remains only weakly informative.

Beliefs about Belief Correction For each of the five receiver-fact pairs, senders predict the share of receivers with the same stance who will identify the fact correctly after disclosure versus without disclosure. Predictions are scored against matched receiver outcomes via a quadratic scoring rule.

Beliefs about Persuasiveness Senders also predict how each fact shifts the assigned receiver's policy position on the five-point scale. They report the Likert position the receiver will adopt after seeing the fact and the position the receiver would take without seeing it; both predictions are scored via a quadratic scoring rule against the receiver's realized outcome.

Beliefs about Congeniality I measure expressive congeniality—how much a fact validates the sender's own position—without researcher-imposed directional classification, using a cross-domain rating task. After the provision task, senders rate two facts from the topic they were *not* assigned to—one hypothesized pro-policy and one hypothesized anti-policy—answering: “*Does this information show that your opinion is correct?*” The instructions specify that senders should not assess whether their abstract opinion is correct, nor whether the fact would convince others. The cross-domain design limits contamination by the provision decisions. Senders also rate whether each cross-domain fact is correct.

Treatments and Hypotheses Senders are assigned between subjects to CONTROL or ACTION. In both arms, receivers re-state their policy opinion after potentially receiving a fact.¹⁶ In the ACTION arm, receivers additionally make a real donation decision (€250 split between the same two organizations as Part 1): one decision per session is implemented with actual money, so receivers' choices carry real financial consequences for these organizations. Additionally, I again exploit the within-sender variation in alignment to the facts.

¹⁶The ACTION treatment may also shift task salience or decision construal beyond political consequentiality.

I pre-registered three hypotheses with respect to the treatments, mirroring Experiment 1:

Hypothesis 4 (Baseline provision). *Provision rates for unaligned facts exceed the indifference benchmark of 50% in each treatment-topic cell.*

Hypothesis 5 (Selective provision). *Senders are more likely to provide a fact aligned with their own political position than one that is not, holding fact and topic fixed.*

Hypothesis 6 (Action amplification). *The alignment gap in provision is larger when the receiver makes a real political donation (ACTION) than when the receiver only restates their opinion (CONTROL).*

4.2 Sample, Incentives, and Protocol

Data collection ran from May 27 to June 4, 2026 on Bilendi with quota-sampling for age, gender, and state (Bundesland) to approximate the German adult population. (See Appendix Table C.30 for details.) This study was approved by the Ethics Council of the Max Planck Society under the Generalized Approval of Experiments Following the Protocol that is Standard in Experimental Economics and was pre-registered on AsPredicted (No. 293252, an update of No. 292898). Senders took a median completion time of 17 minutes and earned €3 on average.

4.3 Policy Preferences and Beliefs

Policy preferences. Figure 4 reports stated and revealed preference distributions for each topic. For immigration, 66.1% of senders state policy preference in favor of more restricted immigration and 26.3% donate exclusively to the organization advocating more restrictions. The correlation between the two immigration measures is strong ($r = 0.63$, $p < 0.001$). For climate, a majority states preferences for increased national climate action (52.8%), while the plurality of donations leans toward the organization advocating fewer and more market-based restrictions. The within-person ranking is nonetheless strongly correlated ($r = 0.50$, $p < 0.001$). As pre-registered,

I combine the two measures, each normalized to $[-1, +1]$ with $+1$ the most pro-policy position: revealed preference $\text{Pref}_i^{\text{rev}} = \frac{\text{Donation}_i^{\text{pro}} - 125}{125}$ from the €250 split, stated preference $\text{Pref}_i^{\text{st}} = \frac{3 - \text{Stance}_i}{2}$ from the five-point Likert stance, and their average $\text{Pref}_i = (\text{Pref}_i^{\text{rev}} + \text{Pref}_i^{\text{st}})/2$. Preferences are distributed across the full scale in both samples within each topic allowing for identification of selectivity.




Figure 4: Distribution of policy preferences among Experiment 2 senders ($N = 1,548$). *Top row (Panels A–B):* stated preferences from the 5-point Likert item; tick labels complete the stance statement shown in the panel title (e.g., *Immigration to Germany should be ... greatly eased*). *Bottom row (Panels C–D):* revealed preferences from the real donation task, grouped into five buckets from All pro (€250 to the pro-policy organisation) to All anti (€0), displayed pro-to-anti left to right. *Left column:* immigration senders ($N = 809$). *Right column:* climate senders ($N = 739$).

Senders' fact congeniality rating. The cross-domain rating task asks senders whether each other-topic fact “shows that your view is correct.” I average the yes-responses within cells of sender stated stances and individual fact direction and transform $\hat{\alpha}_{ij} = \frac{\text{share_yes}_{ij} - 0.5}{0.5} \in [-1, +1]$, where $+1$ denotes universal validation and -1 no validation

by senders with the stance. Panel (a) of Figure 5 plots the resulting pattern. The empirical range across the ten stance $\times$ direction cells is $[-0.04, +0.78]$; only two of the ten cells are negative, and both are negligible (-0.01 and -0.04).¹⁷ Senders almost never regard a fact as purely invalidating their position. Neutral senders (stance 3) also record large variation in perceived congeniality for the different facts, $\hat{\alpha}_{ij} = +0.15$ on anti-direction and $+0.43$ on pro-direction facts on average. Notably, positional alignment assigns zero to all facts for senders with $\text{Pref}_i \approx 0$, erasing their within-sender variation. The measure correlates only moderately with positional alignment ($r = 0.184$): both carry the same directional signal, but congeniality adds gradation that alignment cannot express. Senders also rated factual credibility; trust ratings range 52%–85% and follow the same directional gradient ($\beta = +0.119$, $p < 0.001$). Since all twelve facts are verified true, this alignment-predicted variation in truthfulness judgments is consistent with motivated belief formation.

Senders’ beliefs about factual belief correction. Without disclosure, senders expect receivers to answer correctly roughly half the time: the baseline prior that the receiver answers correctly averages 49.4% across fact–receiver pairs. Disclosure is expected to raise this to 66.5%, a 17-percentage-point improvement. Panel (b) of Figure 5 shows the distribution of the expected factual belief correction gain $\widehat{\Delta b}_{ijk}$ for aligned and unaligned facts, where k indexes the receiver assigned to sender i ’s fact j ; each sender–fact pair faces exactly one receiver, so k is implied by (i, j) and estimation remains at the sender–fact level. The mean expected factual belief correction is positive for all twelve facts (range: 0.139–0.199). Climate facts yield higher expected gains (0.188 vs. 0.155 for immigration). Senders expect partial updating: only 11.9% of fact–receiver pairs receive a post-disclosure prediction of 100%, and 23.7% remain below 50% even after disclosure.

¹⁷The two negative cells belong exclusively to lean senders (stances 2 and 4). Extreme senders (stances 1 and 5) rate counter-directional facts positively: $+0.32$ (stance 1 $\times$ anti) and $+0.27$ (stance 5 $\times$ pro). For a stance-5 sender the within-sender contrast is $+0.78$ congenial versus $+0.27$ counter-directional; this 0.51-unit spread drives identification. Within individual topics lean-sender negative values are more pronounced (stance 2 reaches -0.31 for immigration anti-facts; stance 4 reaches -0.22 for climate pro-facts), confirming the pattern is not a pooling artifact (Appendix Table C.35).




Figure 5: Elicited returns to disclosure, Experiment 2 (Citizens). Panel (a): perceived congeniality $\hat{\alpha}_{ij}$ (Expression) of each fact by the sender’s five-point stance on the rated topic; one marker per fact $\times$ stance cell with at least 10 ratings (6 anti-direction facts, red circles; 6 pro-direction facts, green squares), each a cell mean from the cross-domain rating task (2,308 ratings by 1,147 respondents; all completed rating-task responses across recruitment batches, the sample defining the cell-level Expression measure); lines show the observation-weighted mean per fact direction at each stance. Panels (b) and (c): distributions of the two incentivized elicitations for positionally aligned facts ($\text{Align}_{ij} > 0$; bars) and unaligned facts ($\text{Align}_{ij} < 0$; black outline), as shares within each group; observations of senders with $\text{Pref}_i = 0$ (260 of 7,740) are omitted from the split. Panel (b): expected factual belief correction $\hat{\Delta}b_{ijk}$ (0.05-unit bins; group means 0.173 aligned, 0.175 unaligned). Panel (c): expected persuasion $\hat{\Delta}\pi_{ijk}$ (0.125-unit bins; the bin at $[-0.125, 0)$ contains exact-zero predictions; group means 0.046 aligned, 0.047 unaligned). Shaded regions mark negative values (disclosure expected to move the receiver’s factual belief away from the truth, respectively the receiver away from the sender’s position). Elicitations incentivized via Quadratic Scoring Rule.

About 26% of fact–receiver pairs carry negative expected factual belief correction, i.e., the sender expects disclosure to leave the receiver slightly further from the truth. This rate is flat across receiver stances (23.4%–26.9%), receiver–fact alignment (counter-attitudinal vs. same-side), and sender–receiver distance (all regression coefficients $p > 0.35$). The rate is not politically patterned. The mean expected factual belief correction is similarly unrelated to receiver position and flat in sender–receiver distance (pooled slope $\beta = -0.003$, $p = 0.38$, fact fixed effects, sender-clustered standard errors; Appendix Figure C.6, Panel (a)). Therefore, the factual belief correction return does not vary with the receiver’s position relative to the sender.

The incentivized elicitation shows no evidence of politically motivated over-rating: the aligned and unaligned distributions in Figure 5, Panels (b) and (c), lie nearly on top of one another. Senders assign statistically indistinguishable expected factual belief correction gains to aligned and unaligned facts (0.173 vs. 0.175, $p = 0.62$, fact FE), ruling out the concern that senders shade their predictions toward politically congenial facts. The same null holds for the Persuasion elicitation.¹⁸

Senders’ beliefs about the persuasive effect of disclosure. The persuasion return $\widehat{\Delta\pi}_{ijk} = |\text{Stance}_i - \hat{\pi}_{ijk}^0| - |\text{Stance}_i - \hat{\pi}_{ijk}^1|$ captures expected movement of receiver k ’s stance toward sender i ’s own position from (non-)disclosure of fact j ; $\hat{\pi}_{ijk}^0$ and $\hat{\pi}_{ijk}^1$ are the sender’s incentivized expectations of receiver k ’s stance without and with disclosure. Panel (c) of Figure 5 shows its distribution for positionally aligned and unaligned facts. I find that in 41% of pairs, senders predict the fact will shift the receiver’s stance counter to the fact’s political direction—a distinct question from whether the receiver moves toward the sender’s own position. Computing the rate relative to sender congeniality rather than fact direction leaves it unchanged (40.8%), so it does not reflect a mechanical confound.¹⁹ Receiver extremity does not predict this rate (range: 36.9%–

¹⁸Using Expression ($\alpha_{ij} > 0$) as the conditioning variable yields $\Delta\pi_{ijk} = 0.055$ (congenial) versus 0.022 (uncongenial), with $\beta = +0.033$, $p = 0.25$, fact FE—directional but not statistically significant.

¹⁹Directional backfire does not predict provision choices ($\beta = +0.002$, $p = 0.84$). Facts expected to move the receiver *away from the sender’s own position* ($\widehat{\Delta\pi}_{ijk} < 0$) do predict lower provision ($\beta = -0.044$, $p < 0.001$), but this is the Persuasion motive’s negative tail operating as a selection criterion, not a separate directional effect.

43.9% across the five receiver stances); sender–receiver distance is not statistically distinguishable from zero ($\beta = +0.002$, $p = 0.62$).

Senders' persuasion expectations differ widely across fact–receiver pairs: the standard deviation of $\widehat{\Delta\pi}_{ijk}$ is roughly 0.9 in both alignment groups (0.888 aligned, 0.906 unaligned) against group means of 0.046 and 0.047; about 17% of pairs are exact zeros. In 59% of fact–receiver pairs, senders expect disclosure to bring the receiver's stance weakly closer to their own relative to non-disclosure ($\widehat{\Delta\pi}_{ijk} \geq 0$; 41.7% strictly closer, 17.3% unchanged). Receiver position enters mechanically: a receiver already at the sender's position leaves essentially no gap to close, and the average return at distance zero is essentially zero (-0.067). From distance 2 onward, senders expect disclosure to close a roughly constant share of the sender–receiver gap, about six percent (5.6%, 5.8%, and 5.9% at distances 2, 3, and 4; equality Wald $p = 0.98$, sender-clustered), so the level of the return rises with distance (Appendix Figure C.6, Panel (b)). The aligned and unaligned mean profiles are statistically indistinguishable (distance $\times$ alignment interaction $p = 0.12$). Unlike the factual belief correction return, the persuasion return is positive only where a gap exists to close. What characterizes senders' persuasion beliefs is thus not their level, which is near zero, but their dispersion.

Variation in the elicited returns. The systematic structure just described accounts for little of the variation in the elicited returns: fact identity, receiver position, and sender–receiver distance together explain at most 1.6% of the variance of either return. About half of the factual belief correction variance (53%) and a third of the persuasion–return variance (30%) lies between senders, who differ persistently in how much they expect facts to correct and to persuade. The remainder (roughly half of the factual belief correction return, two-thirds of the persuasion return) is idiosyncratic to the fact–receiver pair within sender. This pair-level component is the within-sender variation on which the estimation relies. Observable sender characteristics—political extremity, left–right position, gender, age, education, party affiliation, and topic—capture little of the persistent sender component: regressing the sender mean returns on the full set yields an R^2 of 5.1% for the factual belief correction return and 3.9% for the persuasion

return ($N = 1,543$ senders).²⁰

Uncorrelatedness of motive proxies. Expression and Persuasion are analytically distinct: Expression captures whether a fact validates the sender’s view, a property of the fact relative to the sender alone, while Persuasion captures whether a specific receiver will move toward the sender’s position—a property of the fact–receiver pair. They diverge whenever a congenial fact faces a receiver who already agrees. A stance-5 sender rates anti-policy facts highly congenial ($\hat{a}_{ij} = +0.78$), but a receiver also at stance 5 is already at the sender’s position; the sender expects near-zero persuasive gain ($\widehat{\Delta\pi}_{ijk} \approx 0$). Positional alignment carries no persuasion-expectation signal either. Conditional on a fact being positionally aligned rather than unaligned, the share of pairs in which the sender expects movement toward her own position is essentially identical (44.2% versus 43.2%), and the mean difference in $\widehat{\Delta\pi}_{ijk}$ is -0.001 . The factual belief correction proxy is uncorrelated with both. I find all pairwise correlations below $|r| = 0.04$: Expression and Factual Belief Correction ($r = -0.005$, $p = 0.69$), Expression and Persuasion ($r = +0.016$, $p = 0.15$), Factual Belief Correction and Persuasion ($r = +0.033$, $p = 0.004$); $VIF \approx 1.001$. The design generates this separation across the three predictors. The ACTION randomization does not sort senders differentially on any of the three proxies or on policy preferences: all four balance tests yield $p > 0.20$.

4.4 Treatment Effects

Figure 6 shows provision rates by arm and alignment. I find that provision rates for unaligned facts exceed the 50% indifference benchmark in all four treatment–topic cells (range: 60.9%–63.8%, one-sample t -tests, all $p < 0.001$), confirming Hypothesis 4.

Result 4 (Baseline provision). *Senders provide unaligned facts at rates significantly above the 50% indifference benchmark in all four treatment–topic cells ($p < 0.001$).*

²⁰A descriptive battery interacting each characteristic with the congeniality rating and with sender–receiver distance yields one interaction of twenty-four below a Benjamini–Hochberg q of 0.10: senders with more extreme issue positions expect larger persuasion returns overall ($\beta = +0.114$, $p < 0.001$), with a flatter distance gradient ($\beta = -0.035$, $p = 0.004$).




Figure 6: Mean provision rates by arm and fact alignment, Experiment 2 (Citizens; $N = 1,548$ quota-sampled German adults, Bilendi, 2026). Panel (A): mean sender-level provision rates for aligned (dark green) and unaligned (dark red) facts, by experimental arm. Panel (B): alignment bias (aligned minus unaligned provision rate) by arm. Sample: senders with a directional preference (revealed donation $\neq$ €125). Alignment: $\text{Pref}_i \times \text{Dir}_j \in [-1, 1]$. Error bars are 95% confidence intervals (sender-level means).

Provision rates are between 61.1% and 65.7%, much lower than in Experiment 1, where unaligned rates ranged from 81.9% to 86.5%. First, sharing is implemented with 50% probability under the Dye-type structure, halving the expected persuasive return per sharing act relative to Experiment 1’s deterministic implementation. Second, receivers know the sender’s political position; in Experiment 1 the receiver was unaware a disclosure decision was occurring at all. Third, senders have no information about receivers’ prior beliefs, removing the ability to target disclosure toward those most in need of correction.

Table 5 reports the pre-registered tests. I estimate

$$y_{i,j} = \beta_0 + \beta_1 \text{Align}_{i,j} + \beta_2 \text{ACTION}_i + \beta_3 (\text{Align}_{i,j} \times \text{ACTION}_i) + \mathbf{S}_j + \gamma_3 \mathbf{X}_i + \varepsilon_{i,j}, \quad (5)$$

with $\text{ACTION}_i \in \{0, 1\}$ an indicator equal to one if sender i is assigned to the ACTION arm. The unit of observation is a sender–fact pair; standard errors are clustered at the sender level throughout. β_1 identifies baseline alignment-based selectivity (Hypothesis 5) and β_3 whether political action amplifies it (Hypothesis 6). Column (1) includes no controls, column (2) adds fact fixed effects $\mathbf{S}_j$, and column (3)—the pre-registered

preferred specification—adds demographic and political controls X_i .²¹ The confirmatory tests use positional alignment Align_{ij} (not the motive proxies), with H1 ($\beta_1 > 0$, one-sided) and H2 ($\beta_3 > 0$, one-sided) as the pre-registered directional tests.

Senders share aligned facts at higher rates ($p = 0.015$, one-sided, column (3)), confirming Hypothesis 5.

Result 5 (Selective provision). *Senders provide aligned facts at higher rates than unaligned facts ($p = 0.015$, one-sided, column (3)).*

The pre-registered action amplification of the alignment gap is not statistically distinguishable from zero ($\beta = -0.010$, one-sided $p = 0.73$, column (3)); I do not confirm Hypothesis 6. This null is informative about the proxy rather than the treatment: positional alignment is uncorrelated with the Persuasion proxy within facts ($r = 0.000$), so the alignment interaction can inherit almost none of an action-conditional persuasion effect. Under the estimated instrumental coefficient, the implied $\text{Align} \times \text{ACTION}$ coefficient is 0.1% of the test's minimum detectable effect (0.046), and even perfect correlation between the measures would leave it below that threshold ($0.024 \times 1.53 = 0.037 < 0.046$). As pre-registered, I examine the motive channels directly in the following subsection.

4.5 Separating the Motive Channels

Unlike Experiment 1, the model's three belief arguments—expected factual belief correction, expected expressive congeniality, and expected receiver shift toward the sender's position—are directly elicited rather than inferred from revealed preference. I estimate a linear probability model at the sender-fact level, with standard errors clustered at the

²¹The pre-registration specifies a perceived believability control jointly with fact fixed effects in column (2). This control was not collected: the cross-domain rating module was shortened after pre-registration to reduce survey length and dropout, a deviation documented in the post-pre-registration decision log. Column (2) therefore reports fact fixed effects without a believability control.

Table 5: Alignment-Based Selective Provision — Experiment 2 (Citizens)

	Probability of Fact Provision		
	(1)	(2)	(3)
Align	0.024** (0.011)	0.024** (0.012)	0.026** (0.012)
Action	0.025** (0.013)	0.026** (0.013)	0.023* (0.012)
Align $\times$ Action	-0.008 (0.016)	-0.007 (0.016)	-0.010 (0.017)
Clustered SE	Sender	Sender	Sender
Fact FE	No	Yes	Yes
Controls	No	No	Yes
Num. obs.	7740	7740	7715
R ²	0.001	0.002	0.016
Adj. R ²	0.001	0.001	0.009

Sample: $N = 1,548$ Experiment 2 senders (Bilendi panel, Germany, 2026), 7,740 sender-fact observations. Each sender makes five provision decisions, one per fact. Column (3) restricts to non-missing demographic controls (7,715 observations). *Outcome:* binary provision indicator ($y_{ij} \in \{0,1\}$). *Alignment:* $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1,1]$ (pre-registered positional alignment measure). *Fact FE:* indicator for each of the 12 facts. *Controls:* age, gender, education, Bundesland, voting intention, left-right self-placement. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$. *Pre-registered tests on column (3):* H1 (Alignment > 0 , one-sided $p = 0.015$); H2 (Alignment $\times$ Action > 0 , coefficient negative, one-sided $p = 0.73$, not confirmed).

sender level throughout:

$$y_{ij} = \beta_0 + \beta_E \alpha_{ij} + \beta_P \Delta \pi_{ijk} + \beta_2 \text{ACTION}_i + \beta_A (\Delta \pi_{ijk} \times \text{ACTION}_i) + \beta_B \Delta b_{ijk} + \mathbf{S}_j + \varepsilon_{ij}, \quad (6)$$

where $y_{ij} \in \{0,1\}$ is sender i 's provision decision for fact j and $\mathbf{S}_j$ are fact fixed effects. The fact fixed effects absorb all time-invariant fact-level characteristics, including political direction, topic-level variation, and fact-specific shareability. Within-sender variation across the five facts assigned to each sender provides identification.²² Columns (1) and (2) isolate the expressive and persuasion channels in turn; column (3) includes both jointly. Column (4) is the specification implied by the utility function in the conceptual framework: Expression main effect plus the instrumental Persuasion $\times$ ACTION inter-

²²Because identification rests on within-sender variation across facts rather than across experimental conditions, experimenter-demand effects are limited (de Quidt et al. 2018).

Table 6: Action Decomposition: Expression and Persuasion — Experiment 2 (Citizens)

	Probability of Fact Provision			
	(1)	(2)	(3)	(4)
Expression	0.083*** (0.031)		0.083*** (0.031)	0.069*** (0.023)
Expression $\times$ Action	-0.027 (0.042)		-0.027 (0.042)	
Persuasion		0.003 (0.009)	0.003 (0.009)	0.003 (0.009)
Persuasion $\times$ Action		0.024* (0.013)	0.024* (0.013)	0.024* (0.013)
Action	0.035* (0.020)	0.025* (0.013)	0.034* (0.020)	0.024* (0.013)
Factual belief correction gain	0.038** (0.017)	0.037** (0.017)	0.037** (0.017)	0.037** (0.017)
Clustered SE	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes
Spec	(1)	(2)	(3)	(4)
Num. obs.	7740	7740	7740	7740
R ²	0.004	0.004	0.005	0.005
Adj. R ²	0.002	0.002	0.003	0.003

Sample: $N = 1,548$ Experiment 2 senders (Bilendi panel, Germany, 2026), 7,740 sender-fact observations; requires non-missing expressive identity, projective persuasion, and factual belief correction gain measures. *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating a_{ij} , elicited post-provision on a different topic. *Persuasion:* incentivized expected preference shift $\Delta\pi_{ijk}$ toward the sender's position (positional, preference-free reading). *Factual belief correction gain:* incentivized expected correction of the receiver's factual belief Δb_{ijk} . *Receiver distance (dist):* reported separately in Table C.20 (pre-registered heterogeneity). Column (1): expressive motive and its action interaction. Column (2): the persuasion channel—projective main effect and instrumental action interaction. Column (3): both interactions jointly. Column (4): expressive main effect plus the instrumental (persuasion $\times$ action) interaction. All specifications include fact fixed effects and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

action, omitting Expression $\times$ ACTION whose null is established in column (1).

All three motive proxies predict provision and survive fact fixed effects, sender-clustered standard errors, and logit estimation (Table 6). Presentation order changes nothing either: fact order is randomized, and although provision declines mildly over the five positions (-0.8 pp per position, $p = 0.037$), adding presentation-position fixed effects leaves every coefficient in the table unchanged to the third decimal.²³ Positional alignment (Align_{ij}), the pre-registered proxy, is distinct from all three: it correlates

²³Appendix Table C.31 re-estimates the confirmatory and motive specifications under alternative sample definitions: relaxing the pre-registered speeder exclusion, adding the 106 participants excluded for instrument comparability (they started before the pre-registration update of 27 May 2026, which shortened the survey after its length caused elevated dropouts), and — the pre-specified robustness check — excluding attention-check failers, who are retained in the primary analysis as pre-registered. Every estimate keeps its sign in every variant. Excluding the attention-check failers, who rush the study and report flatter belief elicitation but share at nearly the same rate, moves the Persuasion $\times$ ACTION point estimate by less than 0.001 and leaves all three channel estimates significant.

with Expression at $r = 0.184$ but is uncorrelated with Persuasion ($r = 0.000$); it therefore proxies Expression alone. This uncorrelatedness is the quantitative basis of the Hypothesis 2 null documented in the previous subsection.

Expression predicts provision at +13.8 pp across its full range (column (4)) and displaces positional alignment in a joint specification (Appendix Table C.25): Expression holds ($\beta = +0.061$, $p = 0.013$) while Alignment loses significance ($\beta = +0.012$, $p = 0.195$). The expressive motive responds to a fact's political direction alone.

Result 6 (Expression predicts baseline provision). *Senders are more likely to share facts with higher expressive congeniality ($\beta = +0.069$, column (4)), adding +13.8 pp across the full range of Expression.*

The Expression $\times$ ACTION interaction is null ($\beta = -0.027$, $p = 0.52$, column (3)), consistent with an action-invariant expressive motive.

Persuasion ($\Delta\pi_{ijk}$) does not predict provision in the CONTROL arm ($\beta = +0.003$, $p = 0.76$; column (4)). It activates under political action: the ACTION interaction is marginally positive ($\beta = +0.024$, $p = 0.064$; exact wild-cluster bootstrap $p = 0.047$; column (4)), and a randomization test that reassigns the ACTION arm across senders (5,000 permutations) yields one-sided $p = 0.030$ in the pre-registered direction.²⁴ Appendix Figure C.5 shows the pattern by arm.

Result 7 (Persuasion is instrumental: it activates under political action). *Facts expected to move the receiver toward the sender's own position are shared at higher rates when the receiver takes a downstream political action ($p = 0.064$; exact wild-cluster bootstrap $p = 0.047$; randomization test, one-sided $p = 0.030$; column (4)). Little evidence of selection on Persuasion in the CONTROL arm ($p = 0.76$).*

The factual belief correction gain Δb_{ijk} predicts provision ($\beta = +0.037$ on $[0, 1]$, +3.7 pp across its full range, $p = 0.033$), with a positive sign across all 12 facts.²⁵ Factual belief correction plays a positive but secondary role in provision decisions.

²⁴The treatment shifts how senders weight beliefs, not the beliefs themselves: the distribution of $\widehat{\Delta\pi}_{ijk}$ is statistically indistinguishable between arms.

²⁵The pooled factual belief correction coefficient ($\beta = +0.037$, $p = 0.033$) is directionally consistent across all twelve facts. The immigration subsample reaches conventional significance ($\beta = +0.052$, $p = 0.042$); the climate coefficient is directionally consistent but loses significance at reduced power

Result 8 (Belief correction predicts provision). *Senders share facts with higher expected factual belief correction at greater rates ($\beta = +0.037$, $p = 0.033$, column (4)); the sign is positive for all 12 facts.*

Expressive selection operates in both arms; persuasion-based selection appears only when the receiver’s beliefs carry consequences.

A direction-only analog $\delta_{ijk}^{\text{dir}} \equiv -\Delta\text{stance}_{ijk} \cdot \text{sign}(\text{pref}_i)$ scales expected receiver movement by the sign rather than the magnitude of the sender’s preference. Appendix C.3 compares the two proxies and their binary codings (Appendix Tables C.33 and C.34): baseline provision responds to the direction of expected persuasion, whereas the instrumental interaction survives only under codings that carry magnitude.

4.6 Heterogeneity

I examine three heterogeneity dimensions: sender-receiver ideological distance, sender political extremity (both pre-registered exploratory), and political topic. Appendix Table C.20 covers the treatment-effects contrast; Appendix Tables C.21 and C.22 cover the motive channels. The analysis estimates twenty interaction terms across these specifications; Benjamini–Hochberg-adjusted q -values within this family are reported in the table notes. Nine of the twenty terms fall below $p = 0.10$, against two expected under the null; no individual term survives false-discovery control at conventional levels (all $q \approx 0.20$), so I read the patterns below as suggestive. Jointly, the ACTION interactions of the audience margins and the persuasion channel (six terms) reject at $p = 0.059$ (randomization $p = 0.061$); this joint test uses the same variation as the main estimate and is corroborative rather than independent evidence.

Treatment effects. *Distance.* Sender–receiver ideological distance is $\text{dist}_{ij} \equiv |\text{Stance}_i - \text{Stance}_{ij}^R|$, the absolute gap on the five-point Likert scale between sender i ’s self-placement

($\beta = +0.023$, $p = 0.338$). Within immigration, five of six facts maintain a positive sign. The sixth is the leveraged observation; without it, $\beta = +0.025$ ($p = 0.160$), though the direction is preserved. Within-arm estimates are directionally consistent throughout ($\beta = +0.056$ in CONTROL, $\beta = +0.020$ in ACTION), and the formal factual belief correction $\times$ ACTION interaction does not reach significance ($\beta = -0.034$, $p = 0.32$).

and the revealed position $\text{Stance}_{ij}^R \in \{1, \dots, 5\}$ of the receiver paired with that decision. Provision falls -9.5 pp per Likert unit of distance ($p < 0.001$; column (1) of Appendix Table C.20). The ACTION frame partially offsets this penalty ($\text{dist} \times \text{ACTION}$, $\beta = +0.019$, $p = 0.049$). The alignment slope is invariant to distance.²⁶

Political extremity. More polarised senders provide more overall ($\beta = +0.010$, $p < 0.10$, column (2)), but neither the alignment slope nor the ACTION treatment effect varies with continuous polarisation. Categorically, extreme-party supporters (AfD, Linke, BSW) provide -4.6 pp less than center voters ($p = 0.019$) and non-affiliates -9.3 pp less ($p < 0.001$; column (3)). No group exhibits a differential alignment effect. ACTION lifts provision for non-affiliates ($\beta = +0.066$, $p = 0.075$).

Topic. Alignment is stronger in climate ($\beta = +0.034$, $p < 0.10$, column (5)) than immigration ($\beta = +0.013$, column (4)). Action and the alignment-action interaction do not differ by topic.

Result 9 (Provision falls with sender-receiver ideological distance). *Provision declines -9.5 pp per Likert unit of sender-receiver ideological distance ($p < 0.001$). The ACTION frame partially offsets the penalty ($\beta = +0.019$, $p = 0.049$); alignment and action do not differ by political group or topic.*

Motive decomposition. *Distance.* The distance penalty persists in the motive specification (-10.5 pp per Likert unit; column (1) of Appendix Table C.21). Each motive slope is invariant to distance. The triple $\text{Persuasion} \times \text{ACTION} \times \text{dist}$ is positive ($\beta = +0.018$, $p < 0.05$): when a distant receiver will act politically, senders select more strongly on expected persuasion direction, consistent with strategic targeting of persuasive facts at distant receivers under high-stakes conditions.

Political extremity. Polarisation is associated with stronger responses on both channels: $\text{Expression} \times \text{pol}$ $\beta = +0.027$ ($p < 0.10$) and $\text{Persuasion} \times \text{pol}$ $\beta = +0.011$ ($p < 0.10$; column (2)). Under ACTION, the instrumental interaction is concentrated

²⁶Senders with extreme preferences face more high-distance receiver cells by construction, potentially conflating behavioral effects with a mechanical alignment-distance correlation. Within-sender variation in distance across the five receivers is the primary identifying variation and is independent of sender extremity.

among center voters (Persuasion $\times$ ACTION $\beta = +0.047$, $p = 0.011$, center reference, column (3)); it is not detected for extreme-party supporters and is muted for non-affiliates (triple $\beta = -0.065$, $p < 0.10$). Center senders exhibit expressive bias at baseline; ACTION adds the instrumental channel on top.

Topic. Appendix Table C.22 points to domain-specific patterns: Expression predicts provision in climate ($\beta = +0.086$, $p = 0.007$) but not immigration in this specification, while Factual Belief Correction ($\beta = +0.053$, $p = 0.036$) and the instrumental interaction (Persuasion $\times$ ACTION $\beta = +0.037$, $p = 0.035$) are detected in immigration only.

Result 10 (Partisan and topic structure of the motives). *Polarisation is associated with stronger expressive and persuasion responses. The instrumental interaction is concentrated among center voters ($\beta = +0.047$, $p = 0.011$) and not detected for extreme-party supporters. Selection on expected persuasion under ACTION is stronger for ideologically distant receivers (Persuasion $\times$ ACTION $\times$ dist, $\beta = +0.018$), consistent with strategic targeting. Expression is concentrated in climate; Factual Belief Correction and the instrumental interaction in immigration.*

5 Experiment 3 (Journalists)

Results below are based on $N = 307$ journalists in the regression sample (1,535 sender-fact observations; $N = 308$ collected), reaching the pre-registered target of 300.

Models of media bias attribute slanted coverage to extrinsic incentives: advertiser revenue, audience demand for confirmation, or owner preferences (Gentzkow and Shapiro 2006, Mullainathan and Shleifer 2005, Besley and Prat 2006). Experiment 3 supplements this literature by isolating the intrinsic component: *some* journalist selectivity persists when those incentives are absent from the task environment, and I ask whether the same motive structure that drives citizen fact-selection operates among professional journalists. Journalists differ from citizen senders along two dimensions that bear on the motive structure. First, they work under explicit professional norms—

codified in Germany’s *Pressekodex* and enforced by the Presserat ([Deutscher Presserat 2025](#), [Steindl et al. 2017](#))—which impose veracity and balance obligations absent in lay information sharing. Second, journalists supply information to mass audiences: a single editorial decision reaches many receivers, not the single known receiver of Experiment 2, which may alter how strongly expressive and persuasion motives enter the provision decision relative to factual belief correction.

5.1 Design and Sample

I apply the ACTION arm of Experiment 2 to professional journalists: the fact set, incentivized elicitations, and receiver matrix are identical, and receivers are matched with 100 % probability. Two features differ: the attention check is omitted,²⁷ and the post-questionnaire replaces citizen items with journalist-specific measures of medium type, geographic reach, and editorial autonomy. Because all journalists are in the ACTION arm, θ_I is not separately identified; I recover $\theta_P + \theta_I$.

Journalists were recruited via the Deutscher Journalisten-Verband (DJV) and Deutsche Journalistinnen- und Journalisten-Union (DJU) mailing lists (pre-registered target $N = 300$; AsPredicted No. 292902).²⁸ This study was approved by the Ethics Council of the Max Planck Society under the same Generalized Approval as Experiment 2. Journalists took a median completion time of 22 minutes and earned €15 on average. Appendix Table D.36 compares the regression sample to Loosen et al. (2023): role composition and sector affiliation are close to the benchmark, but online journalists are substantially overrepresented (45.4 % vs. 3.1 % native-online), reflecting the digital recruitment channel.²⁹ The sample is politically left-leaning (58.5 % Grüne or Linke; mean left-right self-placement ≈ 3.2). Because journalists completed no cross-domain rating task, I merge the Expression measure α_{ij} from citizen cross-domain ratings using (sender per-

²⁷The professional journalist population is small enough that exclusions would be costly and uninformative

²⁸One journalist completed the study seven times, identified through duplicate payout codes; only the first participation is retained.

²⁹The sample is a convenience sample of German journalists recruited via online channels, not a probability sample. Perfect random sampling of German journalists is infeasible; Loosen et al. required eighteen months of multi-source population tracing to assemble their sampling frame.

Table 7: Pre-Registered Treatment Effects: Journalists

	Probability of Fact Provision		
	(1)	(2)	(3)
Alignment	0.017 (0.014)	0.006 (0.019)	0.007 (0.019)
Clustered SE	Journalist	Journalist	Journalist
Fact FE	No	Yes	Yes
Controls	No	No	Yes
Num. obs.	1535	1535	1535
R ²	0.001	0.005	0.014
Adj. R ²	0.000	−0.003	0.003

Sample: $N = 307$ Experiment 3 journalists, 1,535 journalist-fact observations. All journalists in the ACTION arm; no between-subject Action treatment variation. *Alignment:* $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, 1]$ (pre-registered expressive proxy). *Controls (col. 3):* age. Standard errors clustered at journalist level. *Pre-registered H1 test, col. (3):* Alignment > 0, one-sided p reported in text. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

ceived stance, fact direction), extending the stance-equivalence assumption used for positional congeniality.

5.2 Results

Among journalists, the positional alignment gap has the sign H1 predicts: aligned facts are provided more often than unaligned ones ($\beta = +0.007$, one-sided $p = 0.360$, column (3) of Table 7), matching the direction of the citizen result in Experiment 2, but not statistically distinguishable from zero.³⁰ Because all journalists are in the ACTION arm, the between-subject amplification test (H2) is not identified from Experiment 3 alone.

Result 11 (Directionally consistent alignment effect among journalists). *Journalists provide aligned facts more often than unaligned ones, consistent with H1 in direction, though not statistically distinguishable from zero.*

The three elicited proxies remain empirically uncorrelated in the journalist sample (pairwise $|r| \leq 0.05$, within-fact VIF ≈ 1.00), so the channels are separately estimable,

³⁰The coefficient is positive in all three specifications; column (1) omits fact fixed effects and controls, column (2) adds fact fixed effects, and column (3) adds controls.

Table 8: Motive Decomposition: Citizens versus Journalists — Experiments 2 and 3

	Probability of Fact Provision		
	(1)	(2)	(3)
Expression	0.069*** (0.023)	0.058* (0.032)	0.064 (0.058)
Factual belief correction gain	0.037** (0.017)	0.020 (0.024)	0.102*** (0.039)
Persuasion	0.015** (0.007)	0.027*** (0.009)	0.047*** (0.014)
Clustered SE	Sender	Sender	Journalist
Fact FE	Yes	Yes	Yes
Arm	Both	Action only	Action only
Population	Citizens	Citizens	Journalists
Num. obs.	7740	3965	1535
R ²	0.004	0.006	0.020
Adj. R ²	0.002	0.002	0.011

Sample: Experiment 2 citizens (columns 1–2; $N = 1,548$, Bilendi panel, Germany, 2026) and Experiment 3 journalists (column 3; $N = 307$ in the regression sample, recruited via DJV/DJU mailing lists). *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating α_{ij} for journalists, merged from the citizen cross-domain rating cell lookup under the stance-equivalence assumption. *Factual belief correction gain:* incentivized expected correction of the receiver’s factual belief Δb_{ijk} . *Persuasion:* incentivized expected preference shift $\Delta \pi_{ijk}$ toward the sender’s position (positional reading, `delta_gain_2`). All columns include fact fixed effects; standard errors clustered at the sender (journalist) level. All journalists are in the ACTION arm. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

as in Experiment 2. I find Persuasion positive and significant ($\Delta \pi_{ijk}$, $\beta = +0.047$, $p = 0.001$; the journalist–citizen gap widens under the directional specification, $\beta = +0.058^{***}$ vs $+0.026^{**}$ in the citizen ACTION arm; specifications as in Appendix Table C.33, estimated on the journalist sample). The factual belief correction coefficient is $5.1\times$ the citizen ACTION-arm estimate ($\beta = +0.102$, $p = 0.009$ vs. $\beta = +0.020$)—consistent with professional norms that preserve veracity as a selection criterion even when the receiver acts politically, though this cross-sample comparison is not within-study randomization.

Journalists also report roughly twice the expected factual belief correction per fact ($\bar{\Delta b} = 0.322$ vs. 0.172 for citizens): the climate-over-immigration hierarchy and the absence of motivated beliefs appear in both populations. Journalists thus both perceive larger factual belief correction gains per fact and respond more steeply to them—

neither comparison identifies a difference in underlying motives.³¹ As for citizens, the aligned and unaligned distributions of both elicited returns nearly coincide (factual belief correction means 0.324 vs. 0.319; persuasion means 0.030 vs. 0.003), mirroring the uncorrelatedness between congeniality and the elicited returns.

Expression is positive but imprecisely estimated (α_{ij} , $\beta = +0.064$, $p = 0.266$); the point estimate is close to the citizen ACTION-arm estimate ($\beta = +0.058$). Positional alignment is also not significant in the secondary specification ($\beta = +0.006$, $p = 0.75$).³² The three estimates line up with the citizen ACTION arm throughout.³³

Result 12 (The motive structure appears among journalists). *Among professional journalists facing the same task without extrinsic, financial incentives, all three motive proxies are directionally consistent with the citizen ACTION arm. Persuasion and the factual belief correction gain are statistically significant ($p = 0.001$ and $p = 0.009$); Expression is positive and directionally consistent but not significant ($p = 0.266$).*

I close with the pre-registered exploratory heterogeneity battery. Heterogeneity of the alignment gap by sender–receiver distance, polarisation, issue extremity, and the journalist background variables (medium type, broadcaster type, reach, role, and editorial autonomy) is reported in Appendix Table D.37; no interaction survives false-discovery control at the 5 % level (minimum $q = 0.11$ in the 8-test alignment family). Extending the battery to the three motive channels (24 tests), one pattern survives at the 10 % level: journalists at online outlets select less on Expression ($\beta = -0.29$, $p = 0.003$, $q = 0.074$), directionally echoed in the alignment battery ($\beta = -0.070$, $p = 0.017$, $q = 0.11$). I read this pattern as suggestive.

³¹Approximately 21 % of journalist participants provided their names and outlets, enabling a planned match to their published articles over the past twelve months to test whether experimental selectivity predicts real-world reporting direction.

³²For journalists, positional alignment correlates with the merged Expression measure at $r = 0.48$ (citizens: $r = 0.18$): because α_{ij} is merged from the citizen cross-domain rating cells, it shares the sender-distance input with positional alignment. Positional alignment is not part of the channel specification, so this correlation does not affect the separability of the three elicited proxies.

³³The $\sim 2\times$ level gap in expected factual belief correction holds at every distance, and the by-distance shape—flat for Δb_{ijk} , rising for $\Delta \pi_{ijk}$ —matches the citizen profile.

5.3 Pooled Comparison: Citizens and Journalists

I pool the Experiment 2 citizens with the Experiment 3 journalists, adding a journalist indicator and demographic controls; because journalists are observed only in the ACTION arm, the journalist indicator captures their provision level relative to citizens in that arm while alignment and action effects are identified from citizens. Table 9 shows the pre-registered treatment effects. The journalist indicator is large and precisely estimated (+9.5 pp, conditional on alignment and demographics), and the alignment slope does not differ by population ($\text{Align} \times \text{Journalist}$, $\beta = +0.006$, $p = 0.74$): journalists are more forthcoming in level but select along the same dimension and to the same degree.

All three motives are positive and significant in the pooled sample (Table 10, column (2)). Adding the action decomposition (column (3)), the Persuasion $\times$ ACTION interaction is positive and significant ($\beta = +0.030$, $p < 0.01$), identified from the citizens. Column (4) allows each motive to differ for journalists: the factual belief correction slope is larger (Factual belief correction gain $\times$ Journalist, $p = 0.043$) and the expressive slope does not differ ($p = 0.30$), a conservative null given the sample's left-political composition. The implied persuasion slope rises from the citizen CONTROL arm (0.019) to the citizen ACTION arm (0.039) to journalists (0.076); the journalist-specific increment is exploratory and the confidence interval is wide.

5.4 Self-Reported Motives and Norms

Senders' open-ended accounts of their own decisions offer an independent check on the revealed channels. These self-reports are exploratory and were not pre-registered. The two samples differ in how much they say: 27 percent of citizens but 74 percent of journalists give a reason that refers to the sharing decision, and referring senders are somewhat more issue-extreme in both.³⁴ A large language model classifies each reason into a primary motive (Appendix E.1; Figure B.4). Citizens spread across motives—

³⁴Referring senders average higher issue extremity than non-referrers in both samples ($p < 0.01$); the analysis conditions on giving a codable reason.

Table 9: Provision Decisions: Citizens and Journalists Pooled

	Probability of Fact Provision		
	(1)	(2)	(3)
Align	0.025** (0.012)	0.020* (0.012)	0.020* (0.012)
Action	0.026** (0.013)	0.022* (0.013)	0.022* (0.013)
Align $\times$ Action	-0.009 (0.015)	-0.006 (0.015)	-0.008 (0.016)
Journalist	0.116*** (0.016)	0.095*** (0.019)	0.095*** (0.019)
Align $\times$ Journalist			0.006 (0.019)
Clustered SE	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes
Controls	No	Yes	Yes
Sample	Pooled	Pooled	Pooled
Num. obs.	9275	9245	9245
R ²	0.012	0.066	0.066
Adj. R ²	0.010	0.060	0.060

Sample: pooled Experiment 2 citizens ($N = 1,548$; both arms; 7,740 observations) and Experiment 3 journalists ($N = 307$; ACTION arm only; 1,535 observations). *Outcome:* binary provision indicator. *Align:* positional alignment $\text{Pref}_i \times \text{Dir}_j \in [-1, +1]$. *Journalist:* indicator for Experiment 3 participants (all in the ACTION arm). *Controls* (columns (2)–(3)): sender-receiver ideological distance, age, gender, education, Bundesland, left-right self-placement, and closest-party affiliation—the demographic variables measured in both samples. All specifications include fact fixed effects. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Expression (21 percent), Persuasion (19 percent), and factual belief correction (14 percent)—while journalists concentrate on factual belief correction (46 percent) and Persuasion (28 percent).

I now ask whether these stated reasons track revealed selection. For each motive I interact its revealed return—the expressive rating α_{ij} , the persuasion return $\Delta\pi_{ijk}$, and the factual belief correction Δb_{ijk} —with an indicator for whether the sender states that motive (Table 11). Senders who state a persuasion motive select more steeply on the revealed persuasion return, in both samples ($\Delta\pi_{ijk} \times \text{Stated Persuasion}$, +0.121, $p < 0.01$ for citizens; +0.099, $p < 0.01$ for journalists). Senders who state factual belief correction provide more overall (+0.224 and +0.223, $p < 0.01$) without selecting more steeply on it, and the expressive interaction is positive in both samples but not

Table 10: Motive Decomposition: Citizens and Journalists Pooled

	Probability of Fact Provision			
	(1)	(2)	(3)	(4)
Expression	0.070*** (0.021)	0.074*** (0.020)	0.073*** (0.020)	0.066*** (0.022)
Factual belief correction gain	0.046*** (0.016)	0.052*** (0.016)	0.052*** (0.016)	0.041** (0.017)
Persuasion	0.019*** (0.006)	0.036*** (0.006)	0.018** (0.009)	0.019** (0.009)
Persuasion $\times$ Action			0.030*** (0.012)	0.020 (0.012)
Action			0.021* (0.013)	0.021* (0.013)
Journalist	0.123*** (0.015)	0.100*** (0.019)	0.090*** (0.019)	0.046 (0.030)
Expression $\times$ Journalist				0.054 (0.052)
Factual belief correction gain $\times$ Journalist				0.082** (0.040)
Persuasion $\times$ Action $\times$ Journalist				0.037** (0.017)
Clustered SE	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes
Controls	No	Yes	Yes	Yes
Sample	Pooled	Pooled	Pooled	Pooled
Num. obs.	9275	9245	9245	9245
R ²	0.014	0.072	0.073	0.074
Adj. R ²	0.013	0.066	0.067	0.068

Sample: pooled Experiment 2 citizens ($N = 1,548$; both arms; 7,740 observations) and Experiment 3 journalists ($N = 307$; ACTION arm only; 1,535 observations). *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating $a_{ij} \in [-1, +1]$ from the citizen cross-domain rating cell lookup (merged for journalists under the stance-equivalence assumption). *Factual belief correction gain:* incentivized expected correction of the receiver's factual belief $\Delta b_{ijk} \in [0, 1]$. *Persuasion:* incentivized expected movement of the receiver toward the sender's position $\Delta \pi_{ijk}$ (positional reading). *Journalist:* indicator for Experiment 3 participants. Column (3) adds the ACTION main effect and the Persuasion $\times$ ACTION interaction, identified from the citizens (who vary across arms). Column (4) adds journalist heterogeneity: the Expression and Factual belief correction interactions with the journalist indicator, and the persuasion channel's heterogeneity as the Persuasion $\times$ ACTION $\times$ Journalist term; because journalists are observed only in the ACTION arm, the Persuasion $\times$ Journalist double is collinear with this triple and omitted. *Controls* (columns (2)–(4)): sender-receiver ideological distance, age, gender, education, Bundesland, left-right self-placement, and closest-party affiliation—the demographic variables measured in both samples. All specifications include fact fixed effects. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

distinguishable from zero. Words and behavior agree on the persuasion channel: the senders who say they are trying to move the receiver are the ones whose provision responds most to how far a fact is expected to move the receiver.

A second open-text question reaches the same beliefs from a different angle. Separate from the sharing reasons coded above, senders also explained what they based their receiver-belief predictions on, and a large language model codes these rationales into stated beliefs about whether the receiver's position is movable and whether the receiver already knows the facts (Appendix F.2; Figure F.7); rationales are codable for 11 percent of citizens and 54 percent of journalists, and Experiment 1 elicited no receiver-belief predictions. These stated beliefs track the incentivized elicited returns

Table 11: Stated Motives and Revealed Selection Returns (Experiments 2 and 3)

	Provision	
	(1)	(2)
Expression ($\hat{\alpha}_{ijk}$)	0.058** (0.023)	0.035 (0.053)
Stated Expression	-0.081* (0.042)	-0.343** (0.144)
Persuasion ($\hat{\Delta}\pi_{ijk}$)	0.009 (0.007)	0.028* (0.015)
Stated Persuasion	-0.073*** (0.024)	-0.091*** (0.031)
Belief correction ($\hat{\Delta}b_{ijk}$)	0.031* (0.018)	0.044 (0.052)
Stated Belief correction	0.224*** (0.032)	0.223*** (0.032)
$\hat{\alpha}_{ijk} \times$ Stated Expression	0.084 (0.092)	0.422 (0.346)
$\hat{\Delta}\pi_{ijk} \times$ Stated Persuasion	0.121*** (0.033)	0.099*** (0.031)
$\hat{\Delta}b_{ijk} \times$ Stated Belief correction	0.109 (0.067)	0.018 (0.065)
Sample	Citizens (Exp. 2)	Journalists (Exp. 3)
Fact FE	Yes	Yes
Num. obs.	7740	1530
R ²	0.019	0.125
Adj. R ²	0.016	0.114

Sample: decision-level provision, Experiment 2 citizens (column 1) and Experiment 3 journalists (column 2).
Specification: each revealed selection return—the expressive rating α_{ij} (Expression), the elicited persuasion return $\Delta\pi_{ijk}$ (Persuasion), and the elicited factual belief correction Δb_{ijk} (Belief correction)—is interacted with an indicator for whether the sender’s open-ended decision reason states the corresponding motive (Appendix E.1). A positive interaction means senders who *state* a motive select more strongly on the matching revealed return. Fact fixed effects included; standard errors clustered by sender. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

(Table F.44): senders who state that the receiver is unpersuadable have a predicted magnitude of receiver movement about 20 percent smaller among citizens (-0.14 , $p < 0.05$) and about 40 percent smaller among journalists (-0.27 , $p < 0.01$), and senders who state that the receiver does not already know the facts predict a lower baseline share who know (-6.9 percentage points for citizens, -9.9 points for journalists) and a larger factual belief correction. The two independently coded questions also agree at the sender level (Table F.45): senders whose stated motive is Persuasion are far more

likely to state that the receiver is persuadable (25 vs. 3 percent among citizens, Fisher $p < 0.001$; 34 vs. 18 percent among journalists, $p = 0.007$), and senders whose stated motive is factual belief correction are more likely to state that the receiver does not know the facts (Fisher $p = 0.031$ for citizens, $p = 0.011$ for journalists). The incentivized persuasion return $\Delta\pi_{ijk}$ and the factual belief correction Δb_{ijk} thus reflect beliefs senders can articulate: a sender's stated motive, her stated model of the receiver, and her incentivized elicitation cohere.

Senders' stated *norms*, by contrast, predict how much they share rather than how selectively. The post-questionnaire of both experiments elicited agreement (1–5) with three normative statements: that whoever passes on political information should also pass on factually correct information contradicting their own position; that social media should be more transparent about why users see some political information and not others; and that platforms and media should be obligated to surface relevant information from different perspectives. Journalists endorse the two principle statements more strongly than citizens (means 4.60 vs. 4.13 and 4.61 vs. 4.05) but not the obligation statement (4.22 vs. 4.17). Appendix Table F.42 reports the distributions.

All three items positively predict overall provision among citizens ($\beta = +0.015$ to $+0.023$ per scale point, $p \leq 0.042$), as do the two principle items among journalists ($+0.096$ and $+0.044$). I interact each item with the three selection measures—positional alignment, Expression, and Persuasion—and apply a Benjamini–Hochberg correction across the resulting family of 18 tests spanning both samples. No citizen interaction is distinguishable from zero, and a pooled index of the three items is also null (all $p > 0.13$). Among journalists, the transparency item is the exception: journalists who endorse platform transparency select less on positional alignment ($\beta = -0.075$, $p = 0.003$, $q = 0.060$) and less on Expression ($\beta = -0.194$, $p = 0.007$, $q = 0.060$). The correlation is consistent with the professional-norms reading of the factual-belief-correction result above, though journalist answers are compressed at the top of the scale (95 % agreement), so I read the moderation as suggestive. The full battery is in Appendix Table F.43.

6 Experiment 4

6.1 Conceptual Framework: Receiver Updating

I characterize how a receiver's beliefs respond to selective but truthful disclosure under imperfect transmission, and I evaluate those beliefs by their accuracy. I take no stand on which policy is socially desirable but adopt an epistemic notion of welfare: accurate beliefs are valuable.

Environment. Nature draws a binary state of the world $\omega \in \{0, 1\}$. The receiver does not observe the state and holds prior $p = \Pr(\omega = 1)$, with $p \in (0, 1)$. The sender observes ω and chooses whether to disclose it, $d \in \{0, 1\}$; disclosure is truthful, so she can transmit the true state or stay silent but cannot report a state that did not occur. Let

$$\sigma_1 = \Pr(d = 1 \mid \omega = 1), \quad \sigma_0 = \Pr(d = 1 \mid \omega = 0)$$

denote the receiver's beliefs about the sender's state-contingent disclosure probabilities. A disclosed state reaches the receiver with probability $\tau \in (0, 1]$; a withheld state, or a disclosed state that fails to transmit, leaves no message. The receiver observes $m \in \{\omega, \emptyset\}$: the true state after successful transmission, and nothing otherwise. She never observes d , so non-receipt ($m = \emptyset$) pools a withheld state with a disclosed state that failed to transmit.

Receiver beliefs. I separate the information a message carries from the receiver's response to it, and let a Bayesian receiver fix the reference. A received message identifies the state, so Bayes beliefs are $b_1^{\text{Bayes}} = 1$ and $b_0^{\text{Bayes}} = 0$. Non-receipt does not identify the state: it arises in both states, through withholding or failed transmission, and Bayesian belief is

$$b_{\emptyset}^{\text{Bayes}} = \Pr(\omega = 1 \mid m = \emptyset) = \frac{p(1 - \tau\sigma_1)}{p(1 - \tau\sigma_1) + (1 - p)(1 - \tau\sigma_0)}. \quad (7)$$

If $\sigma_1 = \sigma_0$, non-receipt is uninformative and $b_{\emptyset}^{\text{Bayes}} = p$; if $\sigma_1 \neq \sigma_0$, the more-disclosed state is discounted under silence ($\sigma_1 > \sigma_0$ implies $b_{\emptyset}^{\text{Bayes}} < p$); and a smaller τ pulls $b_{\emptyset}^{\text{Bayes}}$ toward p , because non-receipt is then more readily explained by failed transmission than by withholding. These are properties of the reference, not assumptions about behavior. Let $\tilde{b}_1, \tilde{b}_0, \tilde{b}_{\emptyset} \in [0, 1]$ denote the receiver's actual beliefs after each message. I place no restriction on how they relate to the reference: the receiver may update as Bayes prescribes, under- or over-react, ignore silence, or read silence as confirmation of her own position, moving opposite the Bayesian direction.

Epistemic Welfare. I score a belief by its squared distance from the realized state,

$$L(b, \omega) = (b - \omega)^2, \quad (8)$$

the proper metric implied by the quadratic scoring rule used to elicit these beliefs.³⁵ The two message events then differ in kind. Because a received message pins down the state, Bayesians attain the truth ($b_m^{\text{Bayes}} = \omega$), and any loss— $(1 - \tilde{b}_1)^2$ or $\tilde{b}_0^2$ —is due solely to the receiver not adopting the disclosed state. Silence, by contrast, leaves even Bayesians in error, so belief error after non-receipt reflects both its ambiguity and the receiver's reading of it. I leave the direction and magnitude of that reading unrestricted, and its dependence on τ an empirical question.

Expected loss and the cost of selection. Taking expectations over the state and over transmission, the epistemic loss generated by the environment is

$$\mathcal{L} = p\tau\sigma_1(1 - \tilde{b}_1)^2 + (1 - p)\tau\sigma_0\tilde{b}_0^2 + p(1 - \tau\sigma_1)(1 - \tilde{b}_{\emptyset})^2 + (1 - p)(1 - \tau\sigma_0)\tilde{b}_{\emptyset}^2. \quad (9)$$

Evaluated with the receiver's beliefs, $\mathcal{L}$ is the belief-accuracy loss relative to the realized state and requires no assumption on how she updates. Writing $q_1 = p(1 - \tau\sigma_1)$ and $q_0 = (1 - p)(1 - \tau\sigma_0)$ for the silence weights, with $b_{\emptyset}^{\text{Bayes}} = q_1/(q_1 + q_0)$, the

³⁵An absolute-distance metric of the belief-correction return yields the same qualitative conclusions but, unlike the quadratic form, is not minimized by a calibrated probability, so I adopt squared distance here.

quadratic metric decomposes it as

$$\begin{aligned} \mathcal{L} = & \underbrace{q_1(1 - b_{\emptyset}^{\text{Bayes}})^2 + q_0(b_{\emptyset}^{\text{Bayes}})^2}_{\text{ambiguity of non-receipt}} \\ & + \underbrace{p\tau\sigma_1(1 - \tilde{b}_1)^2 + (1 - p)\tau\sigma_0\tilde{b}_0^2 + (q_1 + q_0)(\tilde{b}_{\emptyset} - b_{\emptyset}^{\text{Bayes}})^2}_{\text{departure from Bayes}}. \end{aligned}$$

The first term is the loss a Bayesian still bears because silence overlaps across states—a function of $(p, \sigma_1, \sigma_0, \tau)$ alone, and zero only when non-receipt separates the states. The remaining terms are the receiver’s departures from Bayes, each a squared deviation $(\tilde{b}_m - b_m^{\text{Bayes}})^2$ that vanishes under Bayesian updating and is strictly positive otherwise. The cost attributable to selective disclosure is the excess of the ambiguity term over its value under a non-selective counterfactual $\sigma_1 = \sigma_0$ at fixed disclosure volume.

6.2 Experimental Design and Hypotheses

[PLACEHOLDER — September 2026. Receivers matched to Exp2/3 senders.]

6.3 Sample, Incentives, and Protocol

[PLACEHOLDER.]

6.4 Results

[PLACEHOLDER.]

7 Discussion

This paper makes four contributions. First, I provide causal evidence of a distortion of the political information environment that operates entirely on true content. The observational record documents that citizens, politicians, and fact-checkers pass on congenial material at higher rates; my experiments establish the causal counterpart: senders withhold true political facts because of what those facts say about their own

position, and the selectivity intensifies when the receiver’s beliefs feed into political action. In both designs, silence is unattributable by construction—in Experiment 1 the receiver does not know a disclosure decision was made, and in Experiments 2 and 3 the probabilistic transmission pools a withheld fact with one that failed to arrive—so selective omission operates without the reputational discipline that a visible refusal to share would invite. The distortion differs from misinformation in kind, not degree: every shared fact is true, and the distortion rides entirely on what is missing.

Second, I separate the motive channels behind this selectivity. Incentivized belief elicitation measures the persuasive and belief-correction returns each fact offers; a cross-domain rating task measures its expressive return. The coefficients on these returns identify which motives are active, and the channels differ sharply in when they operate: expressive selection is unconditional, while persuasion-based selection activates when the receiver’s beliefs carry consequences. Belief correction raises provision without tilting it toward the sender’s side. The estimates recover the sign and activation of each motive, not the magnitudes of the underlying weights; the one design-identified magnitude claim is the instrumental activation itself, which follows from the ACTION randomization independently of functional-form assumptions.

Third, I provide evidence for an intrinsic component of media bias, together with its motives. Journalist selectivity persists in an environment stripped of the extrinsic incentives—audience demand, ownership, competition—that media-bias models invoke, and it is organized by the same channels that organize citizen behavior. The factual belief correction coefficient, roughly five times the citizen ACTION-arm estimate, is consistent with professional norms preserving veracity as a selection criterion when beliefs are consequential. This is a cross-sample comparison rather than within-study randomization, and it reflects both a higher perceived correction potential per fact and a steeper response to it.

Fourth, the receiver experiment is designed to measure what curated truth costs its audience: under the epistemic benchmark of Section 6.1, the belief-accuracy loss that selective omission imposes on receivers is [\[PENDING DATA\]](#).

Why Political Action Moves Only the Persuasion Channel. The ACTION treatment could in principle amplify selectivity through any of the three channels; it detectably moves exactly one. The pattern follows the utility function. Expressive utility attaches to the act of disclosing a congenial fact, not to the receiver’s response, so downstream consequences leave it unchanged: the Expression $\times$ ACTION interaction is null ($\beta = -0.027$, $p = 0.52$). The factual belief correction slope does not respond significantly either ($\beta = -0.034$, $p = 0.32$). The persuasion activation, in turn, operates through weights rather than beliefs: the elicited persuasion returns are statistically indistinguishable across arms, so the treatment changes how senders act on their beliefs, not the beliefs themselves. Because the expected persuasion return is near zero on average and widely dispersed, activating the instrumental motive redistributes provision within sender—toward facts expected to move the receiver and away from facts expected to backfire—rather than raising overall provision.

Reconciling Experiments 1 and 2. The two citizen experiments differ in provision levels and in what the positional-alignment proxy detects, and both differences are informative about where curated truth operates. Provision falls from above 80 percent in Experiment 1 to the low 60s in Experiment 2, and the full-range positional alignment gap shrinks from 7–8 pp to about 5 pp; Section 4.4 traces the level difference to three design changes. The substantive reading is that selectivity persists across two very different sharing environments: an anonymous one, in which the receiver holds a false or absent belief and silence leaves no trace, and a socially legible one, in which positions are mutually salient and a missing fact is only partially attributable to the sender. The action results are consistent once the proxies are read for what they measure. In Experiment 1 the receiver is documented to hold a false or absent belief on the fact itself, so the direction of belief correction—and with it the direction of expected persuasion—is pinned down by the environment. Senders’ unincentivized expectations confirm the coincidence: 78.0% toward their own position for aligned facts, 17.2% for unaligned ones. Alignment therefore bundles the expressive and persuasive value of a fact, and the ACTION amplification loads onto the alignment interaction. In Exper-

iment 2 receivers span the full stance scale, and the direction of expected persuasion is governed by where the receiver stands rather than by the fact's congeniality: in the incentivized elicitation the same shares are 44.2% and 43.2%, and positional alignment is uncorrelated with the persuasion return ($r = 0.000$). The alignment interaction is structurally empty as a test of an action-conditional persuasion motive, and the amplification appears where the framework says it must—on the persuasion channel. The Experiment 2 Hypothesis 6 null and the Experiment 1 action doubling are the same finding viewed through proxies of different resolution.

Elicited Beliefs Rather than a Structural Model. I identify motives from incentivized belief elicitation interacted with randomization, not from a structural model of sharing. The choice trades magnitude for robustness. Structural decompositions of sharing motives recover parameter values, but they require equilibrium assumptions on receiver behavior (Guriev et al. 2026); because the belief arguments senders act on are elicited directly, my identification requires none—no parametric assumptions on the disclosure-shock distribution, no functional form for the persuasion-value function, and no equilibrium restrictions on what receivers infer. What the approach gives up is magnitudes and counterfactuals: the coefficients establish which motives are active and when, not how much provision would change under a counterfactual motive structure. The structural counterpart, in which the elicited returns are scored against realized receiver outcomes—still with no equilibrium fixed point, since beliefs are measured rather than solved for—is [PENDING DATA].

The design also carries a measurement lesson for observational work. Positional alignment, the proxy that naturalistic sharing data support, is empirically uncorrelated with the elicited persuasion return within facts ($r = 0.000$); in these data it therefore could not have detected the action-conditional component of persuasion, however large. Studies that rely on alignment-style proxies identify expressive selectivity; the instrumental margin requires a persuasion measure of its own.

Experimenter Demand and Social Desirability. Provision decisions were observed, and observation raises the concern that selectivity reflects what senders believed the study wanted. The direction of the bias argues otherwise. All three sources of social desirability bias—material costs, social-image concerns, and self-image concerns (Bursztyn et al. 2025)—favor sharing true, sourced information over withholding it, so desirability pressure attenuates measured selectivity rather than producing it. A demand account must also explain within-sender gradients: identification comes from variation across facts within the same sender, which a uniform demand effect cannot generate (de Quidt et al. 2018). Nor does ACTION change what is socially desirable—sharing remains the creditable answer in both arms—so differential demand is a poor account of the persuasion activation. The direct measurements point the same way: 87 percent of Experiment 1 senders judged the survey unbiased, and 85.5 and 82.4 percent of Experiment 2 and Experiment 3 senders perceived no political imbalance (Appendix F.1). Only 7 percent of Experiment 1 senders correctly identified the research question; the corresponding shares for Experiments 2 and 3 are [PENDING DATA]. The primary estimates are [PENDING DATA] when senders who guessed the purpose are excluded.

External Validity. The citizen samples are quota-sampled, not representative, and each experiment covers one country and at most two policy domains; the qualitative structure is the same in both domains tested, and I make no cross-national claim. The journalist sample is a convenience sample recruited online, overrepresenting online outlets and left-leaning positions; the political skew is conservative for the expressive test, since a left-skewed sample should exhibit a stronger expressive motive and the Expression coefficient is nonetheless not significant. Each sender faces a single receiver of known position, whereas journalists in the field supply many receivers at once, and the ACTION treatment may shift salience or task construal beyond political consequentiality. Finally, several design features—true, sourced facts, no cost of sharing—push toward provision, suggesting that the selectivity estimates are conservative. Stated norms point the same way as the behavioral estimates: senders who endorse balanced-

information norms share more overall but, with one suggestive exception among journalists, not less selectively (Section 5.4).

Policy Implications. The defining property of curated truth is invisibility. A false claim is an object: it can be detected, labeled, and corrected, and the policy architecture for political information (fact-checking, content labeling, platform moderation) is built around these operations. A withheld truth is the absence of an object. Nothing arrives to be flagged, the beliefs it would have corrected persist undisturbed, and the receiver rarely knows what was available; curated truth is a distortion that fact-checking and labeling interventions do not reach. Reaching it requires tools that target completeness rather than accuracy, such as transparency about selection and omission, or source multiplicity that lets one sender's omissions be filled by another's disclosures. Completeness is the harder dimension to govern, because accuracy can be checked against ground truth while completeness is relative to an open-ended universe of relevant facts. The motive evidence disciplines the target further. Remedies premised on extrinsic incentives cannot reach the intrinsic component Experiment 3 isolates, though the suggestive correlation between journalists' endorsement of transparency norms and lower selectivity (Section 5.4) marks norm-level interventions as a candidate lever. And because the instrumental motive activates where beliefs translate into political action, the distortion should be strongest around elections and referenda—exactly where the information environment matters most. How much policy attention completeness deserves depends on the epistemic cost that curated truth imposes on receivers, which is [PENDING DATA].

8 Concluding Remarks

The political information environment can be distorted without a single false statement. I study the supply side of this distortion: why senders choose which true political facts to pass on and which to withhold. Across three experiments—with quota-sampled U.S. and German citizens and professional German journalists—provision of

true political facts is selective: it favors the sender's own position and sharpens when the receiver's beliefs carry political consequences. The channels behind the selectivity separate cleanly: sharing congenial facts is expressive and operates regardless of consequences; selecting facts for their persuasive effect is instrumental and switches on when the receiver acts; the wish to correct factual beliefs raises how much senders share without tilting selection toward the sender's side. Professional journalists, deciding without the extrinsic pressures that media-bias models invoke, show the same structure, and their response to a fact's expected factual belief correction is, if anything, steeper than citizens': some of the selectivity that media-bias models attribute to markets is intrinsic to the people who select.

A receiver can scrutinize what arrives; she cannot scrutinize what was never sent, and in most information environments she does not know what was available. Curated truth therefore evades the receiver's own defenses as well as the policy architecture built for false content, and it should be strongest around the moments when beliefs turn into collective decisions.

Whether supply-side selectivity translates into distorted beliefs depends on the receiver: one who reads silence correctly undoes part of the curation; one who underreads it inherits the sender's selection. The receiver experiment closes this loop, and its answer to what curated truth costs its audience is [\[PENDING DATA\]](#). What is already clear is the shape of the problem: the supply of political information is curated by the motives of those who hold it, the curation survives professional norms and visible positions, and no amount of verifying what is said will reveal what is not.

References

- Abramitzky, R., Ager, P., Boustan, L., Cohen, E., and Hansen, C. W. (2024). The effect of immigration on the economy: Lessons from the 1920s border closure. *American Economic Review: Insights*.
- Alesina, A., Miano, A., and Stantcheva, S. (2023). Immigration and redistribution. *The Review of Economic Studies*, 90(1):1–39.

- Allcott, H. and Gentzkow, M. (2017). Social media and fake news in the 2016 election. *Journal of Economic Perspectives*, 31(2):211–236.
- Allen, J., Howland, B., Mobius, M., Rothschild, D., and Watts, D. J. (2020). Evaluating the fake news problem at the scale of the information ecosystem. *Science Advances*, 6(14):eaay3539.
- Allen, J., Mobius, M., Rothschild, D. M., and Watts, D. J. (2021). Research note: Examining potential bias in large-scale censored data. *Harvard Kennedy School Misinformation Review*, 2(4).
- Ambuehl, S., Bernheim, B. D., and Ockenfels, A. (2021). What motivates paternalism? an experimental study. *American economic review*, 111(3):787–830.
- Andreottola, G. and Louis-Sidois, C. (2025). Why us and not them? a theory of political fact-checking. SSRN Working Paper No. 5122378.
- Angelucci, C. and Prat, A. (2024). Is journalistic truth dead? measuring how informed voters are about political news. *American Economic Review*, 114(4):887–925.
- Aridor, G., Jiménez-Durán, R., Levy, R., and Song, L. (2024). The economics of social media. *Journal of Economic Literature*, 62(4):1422–1474.
- Barrera, O., Guriev, S., Henry, E., and Zhuravskaya, E. (2020). Facts, alternative facts, and fact checking in times of post-truth politics. *Journal of public economics*, 182:104123.
- Barron, K. (2021). Belief updating: Does the ‘good-news, bad-news’ asymmetry extend to purely financial domains? *Experimental Economics*, 24(1):31–58.
- Bartling, B., Cappelen, A. W., Hermes, H., Tungodden, B., and Skivenes, M. (2023). Free to fail? paternalistic preferences in the united states. *NHH Dept. of Economics Discussion Paper*, (09).
- Bénabou, R. and Tirole, J. (2006). Incentives and prosocial behavior. *American economic review*, 96(5):1652–1678.
- Berger, L. M. (2026). How digital media markets amplify news sentiment. Working paper, University of Cologne.
- Berger, L. M., Kerkhof, A., Mindl, F., and Münster, J. (2025). Debunking “fake news” on social media: Immediate and short-term effects of fact-checking and media literacy interventions. *Journal of Public Economics*, 245:105345. Open access, CC BY. Available online 5 April 2025. Pre-registered at AEA Registry AEARCTR-0008199.

- Besley, T. and Prat, A. (2006). Handcuffs for the grabbing hand? media capture and government accountability. *American Economic Review*, 96(3):720–736.
- Bowen, T. R., Dmitriev, D., and Galperti, S. (2023). Learning from shared news: When abundant information leads to belief polarization. *The Quarterly Journal of Economics*, 138(2):955–1000.
- Boxell, L. and Conway, J. (2022). Journalist ideology and the production of news: Evidence from movers. Working paper.
- Braghieri, L. (2024). Political correctness, social image, and information transmission. *American Economic Review*, 114(12):3877–3904.
- Braghieri, L., Eichmeyer, S., Levy, R., Mobius, M. M., Steinhardt, J., and Zhong, R. (2025). Level slant and polarization of news consumption on social media. *SSRN Working Paper*.
- Budak, C., Nyhan, B., Rothschild, D. M., Thorson, E., and Watts, D. J. (2024). Misunderstanding the harms of online misinformation. *Nature*, 630:45–53.
- Bullock, J. G. and Lenz, G. (2019). Partisan bias in surveys. *Annual Review of Political Science*, 22(1):325–342.
- Burdea, V., Montero, M., and Sefton, M. (2023). Communication with partially verifiable information: An experiment. *Games and Economic Behavior*, 142:113–149.
- Bursztyn, L., Egorov, G., Haaland, I., Rao, A., and Roth, C. (2023a). Justifying dissent. *The Quarterly Journal of Economics*, 138(3):1403–1451.
- Bursztyn, L., Haaland, I. K., Röver, N., and Roth, C. (2025). The social desirability atlas. Technical report, National Bureau of Economic Research.
- Bursztyn, L., Rao, A., Roth, C., and Yanagizawa-Drott, D. (2023b). Opinions as facts. *The Review of Economic Studies*, 90(4):1832–1864.
- Cagé, J., Gallo, N., Hengel, M., Henry, E., and Huang, Y. (2025). Fact-checking and misinformation: Evidence from the market leader. CESifo Working Paper 12319, CESifo.
- Cagé, J., Hengel, E., Hervé, N., and Urvoy, C. (2022). Hosting media bias: Evidence from the universe of French broadcasts 2002–2020. HAL preprint hal-03878119.
- Chopra, F., Haaland, I., and Roth, C. (2022). Do people demand fact-checked news? evidence from us democrats. *Journal of Public Economics*, 205:104549.

- Chopra, F., Haaland, I., and Roth, C. (2024). The demand for news: Accuracy concerns versus belief confirmation motives. *The Economic Journal*, 134(661):1806–1834.
- Crawford, V. P. and Sobel, J. (1982). Strategic information transmission. *Econometrica*, 50(6):1431–1451.
- de Quidt, J., Haushofer, J., and Roth, C. (2018). Measuring and bounding experimenter demand. *American Economic Review*, 108(11):3266–3302.
- Degan, A., Li, M., and Xie, H. (2023). An experimental investigation of persuasion through selective disclosure of evidence. *Canadian Journal of Economics/Revue canadienne d'économie*, 56(4):1490–1516.
- Deutscher Presserat (2025). Publizistische Grundsätze (Pressekodex): Richtlinien für die publizistische Arbeit nach den Empfehlungen des Deutschen Presserats. Fassung vom 19. März 2025; first adopted 1973.
- Durante, R. and Knight, B. (2012). Partisan control, media bias, and viewer responses: Evidence from berlusconis italy. *Journal of the European Economic Association*, 10(3):451–481.
- Dye, R. A. (1985). Disclosure of nonproprietary information. *Journal of accounting research*, pages 123–145.
- Eil, D. and Rao, J. M. (2011). The good news-bad news effect: Asymmetric processing of objective information about yourself. *American Economic Journal: Microeconomics*, 3(2):114–138.
- Ershov, D. and Morales, J. S. (2024). Sharing news left and right: Frictions and misinformation on twitter. *The Economic Journal*, 134:2391–2417.
- Eyster, E. and Rabin, M. (2005). Cursed equilibrium. *Econometrica*, 73(5):1623–1672.
- Farina, A., Fréchette, G. R., Ispano, A., Lizzeri, A., and Perego, J. (2026). The selective disclosure of evidence: An experiment. *The Review of Economic Studies*. Forthcoming.
- Farina, A. and Leccese, M. (2025). Hiding a flaw? experimental evidence on multi-dimensional information disclosure. Working paper, September 2025. Princeton University and Boston University.
- Feess, E., Jost, P.-J., and Ressi, A. (2024). Fake news and the asymmetry between type I and type II errors: Dismissing true information is more prevalent than following false information. Working paper. AEA RCT Registry AEARCTR-0010984.

- Fong, M.-J., Lin, P.-H., and Palfrey, T. R. (2025). Cursed sequential equilibrium. *American Economic Review*, 115(8):2616–2658.
- Freitag, J., Kerkhof, A., and Münster, J. (2021). Selective sharing of news items and the political position of news outlets. *Information Economics and Policy*, 56:100926.
- Gentzkow, M. and Shapiro, J. M. (2006). Media bias and reputation. *Journal of Political Economy*, 114(2):280–316.
- Gentzkow, M. and Shapiro, J. M. (2010). What drives media slant? evidence from us daily newspapers. *Econometrica*, 78(1):35–71.
- Golman, R., Loewenstein, G., Molnar, A., and Saccardo, S. (2022). The demand for, and avoidance of, information. *Management Science*, 68(9):6454–6476.
- González-Bailón, S., Lazer, D., Barberá, P., Zhang, M., Allcott, H., Brown, T., Crespo-Tenorio, A., Freelon, D., Gentzkow, M., Guess, A. M., et al. (2023). Asymmetric ideological segregation in exposure to political news on facebook. *Science*, 381(6656):392–398.
- Grigorieff, A., Roth, C., and Ubfal, D. (2020). Does information change attitudes toward immigrants? *Demography*, 57:1117–1143.
- Grossman, S. J. (1981). The informational role of warranties and private disclosure about product quality. *The Journal of law and Economics*, 24(3):461–483.
- Grossmann, M. R. P. (2024). Knowledge and freedom: Evidence on the relationship between information and paternalism.
- Guess, A., Nagler, J., and Tucker, J. (2019). Less than you think: Prevalence and predictors of fake news dissemination on facebook. *Science advances*, 5(1):eaau4586.
- Guriev, S., Henry, E., Marquis, T., and Zhuravskaya, E. (2026). Curtailing false news, amplifying truth. *Econometrica*. Forthcoming.
- Haaland, I. and Roth, C. (2020). Labor market concerns and support for immigration. *Journal of Public Economics*, 191:104256.
- Haaland, I., Roth, C., and Wohlfart, J. (2023). Designing information provision experiments. *Journal of Economic Literature*, 61(1):3–40.
- Henry, E., Zhuravskaya, E., and Guriev, S. (2022). Checking and sharing alt-facts. *American Economic Journal: Economic Policy*, 14(3):55–86.

- Hillman, A. L. (2010). Expressive behavior in economics and politics. *European Journal of political economy*, 26(4):403–418.
- Jin, G. Z., Luca, M., and Martin, D. (2021). Is no news (perceived as) bad news? an experimental investigation of information disclosure. *American Economic Journal: Microeconomics*, 13(2):141–173.
- Jung, W.-O. and Kwon, Y. K. (1988). Disclosure when the market is unsure of information endowment of managers. *Journal of Accounting research*, pages 146–153.
- Kamenica, E. and Gentzkow, M. (2011). Bayesian persuasion. *American Economic Review*, 101(6):2590–2615.
- Kuziemko, I., Norton, M. I., Saez, E., and Stantcheva, S. (2015). How elastic are preferences for redistribution? evidence from randomized survey experiments. *American Economic Review*, 105(4):1478–1508.
- Lazer, D. M., Baum, M. A., Benkler, Y., Berinsky, A. J., Greenhill, K. M., Menczer, F., Metzger, M. J., Nyhan, B., Pennycook, G., Rothschild, D., et al. (2018). The science of fake news. *Science*, 359(6380):1094–1096.
- Light, M. T., He, J., and Robey, J. P. (2020). Comparing crime rates between undocumented immigrants, legal immigrants, and native-born us citizens in texas. *Proceedings of the National Academy of Sciences*, 117(51):32340–32347.
- Loosen, W., von Garmissen, A., Bartelt, E., and van Olphen, T. (2023). Journalismus in Deutschland 2023: Aktuelle Befunde zu Situation und Wandel. Arbeitspapiere des Hans-Bredow-Instituts 68, Hans-Bredow-Institut, Hamburg.
- Louis-Sidois, C. (2025). Both judge and party? investigating the political unbiasedness of fact-checkers. *Journal of the European Economic Association*, 23(6):2137–2164.
- Manili, J. (2024). Disclosure to an heterogeneous audience: The role of motivated beliefs. Working paper, Northwestern University. SSRN 4696951.
- Marie, A., Altay, S., and Strickland, B. (2023). Moralization and extremism robustly amplify myside sharing. *PNAS Nexus*, 2(4):pgad078.
- Marie, A. and Petersen, M. B. (2025). Motivations to connect with like-minded audiences increase partisan sharing on social media. *PNAS nexus*, 4(7):pgaf197.

- Martin, G. J. and McCrain, J. (2019). Local news and national politics. *American Political Science Review*, 113(2):372–384.
- Milgrom, P. R. (1981). Good news and bad news: Representation theorems and applications. *The Bell Journal of Economics*, pages 380–391.
- Möbius, M. M., Niederle, M., Niehaus, P., and Rosenblat, T. S. (2021). Managing self-confidence: Theory and experimental evidence. Working paper, October 2021.
- Mullainathan, S. and Shleifer, A. (2005). The market for news. *American Economic Review*, 95(4):1031–1053.
- National Academies of Sciences, Engineering, and Medicine (2017). *The Economic and Fiscal Consequences of Immigration*. The National Academies Press, Washington, DC.
- Osmundsen, M., Bor, A., Vahlstrup, P. B., Bechmann, A., and Petersen, M. B. (2021). Partisan polarization is the primary psychological motivation behind political fake news sharing on twitter. *American political science review*, 115(3):999–1015.
- Pennycook, G., Epstein, Z., Mosleh, M., Arechar, A. A., Eckles, D., and Rand, D. G. (2021). Shifting attention to accuracy can reduce misinformation online. *Nature*, 592:590–595.
- Peterson, E. and Iyengar, S. (2021). Partisan gaps in political information and information-seeking behavior: Motivated reasoning or cheerleading? *American Journal of Political Science*, 65(1):133–147.
- Pew Research Center (2024). How americans get news on TikTok, X, Facebook and Instagram. Report, June 12, 2024. <https://www.pewresearch.org/journalism/2024/06/12/>.
- Pew Research Center (2025). 42% of social media users say the sites are important for them getting involved with political or social issues. Report, September 16, 2025; survey of 5,097 U.S. adults, Feb. 10–17, 2025. <https://www.pewresearch.org/short-reads/2025/09/16/>.
- Prat, A. and Strömberg, D. (2013). The political economy of mass media. *Advances in economics and econometrics*, 2:135.
- Shin, J. and Thorson, K. (2017). Partisan selective sharing: The biased diffusion of fact-checking messages on social media. *Journal of Communication*, 67(2):233–255.
- Steindl, N., Lauerer, C., and Hanitzsch, T. (2017). Journalismus in Deutschland: Aktuelle Befunde zu Kontinuität und Wandel im deutschen Journalismus. *Publizistik*, 62(4):401–423.

- Thaler, M. (2021). The supply of motivated beliefs. *arXiv preprint arXiv:2111.06062*.
- Thaler, M. (2024). The fake news effect: Experimentally identifying motivated reasoning using trust in news. *American Economic Journal: Microeconomics*, 16(2):1–38.
- Verrecchia, R. E. (1983). Discretionary disclosure. *Journal of Accounting and Economics*, 5:179–194.
- Vosoughi, S., Roy, D., and Aral, S. (2018). The spread of true and false news online. *Science*, 359(6380):1146–1151.
- Zhu, Y. and Dukes, A. (2015). Selective reporting of factual content by commercial media. *Journal of Marketing Research*, 52(1):56–76.
- Zimmermann, F. (2020). The dynamics of motivated beliefs. *American Economic Review*, 110(2):337–361.

A Experimental Design Details

A.1 Statistic Selection and Validation Survey

I selected statistics for the main experiment from a larger pool of candidates using a pre-registered validation survey ($N = 100$, AsPredicted No. 232064) conducted on Prolific between June 10 and June 21, 2025. The survey evaluated eight candidate statistics—four hypothesized as pro-immigration and four as anti-immigration—against three survey-measured criteria (perceived credibility, policy relevance, and perceived political direction) and one researcher-verified criterion (veracity). The selection rule was pre-specified: within each direction, select the two statistics with the highest policy relevance among those with adequate perceived credibility, subject to the constraint that the final set contains two pro-immigration and two anti-immigration statistics.

Candidate Statistics and Selection Criteria. Table A.1 reports the eight candidate statistics with their survey-measured scores. The four criteria were operationalized as follows:

Veracity was established by the researcher through source verification. Each selected statistic is supported by multiple independent, authoritative sources: peer-reviewed academic publications, federal statistical agencies, or nonpartisan research institutions. Source documentation is provided below.

Policy relevance was measured by asking respondents: “How relevant do others consider the statistic when forming immigration policy positions?” (1 = completely irrelevant, 5 = extremely relevant).

Perceived credibility was measured by asking respondents: “How credible is the statistic?” (1 = not at all, 5 = entirely).

Perceived political direction was measured through two items: *usability* (“For whom would the statistic be most useful in a policy debate?”; 1 = exclusively supporting expansive immigration, 5 = exclusively supporting restrictive immigration) and *opinionability* (“How does learning about the statistic change someone’s immigration policy

views?"; 1 = convince to support expansion, 5 = convince to support limitation). These directional measures serve the classification of statistics as pro- or anti-immigration and the balance criterion.

Table A.1: Candidate Statistics: Survey-Measured Indicators

Statistic	Policy	Perceived	Perceived Direction	
	Relevance	Credibility	Usability	Opinionability
<i>Pro-immigration</i>				
Crime Rate [†]	3.74 (1.18)	3.50 (0.97)	2.23 (1.06)	2.36 (1.11)
Labor [†]	3.64 (1.04)	3.81 (1.02)	2.40 (1.09)	2.57 (1.23)
Entrepreneurship	3.12 (1.08)	3.54 (1.00)	2.09 (1.05)	2.18 (1.15)
Healthcare	3.28 (1.03)	3.59 (1.07)	2.03 (1.05)	2.19 (1.15)
<i>Anti-immigration</i>				
Border Crossings [†]	3.85 (1.14)	3.97 (0.90)	3.94 (1.14)	4.08 (1.11)
Fiscal Impact [†]	3.79 (1.01)	3.55 (0.93)	3.68 (1.15)	3.73 (1.01)
Poverty Rate	3.41 (1.01)	3.93 (0.93)	3.40 (1.13)	3.52 (1.11)
Wage Inequality	3.35 (0.95)	3.53 (0.90)	3.41 (1.06)	3.43 (0.99)

Notes: $N = 100$. Cells report means with standard deviations in parentheses. Selected statistics with [†]. Policy relevance: 1 = completely irrelevant, 5 = extremely relevant. Perceived credibility: 1 = not at all credible, 5 = entirely credible. Usability and opinionability: 1 = supports expansive immigration, 5 = supports restrictive immigration. One-sample t -tests against the neutral midpoint of 3 reject the null for all eight statistics on both directional measures ($p < 0.001$ for all selected statistics; see text).

Selection. The selection rule applied within each direction independently. Among the four pro-immigration candidates, Crime Rate (3.74) and Labor Force Participation (3.64) had the highest policy relevance scores, both with adequate perceived credibility (≥ 3.50). Among the four anti-immigration candidates, Border Crossings (3.85) and Fiscal Impact (3.79) had the highest policy relevance, again with adequate credibility. This rule uniquely determines the four selected statistics reported in the main text (Section 2.1). The one rejected statistic with notably high credibility—Poverty Rate (3.93)—scored lower on policy relevance (3.41) than both selected anti-immigration statistics,

confirming that the selection was not driven by credibility alone.

Perceived directional scores did not vary significantly by respondents' political affiliation (Democrat, Republican, Independent), indicating that the directional classification reflects consensus perception rather than partisan projection.

Direction Validation. I classify Crime Rate and Labor Force Participation as pro-immigration and Border Crossings and Fiscal Impact as anti-immigration. This classification is validated by the survey's directional measures. For usability and opinionability respectively, the pro-immigration statistics score below the neutral midpoint of 3 (Crime: 2.23, 2.36; Labor: 2.40, 2.57) and the anti-immigration statistics score above it (Border: 3.94, 4.08; Fiscal: 3.68, 3.73). One-sample t -tests against the midpoint of 3 reject the null of no directional content for all four selected statistics on both measures: Crime ($t = -7.25$ and $t = -5.74$), Labor ($t = -5.50$ and $t = -3.49$), Border Crossings ($t = 8.21$ and $t = 9.75$), and Fiscal Impact ($t = 5.89$ and $t = 7.20$), all $p < 0.001$.

This classification is independently confirmed by the main experiment. In the expected attitude shift elicitation (Section 2.3.1, Table B.7), senders in the main experiment ($N = 1,200$) expect Crime and Labor to shift receiver preferences toward pro-immigration positions (25 and 17 pp, both $p < 0.001$) and Border and Fiscal to shift preferences toward anti-immigration positions (-34 and -32 pp, both $p < 0.001$). The convergence of two independent samples using two independent methods—survey-based directional perception in the validation study and incentivized attitude-shift predictions in the main experiment—makes the directional classification robust.

Source Documentation. Veracity was established through source verification rather than survey measurement. Each selected statistic is supported by multiple independent, authoritative sources. I report the specific values shown to participants, the primary source for each, and corroborating evidence. All sources were available at the time of the experiment (July–August 2025).

Crime Rate (immigrants have lower crime rates than U.S.-born Americans). Participants were shown violent crime *arrest* rates per 100,000 residents in Texas (2012–2018):

213.0 for U.S.-born citizens, 185.3 for legal immigrants (13% lower), and 96.2 for undocumented immigrants (55% lower). These figures are from [Light et al. \(2020\)](#), who analyze the universe of felony arrests recorded in the Texas Computerized Criminal History database cross-referenced with Department of Homeland Security immigration status checks.³⁶ Other studies use Texas Department of Public Safety conviction data over a longer period (2013–2022) and reports homicide conviction rates of 3.0 per 100,000 for native-born Americans, 2.2 for unauthorized immigrants (26% lower), and 1.2 for legal immigrants (61% lower).³⁷ At the national level, [Abramitzky et al. \(2024\)](#) analyze census-linked incarceration records from 1870–2020 and find that immigrants have been consistently less likely to be incarcerated than the U.S.-born population throughout this period.³⁸ Texas is the only U.S. state that systematically records the immigration status of all arrestees, making it the sole source of population-level crime-rate comparisons by immigration status. The national-level incarceration evidence in [Abramitzky et al. \(2024\)](#) suggests the Texas pattern is not anomalous.

Labor Force Participation (immigrants have higher labor force participation than U.S.-born Americans). Participants were shown 2023 annual average labor force participation rates: 66.6% for the foreign-born population and 61.8% for the native-born population (4.8 percentage points higher). These figures are from Table 1 of the Bureau of Labor Statistics annual report *Foreign-Born Workers: Labor Force Characteristics—2023*.³⁹ The same data are available through the Federal Reserve Bank of St. Louis FRED database, Current Population Survey Table A-7.⁴⁰ The foreign-born participation rate has exceeded the native-born rate in every year of the BLS series.

Border Crossings (unauthorized border crossings have been higher in the past five years relative to historical averages). Participants were shown two figures computed

³⁶Published in *Proceedings of the National Academy of Sciences*; data and replication materials available at <https://www.openicpsr.org/openicpsr/project/124923>.

³⁷Cato Institute Policy Analysis no. 977; available at <https://www.cato.org/policy-analysis/illegal-immigrant-murderers-texas-2013-2022>.

³⁸Published in *American Economic Review: Insights*.

³⁹The annual report is available at <https://www.bls.gov/news.release/pdf/forbrn.pdf>. At the time of writing, the 2024 edition (released May 2025) reports 66.5% versus 61.7%, confirming the same pattern.

⁴⁰<https://fred.stlouisfed.org/release/tables?rid=50&eid=1736>

from U.S. Customs and Border Protection (CBP) enforcement statistics:⁴¹ a historical average of 890,435 apprehensions per year over fiscal years 1970–2019, and a recent five-year average of 1,584,573 encounters per year over fiscal years 2020–2024 (78% higher). The comparison requires combining two CBP reporting series: “apprehensions” (reported through FY2019) and “encounters” (reported from FY2020 onward, which combine apprehensions and expulsions under Title 42 and Title 8 processing).⁴² The statistic reflects data available at experiment time (July–August 2025). Subsequent enforcement changes under the second Trump administration reduced encounters to record lows in FY2025, as documented by Pew Research Center analysis of CBP data.⁴³ The FY2025 decline does not affect the validity of the statistic as presented to participants, which concerned the five-year average through FY2024.

Fiscal Impact (net fiscal costs of first-generation immigrants are higher than U.S.-born Americans at the state and local level). Participants were shown that first-generation immigrant households generate \$57.4 billion more in annual state and local government costs than they contribute in taxes, based on 2011–2013 data. This figure is from the National Academies of Sciences, Engineering, and Medicine comprehensive report on the economic and fiscal consequences of immigration, which remains the most authoritative and comprehensive fiscal analysis of immigration in the United States.⁴⁴ The report notes that this fiscal gap reverses at the federal level and across generations: the children and grandchildren of immigrants are among the strongest net fiscal contributors. I present only the state-and-local first-generation finding because it represents the anti-immigration direction in my balanced design. The Cato Institute provides complementary analysis of immigrants’ fiscal effects over a longer period (1994–2023), confirming that the fiscal picture varies substantially by level of government and

⁴¹CBP historical data available at <https://www.cbp.gov/newsroom/stats/cbp-enforcement-statistics>.

⁴²Beginning in March 2020, CBP shifted from reporting apprehensions to reporting “encounters,” a broader category. The historical series (1970–2019) uses the apprehension metric; the recent series (2020–2024) uses the encounter metric. Since encounters $\geq$ apprehensions by construction, the 78% comparison is an upper bound—i.e., it may slightly overstate the increase by comparing a narrower historical measure to a broader recent measure.

⁴³<https://www.pewresearch.org/short-reads/2026/02/02/migrant-encounters-at-the-us-mexico-border-ar>

⁴⁴National Academies of Sciences, Engineering, and Medicine (2017); the state and local fiscal impact analysis is in Chapter 8. Available at <https://www.nationalacademies.org/read/23550/chapter/14>.

generational horizon.⁴⁵

A.2 Belief Elicitation and Donation Task (Experiment 1)

This section documents two implementation details of the Experiment 1 sender protocol referenced in Section 2.1: the incentivization of the prior-knowledge elicitation and the revealed-preference donation task.

Prior-Knowledge Elicitation. Senders complete each comparative statement by choosing “higher,” “lower,” or “don’t know.” Correct completions receive a bonus, incorrect completions do not, and “don’t know” responses receive the bonus with 25% probability. The 25% payoff for “don’t know” is strictly dominated under expected utility maximization with binary alternatives, where uninformed guessing yields a 50% expected bonus. The option may nonetheless be selected by highly ambiguity-averse participants who prefer an objective 25% lottery over an ambiguous 50% guess. This implies that reported “don’t know” rates likely understate genuine ignorance among expected-utility maximizers.

Donation Task. Senders allocate \$500 between the *American Immigration Council* (AIC), which advocates for increased immigration, and *NumbersUSA*, which advocates for decreased immigration—both well-established organizations representing opposing positions in the immigration debate. I briefly describe the stated goals of each organization and present their missions in their own words, with the display order randomized. Participants are informed that one donation decision from the experiment will be randomly selected and implemented, and they receive a verification link to confirm the donation was placed.

⁴⁵<https://www.cato.org/white-paper/immigrants-recent-effects-government-budgets-1994-2023>

A.3 Fact Selection and Pre-Validation Survey — Experiments 2 and 3

I selected facts for Experiments 2 and 3 from a larger candidate pool using a pre-registered pre-validation survey (AsPredicted #290989, quota target $N = 330$, final analysis sample $N \approx 300$ after speeder and consent exclusions) conducted on the Bilendi panel in Germany in May 2026. The survey evaluated 36 candidate facts—six hypothesized as pro-policy and six as anti-policy in each of three domains (immigration, climate, redistribution)—against four criteria, and additionally validated the domain-specific advocacy organizations used in the ACTION treatment.

Criteria and Selection Rule. Four criteria were applied to each candidate fact. *Prior knowledge* was elicited via an incentivized fill-in-the-blank completion of the fact statement with a confidence slider [50, 100]; the probabilistic prior knowledge score is $PK = \text{confidence}/100$ for correct completions and $1 - \text{confidence}/100$ for incorrect completions. *Directional consistency* was measured by a 5-point rating (1 = strongly pro-policy, 5 = strongly anti-policy); a fact passes if the share of respondents rating it in its assigned direction exceeds the share rating it in the opposite direction. *Policy relevance* and *perceived credibility* were each measured on a 1–5 scale.

The pre-specified selection rule selected the top three facts by policy relevance within each topic $\times$ direction cell among those passing directional consistency, subject to balance constraints across the pro and anti sets: $|\bar{c}_{\text{pro}} - \bar{c}_{\text{anti}}| \leq 1.0$ for credibility, $|\overline{PK}_{\text{pro}} - \overline{PK}_{\text{anti}}| \leq 0.25$, and $|\bar{r}_{\text{pro}} - \bar{r}_{\text{anti}}| \leq 1.0$ for policy relevance.

The ACTION treatment organizations were additionally validated against two criteria: (i) at least 50% of respondents correctly recognize each organization as representing the pro-policy or anti-policy position respectively, and (ii) a positive Spearman correlation between topic-specific attitude and hypothetical donation to the pro-organization, restricted to respondents with directional attitudes.

Domain Selection. The redistribution domain failed on both grounds and is excluded. All six redistribution anti-direction candidate facts failed the directional consistency criterion, and the redistribution anti-organization had a recognition rate of 33.3%, well

below the 50% threshold. Immigration passed all fact and organization criteria (anti-organization recognition: 93.6%; Spearman $r = 0.57$). Climate's facts all passed directionality and the organization correlation criterion ($r = 0.57$); the climate anti-organization fell marginally below the recognition threshold (42.6% vs. 50%). Since climate's failure was marginal and isolated to one metric while redistribution's failures were compounding, climate was retained and this deviation is noted here.

Deviations from Algorithmic Fact Selection. *Immigration.* The algorithm selected *Bildungsabschluss* (PK = 0.81), which produced a pro-anti PK gap of 0.29, exceeding the 0.25 balance threshold. It was replaced by *Illegale Aufenthalte* (PK = 0.51), closing the gap to 0.15 while satisfying all other balance constraints. The label of one anti-fact was updated from *Arbeitslosigkeit* to *Bürgergeld* to reflect the 2026 German policy context; the underlying statistic is unchanged.

Climate. The pro set was revised to *Erneuerbare Energien*, *Stromgestehungskosten*, and *Grüne Beschäftigung* for higher directional consistency. *Netzausbaukosten* (anti) was reformulated between the pre-validation survey and the main experiment; directional ratings in Table A.2 reflect the pre-reformulation version, which failed the consistency criterion (17.5% consistent vs. 42.1% inconsistent). The reformulated version is what was fielded in the main experiment.

Selected Facts and Pre-Validation Ratings. Table A.2 reports the 12 selected facts with their pre-validation survey scores. Table A.3 reports the balance check.

Table A.2: Selected Facts: Pre-Validation Ratings (Experiments 2 and 3)

Fact	Dir.	Relevance	Credibility	PK	% Consistent
<i>Immigration</i>					
Fachkräftemangel	Pro	3.79	3.65	0.52	45.8
Beschäftigungswachstum	Pro	3.43	3.45	0.55	38.8
Integrationskurs	Pro	3.61	3.29	0.37	32.1
Kriminalität	Anti	4.04	3.76	0.70	90.0
Bürgergeld [‡]	Anti	3.98	3.74	0.67	78.7
Illegale Aufenthalte	Anti	3.69	3.45	0.51	76.5
<i>Climate policy</i>					
Erneuerbare Energien	Pro	3.67	3.61	0.52	51.0
Stromgestehungskosten	Pro	3.46	3.51	0.56	61.4
Grüne Beschäftigung	Pro	3.28	3.46	0.57	47.8
Netzausbaukosten [†]	Anti	3.84	3.53	0.72	17.5
Energiekosten Industrie	Anti	4.09	3.60	0.81	35.8
Preisvolatilität	Anti	3.73	3.36	0.66	33.9

Notes: $N \approx 300$. Policy relevance and perceived credibility: 1 = lowest, 5 = highest. PK is the confidence-weighted prior knowledge probability (see text). % Consistent is the share of respondents whose directional rating matches the assigned fact direction. [†] Reformulated after pre-validation; ratings are pre-reformulation and failed the consistency criterion. [‡] Relabeled from *Arbeitslosigkeit*; underlying statistic unchanged.

Table A.3: Balance Check: Selected Facts by Domain (Experiments 2 and 3)

Domain	Δ Relevance	Δ Credibility	Δ PK	Passes
Immigration	0.29	0.19	0.15	Yes
Climate	0.42	0.03	0.18	Yes

Notes: Δ = absolute difference between mean pro-direction and mean anti-direction facts. Pre-specified thresholds: Δ Relevance ≤ 1.0 , Δ Credibility ≤ 1.0 , Δ PK ≤ 0.25 . Climate Δ PK uses pre-reformulation statistics for *Netzausbaukosten*.

A.4 Information Condition Implementation

Part 2 of the experiment proceeds in three stages: study presentation and elicitation (all senders), information feedback (INFO = 1 only), and expected attitude change elicitation (all senders). I describe each stage and reproduce the text shown to participants.

Stage 1: Study Presentation and Elicitation (All Senders). All senders read the following description of the [Haaland and Roth \(2020\)](#) study:

A recent study published in the prestigious academic journal, *Journal of Public Economics*, examined how people form political opinions. The researchers surveyed more than 3,000 Americans that are representative of the US population. People were told about a 1980 event when 55,000 Cuban immigrants unexpectedly arrived in Miami, increasing the city's workforce by 8% almost overnight. They were asked how this affected the wages and employment of U.S.-born workers. Then, half of the participants learned that research found this wave of immigration had virtually no effect on the wages or unemployment for U.S.-born workers, because the new immigrants increased demand for goods and services, which created more jobs. The other half of the participants received no information.

Senders then make two incentivized predictions:

1. *Behavioral effect (continuous)*: "By what percentage did learning about the research findings change people's petition signing rates?" Senders enter a percentage (positive or negative).
2. *Differential updating (categorical)*: "After learning about the research findings, who do you think changed their attitudes toward immigration more?" Senders choose among: (a) people who initially believed immigration harmed U.S.-born workers,

- (b) people who initially believed immigration did not harm U.S.-born workers, or
- (c) both groups changed about the same amount.

Each question is equally likely to be selected for the bonus payment. For the continuous estimate, the probability of receiving the bonus is 100% if the estimate equals the true value (69%) and decreases linearly to 0% at a deviation of 20 percentage points or more.⁴⁶ For the categorical estimate, the probability is 100% if correct and 0% otherwise. Senders are instructed not to search for the answers and tab-switching is monitored to enforce compliance; senders who switched tabs during this stage were excluded as pre-registered.

Stage 2: Information Feedback (INFO = 1 Only). Senders in the INFO condition receive personalized feedback comparing their predictions to the actual results. Senders in the no-INFO condition proceed directly to Stage 3. The feedback screen displays the following:

Petition Signatures

The percentage of people who signed the White House petition actually increased by **69%**.

You estimated that the percentage of people who signed the White House petition changed by *[participant's estimate]%*.

That means *you [underestimated/overestimated] the change by [participant-specific gap] percentage points.*

Effect on Attitudes towards Immigration

Actual Result: People who initially believed immigration harmed U.S.-born workers changed their attitudes towards immigration **significantly more** than those who already held correct beliefs.

Your Answer: *[participant's answer]*.

The feedback is accompanied by a bar chart that displays the sender's estimate and

⁴⁶This implements a piecewise-linear proper scoring rule that is incentive-compatible under risk neutrality and approximately so under moderate risk aversion.

the true value (69%) side by side, with bar heights scaled proportionally. The sender’s categorical prediction is marked as correct or incorrect with a color-coded indicator. A summary reiterates the two key findings: (i) receiving information strongly changed behavior and (ii) the effect was strongest for participants who initially held false beliefs.⁴⁷

Stage 3: Expected Attitude Change Elicitation (All Senders). All senders then predict how receiving each of the four immigration statistics would shift the attitudes of hypothetical receivers. For each statistic j , senders allocate 10 hypothetical participants—described as holding the receiver belief state realized for that statistic (false belief or “don’t know”)—among three categories: those who become *more supportive* of immigration, those who become *less supportive* of immigration, and those whose opinions *do not change*. The allocations must sum to 10. This procedure generates the expected attitude change measure $\Delta\text{Attitude}_{i,j}$ defined in Section 2.3.1, which provides within-sender variation in perceived persuasiveness across statistics, and between-sender variation across receiver belief states.

You will estimate how ten participants who hold **false beliefs** or **do not know the answer** would change their attitudes about immigration if they **are provided with research findings**.

In particular, you will estimate how many of them. . .

- become less supportive of immigration.
- don’t change their opinion.
- become more supportive of immigration.

The interface displays 10 person icons to make the allocation task concrete. Senders complete this task four times: once for each of the four statistics, with the hypothetical receivers described by the belief state of the receiver group matched to that statistic. The order of statistics is randomized across senders.

⁴⁷The gap between the participant’s estimate and the true value is computed individually and displayed dynamically in the experimental interface. Senders who estimated the effect exactly see a confirmation message rather than a deviation.

A.5 Experimental Instructions and Interface

Available upon request.

Preliminary

B Experiment 1 Supplementary Tables and Figures

B.1 Sample Characteristics

Table B.4: Sample characteristics and treatment balance

	Population (1)	Sample (2)	BASE (3)	ACTION (4)	INFO (5)	ACTION $\times$ INFO (6)
<i>Gender</i>						
Female	0.509	0.502	0.507	0.490	0.503	0.510
Male	0.491	0.478	0.470	0.477	0.477	0.487
<i>Ethnicity</i>						
White	0.578	0.587	0.577	0.590	0.587	0.593
Black	0.121	0.129	0.137	0.120	0.140	0.120
Hispanic	0.187	0.182	0.180	0.183	0.177	0.187
Asian	0.059	0.058	0.063	0.063	0.057	0.050
Other	0.055	0.044	0.043	0.043	0.040	0.050
<i>Political Affiliation</i>						
Democrat	0.386	0.388	0.377	0.393	0.393	0.387
Independent	0.310	0.308	0.323	0.303	0.303	0.300
Republican	0.304	0.298	0.297	0.297	0.300	0.300

Sample: $N = 1,200$ U.S. adults recruited on Prolific (quota-sampled), Experiment 1. Column (1): 2020 U.S. Census population proportions. Column (2): full sample proportions. Columns (3)–(6): proportions by treatment cell (BASE = no-Action, no-Info; ACTION = Action only; INFO = Info only; ACTION $\times$ I = Action and Info). Population benchmarks from the 2020 U.S. Census.

B.2 Additional Summary Statistics

Table B.5: Prior knowledge of the four immigration statistics

	Share of Responses		
	Correct	False	Don't know
Crime	0.700	0.223	0.076
Labor	0.794	0.141	0.065
Border	0.612	0.317	0.071
Fiscal	0.304	0.553	0.143

Sample: $N = 1,200$ senders, Experiment 1. Elicited before senders were informed they would make provision decisions. *Correct:* stated belief matches the verified fact direction; *False:* stated belief contradicts the fact. *Don't know:* explicit non-answer (senders who select “Don't know” receive a 25% probability of being scored as correct, incentivizing truthful reporting of ignorance).




Figure B.2: Distribution of senders donation allocations to the pro-immigration advocacy group ($N = 1,200$, Experiment 1). Horizontal axis: allocation to the pro-immigration group collapsed into five categories ("All pro = \$500, "Pro-leaning = \$251–499, "Equal-split = \$250, "Anti-leaning = \$1–249, "All anti = \$0) from a total \$500 stake. Vertical axis: proportion of senders. The incentive-compatible preference measure is $\text{Pref}_i = (AIC_i - 250)/250 \in [-1, 1]$; senders with $AIC_i = 250$ ("Equal-split") are excluded from alignment-based analyses ($N = 86, 7.2\%$).




Figure B.1: Senders' beliefs about the persuasive effect of information (pre-treatment). *Notes:* The green line marks the true effect of 69% from [Haaland and Roth \(2020\)](#).

B.3 Robustness of Main Results

Logistic Regression. Table B.10 re-estimates specification (2) using logistic regression, which is the correctly specified model if the latent propensity to disclose follows a logistic distribution—as implied by the cumulative logistic F assumed in Section 3.2. All results replicate: the $\text{Align} \times \text{ACTION}$ interaction is positive and significant across all

Table B.6: Perceived Credibility of Statistics by Political Affiliation

	Percentage Perceiving Statistic as Credible			
	All	Democrats	Independents	Republicans
Crime	0.962	0.996	0.949	0.933
Labor	0.978	0.998	0.965	0.964
Border	0.978	0.985	0.976	0.972
Fiscal	0.968	0.981	0.962	0.958

Sample: $N = 1,200$ senders, Experiment 1. Credibility assessed post-provision (potentially endogenous to the provision decision). Values shown are the share of senders rating each statistic as credible (1 = credible, 0 = not credible).

Table B.7: Expected effects of information on political preferences

Statistic	Expect Effect of Information on Political Preferences			
	Crime (1)	Labor (2)	Border (3)	Fiscal (4)
Constant	0.248*** (0.016)	0.168*** (0.016)	-0.338*** (0.018)	-0.322*** (0.020)
INFO	0.058*** (0.019)	0.038** (0.019)	-0.033 (0.021)	0.031 (0.023)
Estimated Effect Size	0.000 (0.000)	0.001*** (0.000)	0.001*** (0.000)	0.002*** (0.000)
Clustered SE	No	No	No	No
Controls	No	No	No	No
R ²	0.009	0.018	0.017	0.022
Adj. R ²	0.007	0.016	0.015	0.020
Num. obs.	1200	1200	1200	1200

Sample: $N = 1,200$ senders, Experiment 1. *Outcome:* expected effect of each statistic on receivers' immigration policy preferences, $[-1, 1]$ (positive = pro-immigration shift). *Info:* binary indicator for the between-subject feedback treatment. *Estimated Effect Size:* sender's incentivized estimate of the treatment effect magnitude in Haaland and Roth (2020), included as a covariate. Standard errors are not clustered. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

specifications ($p < 0.05$), and the Align $\times$ INFO interaction remains null. I report the linear probability model in the main text for interpretability of coefficients as percentage-point effects.

Table B.8: Expected effects of information on political preferences with controls

Statistic	Expect Effect of Information on Political Preferences			
	Crime (1)	Labor (2)	Border (3)	Fiscal (4)
<i>Panel (A): INFO</i>				
Constant	0.236*** (0.067)	0.033 (0.071)	−0.336*** (0.073)	−0.259*** (0.080)
INFO	0.055*** (0.019)	0.035* (0.019)	−0.034 (0.021)	0.028 (0.023)
Estimated Effect Size	0.000 (0.000)	0.001*** (0.000)	0.001*** (0.000)	0.001*** (0.000)
Belief Accuracy	0.116*** (0.023)	0.133*** (0.027)	0.003 (0.024)	0.129*** (0.026)
<i>Panel (B): Receiver Belief State</i>				
Constant	0.249*** (0.068)	0.043 (0.071)	−0.347*** (0.074)	−0.255*** (0.080)
Receiver Belief False	0.014 (0.019)	0.006 (0.019)	−0.000 (0.021)	0.018 (0.023)
Estimated Effect Size	0.000 (0.000)	0.001*** (0.000)	0.001*** (0.000)	0.001*** (0.000)
Belief Accuracy	0.116*** (0.023)	0.134*** (0.027)	0.001 (0.024)	0.129*** (0.026)
Clustered SE	No	No	No	No
Controls	Yes	Yes	Yes	Yes
<i>Panel (A): R²</i>	0.049	0.054	0.047	0.076
<i>Panel (B): R²</i>	0.042	0.051	0.044	0.075
<i>Panel (A) Adj. R²</i>	0.031	0.036	0.029	0.058
<i>Panel (B) Adj. R²</i>	0.025	0.033	0.027	0.058
Num. obs.	1200	1200	1200	1200

Sample: $N = 1,200$ senders, Experiment 1. *Outcome:* expected effect of each statistic on receivers' immigration policy preferences, $[-1, 1]$ (positive = pro-immigration shift). *Panel (A):* regressor is the INFO indicator. *Panel (B):* regressor is Receiver Belief False (= 1 if the matched receiver holds an incorrect prior; = 0 if the receiver reports not knowing). *Estimated Effect Size:* sender's incentivized estimate of the Haaland & Roth (2020) treatment effect. *Belief Accuracy:* incentivized prior knowledge indicator. Controls: demographic covariates. Standard errors are not clustered. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Alternative Information Provision Specification. Table B.11 allows asymmetric INFO effects by fact direction to accommodate the asymmetric belief-shifting effect documented in Table B.7. The additional terms $\text{INFO}_i \times \text{PROSTAT}_j$ and $\text{Align}_{i,j} \times \text{INFO}_i \times \text{PROSTAT}_j$ are not statistically significant in the specifications with fact fixed effects. The core result— $\text{Align} \times \text{ACTION}$ —is unchanged.

Table B.9: Treatment Effects on Alignment Bias

	Alignment Bias	
	(1)	(2)
Intercept	0.067*** (0.021)	−0.026 (0.071)
<i>Treatment: A</i>	0.065** (0.029)	0.066** (0.029)
<i>Treatment: I</i>	0.008 (0.029)	−0.000 (0.029)
<i>Treatment: A × I</i>	0.029 (0.029)	0.024 (0.029)
Controls	No	Yes
R ²	0.005	0.048
Adj. R ²	0.003	0.029
Num. obs.	1114	1114

Sample: $N = 1,114$ senders with revealed directional policy preference ($AIC_i \neq 250$), Experiment 1. *Outcome:* sender-level alignment bias (mean aligned minus mean unaligned provision rate). *Reference:* BASE (no-Action, no-Info). *Treatment A:* Action only. *Treatment I:* Info only. *Treatment A × I:* Action and Info. Column (2) adds age, gender, education, ethnicity, and party identification. Standard errors are not clustered. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Leave-One-Statistic-Out Robustness. Table B.12 re-estimates the preferred specification (Table 2, column 3) four times, each time dropping one of the four statistics. Both coefficients are stable in magnitude across all four runs—Align stays within 0.031–0.043 (full sample: 0.037) and remains significant at $p < 0.05$ in every run; Align × ACTION stays within 0.039–0.044 (full sample: 0.041) and is significant at $p < 0.05$ in three of four runs and at $p < 0.10$ in the fourth. No single statistic drives either result.

B.4 Heterogeneous Treatment Effects

B.5 Mechanism-Consistent Evidence

Table B.10: Treatment Effects on Information Provision Logistic

	Probability of Statistic Provision		
	(1)	(2)	(3)
Intercept	1.967*** (0.148)	1.800*** (0.163)	2.120*** (0.502)
Align	0.460*** (0.133)	0.416*** (0.134)	0.430*** (0.137)
INFO	-0.126 (0.183)	-0.126 (0.184)	-0.126 (0.188)
ACTION	0.173 (0.195)	0.172 (0.196)	0.167 (0.200)
Belief Accuracy	0.399*** (0.103)	0.308*** (0.114)	0.317*** (0.113)
Align $\times$ INFO	0.005 (0.177)	-0.009 (0.177)	-0.010 (0.183)
Align $\times$ ACTION	0.535** (0.210)	0.538** (0.209)	0.553*** (0.214)
INFO $\times$ ACTION	-0.234 (0.259)	-0.233 (0.260)	-0.241 (0.268)
Align $\times$ INFO $\times$ ACTION	-0.434 (0.270)	-0.425 (0.270)	-0.433 (0.277)
Clustered SE	Subject	Subject	Subject
Statistics FE	No	Yes	Yes
Controls	No	No	Yes
No. of Senders	1200	1200	1200
AIC	3258.054	3247.312	3248.924
BIC	3316.341	3325.029	3449.691
Log Likelihood	-1620.027	-1611.656	-1593.462
Deviance	3240.054	3223.312	3186.924
Num. obs.	4800	4800	4800

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. Logistic regression; coefficients are log-odds (not percentage-point effects). Column structure matches Table 2: column (1) no fixed effects or controls; column (2) adds statistic fixed effects; column (3) adds statistic fixed effects and demographic controls. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.11: Treatment Effects on Information Provision

	Probability of Statistic Provision		
	(1)	(2)	(3)
Intercept	0.871*** (0.015)		
Align	0.041*** (0.012)	0.037*** (0.012)	0.037*** (0.013)
INFO	-0.034* (0.019)	-0.018 (0.020)	-0.018 (0.020)
ACTION	-0.002 (0.016)	-0.002 (0.016)	-0.003 (0.016)
Belief Accuracy	0.038*** (0.012)	0.035*** (0.012)	0.034*** (0.012)
Align $\times$ INFO	-0.015 (0.021)	-0.009 (0.021)	-0.011 (0.022)
Align $\times$ ACTION	0.041** (0.018)	0.041** (0.018)	0.041** (0.018)
INFO $\times$ ACTION	-0.007 (0.024)	-0.007 (0.024)	-0.008 (0.024)
Align $\times$ INFO $\times$ ACTION	-0.027 (0.026)	-0.027 (0.026)	-0.027 (0.026)
INFO $\times$ PROSTAT	0.038*** (0.015)	0.005 (0.020)	0.005 (0.020)
Align $\times$ INFO $\times$ PROSTAT	0.025 (0.023)	0.025 (0.023)	0.029 (0.025)
Clustered SE	Subject	Subject	Subject
Statistics FE	No	Yes	Yes
Controls	No	No	Yes
No. of Senders	1200	1200	1200
Num. obs.	4800	4800	4800
R ²	0.032	0.034	0.041
Adj. R ²	0.030	0.031	0.034

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. Extends Table 2 with asymmetric INFO effects by statistic direction. *ProStat*: binary indicator for pro-immigration statistics (Crime and Labor). *INFO* $\times$ *ProStat* tests whether the INFO treatment shifts provision differentially by statistic direction; *Align* $\times$ *INFO* $\times$ *ProStat* tests whether the alignment gradient shifts differentially. Column structure matches Table 2. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.12: Leave-One-Statistic-Out Robustness of the Alignment and Action-Amplification Effects — Experiment 1

Coefficient	Sign-positive	Min β	Max β	Max p
Align	4/4	0.031	0.043	0.028
Align $\times$ Action	4/4	0.039	0.044	0.068
Statistics FE	Yes			
Controls	Yes			
Clustered SE	Sender			

Sample: Experiment 1 senders. Each coefficient is re-estimated four times in the preferred specification (Table 2, column 3; statistic fixed effects, full demographic and political controls, sender-clustered standard errors), each time dropping one of the four statistics (crime, labor, border, fiscal). Columns report the number of leave-one-out estimates with a positive coefficient, the minimum and maximum coefficient across the four runs, and the largest p -value observed. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.13: Sender-level Heterogeneous Treatment Effects

Heterogeneity	Alignment Bias					
	Partisanship		Political Party		Issue Polarization	
	(1)	(2)	(3)	(4)	(5)	(6)
Intercept	−0.005 (0.036)	−0.105 (0.096)	0.132*** (0.032)	0.042 (0.095)	0.008 (0.045)	−0.102 (0.100)
A	0.118** (0.052)	0.120** (0.052)	0.030 (0.046)	0.028 (0.046)	−0.006 (0.067)	0.003 (0.068)
Partisanship	0.140** (0.058)	0.142** (0.059)				
A $\times$ Partisanship	−0.102 (0.082)	−0.101 (0.083)				
Independent			−0.132*** (0.048)	−0.130*** (0.048)		
Republican			−0.075 (0.050)	−0.069 (0.050)		
A $\times$ Independent			0.146** (0.068)	0.150** (0.069)		
A $\times$ Republican			−0.039 (0.069)	−0.025 (0.071)		
Alignment Intensity					0.088 (0.061)	0.104* (0.061)
A $\times$ Alignment Intensity					0.088 (0.087)	0.079 (0.088)
Controls	No	Yes	No	Yes	No	Yes
R ²	0.020	0.058	0.036	0.069	0.027	0.067
Adj. R ²	0.015	0.026	0.028	0.035	0.022	0.035
Num. obs.	554	554	554	554	554	554

Sample: $N = 554$ senders with revealed directional policy preference ($AIC_i \neq 250$), restricted to BASE and A treatment cells, Experiment 1. *Outcome:* alignment bias = sender's mean aligned provision rate minus mean unaligned provision rate. *Partisanship:* absolute deviation of self-reported political leaning from center (0 = pure Independent, 1 = strong partisan). *Political Party:* Democrat is the reference category; Independent and Republican are indicator contrasts. *Alignment Intensity:* $|AIC_i - 250|/250$, immigration-specific preference intensity. Controls (even columns): age, gender, education, ethnicity, party identification. Standard errors are not clustered. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.14: Sender-level heterogeneous treatment effects, pooling across INFO conditions

Heterogeneity	Alignment Bias					
	Partisanship		Political Party		Issue Polarization	
	(1)	(2)	(3)	(4)	(5)	(6)
Intercept	−0.003 (0.027)	−0.054 (0.068)	0.114*** (0.023)	0.062 (0.069)	0.008 (0.033)	−0.054 (0.070)
ACTION	0.073* (0.037)	0.070* (0.038)	0.027 (0.033)	0.030 (0.033)	−0.017 (0.049)	−0.018 (0.049)
Partisanship	0.137*** (0.041)	0.139*** (0.041)				
ACTION × Partisanship	−0.055 (0.058)	−0.046 (0.058)				
Independent			−0.086** (0.035)	−0.081** (0.035)		
Republican			−0.052 (0.036)	−0.039 (0.036)		
ACTION × Independent			0.088* (0.049)	0.085* (0.049)		
ACTION × Republican			−0.040 (0.050)	−0.044 (0.050)		
Alignment Intensity					0.093** (0.043)	0.107** (0.043)
ACTION × Alignment Intensity					0.075 (0.062)	0.078 (0.062)
Controls	No	Yes	No	Yes	No	Yes
R ²	0.017	0.042	0.017	0.038	0.020	0.047
Adj. R ²	0.015	0.026	0.013	0.021	0.018	0.031
Num. obs.	1114	1114	1114	1114	1114	1114

Sample: $N = 1,114$ senders with revealed directional policy preference, pooling all four treatment cells, Experiment 1. Outcome: alignment bias (sender-level). ACTION: binary indicator equal to 1 for senders in either A or A×I cell (both Action conditions). Moderator definitions, controls, and specification structure match Table B.13. Standard errors are not clustered. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.15: Receiver belief state: perceived effect of providing information

Statistic	Expected Effect of Information on Political Preferences			
	Crime (1)	Labor (2)	Border (3)	Fiscal (4)
<i>Panel (A): Receiver Belief State</i>				
Constant	0.265*** (0.016)	0.185*** (0.017)	−0.352*** (0.018)	−0.313*** (0.020)
Receiver Belief False	0.020 (0.019)	0.004 (0.019)	−0.003 (0.021)	0.011 (0.023)
Estimated Effect Size	0.000 (0.000)	0.001*** (0.000)	0.001*** (0.000)	0.002*** (0.000)
R ²	0.002	0.015	0.015	0.021
Adj. R ²	0.000	0.013	0.013	0.019
Num. obs.	1200	1200	1200	1200

*** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$. Panel (A): Dependent variable is the sender's expected effect of each statistic on receivers' immigration preferences (range $[-1, 1]$). Receiver Belief False equals one if the matched receiver holds a false belief, zero if the receiver reports not knowing. The null effect confirms that the Δb variation (receiver belief state) does not contaminate senders' perceived δ : the two channels are empirically separable.

Table B.16: Heterogeneous Treatment Effects of Receiver Beliefs

	Probability of Statistic Provision	
	(1)	(2)
Intercept	0.863*** (0.018)	0.894*** (0.045)
Align	0.026* (0.014)	0.025* (0.014)
INFO	-0.034 (0.021)	-0.034 (0.022)
ACTION	-0.021 (0.020)	-0.022 (0.021)
Receiver Belief False	-0.025 (0.018)	-0.024 (0.018)
Accuracy Prior Belief	0.035*** (0.012)	0.034*** (0.012)
Align × INFO	0.025 (0.023)	0.025 (0.023)
Align × ACTION	0.053** (0.024)	0.054** (0.024)
INFO × ACTION	0.014 (0.030)	0.014 (0.031)
Align × Receiver Belief False	0.019 (0.022)	0.020 (0.022)
INFO × Receiver Belief False	0.043 (0.027)	0.043 (0.027)
ACTION × Receiver Belief False	0.034 (0.026)	0.034 (0.026)
Align × INFO × ACTION	-0.050 (0.035)	-0.051 (0.035)
Align × INFO × Receiver Belief False	-0.042 (0.032)	-0.042 (0.032)
Align × ACTION × Receiver Belief False	-0.020 (0.033)	-0.022 (0.033)
INFO × ACTION × Receiver Belief False	-0.041 (0.038)	-0.043 (0.038)
Align × INFO × ACTION × Receiver Belief False	0.047 (0.045)	0.048 (0.045)
Clustered SE	Subject	Subject
Statistics FE	Yes	Yes
Controls	No	Yes
No. of Senders	1200	1200
R ²	0.034	0.041
Adj. R ²	0.030	0.033
Num. obs.	4800	4800

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. Outcome: binary provision indicator. Full four-way interaction: Align, ACTION, INFO, and Receiver Belief False (= 1 if the matched receiver holds a factually incorrect prior belief; = 0 if the receiver reports not knowing). Column (2) adds demographic controls. Summary statistics from this specification are reported in Appendix Table B.15, Panel (B). Statistic fixed effects and sender-clustered standard errors throughout. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.17: Individual-level Heterogeneous Treatment Effects

Heterogeneity	Probability of Statistic Provision			
	<i>Partisanship</i>		<i>Political Party</i>	
	(1)	(2)	(3)	(4)
Align	0.001 (0.016)	0.001 (0.016)	0.057*** (0.019)	0.057*** (0.019)
ACTION	0.026 (0.024)	0.026 (0.025)	0.018 (0.023)	0.014 (0.023)
Partisanship	-0.003 (0.026)	-0.001 (0.027)		
Accuracy Prior Belief	0.031** (0.012)	0.031** (0.012)	0.033*** (0.012)	0.034*** (0.012)
Align $\times$ ACTION	0.057** (0.027)	0.057** (0.027)	0.019 (0.025)	0.019 (0.025)
Align $\times$ Partisanship	0.065*** (0.023)	0.065*** (0.023)		
ACTION $\times$ Partisanship	-0.053 (0.036)	-0.054 (0.037)		
Align $\times$ ACTION $\times$ Partisanship	-0.029 (0.037)	-0.029 (0.037)		
Independent			0.009 (0.021)	0.005 (0.021)
Republican			0.005 (0.021)	0.001 (0.021)
Align $\times$ Independent			-0.043** (0.020)	-0.043** (0.020)
Align $\times$ Republican			-0.022 (0.023)	-0.022 (0.023)
ACTION $\times$ Independent			-0.034 (0.029)	-0.031 (0.029)
ACTION $\times$ Republican			-0.032 (0.029)	-0.027 (0.029)
Align $\times$ ACTION $\times$ Independent			0.061** (0.030)	0.061** (0.030)
Align $\times$ ACTION $\times$ Republican			0.009 (0.032)	0.009 (0.032)
Clustered SE	Subject	Subject	Subject	Subject
Statistics FE	Yes	Yes	Yes	Yes
Controls	No	Yes	No	Yes
R ²	0.037	0.042	0.035	0.040
Adj. R ²	0.034	0.036	0.032	0.033
Num. obs.	4800	4800	4800	4800

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. Outcome: binary provision indicator. Columns (1)–(2): three-way interaction Align $\times$ ACTION $\times$ Partisanship (Democrat reference for ACTION interaction). Columns (3)–(4): three-way interaction Align $\times$ ACTION $\times$ Political Party (Democrat is reference). Partisanship: absolute deviation of self-reported leaning from center. Statistic fixed effects and sender-clustered standard errors throughout. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.




Figure B.3: Provision rates by expected effect on receiver's immigration preferences, Experiment 1 ($N = 1,200$). Horizontal axis: *AlignPersuasion* = expected preference shift toward receiver's policy position, signed by sender direction, $\in [-1, 1]$. Vertical axis: share of statistics disclosed. Red dots: sender expects a preference-unaligned shift ($\delta < 0$); gray squares: no expected shift ($\delta = 0$); blue dots: sender expects a preference-aligned shift ($\delta > 0$). Dashed horizontal lines indicate the within-category mean provision rate (red for unaligned, blue for aligned category). Error bars are 95% confidence intervals.

Table B.18: Mechanisms

Mechanism Evidence	Probability of Statistic Provision			
	<i>Belief correction</i>		<i>Persuasiveness</i>	
	(1)	(2)	(3)	(4)
Intercept	0.600*** (0.062)	0.592*** (0.091)	0.581*** (0.062)	0.566*** (0.065)
Align	0.030** (0.013)	0.031** (0.012)		
INFO	-0.014 (0.017)	-0.014 (0.017)	-0.014 (0.017)	-0.014 (0.017)
ACTION	0.010 (0.023)	0.023 (0.122)	0.001 (0.016)	0.012 (0.033)
Belief Accuracy	0.032* (0.017)	0.025** (0.012)	0.039*** (0.012)	0.038*** (0.012)
Credibility	0.256*** (0.060)	0.269*** (0.090)	0.264*** (0.060)	0.260*** (0.060)
Align × INFO	0.007 (0.017)	0.007 (0.017)		
Align × ACTION	0.047** (0.019)	0.045** (0.018)		
INFO × ACTION	-0.008 (0.023)	-0.008 (0.023)	-0.008 (0.023)	-0.008 (0.024)
ACTION × Belief Accuracy	-0.014 (0.022)			
Align × INFO × ACTION	-0.032 (0.025)	-0.032 (0.026)		
ACTION × Credibility		-0.022 (0.122)		
AlignPersuasion			0.021 (0.016)	
AlignPersuasion × ACTION			0.042* (0.024)	
AlignPersuasion(-)				0.006 (0.025)
AlignPersuasion(+)				0.040* (0.024)
AlignPersuasion(-) × ACTION				-0.022 (0.036)
AlignPersuasion(+) × ACTION				-0.004 (0.034)
Clustered SE	Subject	Subject	Subject	Subject
Statistics FE	Yes	Yes	Yes	Yes
Controls	No	No	No	No
No. of Senders	1200	1200	1200	1200
R ²	0.051	0.051	0.041	0.041
Adj. R ²	0.048	0.048	0.039	0.039
Num. obs.	4800	4800	4800	4800

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. *Outcome:* binary provision indicator. Columns (1)–(2) test belief-correction channels: column (1) adds Action×Belief Accuracy to the main specification; column (2) replaces it with Action×Credibility. *Credibility:* sender's unincentivized posterior assessment of statistic credibility (binary), elicited post-provision. *Belief Accuracy:* incentivized prior knowledge indicator, elicited pre-treatment. Columns (3)–(4) replace Align with persuasion-based measures. *AlignPersuasion:* expected preference shift signed by sender direction, $[-1, 1]$. Column (4) decomposes into AlignPersuasion(+) and AlignPersuasion(-) binary indicators (zero = reference). All specifications include statistic fixed effects and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table B.19: Treatment Effects on Information Provision

Statement	Probability of Statistic Provision			
	<i>False factual beliefs should be corrected</i> (1)	<i>Facts are important for forming opinions</i> (2)	<i>Social media should fact check</i> (3)	<i>Disagreement stems from facts not views</i> (4)
Intercept	0.538*** (0.056)	0.415*** (0.083)	0.734*** (0.033)	0.898*** (0.027)
Align	0.197*** (0.052)	0.213** (0.088)	0.109*** (0.028)	0.007 (0.021)
INFO	-0.009 (0.016)	-0.007 (0.016)	-0.011 (0.017)	-0.014 (0.017)
ACTION	-0.003 (0.016)	-0.005 (0.016)	-0.004 (0.016)	-0.003 (0.016)
Accuracy Belief	0.034*** (0.012)	0.032*** (0.012)	0.035*** (0.012)	0.034*** (0.012)
Statement_Agree	0.069*** (0.011)	0.091*** (0.016)	0.026*** (0.006)	-0.013** (0.006)
Align × INFO	0.002 (0.017)	0.001 (0.017)	0.002 (0.017)	0.004 (0.017)
Align × ACTION	0.039** (0.018)	0.039** (0.018)	0.043** (0.018)	0.041** (0.018)
INFO × ACTION	-0.011 (0.023)	-0.010 (0.023)	-0.010 (0.024)	-0.006 (0.024)
Align × Statement_Agree	-0.036*** (0.011)	-0.036** (0.018)	-0.018*** (0.006)	0.009 (0.006)
Align × INFO × ACTION	-0.024 (0.026)	-0.023 (0.026)	-0.024 (0.026)	-0.027 (0.026)
Clustered SE	Subject	Subject	Subject	Subject
Statistics FE	Yes	Yes	Yes	Yes
Controls	Yes	Yes	Yes	Yes
No. of Senders	1200	1200	1200	1200
R ²	0.065	0.066	0.044	0.036
Adj. R ²	0.063	0.063	0.042	0.033
Num. obs.	4800	4800	4800	4800

Sample: $N = 1,200$ senders, 4,800 sender-statistic observations, Experiment 1. Outcome: binary provision indicator. Each column corresponds to a different pro-accuracy norm statement (column headers). *Statement_Agree*: sender's agreement with the column-header statement, measured post-provision on a five-point scale (1 = strongly disagree; 5 = strongly agree). *Align* × *Statement_Agree* tests whether norm endorsement moderates the alignment gradient. All specifications include statistic fixed effects, demographic controls, and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.




Figure B.4: Distribution of senders' stated decision motives across experiments, among senders whose open-ended decision reason refers to the disclosure decision (referral rates: Experiment 1 48.6%, Experiment 2 27.0%, Experiment 3 73.9%). Panel (a): share for whom each motive is the *primary* stated reason. Panel (b): share mentioning each motive at all (multiple motives may be coded per response). Motives coded from open-ended text by a large language model (claude-sonnet-5; Appendix E.1).

C Experiment 2 Supplementary Tables and Figures

C.1 Heterogeneity, Controls, and Provision Figures

I omit demographic controls from the main-text provision tables because the alignment effect is identified within sender and the ACTION effect is randomized. Adding the full set of sender demographics to the pre-registered tests and to the action decomposition leaves the coefficients stable (see Tables C.23 and C.24).

C.2 Robustness of the Motive-Channel Estimates

The action decomposition of Table 6 is robust to alternative inference, ordering, functional-form, and sample-restriction choices; all checks use the canonical positional persuasion measure $\Delta\pi_{ijk}$. Table C.26 re-estimates the preferred specification under two-way clustering by sender and fact and under an exact wild-cluster bootstrap over all $2^{12} = 4,096$ fact-level sign patterns. Table C.27 adds the within-sender presentation order and a lagged provision indicator. Table C.28 probes the shape of $w(\Delta\pi)$ by replacing $\Delta\pi_{ijk}$




Figure C.5: Mean provision rates by expected preference-aligned belief shift and arm, Experiment 2 (Citizens). Dark green: facts for which the sender expects the receiver’s policy position to move toward the sender’s own position ($\Delta\pi_{ijk} > 0$). Dark red: facts for which no such shift is expected ($\Delta\pi_{ijk} \leq 0$). x -axis: experimental arm (CONTROL or ACTION). Error bars are 95% confidence intervals (sender-level means). Sample: Experiment 2 senders with non-missing $\Delta\pi_{ijk}$ elicitation; $\Delta\pi_{ijk}$ is incentivized via Quadratic Scoring Rule.

Table C.20: Heterogeneity by Sender Profile: Crude Alignment and Action — Experiment 2 (Citizens)

Moderator Topic	Probability of Fact Provision				
	(1) Distance Pooled	(2) polarization Pooled	(3) Pol. group Pooled	(4) Immigration	(5) Climate
Alignment	0.024 (0.018)	0.023 (0.018)	0.032* (0.019)	0.013 (0.015)	0.034* (0.019)
Action	-0.006 (0.018)	0.039** (0.018)	0.005 (0.019)	0.027 (0.018)	0.025 (0.017)
Alignment × Action	0.020 (0.026)	0.003 (0.025)	-0.027 (0.028)	-0.023 (0.021)	0.016 (0.027)
Receiver distance	-0.095*** (0.007)				
Alignment × Distance	-0.005 (0.008)				
Action × Distance	0.019** (0.010)				
Alignment × Action × Distance	-0.013 (0.012)				
Sender polarization		0.010* (0.006)			
Alignment × polarization		0.001 (0.007)			
Action × polarization		-0.009 (0.009)			
Alignment × Action × polarization		-0.006 (0.011)			
Extremes (ref: center)			-0.046** (0.020)		
Others (ref: center)			-0.093*** (0.026)		
Alignment × Extremes			0.002 (0.026)		
Alignment × Others			-0.049 (0.035)		
Action × Extremes			0.023 (0.027)		
Action × Others			0.066* (0.037)		
Alignment × Action × Extremes			0.015 (0.036)		
Alignment × Action × Others			0.068 (0.053)		
Clustered SE	Sender	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes	Yes
Num. obs.	7740	7740	7715	4045	3695
R ²	0.050	0.003	0.005	0.002	0.004
Adj. R ²	0.048	0.001	0.003	0.000	0.001

Sample: $N = 1,548$ Experiment 2 citizens (Bilendi panel, Germany, 2026). Columns (1)–(2): 7,740 sender-fact observations. Column (3): 7,715 (non-missing political group). Columns (4)–(5): topic subsamples (immigration: $N = 809$ senders; climate: $N = 739$ senders). Outcome: binary provision indicator. Alignment: $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, 1]$ (pre-registered confirmatory measure). Receiver distance: $|\text{sender_likert} - \text{recv_stance}| \in \{0, 1, 2, 3, 4\}$. Sender polarization: absolute left-right distance from centre. Political group: Extremes = AfD, Linke, BSW; Others = other/no vote; Center = CDU/CSU, SPD, Grüne, FDP (reference). Columns (1)–(3) include the moderator and all interactions; columns (4)–(5) are subsample estimates with no moderator. All specifications include fact fixed effects and sender-clustered standard errors. Multiple testing: the exploratory heterogeneity family spans twenty interaction terms across this table and Table C.21; Benjamini-Hochberg-adjusted q -values within that family are approximately 0.20 for every interaction significant at raw $p < 0.10$. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.21: Heterogeneity by Sender Profile: Motive Decomposition — Experiment 2 (Citizens)

	Probability of Fact Provision		
	(1) Distance	(2) Polarization	(3) Political group
Expression	0.038 (0.033)	0.028 (0.031)	0.071* (0.037)
Persuasion	0.021 (0.014)	-0.014 (0.013)	-0.011 (0.013)
Action	-0.007 (0.018)	0.037** (0.018)	0.001 (0.019)
Persuasion $\times$ Action	-0.018 (0.020)	0.025 (0.018)	0.047** (0.018)
Receiver distance	-0.105*** (0.009)		
Expression $\times$ Distance	0.021 (0.017)		
Persuasion $\times$ Distance	-0.001 (0.007)		
Action $\times$ Distance	0.018* (0.010)		
Persuasion $\times$ Action $\times$ Distance	0.018** (0.009)		
Sender polarisation		-0.000 (0.008)	
Expression $\times$ Polarisation		0.027* (0.014)	
Persuasion $\times$ Polarisation		0.011* (0.006)	
Action $\times$ Polarisation		-0.009 (0.009)	
Persuasion $\times$ Action $\times$ Polarisation		-0.001 (0.008)	
Extremes (ref: center)			-0.065** (0.026)
Others (ref: center)			-0.072** (0.034)
Expression $\times$ Extremes			0.037 (0.048)
Expression $\times$ Others			-0.071 (0.061)
Persuasion $\times$ Extremes			0.014 (0.019)
Persuasion $\times$ Others			0.043* (0.026)
Action $\times$ Extremes			0.025 (0.027)
Action $\times$ Others			0.070* (0.037)
Persuasion $\times$ Action $\times$ Extremes			-0.028 (0.028)
Persuasion $\times$ Action $\times$ Others			-0.065* (0.038)
Clustered SE	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes
Spec	(1)	(2)	(3)
Num. obs.	7740	7740	7715
R ²	0.056	0.007	0.009
Adj. R ²	0.053	0.004	0.005

Sample: $N = 1,548$ Experiment 2 citizens (Bilendi panel, Germany, 2026); 7,740 sender-fact observations in columns (1)–(2), 7,715 in column (3) (non-missing political group). Outcome: binary provision indicator. Expression: cross-domain behavioral rating a_{ij} , elicited post-provision on a different topic. Persuasion: incentivized expected preference shift $\Delta\pi_{ijk}$ toward the sender's position. Receiver distance: $|\text{sender_likert} - \text{rcv_stance}| \in \{0, 1, 2, 3, 4\}$. Sender polarisation: absolute left-right distance from centre. Political group: Extremes = AfD, Linke, BSW; Others = other/no vote; Center = CDU/CSU, SPD, Grüne, FDP (reference). All specifications include fact fixed effects and sender-clustered standard errors. Multiple testing: the exploratory heterogeneity family spans twenty interaction terms across this table and Table C.20; Benjamini–Hochberg-adjusted q -values within that family are approximately 0.20 for every interaction significant at raw $p < 0.10$. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.22: Motive Decomposition by Topic — Experiment 2 (Citizens)

	Probability of Fact Provision	
	(1)	(2)
Expression	0.049 (0.033)	0.086*** (0.032)
Factual belief correction	0.053** (0.025)	0.023 (0.024)
Persuasion	0.002 (0.013)	0.003 (0.013)
Action	0.025 (0.018)	0.023 (0.017)
Persuasion $\times$ Action	0.037** (0.017)	0.010 (0.019)
Clustered SE	Sender	Sender
Fact FE	Yes	Yes
Topic	Immigration	Climate
Spec	(1)	(2)
Num. obs.	4045	3695
R ²	0.007	0.004
Adj. R ²	0.004	0.001

Sample: $N = 1,548$ Experiment 2 citizens (Bilendi panel, Germany, 2026). Column (1): immigration subsample ($N = 809$ senders). Column (2): climate subsample ($N = 739$ senders). *Outcome:* binary provision indicator. Each column applies the preferred motive specification (col. (4) of Table 6) within topic subsample: Expression (α_{ij} , cross-domain behavioral), Factual belief correction (Δb_{ijk} , incentivized), Persuasion ($\Delta \pi_{ijk}$, incentivized positional), and the Persuasion $\times$ ACTION interaction. All specifications include fact fixed effects and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

with its sign and with a version winsorized at the ninetieth percentile of $|\Delta\pi|$: the sign coding carries the Persuasion main effect ($\beta = +0.018$, $p < 0.05$) but not its ACTION interaction ($\beta = -0.005$). Table C.29 drops low- $|\Delta\pi|$ senders, perfect sharers, and low- α facts in turn; the instrumental interaction stays positive throughout, though it is not separately significant in these restricted subsamples under the positional measure. Table C.32 reports leave-one-fact-out estimates for each motive; the within-fact variance inflation factors are near unity ($VIF \approx 1.00$) for all three channels.

Table C.31 varies the sample definition. Column (2) relaxes the pre-registered speeder exclusion (completion time below $0.25 \times$ the realised-sample median). Column (3) adds the 106 participants who started before the pre-registration update of 27 May 2026, which dropped a non-essential elicitation module after the survey's length caused

Table C.23: Provision Decisions with Demographic Controls — Experiment 2 (Citizens)

	Probability of Fact Provision	
	(1)	(2)
Align	0.019* (0.011)	0.026** (0.012)
Action	0.021* (0.012)	0.023* (0.012)
Align $\times$ Action	-0.008 (0.016)	-0.010 (0.017)
Clustered SE	Sender	Sender
Fact FE	No	Yes
Controls	Yes	Yes
Num. obs.	7715	7715
R ²	0.061	0.016
Adj. R ²	0.056	0.009

Sample: $N = 1,548$ Experiment 2 citizens, both arms; 7,715 observations with non-missing controls. *Outcome:* binary provision indicator. *Align:* positional alignment $\text{Pref}_i \times \text{Dir}_j \in [-1, +1]$. *Controls:* sender-receiver ideological distance, age, gender, education, Bundesland, closest-party affiliation, and left-right self-placement. Column (1) omits fact fixed effects; column (2) includes them and reproduces the pre-registered preferred specification. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

elevated dropouts; they completed the longer instrument (median 29 vs. 17 minutes) and are excluded from the canonical sample for instrument comparability. Column (4) adds both groups. Column (5) implements the pre-specified attention-check robustness: failers are retained in the primary analysis per the pre-registration. Excluding them leaves the Persuasion $\times$ ACTION point estimate unchanged (+0.024; the p -value rises to 0.11 with the smaller sample) and strengthens the factual belief correction slope (+0.054 vs. +0.037). Failers complete the study faster (median 12 vs. 18 minutes) and report flatter belief elicitation (mean Δb_{ijk} of 0.08 vs. 0.20 for passers), but their provision rate is nearly identical (0.641 vs. 0.638). The sign of every estimate is unchanged in every variant; the instrumental interaction attenuates and loses precision as the excluded groups are added back (+0.017 to +0.024, p from 0.064 to 0.179).

Table C.24: Action Decomposition with Demographic Controls — Experiment 2 (Citizens)

	Probability of Fact Provision	
	(1)	(2)
Expression	0.069*** (0.023)	0.069*** (0.023)
Persuasion	0.003 (0.009)	0.004 (0.009)
Persuasion $\times$ Action	0.024* (0.013)	0.023* (0.013)
Action	0.024* (0.013)	0.021* (0.012)
Factual belief correction gain	0.037** (0.017)	0.046*** (0.017)
Clustered SE	Sender	Sender
Fact FE	Yes	Yes
Controls	No	Yes
Num. obs.	7740	7715
R ²	0.005	0.018
Adj. R ²	0.003	0.012

Sample: $N = 1,548$ Experiment 2 citizens, both arms; column (1) 7,740 observations, column (2) 7,715 with non-missing controls. *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating $\alpha_{ij} \in [-1, +1]$. *Persuasion:* incentivized expected movement of the receiver toward the sender's position $\Delta\pi_{ijk}$ (positional reading). *Factual belief correction gain:* incentivized expected correction of the receiver's factual belief $\Delta b_{ijk} \in [0, 1]$. Column (1) reproduces the preferred action-decomposition specification; column (2) adds demographic controls (age, gender, education, Bundesland, closest-party affiliation, left-right self-placement). Sender-receiver ideological distance is excluded as a heterogeneity variable. All specifications include fact fixed effects. Standard errors clustered at sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

C.3 Directional vs. Positional Persuasion Proxy

The canonical persuasion proxy $\Delta\pi_{ijk}$ measures reduction in Likert distance between the receiver's predicted post-disclosure stance and the sender's own position—a positional measure. A direction-only analog, $\delta_{ijk}^{\text{dir}}$ (defined in Section 4.5), strips preference magnitude from the signal: it is positive whenever the receiver moves toward the sender's preferred side, regardless of how far the receiver is from the sender's exact position or how strong the sender's preference is. Because $\delta_{ijk}^{\text{dir}}$ is undefined for neutral senders, 260 observations from 52 neutral senders (3.4 % of observations) are excluded from the directional columns.

Table C.33 re-estimates the specification replacing $\Delta\pi_{ijk}$ with $\delta_{ijk}^{\text{dir}}$, pooled (columns 1–

Table C.25: Expression vs. Positional Alignment: Horse Race — Experiment 2 (Citizens)

	Probability of Fact Provision		
	(1)	(2)	(3)
Alignment	0.020** (0.009)		0.012 (0.009)
Expression		0.070*** (0.023)	0.061** (0.025)
Clustered SE	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes
Expressive measure	Positional	Behavioral	Both
Num. obs.	7740	7740	7740
R ²	0.002	0.002	0.002
Adj. R ²	0.000	0.001	0.001

Sample: $N = 1,548$ Experiment 2 senders (Bilendi panel, Germany, 2026), 7,740 sender-fact observations. Outcome: binary provision indicator. Alignment: positional alignment $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, +1]$, the pre-registered confirmatory proxy. Expression: cross-domain behavioral rating α_{ij} , elicited post-provision on a different topic. Column (1): Alignment only. Column (2): Expression only. Column (3): horse race including both. All specifications include fact fixed effects and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.26: Inference Robustness: Clustering and Wild Bootstrap — Experiment 2 (Citizens)

	(1) Sender SE (preferred)	(2) Two-way SE (Sender + Fact)
Expression	0.069*** (0.023)	0.069** (0.026)
Factual belief correction gain	0.037** (0.017)	0.037* (0.018)
Persuasion	0.003 (0.009)	0.003 (0.009)
Persuasion $\times$ Action	0.024* (0.013)	0.024* (0.011)
Action	0.024* (0.013)	0.024** (0.010)
Clustered SE	Sender	Sender + Fact
Fact FE	Yes	Yes
Wild bootstrap p (fact, exact 2 ¹²)	—	0.047
Num. obs.	7,740	7,740

Sample: Experiment 2 citizens; 7,740 sender-fact observations. Both columns estimate the preferred action-decomposition specification (Expression, Factual belief correction gain, Persuasion, Persuasion $\times$ Action, Action) with fact fixed effects. Column (1): standard errors clustered at the sender level (preferred). Column (2): two-way clustering by sender and fact (Cameron–Gelbach–Miller; $K = 12$ fact clusters, asymptotic standard errors indicative only). Wild bootstrap: exact enumeration of all $2^{12} = 4,096$ Rademacher sign patterns at the fact level (impose-null variant); p -value for Persuasion $\times$ Action. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.27: Order and Serial-Autocorrelation Robustness — Experiment 2 (Citizens)

	(1) + Order	(2) + Lag	(3) + Both
Expression	0.068*** (0.023)	0.089*** (0.025)	0.089*** (0.025)
Factual belief correction gain	0.036** (0.017)	0.042** (0.019)	0.043** (0.019)
Persuasion	0.003 (0.009)	0.009 (0.011)	0.010 (0.011)
Persuasion $\times$ Action	0.024* (0.013)	0.029* (0.015)	0.029* (0.015)
Action	0.024* (0.013)	0.024* (0.013)	0.024* (0.013)
Presentation order	-0.007* (0.004)	—	0.006 (0.005)
Lag(provision)	—	0.075*** (0.014)	0.076*** (0.014)
Clustered SE	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes
Num. obs.	7,740	6,192	6,192

Sample: Experiment 2 citizens. Column (1): 7,740 observations (all five slots). Columns (2)–(3): slots 2–5 only (slot 1 has no lag). All columns add the listed control(s) to the preferred action-decomposition specification with fact fixed effects and sender-clustered standard errors. *Presentation order*: linear slot index $\in \{1, \dots, 5\}$. *Lag(provision)*: the sender's provision decision for the immediately preceding fact. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

3) and with the ACTION interaction (columns 4–6). Both proxies survive standalone specifications with similar coefficients (positional $\beta = +0.015$, $p = 0.025$; directional $\beta = +0.016$, $p < 0.01$; columns 1–2). When they enter jointly, the directional proxy absorbs the shared variance ($\beta_{\text{dir}} = +0.018$, $p < 0.05$; $\beta_{\text{pos}} = -0.004$, not statistically distinguishable from zero; column 3). The two measures correlate at $r = 0.74$ after fact fixed effects, so the horse race partitions shared variance; it is descriptive and does not adjudicate between the proxies. A common-sample check rules out sample composition as the driver of the joint-specification pattern: positional estimates on the non-neutral subsample remain close to their full-sample values (pooled $\beta = +0.012$, $p = 0.079$; $\times$ ACTION $\beta = +0.022$, $p = 0.095$). The θ_I interaction is directionally consistent under both codings ($\beta = +0.024$, $p = 0.064$ positional; $\beta = +0.019$, $p = 0.103$ directional; columns 4–5).

Table C.34 isolates the functional form: it repeats the ACTION-interaction specifi-

Table C.28: Functional Form of $w(\delta)$: Direction versus Magnitude — Experiment 2 (Citizens)

	(1) $\text{sign}(\delta)$	(2) Winsorized δ
Expression	0.069*** (0.023)	0.069*** (0.023)
Factual belief correction gain	0.037** (0.017)	0.037** (0.017)
Persuasion (transform)	0.018** (0.009)	0.010 (0.011)
Persuasion (transform) $\times$ Action	-0.004 (0.012)	0.020 (0.016)
Action	0.025** (0.013)	0.025* (0.013)
Persuasion transform	$\text{sign}(\delta)$	Winsorized δ
Fact FE	Yes	Yes
Num. obs.	7,740	7,740

Sample: Experiment 2 citizens; 7,740 observations. Both columns replace *Persuasion* ($\Delta\pi_{ijk}$, positional reading) in the preferred action-decomposition specification with a transformed version, keeping fact fixed effects and sender-clustered standard errors. Column (1): $\text{sign}(\delta) \in \{-1, 0, +1\}$ (direction only). Column (2): δ winsorized at the 90th percentile of $|\delta|$ (magnitude, extremes capped; N unchanged). A significant direction main effect with a null direction $\times$ Action interaction, alongside a surviving winsorized magnitude $\times$ Action interaction, indicates w is strictly increasing rather than a step function. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.29: Sample-Exclusion Robustness for the Instrumental Motive — Experiment 2 (Citizens)

	(1) Drop low- $ \delta $	(2) Drop perfect sharers	(3) Drop low- α facts
Expression	0.067** (0.026)	0.067*** (0.025)	0.087*** (0.027)
Factual belief correction gain	0.027 (0.019)	0.047*** (0.018)	0.043** (0.021)
Persuasion	0.003 (0.009)	0.006 (0.010)	0.008 (0.011)
Persuasion $\times$ Action	0.023* (0.013)	0.021 (0.014)	0.022 (0.015)
Action	0.029** (0.014)	0.015 (0.011)	0.022 (0.015)
Fact FE	Yes	Yes	Yes
Num. obs.	5,805	6,355	5,233

Sample: Experiment 2 citizens; preferred action-decomposition specification, fact fixed effects, sender-clustered standard errors. Column (1) drops senders in the bottom quartile of mean $|\delta|$ (expect near-zero persuasion across facts). Column (2) drops perfect sharers (provision rate = 1.0, all five facts shared). Column (3) drops the bottom tertile of facts by mean expressive magnitude $|\alpha|$. The instrumental motive (*Persuasion* $\times$ *Action*) is positive in every restriction. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.30: Sample characteristics and quota targets, Experiment 2

	Population (1)	Sample (2)	CONTROL (3)	ACTION (4)
<i>Gender</i>				
Female	0.499	0.481	0.501	0.462
Male	0.501	0.517	0.497	0.536
Diverse	—	0.003	0.003	0.003
<i>Age band</i>				
18–29	0.179	0.197	0.199	0.195
30–39	0.178	0.176	0.177	0.175
40–49	0.170	0.160	0.150	0.169
50–59	0.193	0.200	0.188	0.212
60–75	0.279	0.267	0.286	0.248
<i>Bundesland</i>				
Baden-Württemberg	0.134	0.129	0.135	0.122
Bayern	0.159	0.156	0.144	0.166
Berlin	0.045	0.042	0.040	0.044
Brandenburg	0.030	0.032	0.036	0.029
Bremen	0.008	0.010	0.011	0.010
Hamburg	0.023	0.023	0.021	0.024
Hessen	0.075	0.072	0.068	0.076
Mecklenburg-Vorpommern	0.019	0.023	0.023	0.024
Niedersachsen	0.095	0.099	0.102	0.096
Nordrhein-Westfalen	0.216	0.215	0.204	0.226
Rheinland-Pfalz	0.050	0.050	0.046	0.054
Saarland	0.012	0.010	0.016	0.005
Sachsen	0.047	0.048	0.050	0.047
Sachsen-Anhalt	0.025	0.032	0.029	0.034
Schleswig-Holstein	0.035	0.031	0.041	0.021
Thüringen	0.025	0.028	0.034	0.021
<i>School-leaving qualification</i>				
Hauptschulabschluss or none	0.297	0.265	0.277	0.253
Mittlerer Abschluss (Realschule)	0.314	0.293	0.270	0.315
Fachhochschul-/Hochschulreife or higher	0.390	0.442	0.453	0.431

Sample: $N = 1,548$ German adults recruited on Bilendi (quota-sampled), Experiment 2. Column (1): German population proportions. Gender, age band, and Bundesland: population aged 18–75 as of 31 December 2024 (Bevölkerungsfortschreibung on the Zensus-2022 basis; Statistisches Bundesamt, Statistischer Bericht 12411, tables 12411-06 and 12411-09). No registry benchmark exists for “Diverse.” School-leaving qualification: Mikrozensus 2023 (Statistisches Bundesamt, Bildungsstand der Bevölkerung), persons aged 15 and over excluding those still in school—a broader base than the sample’s 18–75. Column (2): full-sample proportions. Columns (3)–(4): proportions by experimental arm.

cation under binary codings of both proxies, replacing each continuous measure with an indicator for a strictly positive value. The two codings separate the persuasion response by arm. In the CONTROL arm, senders respond to the *direction* of expected persuasion: a binary indicator for movement toward the sender’s preferred side predicts provision on its own ($\beta = +0.035$, $p < 0.05$), consistent with a sign-like w when $A = 0$. The ACTION amplification—the instrumental motive θ_I —requires *magnitude*:

	Canonical	+ Speeders	+ Early starts	+ Both	Attention passers
Alignment (H1)	0.026** (0.012)	0.030** (0.012)	0.020* (0.011)	0.023** (0.011)	0.026* (0.014)
Alignment $\times$ ACTION (H2)	-0.010 (0.017)	-0.012 (0.016)	-0.007 (0.016)	-0.009 (0.016)	-0.019 (0.019)
Expression (α_{ij})	0.069*** (0.023)	0.069*** (0.023)	0.068*** (0.022)	0.067*** (0.022)	0.058** (0.027)
Factual belief correction gain (Δb_{ijk})	0.037** (0.017)	0.034** (0.017)	0.036** (0.017)	0.034** (0.017)	0.054*** (0.020)
Persuasion ($\Delta \pi_{ijk}$)	0.015** (0.007)	0.015** (0.006)	0.013** (0.006)	0.014** (0.006)	0.019** (0.008)
Persuasion $\times$ ACTION	0.024* (0.013)	0.021* (0.013)	0.019 (0.012)	0.017 (0.012)	0.024 (0.015)
Senders	1548	1588	1654	1694	1156
Observations (confirmatory)	7715	7910	8245	8440	5760
Observations (motive specifications)	7740	7940	8270	8470	5780

Dependent variable: binary provision decision ($y_{ij} \in \{0, 1\}$), Experiment 2 citizens, under alternative sample definitions. Canonical: the main analysis sample (no speeders, no pre-launch starts). + Speeders: adds completers faster than $0.25 \times$ the median completion time. + Early starts: adds participants who started before the launch of the main collection on 27 May 2026, 10:00 CET. + Both: adds both groups. Attention passers: restricts the canonical sample to senders answering the attention-check item correctly (the attention check was administered in Experiment 2 only). Rows (1)–(2) are from the pre-registered confirmatory specification: provision on positional alignment, ACTION, their interaction, demographic controls (age, gender, education, Bundesland, voting intention, left-right), and fact fixed effects. Rows (3)–(5) are from the three-channel motive specification (Expression, factual belief correction gain, Persuasion, fact fixed effects); row (6) adds the Persuasion $\times$ ACTION interaction to that specification. The two observation rows differ because the confirmatory specification drops observations with missing demographic controls (25 observations from five senders in the canonical column). Standard errors clustered at the sender level (in parentheses). *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.31: Sample-inclusion robustness of the Experiment 2 estimates.

Table C.32: Leave-One-Fact-Out Robustness of the Motive Decomposition — Experiment 2 (Citizens)

Motive	Sign-positive	Min β	Max β	Max p
Expression	12/12	0.053	0.082	0.031
Factual belief correction gain	12/12	0.025	0.043	0.160
Persuasion	12/12	0.011	0.018	0.102
Fact FE	Yes			
Clustered SE	Sender			

Sample: Experiment 2 citizens. Each motive is re-estimated twelve times in the three-motive specification (*Expression*, *Factual belief correction gain*, *Persuasion*), each time dropping one of the twelve facts; fact fixed effects, sender-clustered standard errors. Columns report the number of leave-one-out estimates with a positive coefficient, the minimum and maximum coefficient across the twelve runs, and the largest p -value observed.

replacing either continuous proxy with its binary version eliminates the interaction entirely (positional binary $\times$ ACTION $\beta = +0.001$, $p = 0.967$; directional binary $\times$ ACTION $\beta = +0.013$, $p = 0.586$), while both continuous codings retain it (positional $p = 0.064$; directional $p = 0.103$).

	Pooled			With Action Interaction		
	Positional	Directional	Horse race	Positional	Directional	Horse race
Expression α_{ij}	0.069*** (0.023)	0.065*** (0.023)	0.065*** (0.023)	0.069*** (0.023)	0.064*** (0.023)	0.064*** (0.023)
Factual belief correction Δb_{ijk}	0.037** (0.017)	0.036** (0.018)	0.036** (0.018)	0.037** (0.017)	0.036** (0.018)	0.036** (0.018)
Persuasion-positional $\Delta\pi_{ijk}^{\text{pos}}$	0.015** (0.007)		-0.004 (0.009)	0.003 (0.009)		-0.010 (0.013)
Persuasion-directional $\Delta\pi_{ijk}^{\text{dir}}$		0.016*** (0.006)	0.018** (0.008)		0.006 (0.008)	0.013 (0.012)
ACTION				0.024* (0.013)	0.024* (0.013)	0.024* (0.013)
Persuasion-pos. $\times$ ACTION				0.024* (0.013)		0.013 (0.019)
Persuasion-dir. $\times$ ACTION					0.019 (0.011)	0.010 (0.017)
Clustered SE	Sender	Sender	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes	Yes	Yes
Action level + interaction	No	No	No	Yes	Yes	Yes
Sample	All senders	Non-neutral	Non-neutral	All senders	Non-neutral	Non-neutral
Num. obs.	7740	7480	7480	7740	7480	7480
R ²	0.004	0.004	0.004	0.005	0.005	0.005
Adj. R ²	0.002	0.002	0.002	0.003	0.003	0.003

Dependent variable: binary provision decision ($y_{ij} \in \{0,1\}$). Sample: Experiment 2, $N = 1,548$ senders. Directional columns exclude 260 observations from senders with $\text{pref}_i = 0$ (neutral donors; $N = 7,480$); positional columns retain the full sample ($N = 7,740$). Persuasion-positional ($\Delta\pi_{ij}^{\text{pos}}$) is `delta_gain_2`: reduction in Likert distance to the sender's own position. Persuasion-directional ($\Delta\pi_{ij}^{\text{dir}}$) is $\delta_{ij}^{\text{dir}} \equiv -\Delta\text{stance}_{ij} \cdot \text{sign}(\text{pref}_i)$: whether the receiver moves toward the sender's preferred side, preference-magnitude-free. The two proxies correlate at $r = 0.74$ after fact fixed effects; horse-race columns partition shared variance and are descriptive, not an adjudication between the proxies. Standard errors clustered at the sender level. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.33: Positional vs. directional persuasion proxies, Experiment 2.

	Positional		Directional	
	Continuous	Binary	Continuous	Binary
Expression α_{ij}	0.069*** (0.023)	0.069*** (0.023)	0.064*** (0.023)	0.065*** (0.023)
Factual belief correction Δb_{ijk}	0.037** (0.017)	0.038** (0.017)	0.036** (0.018)	0.036** (0.018)
Persuasion-pos. (continuous)	0.003 (0.009)			
Persuasion-pos. (binary)		0.023 (0.016)		
Persuasion-dir. (continuous)			0.006 (0.008)	
Persuasion-dir. (binary)				0.035** (0.017)
ACTION	0.024* (0.013)	0.025 (0.016)	0.024* (0.013)	0.020 (0.017)
Persuasion-pos. (cont.) $\times$ ACTION	0.024* (0.013)			
Persuasion-pos. (bin.) $\times$ ACTION		0.001 (0.023)		
Persuasion-dir. (cont.) $\times$ ACTION			0.019 (0.011)	
Persuasion-dir. (bin.) $\times$ ACTION				0.013 (0.024)
Clustered SE	Sender	Sender	Sender	Sender
Fact FE	Yes	Yes	Yes	Yes
Sample	All senders	All senders	Non-neutral	Non-neutral
Num. obs.	7740	7740	7480	7480
R ²	0.005	0.004	0.005	0.005
Adj. R ²	0.003	0.002	0.003	0.003

Dependent variable: binary provision decision ($y_{ij} \in \{0, 1\}$). Sample: $N = 1,548$ Experiment 2 senders (7,740 observations); directional columns exclude 260 observations from senders with $\text{pref}_i = 0$ (neutral donors; $N = 7,480$). Each column re-estimates the ACTION-interaction specification with one coding of the persuasion proxy. Continuous codings: positional $\Delta\pi_{ijk}^{\text{pos}}$ (delta_gain_2) and directional $\Delta\pi_{ijk}^{\text{dir}}$. Binary codings: an indicator equal to one when the corresponding continuous measure is strictly positive. All specifications include Expression, Factual belief correction, fact fixed effects, and sender-clustered standard errors. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table C.34: Persuasion proxy functional form: continuous vs. binary codings, Experiment 2.

Table C.35: Expression Measure (α_{ij}) by Sender Stance and Individual Fact — Experiment 2

Panel A: Immigration facts

Sender stance	Anti-policy direction (–)			Pro-policy direction (+)		
	Violent crime	Bürgergeld	Res. permits	Shortage fields	Employment	Integration
1 (most pro)	+0.00	+0.29	+0.00	+0.43	+0.00	+0.67
2	–0.33	–0.16	–0.43	+0.44	+0.37	+0.77
3 (neutral)	+0.29	–0.07	+0.00	+0.64	+0.33	+0.59
4	+0.66	+0.83	+0.54	+0.00	+0.28	–0.16
5 (most anti)	+0.82	+0.82	+0.79	+0.38	+0.34	+0.24

Panel B: Climate facts

Sender stance	Anti-policy direction (–)			Pro-policy direction (+)		
	Grid fees	Business surv.	Price fluct.	Renewable share	Solar cost	Renew. jobs
1 (most pro)	+0.23	+0.45	+0.38	+0.49	+0.71	+0.55
2	–0.15	+0.23	+0.26	+0.46	+0.59	+0.46
3 (neutral)	+0.14	+0.10	+0.30	+0.31	+0.50	+0.30
4	+0.64	+0.58	+0.61	+0.00	–0.40	–0.33
5 (most anti)	+0.78	+0.79	+0.54	+0.26	+0.20	–0.10

Notes: Expression (α_{ij}) = $2(\bar{g}_{ij} - 0.5)$, where $g = 1$ if the sender rates the cross-domain fact as “supporting my opinion.” Per-cell sizes are approximately 30–100 observations (individual fact $\times$ stance); values in small cells (fewer than 40 observations) should be interpreted with caution. Fact identifiers in brackets. Sender stance = stated Likert position on the provision topic (1 = most pro-policy; 5 = most anti-policy). Cross-domain ratings are elicited on the *other* topic; all senders in the cell have the same stated stance.




Figure C.6: Elicited returns to disclosure by sender-receiver distance, Experiment 2 (Citizens). One marker per fact $\times$ sender-receiver distance $\times$ alignment (cell mean over ≥ 20 elicitations), positionally aligned facts ($\text{Align}_{ij} > 0$) as filled black circles, unaligned ($\text{Align}_{ij} < 0$) as open gray squares, horizontally offset for legibility; facts of senders with $\text{Pref}_i = 0$ are omitted from the split. Solid (dashed) lines connect the observation-weighted means of aligned (unaligned) elicitations at each distance. Panel (a): expected factual belief correction $\widehat{\Delta b}_{ijk}$ —flat in distance (pooled slope $\beta = -0.003$, $p = 0.38$, fact fixed effects, sender-clustered standard errors). Panel (b): expected persuasion $\widehat{\Delta \pi}_{ijk}$ —rising in distance ($\beta = +0.072$, $p < 0.001$); the aligned and unaligned mean profiles are statistically indistinguishable (distance $\times$ alignment interaction $p = 0.12$). Dashed lines mark zero. Elicitations incentivized via Quadratic Scoring Rule.

D Experiment 3 Supplementary Tables and Figures

D.1 Sample Composition

Table D.36 compares the demographic and occupational characteristics of the Experiment 3 journalist sample to Loosen et al. (2023), the most recent representative survey of the German journalist population. The post-questionnaire collected gender, age, education, medium type, conditional broadcaster affiliation, geographic reach, role (multiple selection), and editorial autonomy. Fields not elicited in Experiment 3—years of experience, full-time versus freelance status, and topics covered—cannot be compared to the benchmark.

D.2 Heterogeneity

Table D.37 reports the pre-registered exploratory heterogeneity of the alignment gap among journalists—by sender–receiver distance, polarisation, issue extremity, and the journalist background variables (medium type, broadcaster type, reach, role, editorial autonomy)—together with the corresponding interactions for the three motive channels.

Table D.36: Experiment 3 Sample Characteristics: Comparison to Population Benchmark

	Experiment 3 (<i>N</i> = 307)	Loosen et al. (2023) (<i>N</i> = 1,221)
<i>Gender</i>		
Female	53.4	44.0
Male	44.6	55.9
<i>Age</i>		
Under 50	77.3	62.4
<i>Education</i>		
University degree (B.A.+)	84.6	67.1
<i>Medium type</i>		
Print	28.4	57.0
Online	45.4	3.1
Television	7.2	17.3
Radio/podcast	10.8	17.2
News agency	3.6	3.5
<i>Broadcaster type</i> (TV/radio respondents only)		
Public broadcaster	80.0	61.4
<i>Role</i> (multiple selection; shares exceed 100 %)		
Editor	57.2	≈ 61
<i>Editorial autonomy</i>		
High ($\geq 4/5$)	75.8	62.0
<i>Political affiliation</i> (no population benchmark)		
Grüne or Linke (closest party)	58.5	—
Mean left-right self-placement (0–10)	3.2	—

Sample: *N* = 307 professional German journalists recruited via DJV and DJU mailing lists, Experiment 3 (convenience sample, not a probability sample). *Benchmark:* Loosen et al. (2023), stratified random sample of active German journalists (*N* = 1,221, September 2022–February 2023, HBI Hamburg; [Loosen et al.](#)). All figures are percentages. *Medium type:* Experiment 3 uses a single online category for all digital-native outlets; the Loosen et al. benchmark reports native-online journalists (3.1%) as a distinct population segment; print aggregates newspapers and magazines (57.0%). *Broadcaster type* is conditional on TV or radio respondents only. *Role* allows multiple selection; the Loosen et al. editor figure is approximate. *Age:* Loosen et al. report a population mean of 45.3 years; Experiment 3 elicited age as a categorical variable. *Political affiliation* and left-right self-placement were not elicited in Loosen et al. and appear without a benchmark.

Moderator	Interaction of moderator with:				Obs.
	Align _{ij}	α_{ij}	Δb_{ijk}	$\Delta \pi_{ijk}$	
Distance Stance _i – Stance _k	–0.009 (0.011)	–0.082** (0.038)	0.021 (0.029)	0.016 (0.012)	1535
Polarisation	–0.010 (0.010)	–0.029 (0.036)	–0.028 (0.027)	0.012 (0.012)	1530
Issue extremity	–0.023 (0.030)	–0.126* (0.072)	0.006 (0.057)	0.020 (0.023)	1535
Online medium	–0.070** (0.029)	–0.288*** (0.096)	0.013 (0.080)	0.006 (0.029)	1530
Public broadcaster	–0.013 (0.077)	–0.302 (0.236)	0.075 (0.243)	–0.069 (0.071)	275
National/international reach	0.068** (0.031)	0.082 (0.102)	0.035 (0.080)	0.027 (0.032)	1535
Editor role	–0.027 (0.030)	–0.063 (0.099)	–0.068 (0.077)	0.012 (0.029)	1535
Editorial autonomy	0.005 (0.015)	–0.009 (0.055)	0.095** (0.040)	0.017 (0.012)	1530

Dependent variable: binary provision decision ($y_{ij} \in \{0, 1\}$), Experiment 3 journalists. Each row reports interaction coefficients between one moderator and the selection measures: column (1) is from a regression of provision on positional alignment, the moderator, and their interaction; columns (2)–(4) are from a single regression interacting the moderator with Expression (α_{ij}), the factual belief correction gain (Δb_{ijk}), and Persuasion ($\Delta \pi_{ijk}$) jointly. All specifications include fact fixed effects and journalist-clustered standard errors (in parentheses); moderators are centered at the estimation-sample mean. Moderators: sender–receiver distance, polarisation (|left-right – 5|), issue extremity (|stance – 3|), and the pre-registered journalist background variables, coarsened as online medium (vs. all other medium types), public broadcaster (defined only for television and radio senders), national or international reach (vs. local or regional), any editor role in the multi-select role item, and editorial autonomy (Likert 1–5). Benjamini–Hochberg correction is applied within two families: the pre-registered alignment battery (column 1; 8 tests, minimum $q = 0.11$) and the channel battery (columns 2–4; 24 tests), in which only Expression $\times$ Online medium survives at the 10 % level ($q = 0.074$). Stars mark unadjusted p -values: *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table D.37: Heterogeneity of alignment- and channel-based selection among journalists, Experiment 3.

E LLM Classification Methodology

E.1 Classification Procedure

I classify the open-ended decision-reason responses from all three experiments—Experiment 1 (1,200 US citizens, in English), Experiment 2 (1,548 German citizens), and Experiment 3 (306 German journalists)—into theory-derived motive categories using a single large language model, Anthropic’s Claude Sonnet 5 (claude-sonnet-5). Each response is classified in a separate API call that returns one structured JSON object; I change no temperature or sampling settings, so that the classification is deterministic given the pinned model and structured outputs. This is a genuine reclassification under a new category scheme, not a relabeling of an earlier coding: the categories, the gating step, and the boundary rules below all differ from the scheme used before.

The scheme is applied in two stages. A first-step gate asks whether the response refers to the sharing decision, the receiver, or the information provision—explicitly or elliptically. Responses that describe only the sender’s own epistemic process, or state a non-directional opinion with no link to the decision, fail the gate and are recorded as *did not refer to disclosure* with no further coding. Responses that pass are coded into the categories in Table E.38: belief correction, expression, persuasion, receiver tailoring, disbelief, balance, relevance, futility, and miscellaneous, with role neutrality added for journalists. Categories are not mutually exclusive; a label is assigned only where the text gives distinct support for it, and no motive is inferred beyond the text. Each response receives exactly one primary motive, the most central justification. For persuasion I additionally record a think/do subtype; the subtype is not used in the analysis.

Three categories carry the theory. Expression marks the fact’s fit with the sender’s own position as the stated decision factor, with no receiver-directed goal; a generic reference to “my beliefs” without a political direction fails the gate rather than counting as expression. Persuasion marks a receiver-directed goal—moving what the receiver thinks or does—and any reference to the receiver changing is coded here rather than as

receiver tailoring. Belief correction marks truth-telling or the receiver's right to accurate information and is direction-blind. These three map onto the channels in the main text: expression onto the expressive channel, persuasion onto the persuasion channel, and belief correction onto factual belief correction.

The complete classification prompt, including all decision rules, boundary conditions, and worked examples, is available upon request.

E.2 Validation

Each classification carries the model's self-reported confidence, which indicates how firmly the category definitions and boundary rules pin down the coding. The median confidence is 0.85. The model reports confidence at or above 0.7 for 82 percent of Experiment 1 responses, 86 percent for Experiment 2, and 81 percent for Experiment 3. The three categories central to the analysis—belief correction, expression, and persuasion—carry explicit boundary rules in the prompt.

E.3 Coding Examples

Table E.39 presents representative examples of open-ended responses classified into each motive category, illustrating the classification boundaries. Tables E.40 and E.41 present the corresponding examples for the Experiment 2 general-population sample and the Experiment 3 journalist sample, translated from German.

Table E.38: Coding Scheme and Decision Rules

Category	Definition	Key boundary rule
Belief correction	Truth-telling, honesty, or the receiver's right to accurate information; direction-blind	Does not depend on the fact's political direction
Expression	The fact's fit with the sender's own position is the stated decision factor, with no receiver goal	No receiver-directed goal; a generic "my beliefs" without a political direction fails the gate
Persuasion	A receiver-directed goal: moving what the receiver thinks or does	Any reference to the receiver changing is coded here; a think/do subtype is recorded but not used
Receiver tailoring	Decision conditioned on the receiver's beliefs or characteristics, without reference to attitude change	Any reference to the receiver changing → <i>persuasion</i>
Disbelief	Credibility doubts about the fact or source as the reason for withholding	Politically motivated skepticism → <i>expression</i> or <i>persuasion</i>
Balance	Deliberately sharing counter-attitudinal or both-sides content	The receiver is meant to see the other side
Relevance	The fact is judged irrelevant or unimportant for the receiver, without a directional stake	Misleadingness with a directional stake → <i>persuasion</i>
Futility	"Facts do not change minds" offered as the stated reason	A general claim about receivers, not about one fact
Role neutrality (journalists only)	Professional role integrity: it is not the journalist's place to filter what the reader sees	Distinct from <i>balance</i> ; concerns the sender's role, not both-sidedness
Miscellaneous	Refers to the decision but gives no interpretable reasoning	Residual category; enters no regression
Did not refer to disclosure	Describes only the sender's own epistemic process, or a non-directional opinion with no decision link	Gate outcome; receives no further coding

Categories coded by Claude Sonnet 5 from the open-ended decision-reason responses of Experiments 1–3. Categories are not mutually exclusive; a label is assigned only where the text gives distinct support. Each response receives one primary motive. Expression, persuasion, and belief correction map onto the expressive, persuasion, and factual-belief-correction channels in the main text.

Table E.39: Coding Examples

Primary motive	Example
Belief correction	<i>"People should know the truth." "I wanted the people who got the questions wrong to know the correct answer. People should be more knowledgeable about things."</i>
Expression	<i>"My belief that immigrants are beneficial to the united states." "My personal political and socioeconomic stand towards immigration."</i>
Persuasion	<i>"I considered how likely each participant was to meaningfully change their attitude after receiving the findings." "When I thought that the fact was unfavorable towards immigration, I chose to not share the fact."</i>
Disbelief	<i>"Whether the facts seemed like they were actually facts, or just made up." "The potential accuracy, or intentional misconstruing of data."</i>
Miscellaneous	<i>"It was dependent on the topic." "The data available influenced my decision."</i>
Did not refer to disclosure	<i>"My own beliefs and also taking into consideration the information from research." "I think that immigration is a valuable thing for the US. I think immigrants help the economy and do important jobs US citizens do not want to do."</i>

Representative open-ended responses from Experiment 1 senders, illustrating the classification boundaries between primary motives. Responses are lightly excerpted for length; full text available upon request. Primary motive classification by Claude Sonnet 5.

Table E.40: Coding Examples — Experiment 2 (Citizens)

Primary motive	Example
Belief correction	<i>"Since the information is always truthful, I wanted every person to receive it—regardless of whether it matched their view."</i> <i>"I shared all the information. Everyone has the right to be informed."</i>
Expression	<i>"I checked whether it fit my own attitude."</i> <i>"As a rule, I sent things that support my opinion."</i>
Persuasion	<i>"I wanted to influence the person so that they would come around to my opinion."</i> <i>"I tried to reinforce views that favour climate protection and to convince the others."</i>
Receiver tailoring	<i>"If the person had thought positively, send the message; if they thought negatively, don't send it."</i> <i>"I looked at whether the information matched the person's opinion and then decided whether to send it."</i>
Balance	<i>"I looked at what might broaden their opinion, so that they would also see the other side once."</i> <i>"I tried to be impartial, so also to send information favourable to migration, as long as they were correct facts."</i>

Representative open-ended responses from Experiment 2 senders (German general-population sample), illustrating the classification boundaries between primary motives. Responses are translated from German and lightly excerpted for length; original text available upon request. Primary motive classification by Claude Sonnet 5.

Table E.41: Coding Examples — Experiment 3 (Journalists)

Primary motive	Example
Belief correction	<i>"All readers should have all the information so that they can form their own opinion. That is why I sent everything."</i> <i>"Everyone has a right to the facts, whatever their opinion. So I always sent them."</i>
Expression	<i>"I checked whether the information supported my opinion or not, and sent it accordingly."</i> <i>"Weighing the reader's existing opinion against the slant of the information—admittedly in favour of my own view."</i>
Persuasion	<i>"You try to nudge people a little towards your own opinion."</i> <i>"Starting from my own opinion, I tried to gauge in which case the person would move closer to my view."</i>
Role neutrality	<i>"I always send information. It is my job."</i> <i>"As a journalist I send facts neutrally—independent of my personal opinion. The readers' opinion was really secondary."</i>
Balance	<i>"Balance was important to me."</i> <i>"I tried to create a counterweight to the constant political propaganda."</i>

Representative open-ended responses from Experiment 3 senders (German journalists), illustrating the classification boundaries between primary motives. "Role neutrality" captures the professional-norm rationale of transmitting facts irrespective of the sender's or receiver's stance. Responses are translated from German and lightly excerpted for length; original text available upon request. Primary motive classification by Claude Sonnet 5.

F Self-Reported Motives, Norms, and Beliefs

F.1 Stated Norms and Perceived Design Balance

Both experiments closed with the question whether respondents perceived a political imbalance in the study. A large majority perceived none (85.5 % of citizens, 82.4 % of journalists), and the shares are flat across treatment arms and topics. Among the minority perceiving an imbalance, a left-favoring reading outnumbers a right-favoring one roughly two to one in both samples (9.5 % vs. 5.0 % among citizens; 11.7 % vs. 5.9 % among journalists), against a fact menu balanced by construction (two pro and two anti facts, plus one randomly drawn). Table F.42, Panel A, reports the shares.

Table F.42, Panel B, reports the distributions of the three normative items, and Table F.43 the item $\times$ selection interaction battery discussed in Section 5.4.

F.2 Stated Receiver-Model Beliefs

Beyond their sharing decisions, senders explained what they based their receiver-belief predictions on; a large language model codes these prediction rationales into stated be-

	Citizens (Exp. 2)			Journalists (Exp. 3)		
	Mean	SD	% Agree	Mean	SD	% Agree
<i>Panel A: Perceived political imbalance (% of valid answers)</i>						
No imbalance perceived		85.5			82.4	
Imbalance favoring the left		9.5			11.7	
Imbalance favoring the right		5.0			5.9	
<i>Panel B: Normative and policy items (1–5 agreement)</i>						
Share opposing information	4.13	0.90	80.0	4.60	0.68	95.4
Social-media transparency	4.05	0.96	75.7	4.61	0.65	95.1
Viewpoint-diversity obligation	4.17	0.90	80.4	4.22	0.91	82.7

Post-questionnaire self-reports, one per sender. Panel A: answers to “Did you perceive a political imbalance in the study?” (options: no imbalance, imbalance favoring the left, imbalance favoring the right), in percent of valid answers; Experiment 2 $N = 1,548$ (no nonresponse), Experiment 3 $N = 307$ (one nonresponse). Panel B: agreement on a 1–5 scale (1 = fully disagree, 5 = fully agree) with three items: whoever passes on political information should also pass on factually correct information that contradicts their own position (worded for journalists in Experiment 3); social media should be more transparent about why users see some political information and not others; platforms and media should be obligated to surface relevant information from different perspectives. % Agree is the share answering 4 or 5. Experiment 2 $N = 1,548$; Experiment 3 $N = 306$ (two nonresponses).

Table F.42: Perceived design balance and normative items, Experiments 2 and 3.

Item	Interaction of item with:			Obs.
	Align _{ij}	α_{ij}	$\Delta\pi_{ijk}$	
<i>Panel A: Citizens (Experiment 2)</i>				
Share opposing information	−0.008 (0.009)	−0.014 (0.022)	0.006 (0.007)	7740
Social-media transparency	−0.007 (0.008)	0.029 (0.022)	0.007 (0.006)	7740
Viewpoint-diversity obligation	−0.004 (0.009)	−0.008 (0.023)	0.011 (0.007)	7740
<i>Panel B: Journalists (Experiment 3)</i>				
Share opposing information	−0.041 (0.026)	−0.095 (0.078)	−0.038 (0.029)	1530
Social-media transparency	−0.075*** (0.025)	−0.194*** (0.071)	−0.041 (0.026)	1530
Viewpoint-diversity obligation	0.002 (0.019)	−0.008 (0.060)	−0.014 (0.017)	1530

Dependent variable: binary provision decision ($y_{ij} \in \{0, 1\}$). Each cell reports the interaction coefficient between the row item (centered at 3, the scale midpoint) and one selection measure, from a separate regression of provision on the measure, the item, and their interaction, with fact fixed effects and sender-clustered standard errors (in parentheses). Selection measures: positional alignment (Align_{ij}), Expression (α_{ij}), and Persuasion ($\Delta\pi_{ijk}$). Items as in Table F.42, Panel B. Benjamini–Hochberg correction is applied across the 18-test family (3 items $\times$ 3 measures $\times$ 2 samples): the two Experiment-3 transparency interactions in columns (1)–(2) survive at the 10% level ($q = 0.060$); no other test has $q < 0.40$. Stars mark unadjusted p -values: *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table F.43: Normative items and selection: channel $\times$ item interactions, Experiments 2 and 3.

liefs about whether receivers update their position (persuadable, unpersuadable, and related claims) and whether they already know the facts (Appendix E.1). Experiment 1 elicited no such predictions and is omitted. Figure F.7 reports the distribution of these stated beliefs. Table F.44 relates them to the incentivized elicited returns, and Table F.45 to the independently coded stated motives of Section 5.4.




Figure F.7: Distribution of senders' stated beliefs about how receivers update, among senders whose open-ended prediction rationale yields a codable claim (codable rates: Experiment 2 11.2%, Experiment 3 54.2%). Panel (a): beliefs about whether receivers' policy position moves. Panel (b): beliefs about whether receivers already know the facts. Shares sum to more than one because a rationale may state several claims. Beliefs coded from open-ended text by a large language model (claude-sonnet-5, delta_reason scheme; Appendix E.1). Experiment 1 elicited no receiver-updating predictions and is omitted.

Table F.44: Convergent Validity of the Incentivized Belief Elicitations (Experiments 2 and 3)

	Persuasiveness $ \Delta\pi $		Belief correction Δb	
	(1)	(2)	(3)	(4)
Intercept	0.649*** (0.013)	0.656*** (0.032)	0.170*** (0.006)	0.303*** (0.013)
Stated <i>persuadable</i>	0.191*** (0.058)	−0.005 (0.059)		
Stated <i>unpersuadable</i>	−0.138** (0.055)	−0.267*** (0.049)		
Stated <i>not-know</i>			0.073* (0.040)	0.159*** (0.041)
Stated <i>know</i>			0.111** (0.054)	0.078** (0.036)
Sample	Citizens	Journalists	Citizens	Journalists
Num. obs.	1548	306	1548	306
R ²	0.008	0.060	0.004	0.054

Sample: sender-level, Experiment 2 citizens ($N = 1,548$) and Experiment 3 journalists ($N = 306$). *Specification:* OLS of each sender's mean incentivized elicited belief on indicators for the receiver-updating claims stated in the sender's open-ended prediction rationale (Appendix E.1). Columns 1–2: dependent variable is the mean predicted movement magnitude $|\Delta\pi| = |\text{delta_shift}|$. Columns 3–4: dependent variable is the mean elicited factual belief correction Δb (acc_gain). Stating that receivers are *unpersuadable* predicts smaller elicited movement; stating that receivers do *not know* the facts predicts larger elicited belief correction. The knowledge axis is corroborated by the predicted baseline-knowledge share (delta_pct_none), on which *not-know* is negative (−6.9pp in column 3, −9.9pp in column 4) and *know* is null. Heteroskedasticity-robust standard errors in parentheses. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

Table F.45: Cross-Source Coherence: Stated Motive Predicts Stated Receiver-Model (Experiments 2 and 3)

	Stated <i>persuadable</i>		Stated <i>not-know</i>	
	(1)	(2)	(3)	(4)
Intercept	0.033*** (0.005)	0.176*** (0.025)	0.010*** (0.003)	0.052*** (0.016)
Stated persuasion motive	0.214*** (0.048)	0.162** (0.063)		
Stated belief-correction motive			0.039 (0.028)	0.087** (0.036)
Sample	Citizens	Journalists	Citizens	Journalists
Num. obs.	1548	306	1548	306
R ²	0.054	0.027	0.005	0.023

Sample: sender-level, Experiment 2 citizens ($N = 1,548$) and Experiment 3 journalists ($N = 306$). *Specification:* linear-probability regressions of an indicator for a stated receiver-updating *belief* (from the `delta_reason` coding of the prediction rationale) on an indicator for the corresponding stated *motive* (from the independent `decision_reason` coding of the sharing rationale). The two open-text questions are coded in separate classification runs. The intercept is the base rate of stating the belief among senders who do not state the motive; the slope is the enrichment among senders who do. Senders who state a persuasion motive are markedly more likely to state that receivers are persuadable, and senders who state a belief-correction motive more likely to state that receivers do not know the facts. Because the stated-belief cells are rare, association is best assessed by Fisher's exact test: $p < 0.001$ (column 1), $p = 0.007$ (column 2), $p = 0.031$ (column 3), $p = 0.011$ (column 4); all four coherences are significant. Heteroskedasticity-robust standard errors in parentheses. *** $p < 0.01$; ** $p < 0.05$; * $p < 0.1$.

G Proofs

This appendix collects the proofs of the propositions in Section 3.4 and develops the magnitude prediction in full.

Proof of Proposition 0 (Belief correction). From equation (4),

$$\frac{\partial \Pr(d = 1)}{\partial \Delta b} = f(U(d = 1)) \theta_B.$$

Since $f > 0$ everywhere and $\theta_B > 0$, the derivative is strictly positive. $\square$

Proof of Proposition 1a (Expressive selective disclosure). From equation (4),

$$\frac{\partial \Pr(d = 1)}{\partial \alpha} = f(U(d = 1)) \theta_E.$$

Since $f > 0$ everywhere and $\theta_E > 0$, the derivative is strictly positive. $\square$

Proof of Proposition 1b (Projective selective disclosure). From equation (4),

$$\frac{\partial \Pr(d = 1)}{\partial w(\Delta \pi)} = f(U(d = 1)) c(A).$$

Since $f > 0$ everywhere and $c(A) > 0$, this derivative is strictly positive.

If w is differentiable and strictly increasing, then $w'(\Delta \pi) > 0$ on the relevant domain. Therefore,

$$\frac{\partial \Pr(d = 1)}{\partial \Delta \pi} = f(U(d = 1)) c(A) w'(\Delta \pi) > 0.$$

$\square$

Proof of Proposition 2 (Political action amplification). Differentiating equation (3) with respect to $w(\Delta\pi)$ gives

$$\frac{\partial U(d=1)}{\partial w(\Delta\pi)} = \theta_P + \theta_I A.$$

The difference in this marginal utility between the action and no-action states is

$$\left[\frac{\partial U(d=1)}{\partial w(\Delta\pi)} \mid A=1 \right] - \left[\frac{\partial U(d=1)}{\partial w(\Delta\pi)} \mid A=0 \right] = (\theta_P + \theta_I) - \theta_P = \theta_I.$$

Since $\theta_I > 0$, political action strictly increases the marginal utility of persuasion; this increment is the instrumental motive. $\square$

Persuasiveness amplification. The framework also yields a magnitude prediction, stated here in full.

Proposition 3 (Persuasiveness amplification). *Fix $A \in \{0, 1\}$ and suppose $c(A) > 0$. If w is differentiable and strictly increasing, then stronger perceived persuasion increases disclosure when persuasion moves the receiver toward the sender's own position and decreases disclosure when persuasion moves the receiver away from it. Specifically, for $\Delta\pi > 0$,*

$$\frac{\partial \Pr(d=1)}{\partial \Delta\pi} > 0,$$

while for $p < 0$, disclosure is strictly decreasing in $|\Delta\pi|$:

$$\frac{\partial \Pr(d=1)}{\partial |\Delta\pi|} < 0.$$

Proof. For $\Delta\pi > 0$, equation (4) implies

$$\frac{\partial \Pr(d=1)}{\partial \Delta\pi} = f(U(d=1)) c(A) w'(\Delta\pi).$$

Since $f > 0$, $c(A) > 0$, and $w'(\Delta\pi) > 0$, this derivative is strictly positive.

For $p < 0$, by the chain rule,

$$\frac{\partial \Pr(d = 1)}{\partial |\Delta\pi|} = \frac{\partial \Pr(d = 1)}{\partial \Delta\pi} \cdot \frac{\partial \Delta\pi}{\partial |\Delta\pi|} = f(U(d = 1)) c(A) w'(\Delta\pi) \cdot (-1).$$

Since $f > 0$, $c(A) > 0$, and $w'(\Delta\pi) > 0$, this derivative is strictly negative. $\square$

Proposition 3 adds a magnitude prediction to the directional prediction in Proposition 1b. When the state of the world is perceived to move the receiver toward the sender's position, disclosure increases with the perceived strength of that movement. When it is perceived to move the receiver away, disclosure decreases with the perceived strength of that movement. The boundary case is a function such as $w(\Delta\pi) = \text{sign}(\Delta\pi)$, where the sender responds only to the direction of persuasion and not to its magnitude.