

# Who Should Know What and Why? Curated Supply of Political Information

Yero Samuel Ndiaye\*

July 8, 2026

[EARLY DRAFT — COMMENTS WELCOME]

[Click here for the most recent version](#)

## Abstract

The information environment is distorted not only by false content but by the selective supply of true content. I study why senders choose which accurate political facts to share, and decompose the motives behind it. Among 1,200 U.S. adults, senders are 8 percentage points more likely to share aligned than unaligned facts; the gap roughly doubles when the receiver acts politically, and operates primarily through withholding unaligned facts. Incentivized belief elicitation across 1,548 German adults reveal the mechanism: an expressive motive drives baseline selectivity; a persuasion motive—shifting the receiver’s beliefs toward one’s own position—activates specifically when the receiver acts. Professional journalists show the same motive pattern, and their selectivity persists even absent the extrinsic incentives the media-bias literature invokes. The *curated truth* distortion operates entirely on true content, a margin that current policies such as fact-checking and labeling do not reach.

*JEL classification:* D83, D82, C91, D72, L82

*Key words:* Information Provision, Selective Disclosure, Political Information, Media Bias

---

\*Ndiaye: University of Cologne and Max Planck Institute for Behavioral Economics; ndiaye@wiso.uni-koeln.de. I am grateful for valuable comments and suggestions from Arno Appfelstaedt, Lara Maria Berger, Khanh Hoang, Christoph Feldhaus, Markus Dertwinkel-Kalt, Max Rainer Pascal Grossmann, Melisa Kurtis, Felix Mindl, Johannes Münster, Shakked Noy, Axel Ockenfels, Eva Maria Oess, Malte Rattenborg, Christopher Roth, Nicolas Röver, Christoph Schottmüller, Sebastian Siegloch, Till Stange, Johannes Wohlfahrt, and Tom Zimmermann, as well as by participants of seminars and conferences at the C-SEB Early Ideas Workshop and IMPRS BeSmart Thesis Workshop. I thank Jan-Luca Bauss, Frauke Nora Gerster, Linda Plehn, Lena Schmelig, Anselm Bührlein, and Moritz Lakenbrink for their excellent research assistance. Support from the Center for Social and Economic Behavior (C-SEB) and the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation) under Germany’s Excellence Strategy EXC 2126/1390838866 is gratefully acknowledged.

# 1 Introduction

People decide which political information to share everywhere and all the time. For example, everybody decides what to share on social media and what to mention in conversations, journalists and editors decide what to include in their articles, and fact-checkers decide which statements to verify.<sup>1</sup> What is shared matters: information shapes beliefs, beliefs shape political preferences, and in democracies these preferences translate into collective decisions.

Distortions in the supply of political information are widely seen as a threat to democratic decision-making, as evidenced by the policy and research focus on the topic. Both have mostly concentrated on fake news as the culprit. Policy targets it directly: the EU Digital Services Act mandates content moderation, while platform fact-checking programs, content labeling, accuracy nudges, and digital-literacy interventions aim to detect, correct, or pre-empt false content. Research has moved in parallel, studying how false content shapes political beliefs ([Allcott and Gentzkow 2017](#), [Barrera et al. 2020](#)), why senders supply it ([Thaler 2021](#), [Marie et al. 2023](#), [Gurieva et al. 2026](#)), and which remedies curb its spread and effects ([Cagé et al. 2025](#), [Aridor et al. 2024](#), [Berger et al. 2025](#)). Yet false content is a small share of what people consume. It is roughly 0.15 percent of Americans' daily media diet and about 1 percent of news consumption ([Allen et al. 2020](#)),<sup>2</sup> exposure is concentrated among a narrow fringe with the strongest motivation to seek it out ([Budak et al. 2024](#)), and sharing it is rare—in one study linking survey and behavioral data, over 90 percent of users shared no articles from fake-news domains at all ([Guess et al. 2019](#)). Where its effects on political beliefs have been measured, they appear limited and concentrated among the already committed ([Allcott and Gentzkow 2017](#), [Budak et al. 2024](#)). The supply of false information that policy targets and research focuses on is therefore only part of what distorts

---

<sup>1</sup>Ordinary individuals are now routine suppliers of political information: about a third of U.S. adults use social media to voice political views ([Pew Research Center 2025](#)), and on Facebook and Instagram more than 70 percent of those who get news on the platform get it from friends, family, and acquaintances ([Pew Research Center 2024](#)).

<sup>2</sup>[Allen et al. \(2021\)](#) show that Facebook's public-sharing dataset inflates the fake-news share by roughly 4× relative to a representative panel due to URL censoring; the corrected figure is approximately 2.5 percent of Facebook news clicks.

political beliefs.

Selective provision of *true* information is the other part, and it is well-documented observationally. Citizens pass on the fact-checks that favor their side (Shin and Thorson 2017) and politicians amplify the coverage that favors theirs (Freitag et al. 2021); fact-checkers, too, selectively choose which claims to scrutinize (Louis-Sidois 2025), and the pattern holds for true and false content alike (Marie et al. 2023). Selecting among true facts can mislead as much as fabrication: when senders share only the congenial findings and receivers read too little into the silence, beliefs polarize (Bowen et al. 2023), and a true message left unheeded can do as much harm as a false one believed (Feess et al. 2024). Worse, falsehood is at least verifiable and detectable, while withheld truth is *invisible*.

I provide the first causal evidence on the selective provision of true political information and decompose the motives behind it. I cast provision as the outcome of four separable motives. A sender may value correcting the receiver’s factual beliefs; she may gain expressively from passing on facts that fit her own political position; she may want to move the receiver’s beliefs toward that position (projective persuasion); and this persuasive motive gains weight when the receiver’s beliefs may translate into political action (the instrumental motive). Identification rests on two sources of variation: within-sender variation across facts—each sender evaluates aligned and unaligned facts about the same topic, holding sender-level heterogeneity constant—and between-subject assignment to a political-action condition that shifts the return to persuasion. I proceed in three steps: I document selective provision and show that political action amplifies it (Experiment 1), decompose the motives that drive it (Experiments 2 and 3), and target its downstream cost with a receiver experiment (Experiment 4).

Experiment 1 establishes the phenomenon. Among 1,200 quota-sampled U.S. adults, each holding a true immigration fact, provision exceeds 80 percent in every condition, but it is selective: senders are about 8 percentage points more likely to share aligned than unaligned facts ( $p < 0.01$ ), controlling for prior beliefs. The gap roughly dou-

bles, to 16 percentage points, when the receiver makes a downstream political donation ( $p < 0.05$ ), and it widens primarily through the withholding of unaligned facts ( $-4.6$  pp) rather than additional sharing of aligned facts ( $+1.9$  pp). Because the receiver never learns that a disclosure decision was made, the pattern is not driven by reputation or signaling. Experiment 1 cannot, however, say *why* senders select: alignment bundles the expressive value of sharing a congenial fact with its persuasive value, and the two cannot be separated within this design.

Experiment 2 separates them. With 1,548 quota-sampled German adults across two policy domains, I add incentivized belief elicitation that distinguish a fact's expected factual belief correction from its expected persuasive effect, and receivers of known political position that make audience-conditional provision observable. I find that expressive identity drives selective provision absent political action ( $p < 0.01$ ), whereas instrumental persuasion activates specifically when the receiver acts politically ( $p < 0.10$ )—the motive that positional alignment alone cannot isolate. Belief-correction raises provision without making it selective ( $p < 0.05$ ).

Experiment 3 takes the same design to 307 professional German journalists. I find their provision is driven by belief correction and persuasion (both  $p < 0.01$ ; the expressive component is positive but not significant), with larger factual belief correction and persuasion coefficients than the citizen action arm, and it persists once the extrinsic incentives that media-bias models invoke are held constant. *Some* journalistic selectivity is therefore intrinsic. Experiment 4 turns to the receiver: when audiences under-read the silence that selective sharing leaves behind, a diet of entirely true facts can still move beliefs apart, and the experiment is designed to quantify this downstream cost ...

I provide causal evidence on a distortion carried entirely by true facts, and separate the motives behind it: belief correction, expressive value, projective persuasion, and instrumental persuasion. The decomposition points to two policy implications, and existing tools reach neither. Accuracy-based content policy (fact-checking and content-labeling) misses *curated truth* because the distortion rides on the truth of individual

facts rather than on their falsehood. Information quality depends on completeness as well as accuracy, and completeness is the dimension current tools leave untouched; this is a distortion that fact-checking and labeling do not reach. Media-bias remedies premised on extrinsic incentives, such as competition, ownership, and consumer demand, face a separate limit: they cannot reach the intrinsic component. Experiment 3 isolates, the selectivity that survives once those incentives are held constant. The decomposition sharpens the concern, because the instrumental motive is the component policy can least detect, and it grows strongest where beliefs translate into collective decisions—the stakes at which the gap in policy is most costly.

## 1.1 Related Literature

I connect to five strands of literature.

*Individual selective supply.* Citizens selectively pass on the fact-checks that favor their side (Shin and Thorson 2017) and politicians amplify coverage that favors theirs (Freitag et al. 2021); fact-checkers, too, are selective in which claims they scrutinize, with the bias residing in target selection rather than in verdict accuracy (Louis-Sidois 2025, Andreottola and Louis-Sidois 2025). The pattern extends to true and false content alike (Marie et al. 2023), is driven by partisanship (Osmundsen et al. 2021), and intensifies before like-minded audiences (Marie and Petersen 2025). This literature documents that selective provision of true information exists but cannot separate its underlying motives: naturalistic sharing data conflate the expressive benefit of sharing politically congenial facts with the persuasive benefit of shifting receivers' beliefs, and both are confounded with belief-correction concerns in real-world settings. A parallel literature shows that senders supply opinions and stated beliefs strategically for expressive and social-image reasons (Braghieri 2024, Bursztyn et al. 2023b,a, Hillman 2010); I identify the corresponding selective-disclosure mechanism when the content is factual and verifiable. I provide the first motive decomposition, using controlled within-sender variation and incentivized belief elicitation to separate what observational work could only document in aggregate.

*Fake news: effects and decomposition.* A body of work establishes that false content is a small share of consumption whose measured effects on political beliefs are limited and concentrated among the already committed (Allen et al. 2020, Allcott and Gentzkow 2017, Budak et al. 2024, Guess et al. 2019, Feess et al. 2024), that it nonetheless spreads faster and farther than true content (Vosoughi et al. 2018, Lazer et al. 2018), and that voters vary widely in how well they tell true news from false (Angelucci and Prat 2024). The remedies policy deploys reach only adjudicable falsehoods: fact-checks correct false claims imperfectly and incompletely, media literacy interventions show limited durable effects (Barrera et al. 2020, Berger et al. 2025, Cagé et al. 2025), and accuracy nudges that redirect attention to truth reduce false sharing (Pennycook et al. 2021)—none reaching the selective omission of true facts. Guriev et al. (2026) structurally decompose the motives for sharing political content—reputation, partisan persuasion, and partisan signaling—in a setting where the receiver observes that a sharing decision has been made and where content may be true or false; related survey-experimental work studies the checking and sharing of alternative facts (Henry et al. 2022). The two designs partition the information environment: their setting covers visible sharing of true and false content, with motives recovered via a structural model; mine covers verified-true content with unattributable withholding, identifying motives from elicited beliefs without a structural model, and isolating the instrumental persuasion motive through an exogenous political-action treatment rather than recovering it structurally. Thaler (2021) shows that incentivizing senders to be rated truthful can backfire into the supply of *false* motivated beliefs; my object is the mirror image, the selective withholding of verified-true content.

*Institutional media bias.* Outlet-level selectivity arises from extrinsic incentives. Gentzkow and Shapiro (2006) show that profit-maximizing outlets slant toward consumers' prior beliefs to build a reputation for accuracy, with competition disciplining slant by raising the probability of ex-post verification of the true state. Gentzkow and Shapiro (2010) establish, across 433 U.S. daily newspapers, that consumer demand is the dominant driver of slant—accounting for roughly 20 percent of cross-newspaper variance—

while owner ideology contributes negligibly and political capture through advertising is a null (Besley and Prat 2006, Prat and Strömberg 2013, Martin and McCrain 2019). Political control nonetheless shapes what outlets report (Durante and Knight 2012), and exposure is segregated asymmetrically across the ideological spectrum, with platform architecture amplifying partisan disparities in news sharing (González-Bailón et al. 2023, Ershov and Morales 2024). Evidence from journalist movers establishes that individual ideology leaves a fingerprint even within these settings: journalist fixed effects explain around 16 percent of between-outlet slant variation and travel with journalists across outlets (Boxell and Conway 2022); ownership transitions amplify this through ideological sorting of journalists into compatible outlets (Cagé et al. 2022). Yet in naturalistic settings all extrinsic channels operate simultaneously, so a controlled setting is needed to isolate the intrinsic component. Most slant, moreover, lives within outlets rather than across them (Braghieri et al. 2025), so outlet-level accounts—whether from journalist ideology or, as theory shows, from competitive incentives that reduce factual provision in equilibrium (Zhu and Dukes 2015)—miss the margin at which individuals select. Experiment 3 finds that selectivity persists once these extrinsic incentives are held constant—including the financial-incentive channel Berger (2026) documents for a similar population of German online journalists. *Some* journalist selectivity is therefore intrinsic: it survives absent the incentives these models invoke.

*Information demand and belief updating.* Individuals disproportionately seek and value congenial information and avoid disconfirming information (Golman et al. 2022, Peterson and Iyengar 2021, Chopra et al. 2022, Grigorieff et al. 2020), partisans diverge even in their stated factual beliefs (Bullock and Lenz 2019), and the platforms that mediate this demand have their own economics (Aridor et al. 2024); provision experiments show that factual information nonetheless shifts beliefs and behavior (Haaland et al. 2023, Alesina et al. 2023, Barrera et al. 2020, Kuziemko et al. 2015). These papers study the demand side. Chopra et al. (2024) decompose news demand into accuracy and belief-confirmation motives—the demand analog of my supply-side decomposition. Bowen et al. (2023) supplies the downstream theory: when senders pass along

only congenial true news and receivers misread the resulting silence, beliefs polarize even though no signal is false. His “dogmatic” sharer is the behavior my expressive motive micro-founds, and the setting his result needs—selective sharing whose omissions go unnoticed—is the one my design creates.

*Voluntary disclosure.* The disclosure literature shows that informed senders reveal verifiable evidence in equilibrium because receivers can infer bad news from silence—the unraveling principle (Grossman 1981, Milgrom 1981). Unraveling attenuates when receivers are uncertain whether the sender is informed (Dye 1985, Jung and Kwon 1988, Verrecchia 1983) or engage in motivated belief formation (Manili 2024), and laboratory evidence confirms that non-disclosure is common even when unraveling should hold (Jin et al. 2021, Degan et al. 2023, Burdea et al. 2023, Farina and Leccese 2025). Experiments 2 and 3 implement the Dye probabilistic-informed-sender structure directly: a 50% transmission probability means the receiver cannot attribute a missing fact to the sender’s decision—silence is weakly informative but not perfectly so, because non-receipt is equally consistent with a failed transmission. Experiment 1 adds a stronger suspension: the receiver does not know a disclosure decision was made at all, so silence carries no inferential content. Unlike cheap talk (Crawford and Sobel 1982), both designs use verifiable content; unlike full-disclosure settings, silence never perfectly reveals selective omission in either.

Farina et al. (2026) shows that selection (choosing which favorable signals to reveal from a larger pool) is the dominant distortion, with receivers failing to fully discount it. My evidence is true political facts rather than abstract urn draws, enabling direct elicitation of politically motivated sender preferences; non-disclosure is unattributable to the sender rather than revealed by an empty message, eliminating full disclosure as the equilibrium benchmark; the receiver holds a known political position rather than being anonymous, making audience-conditional provision directly observable; and I decompose four sender motives from incentivized belief elicitations rather than recovering behavioral types ex post. My framework also relates to Bayesian persuasion (Kamenica and Gentzkow 2011)—I impose no commitment and study motives for selection rather

than optimal signal design—and the selective-omission model of [Mullainathan and Shleifer \(2005\)](#) is an outlet-level mechanism related to the individual-level mechanism I identify experimentally. Projective paternalism ([Ambuehl et al. 2021](#)), with related work on substituting information for choice restriction ([Grossmann 2024](#), [Bartling et al. 2023](#)), is conceptually related to the projective persuasion motive, which my design isolates at the stage of selecting which facts to present rather than restricting choices.

I structure the rest of the paper as follows. Section 2 presents Experiment 1 as motivating evidence. Section 3 develops a conceptual framework for sender behavior. Section 4 presents Experiment 2. Section 5 extends the design to professional journalists. Section 6 develops the receiver-updating framework and the receiver experiment. Section 7 discusses the implications, and Section 8 concludes.

## 2 Experiment 1

Experiment 1 documents a basic phenomenon: people who know a true fact share it selectively, and the selectivity intensifies when the recipient makes a downstream political decision.

### 2.1 Design, Treatment, and Hypotheses

I run a pre-registered online experiment with  $N = 1,200$  U.S. adults recruited on Prolific using quota-sampling on age, gender, ethnicity, and party identification. I choose policy domain U.S. immigration as (i) opinions are polarized and (ii) misconceptions about factual information are common.

**Experimental Timeline** The experiment proceeds in three parts (Figure 1). In Part 1, I elicit senders' prior knowledge about four immigration facts and measure their policy preferences through a donation decision. In Part 2, I elicit senders' beliefs about the persuasive effect of each fact. In Part 3, senders make one provision decision per fact, each directed at a specific receiver, and a closing questionnaire records motivations,




Figure 1: Experimental timeline, Experiment 1 ( $N = 1,200$  U.S. adults, Prolific, quota-sampled, 2025). Shaded boxes indicate between-subject treatment conditions. The design comprises four cells ( $n = 300$  each) crossing INFO (0/1) and ACTION (0/1). Each sender evaluates two aligned and two unaligned facts across both receiver belief states (within-subject), yielding four provision decisions per sender.

credibility assessments, and demographics. Identification rests on two dimensions varied within sender (alignment between sender and fact; receiver belief state) and two varied between sender (whether the receiver takes a political action; whether senders learn about the effect of information). The primary within-sender dimension is alignment: each sender faces two aligned and two unaligned facts, so the design holds all sender-level heterogeneity fixed. A fact is *aligned* for sender  $i$  if its directional implication favors  $i$ 's political position—that is, if the fact affirms what the sender prefers to be true—and *unaligned* otherwise. Alignment is simultaneously an expressive property (sharing a politically congenial fact) and a persuasive one (an aligned fact also tends

to move a receiver toward the sender's position). The primary between-sender dimension is ACTION: receivers make the same incentivized donation decision as senders in the ACTION condition and complete the experiment without further decisions otherwise. Receivers are unaware that a disclosure decision was made: provided information appears as part of the survey flow, withheld information leaves no trace, and non-disclosure therefore carries no inferential content. The INFO treatment provides exogenous variation in senders' beliefs about how persuasive each fact is; the INFO manipulation is analyzed in Appendix B.4.

**Facts about Immigration** I select four verified-true facts from authoritative sources (Bureau of Justice Statistics, Bureau of Labor Statistics, Customs and Border Protection, National Academies of Sciences), balanced across policy directions, and validate them in a pre-registered survey ( $N = 100$ ) that tested a larger candidate pool against criteria of veracity, perceived credibility, policy relevance, and directional balance (see Appendix B.1). Each is framed as an invertible binary comparison (e.g., "higher" versus "lower"), so beliefs are correct or incorrect rather than varying in degree. I classify crime and labor as pro-immigration and border and fiscal as anti-immigration, and verify the classification using senders' expected attitude shifts (Section 2.3.1). The four statements are:

1. *Crime*: The crime rate of immigrants is *lower* than that of U.S.-born Americans.
2. *Labor*: The labor force participation of immigrants is *higher* than that of U.S.-born Americans.
3. *Border*: The average number of unauthorized border crossings is *higher* in the last five years compared to historical averages.
4. *Fiscal*: When comparing what first-generation immigrants use in government services versus what they pay in taxes, the net cost is *higher* than the same calculation for U.S.-born Americans at the state and local levels.

**Beliefs about Facts** I first elicit senders’ prior knowledge by having them complete each comparative statement, in random order, as “higher,” “lower,” or “don’t know.” Responses are incentivized: correct completions receive a bonus, incorrect completions do not, and “don’t know” responses receive the bonus with 25% probability (Appendix B.2). Senders are then told the correct answer for each. Because senders who already knew a fact may hold it with greater conviction, I control for prior belief accuracy throughout.

**Policy Preferences** I measure senders’ policy preferences through an incentive-compatible donation task: each sender allocates \$500 between the *American Immigration Council* (AIC), which advocates for increased immigration, and *NumbersUSA*, which advocates for decreased immigration. Senders are shown each organization’s position in their own words before deciding. The split yields a continuous measure of policy preference, and therefore of alignment, that captures both the direction and the intensity of each sender’s position. Appendix B.2 reports the implementation details.

**Receiver Belief States** Each fact is assigned to one receiver who either completed the comparative statement incorrectly (false belief) or reported not knowing; the sender observes which case applies before deciding. Both types benefit from disclosure, but correcting a false belief arguably yields a larger expected accuracy gain than updating from ignorance.

**Treatments** Table 1 summarizes the four design dimensions.

Table 1: Experimental design and identifying variation

| Dimension                    | Variation       | Levels                                    |
|------------------------------|-----------------|-------------------------------------------|
| Receiver belief state        | Within-subject  | <i>False belief</i> and <i>Don’t know</i> |
| Alignment with statistic     | Within-subject  | <i>Aligned</i> and <i>Unaligned</i>       |
| Political action (ACTION)    | Between-subject | ACTION = 1 and ACTION = 0                 |
| Effect of information (INFO) | Between-subject | INFO = 1 and INFO = 0                     |

Experiment 1 ( $N = 1,200$  U.S. adults, Prolific, quota-sampled). Within-subject dimensions: each sender evaluates both receiver belief states and all four statistics (two aligned, two unaligned), yielding 4 decisions per sender. Between-subject dimensions: senders are randomly assigned to one of four cells ( $n = 300$  each) crossing ACTION and INFO. The Identification column indicates the utility-function proposition each design dimension is intended to separately identify.

In the ACTION condition, receivers make the same incentive-compatible \$500 donation decision as senders; in the no-ACTION condition, they complete the experiment without it.<sup>3</sup> The INFO treatment provides half of senders with the actual results of Haaland and Roth (2020), generating exogenous variation in beliefs about persuasiveness; the manipulation is analyzed in Appendix B.4. Crossing the two between-sender treatments gives a  $2 \times 2$  factorial design with  $n = 300$  per cell: CONTROL (neither), A (ACTION only), I (INFO only), and A  $\times$  I (both).

**Hypotheses** I pre-registered four hypotheses about how senders provide information.

**Hypothesis 0** (Baseline provision). *Provision rates for both aligned and unaligned information exceed the indifference benchmark of 50% in all treatments.*

**Hypothesis 1** (Alignment-based selectivity). *Senders provide aligned information at strictly higher rates than unaligned information in treatment CONTROL.*

**Hypothesis 2** (Political action amplifies selectivity). *The alignment gap in provision increases when receivers take political actions.*

**Hypothesis 3** (Magnitude of preference shift amplifies selectivity). *The alignment gap in provision is increasing in senders' beliefs about the magnitude of the attitude shift that disclosure induces.*

I test Hypothesis 3 two ways: through the between-sender INFO treatment, which shifts beliefs about persuasiveness, and through within-sender variation in expected attitude change, which provides a correlational test.

## 2.2 Sample, Incentives, and Protocol

The sender experiment was run on Prolific between July 16 and August 12, 2025, quota-sampled on age, gender, ethnicity, and party identification to approximate the U.S.

---

<sup>3</sup>Through randomization any difference in alignment-based selectivity is attributable to the downstream political consequence, under the maintained assumption that ACTION operates only through that consequence. ACTION may also shift the salience of the decision and the sender's construal of the task.

adult population. (See Table A.1 for details.) This study was approved by the Ethics Council of the Max Planck Society under the Generalized Approval of Experiments Following the Protocol that is Standard in Experimental Economics (approval no. 2018\_3 Renewal 2024\_28) and was pre-registered on AsPredicted (No. 238343). Of 1,257 starters, I stopped collection at the pre-registered target of 300 per treatment passing the attention check, not switching tabs during the fact questions, and not timing out. Senders took a median of 17.26 minutes and earned \$3.32 on average. Receiver data were collected on Prolific between September 1 and September 10, 2025; 96 receivers took a median of 3 minutes and earned \$0.80 on average.

## 2.3 Results

### 2.3.1 Senders' Beliefs and Policy Preferences

Before turning to treatment effects, I test the prerequisite for my hypotheses: senders have policy preferences. Additionally, I show that senders believe information can shift receivers' policy preferences and that they hold heterogeneous prior beliefs about the facts with partisan sorting.

**Senders' Beliefs about Facts** I find substantial heterogeneity in senders' factual knowledge.<sup>4</sup> A majority correctly identify facts about immigrants' crime rates (70%), labor force participation (79%), and recent border crossings (61%), but only 30% correctly assess the net fiscal impact of first-generation immigrants. Democrats and immigration supporters are more accurate on pro-immigration facts (crime and labor) and Republicans and immigration opponents on anti-immigration facts (border and fiscal; all differences  $p < 0.001$ ). This partisan sorting validates my classification of facts as pro- or anti-immigration and motivates controlling for senders' prior beliefs throughout.

---

<sup>4</sup>See Table A.2 for details. The fiscal fact, where prior knowledge is lowest, is rated as credible ex post as the other three (Table A.3) and provided at comparable rates; results are robust to leaving out each fact individually (Appendix Table A.9).

**Policy Preferences** The distribution of senders' donation decisions is markedly bimodal (Figure A.2): 30.2% allocate their entire endowment to the pro-immigration group, 14.1% give everything to the anti-immigration group, and only 7% split equally. Donations correlate strongly with stated policy preferences ( $r = 0.56$ ,  $p < 0.001$ ) and political affiliation ( $p < 0.001$ ), providing the variation in alignment needed to test for selective provision.

**Senders' Beliefs about the Effect of Information on Political Preferences** Senders already believe that information strongly affects policy preferences (Figure A.1): roughly 90% expect information to increase petition signatures, and 73% correctly identify that individuals with initially false beliefs update more strongly. They do, however, underestimate the magnitude—predicting a 38% increase on average against a true effect of 69%—leaving room for the INFO condition to shift beliefs upward.

**Expected Attitude Effects of Information** Alignment-based selectivity requires that senders expect each fact to shift receivers' preferences in a predictable direction. For each sender  $i$  and fact  $j$ , I define the expected attitude change  $\Delta\text{Attitude}_{i,j} \in [-1, 1]$  as the net share of ten hypothetical receivers whom the sender expects to become more (rather than less) supportive of immigration after receiving fact  $j$ , with 1 (−1) implying all receivers become more (less) supportive of immigration when seeing fact  $j$ .

Senders expect information to shift attitudes in the predicted directions. Providing the crime fact is expected to shift net attitudes 25 percentage points toward pro-immigration and labor 17 pp; the border fact is expected to shift net attitudes 34 pp toward anti-immigration and fiscal 32 pp (all  $p < 0.001$ ; Appendix Table A.4). These expected attitude shifts establish that senders believe disclosure moves receivers' policy attitudes in the hypothesized directions and corroborate the classification. The INFO treatment, however, generates only modest and one-sided variation in expected effect on attitudes, raising it for the pro-immigration facts alone (Appendix Table A.4; fuller detail in Appendix B.4).




Figure 2: Mean provision rates by treatment and statistic alignment, Experiment 1 (U.S. adults,  $N = 1,200$ , Prolific, quota-sampled). Panel (A): mean sender-level provision rates for aligned (dark green) and unaligned (dark red) statistics, by treatment cell (CONTROL, A=Action, I=Info, A×I=Action×Info). Panel (B): alignment bias (aligned minus unaligned provision rate) by treatment cell. Restricted to senders with revealed directional preference ( $N = 1,114$ , 93%;  $AIC_i > 250$  = pro-immigration,  $AIC_i < 250$  = anti-immigration,  $AIC_i = 250$  excluded). Error bars are 95% confidence intervals.

### 2.3.2 Treatment Effects on Selective Provision of Information

Let  $y_{i,j}$  equal 1 if sender  $i$  provides fact  $j$ , and let  $AIC_i \in [0, 500]$  denote  $i$ 's donation to the pro-immigration group (AIC). Sender  $i$  is pro-immigration if  $AIC_i > 250$  and anti-immigration if  $AIC_i < 250$ ; the few with  $AIC_i = 250$  reveal no directional preference and are excluded from the sender-level analysis. I classify  $\mathcal{J}^+ = \{\text{crime, labor}\}$  as pro-immigration and  $\mathcal{J}^- = \{\text{border, fiscal}\}$  as anti-immigration. Sender  $i$ 's aligned set is  $\mathcal{A}_i = \mathcal{J}^+$  if pro-immigration and  $\mathcal{J}^-$  otherwise, with the unaligned set  $\mathcal{U}_i$  defined analogously.

I define the sender's alignment bias as:

$$\text{Alignment Bias}_i = \frac{1}{|\mathcal{A}_i|} \sum_{j \in \mathcal{A}_i} y_{i,j} - \frac{1}{|\mathcal{U}_i|} \sum_{j \in \mathcal{U}_i} y_{i,j}. \quad (1)$$

A positive alignment bias indicates selective provision favoring the sender's own policy position. A zero alignment bias indicates no selectivity, which includes both sharing all information and sharing no information.<sup>5</sup>

<sup>5</sup>At the sender level, I test Hypothesis 0 using a one-sample  $t$ -test of the unaligned provision rate

I now test whether senders selectively provide accurate information and whether political action amplifies this selectivity. Figure 2 summarizes senders' provision decisions, restricting the sender-level analysis to the 1,114 senders (93%) who revealed a directional policy preference.

Panel (A) shows that senders overwhelmingly share accurate information across all four treatment conditions. Unaligned provision rates range from 81.9% (A) to 86.5% (CONTROL), all significantly above 50% ( $p < 0.001$ , one-sample  $t$ -tests).

**Result 0** (Information provision is high). *Senders share accurate information at high rates, including information unaligned with their policy preferences ( $p < 0.001$  against the 50% benchmark in all treatments).*

Panel (B) reveals that provision is selective. The alignment bias is significantly positive in all four conditions ( $p < 0.01$ ), ranging from 6.7 to 13.3 pp. In both pairwise comparisons holding the INFO condition constant, the alignment bias is larger under ACTION (Appendix Table A.6).

**Decision-Level Analysis** I capture within-sender variation in preference intensity with a continuous alignment measure:

$$\text{Align}_{i,j} = \begin{cases} \frac{AIC_i - 250}{250} & \text{if } j \in \mathcal{J}^+ \\ \frac{AIC_i - 250}{250} \cdot (-1) & \text{if } j \in \mathcal{J}^- \end{cases}$$

This measure is bounded in  $[-1, 1]$ : positive values indicate alignment with sender  $i$ 's preference (relative to the midpoint  $AIC_i = 250$ ), negative values misalignment, and  $|\text{Align}_{i,j}|$  captures preference intensity. Because  $\text{Align}_{i,j} = 0$  for indifferent senders, it incorporates all senders. Equivalently,  $\text{Align}_{i,j} = \text{Pref}_i \times \text{Dir}_j$ , where  $\text{Pref}_i = (AIC_i - 250)/250$  is revealed preference intensity and  $\text{Dir}_j = +1$  if  $j \in \mathcal{J}^+$  and  $-1$  if  $j \in \mathcal{J}^-$ .<sup>6</sup>

against 0.5; Hypothesis 1 using a one-sample  $t$ -test of Alignment Bias $_i > 0$  in CONTROL; and Hypotheses 2 and 3 using directed two-sample  $t$ -tests comparing Alignment Bias $_i$  across the relevant treatment conditions. I Holm-correct tests for Hypotheses 2 and 3.

<sup>6</sup>The pre-registration defined alignment on  $[0, 1]$ . Centering at the indifference point to ease interpretation is a one-to-one transformation that leaves all hypothesis tests invariant, so results under the pre-registered coding are identical.

Let  $\text{Belief}_{i,j}$  denote sender  $i$ 's prior accuracy about fact  $j$  (1 if correct, 0.5 if “don't know,” and 0 if incorrect as pre-registered), capturing the stronger accuracy norm faced by senders who already knew a fact was true; without this control, the uniformly higher provision rates for pre-known facts would conflate the belief-correction motive with the alignment effect. I estimate the pre-registered linear probability model:

$$\begin{aligned} y_{i,j} = & \beta_0 + \beta_1 \text{Align}_{i,j} + \beta_2 \text{ACTION}_i + \beta_3 \text{INFO}_i + \beta_4 (\text{Align}_{i,j} \times \text{ACTION}_i) \\ & + \beta_5 (\text{Align}_{i,j} \times \text{INFO}_i) + \beta_6 (\text{ACTION}_i \times \text{INFO}_i) \\ & + \beta_7 (\text{Align}_{i,j} \times \text{ACTION}_i \times \text{INFO}_i) + \gamma_1 \text{Belief}_{i,j} + \gamma_2 \mathbf{S}_j + \gamma_3 \mathbf{X}_i + \eta_{i,j}, \end{aligned} \quad (2)$$

with standard errors clustered at the sender level, fact fixed effects  $\mathbf{S}_j$ , and demographic and political controls  $\mathbf{X}_i$ . Column (1) includes no fixed effects or controls; column (2) adds  $\mathbf{S}_j$ ; column (3) additionally adds demographic and political controls  $\mathbf{X}_i$  and is the preferred specification.  $\beta_1$  captures baseline alignment-based selectivity (Hypothesis 1),  $\beta_4$  whether alignment matters more when receivers act politically (Hypothesis 2), and  $\beta_5$  whether exogenously shifting beliefs about persuasiveness amplifies selectivity (Hypothesis 3). The control  $\gamma_1$  estimates whether senders who knew a fact provide it at higher rates.

Table 2 reports the results.<sup>7</sup> Across all specifications, provision increases significantly in political alignment. A shift from maximally unaligned to maximally aligned increases provision by 7.2 to 8.2 pp ( $p < 0.003$ ).

**Result 1** (Provision is higher for aligned than unaligned information). *Provision is selectively biased: senders are 7 to 8 pp more likely to share aligned information ( $p < 0.003$ ).*

The interaction  $\text{Align}_{i,j} \times \text{ACTION}_i$  is positive and statistically significant across all specifications ( $p < 0.05$ ). Political action widens the alignment gap by an additional

<sup>7</sup>Appendix Table A.8 reports specifications allowing asymmetric INFO effects by fact direction to accommodate the asymmetric first-stage effect. The additional terms  $\text{INFO}_i \times \text{PROSTAT}_j$  and  $\text{Align}_{i,j} \times \text{INFO}_i \times \text{PROSTAT}_j$  are not statistically significant in the specifications with fact fixed effects. Wald tests fail to reject that INFO changes the alignment gradient for pro-immigration facts ( $p = 0.403$ ) or anti-immigration facts ( $p = 0.606$ ), and fail to reject that INFO changes overall provision for pro-immigration facts ( $p = 0.390$ ).

Table 2: Treatment Effects on Information Provision

|                       | Probability of Statistic Provision |                     |                     |
|-----------------------|------------------------------------|---------------------|---------------------|
|                       | (1)                                | (2)                 | (3)                 |
| Intercept             | 0.869***<br>(0.015)                |                     |                     |
| Align                 | 0.041***<br>(0.012)                | 0.036***<br>(0.012) | 0.037***<br>(0.012) |
| INFO                  | −0.013<br>(0.017)                  | −0.012<br>(0.017)   | −0.012<br>(0.017)   |
| ACTION                | −0.002<br>(0.016)                  | −0.002<br>(0.016)   | −0.003<br>(0.016)   |
| Belief Accuracy       | 0.043***<br>(0.011)                | 0.034***<br>(0.012) | 0.033***<br>(0.012) |
| Align × INFO          | 0.006<br>(0.017)                   | 0.004<br>(0.017)    | 0.004<br>(0.017)    |
| Align × ACTION        | 0.041**<br>(0.018)                 | 0.041**<br>(0.018)  | 0.041**<br>(0.018)  |
| INFO × ACTION         | −0.007<br>(0.024)                  | −0.007<br>(0.024)   | −0.008<br>(0.024)   |
| Align × INFO × ACTION | −0.028<br>(0.026)                  | −0.027<br>(0.026)   | −0.027<br>(0.026)   |
| Clustered SE          | Subject                            | Subject             | Subject             |
| Statistics FE         | No                                 | Yes                 | Yes                 |
| Controls              | No                                 | No                  | Yes                 |
| No. of Senders        | 1200                               | 1200                | 1200                |
| Num. obs.             | 4800                               | 4800                | 4800                |
| R <sup>2</sup>        | 0.030                              | 0.033               | 0.040               |
| Adj. R <sup>2</sup>   | 0.028                              | 0.031               | 0.034               |

*Sample:*  $N = 1,200$  U.S. adults recruited on Prolific (quota-sampled), 4,800 sender-statistic observations (4 statistics per sender). *Outcome:* binary provision indicator. *Align:*  $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, 1]$ , where Pref is derived from an incentive-compatible donation split. *Belief Accuracy:* incentivized prior belief correctness, elicited before senders knew they would make a provision decision. *Info:* between-subject treatment sender receives personalized feedback comparing their prediction to the actual results of Haaland and Roth (2020) (see Appendix B.4). *Action:* between-subject treatment receiver makes a downstream political donation. *Controls:* age, gender, education, party affiliation, political knowledge. Standard errors clustered at sender level. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

8.2 pp, approximately doubling the total effect from 7.2–8.2 pp to 15.4–16.4 pp.<sup>8</sup>

The action effect operates primarily by reducing provision of *unaligned* information rather than by increasing provision of aligned information. Moving from CONTROL to A reduces unaligned provision by 4.6 pp (from 86.5% to 81.9%) while increasing aligned provision by only 1.9 pp (from 93.3% to 95.2%). This asymmetry reflects, in

<sup>8</sup>Table A.7 shows that the results are robust to logistic regression estimation.

part, a mechanical ceiling: aligned provision already exceeds 93%, leaving limited margin for further increase.

**Result 2** (Selectivity is higher when receivers take a downstream political action). *Alignment-based selectivity approximately doubles when receivers make a downstream political action ( $p < 0.05$ ). The action effect operates primarily through reduced provision of unaligned information rather than increased provision of aligned information.*

Senders who directly know a fact to be true provide it at higher rates (3.3 to 4.3 pp,  $p < 0.01$ ). This effect is smaller in magnitude than the alignment effect and does not attenuate the Align  $\times$  ACTION interaction.

Neither the main effect of INFO nor its interactions with alignment are statistically distinguishable from zero, including in specifications that allow asymmetric effects for the pro-immigration facts where INFO successfully shifted beliefs about persuasiveness. This null has two informative interpretations. First, the first-stage results in Section 2.3.1 show that INFO generated only modest and asymmetric shifts in expected persuasiveness, so the manipulation had little to amplify. Second, conditional on the first stage, the pattern indicates that senders respond to the *direction* of expected persuasion, not its magnitude: provision is unchanged when the direction is fixed and only the expected degree of attitude change moves.

**Result 3** (Beliefs about the effect of information do not affect selectivity). *Information about the effect of information does not affect the likelihood of overall information provision nor the selectivity of information provision.*

### 2.3.3 Heterogeneous Treatment Effects

Having established that alignment-based selectivity exists at baseline and approximately doubles under political action, I now ask who drives each component, examining heterogeneity by political polarization as pre-registered. Two distinct groups drive the two effects. The baseline alignment bias is concentrated among strong partisans: moving from non-partisanship to full partisanship raises it by 14 percentage points

( $p = 0.017$ ), while Independents show no baseline bias. The action-induced amplification, by contrast, is driven by Independents: political action raises their alignment bias by 14.6 pp ( $p < 0.05$ ) and insignificantly lowers it for strong partisans ( $-10.2$  pp,  $p = 0.214$ ). Appendix Table A.10 reports the full specification by three measures of polarization.<sup>9</sup>

Receiver belief state does not drive provision decisions. Each fact is matched to a receiver who either holds a false belief or reports not knowing; a belief-correction motive would predict higher provision toward false-belief receivers, whose accuracy gap is larger.<sup>10</sup> Table A.13 shows this does not obtain: senders are not more likely to share with false-belief receivers ( $-2.5$  pp,  $p = 0.17$ ), and the alignment gap is comparable across both receiver types (5.2 pp for don't-know receivers, 9.0 pp for false-belief receivers; difference  $p = 0.39$ ). Senders respond to the *direction* of the fact, not the epistemic state of the recipient.

### 2.3.4 Mechanism-Consistent Correlational Evidence

The observed selective provision could reflect veracity concerns: senders may share what they find credible and withhold what they doubt, and they may find aligned facts more credible. Belief-correction measures predict the *level* of provision but not its alignment-based selectivity. Facts that are deemed credible ex-post are 25.6 to 26.9 pp more likely to be shared ( $p < 0.01$ ), yet adding ex-post credibility and prior belief accuracy and their interactions with ACTION leaves the alignment effect (6 to 6.2 pp,  $p < 0.05$ ) and the Align  $\times$  ACTION interaction (9 to 9.4 pp,  $p < 0.05$ ) intact, and neither veracity concern measure interacts with ACTION (both null; Appendix Table A.15). When receivers make political decisions, senders do not raise their epistemic standards; they become more selective along the alignment dimension. The selectivity is thus not an epistemic-standards story. Senders' own self-reported motivations, cor-

<sup>9</sup>The sender-level specification is Alignment Bias<sub>*i*</sub> =  $\pi_0 + \pi_1 A_i + \pi_2 Z_i + \pi_3 A_i Z_i + \gamma' X_i + \eta_i$ , with moderator  $Z_i$  and controls  $X_i$ , estimated on CONTROL and treatment A. Appendix Table A.11 pools across the INFO condition and Appendix Table A.14 reports the decision-level specification; conclusions are identical in direction and significance.

<sup>10</sup>Senders also do not expect larger attitude shifts for false-belief than for don't-know receivers (Appendix Table A.12).

related within sender with their alignment bias, line up with the behavioral pattern: the reasons senders give for sharing and withholding track the direction of their selectivity, and the expected persuasion only marginally predicts selectivity (Tables A.15 and A.16).

### 3 Conceptual Framework

I develop a conceptual framework of selective disclosure of true information. A sender who knows a true fact decides whether to disclose it to a receiver. She may value improving the receiver's accuracy, derive expressive utility from sharing a politically congenial state of the world, value moving the receiver's beliefs or preferences toward her own position, or—when the receiver may act politically—care about the political consequences of that movement. I nest these four motives in a single sender-side utility function and derive behavioral predictions. The framework characterizes the sender's private utility rather than imposing a full receiver-inference structure in the main text, departing in purpose from classical disclosure models (Grossman 1981, Milgrom 1981) that study equilibrium inference.<sup>11</sup>

#### 3.1 Environment

Consider a sender, a receiver, and a state of the world. First, nature determines the true state of the world  $\omega \in \{0,1\}$ , the sender's political position, and whether the receiver's policy preference may translate into political action,  $A \in \{0,1\}$ . Without loss of generality, I normalize the true state of the world to  $\omega = 1$ .

Second, the sender observes the true state of the world and chooses whether to disclose it,  $d = 1$ , or withhold it,  $d = 0$ . Finally, if  $A = 1$ , the receiver's resulting policy-relevant position may translate into political action.

The sender evaluates disclosure using four objects. First, let  $b^0$  denote the belief that the sender expects the receiver to hold if she withholds the state of the world, and

---

<sup>11</sup>Appendix C embeds this utility representation in a receiver-inference environment and characterizes equilibrium behavior.

let  $b^1$  denote the belief that she expects the receiver to hold if she discloses it. Both are the sender's subjective forecasts of the receiver's posterior probability that  $\omega = 1$ , with  $b^d \in [0, 1]$ . Let  $\Delta b \in [-1, 1]$  denote the factual belief correction the sender expects disclosure to produce:

$$\Delta b = |\omega - b^0| - |\omega - b^1|.$$

As  $\omega = 1$ , larger values of  $\Delta b$  correspond to larger expected movements of the receiver's belief toward the truth.

Second, let  $\alpha \in \mathbb{R}$  denote the sender's assessment of the state of the world's congeniality to her political position. Positive values of  $\alpha$  indicate that the sender regards the state of the world as affirming her political position. Negative values indicate that she regards the state of the world as not affirming her political position. A value of zero denotes a state of the world she regards as politically neutral. Expressive utility arises from the act of disclosing the state of the world, not from privately knowing it; the sender receives no expressive payoff from withholding.

Third, let  $\Delta\pi \in \mathbb{R}$  denote the movement in the receiver's policy-relevant position toward the sender's own position that the sender expects disclosure to produce. Thus,  $\Delta\pi > 0$  means that the sender expects disclosure to move the receiver toward her own position, while  $\Delta\pi < 0$  means that she expects disclosure to move the receiver away from it. The distinction between  $\alpha$  and  $\Delta\pi$  is important. The former captures how the sender regards the state of the world's political congeniality to her own position; the latter captures what the sender expects disclosure to do to the receiver.

Finally,  $A \in \{0, 1\}$  indicates whether the receiver's policy-relevant position may translate into political action. When  $A = 0$ , the sender may still value moving the receiver's beliefs or preferences toward her own position. When  $A = 1$ , such movement also has potential instrumental value because it may affect political outcomes.

### 3.2 Sender Utility

The sender's utility from disclosure combines four motives. First, a *belief-correction motive*,  $\theta_B \geq 0$ : the sender values correcting the receiver's factual beliefs, independent of political content. Second, an *expressive motive*,  $\theta_E \geq 0$ : the sender derives utility from disclosing the state of the world when it affirms her own political position, independent of its expected effect on the receiver. Third, a *projective persuasion motive*,  $\theta_P \geq 0$ : the sender values moving the receiver's policy-relevant beliefs or preferences toward her own position, even absent downstream action. Fourth, an *instrumental persuasion motive*,  $\theta_I \geq 0$ : when the receiver's beliefs or preferences may translate into political action, the sender places additional value on moving the receiver toward her own position.

I normalize the utility of withholding to zero,  $U(d = 0) \equiv 0$ , so that all terms represent gains from disclosure relative to withholding. Combining these motives, the sender's net utility from disclosure is

$$U(d = 1) = \theta_B \Delta b + \theta_E \alpha + (\theta_P + \theta_I A) w(\Delta \pi), \quad (3)$$

where  $w : \mathbb{R} \rightarrow \mathbb{R}$  maps the sender's perceived movement in the receiver's policy-relevant position into utility. I assume that  $w$  is sign-preserving:  $w(0) = 0$ ,  $w(\Delta \pi) > 0$  for  $\Delta \pi > 0$ , and  $w(\Delta \pi) < 0$  for  $\Delta \pi < 0$ . For derivative-based predictions, I additionally assume that  $w$  is differentiable and strictly increasing on the relevant domain. Define  $c(A) = \theta_P + \theta_I A$  as the composite persuasion weight on  $w(\Delta \pi)$ , combining the projective sub-motive ( $\theta_P$ , present whether or not the receiver may act) and the instrumental sub-motive ( $\theta_I A$ , present only when  $A = 1$ ).

The additive specification separates the four motives conceptually. The belief-correction term depends on the perceived factual belief correction from disclosure. The expressive term depends on whether the state of the world itself affirms the sender's political position. The projective term depends on the perceived movement of the receiver toward or away from the sender's position. The instrumental term places additional

value on that same movement when the receiver's beliefs or preferences may translate into political action; it is a distinct motive operating on the same persuasion channel, not an amplification of the projective persuasion motive.<sup>12</sup>

### 3.3 Disclosure Decision

The sender discloses if and only if the utility from disclosure exceeds an idiosyncratic disclosure shock:

$$d = 1 \quad \text{iff} \quad U(d = 1) > \varepsilon.$$

The shock  $\varepsilon$  is drawn from a distribution  $F$  with full support on  $\mathbb{R}$  and strictly positive density  $f = F'$ . It captures unmodeled determinants of disclosure, such as attention, salience, idiosyncratic features of the sender-receiver interaction, or idiosyncratic inclinations toward disclosure.

The probability of disclosure is therefore

$$\Pr(d = 1) = F(\theta_B \Delta b + \theta_E \alpha + (\theta_P + \theta_I A)w(\Delta \pi)). \quad (4)$$

### 3.4 Predictions

The framework yields four behavioral predictions. The first is a non-political belief-correction benchmark. The next two distinguish two forms of political selectivity: expressive selectivity and projective selectivity. The fourth identifies the instrumental motive, the component of persuasion that operates only when the receiver may act politically. All proofs are in Appendix C.

**Proposition 0.** *If  $\theta_B > 0$ , disclosure is strictly increasing in the perceived factual belief correc-*

---

<sup>12</sup>The additive structure follows [Bénabou and Tirole \(2006\)](#) in nesting distinct motives in a single utility representation. The model nests pure belief-correction, pure expressive disclosure, pure projective disclosure, and pure instrumental disclosure as special cases. The expressive motive draws on the social-image account of communication, in which senders gain utility from transmitting position-affirming content ([Braghieri 2024](#)).

tion from disclosure:

$$\frac{\partial \Pr(d = 1)}{\partial \Delta b} > 0.$$

Proposition 0 is the belief-correction benchmark. Holding political content and expected persuasion fixed, senders with a belief-correction motive are more likely to disclose the state of the world when they perceive disclosure to correct the receiver's factual beliefs more.

**Proposition 1a** (Expressive selective disclosure). *If  $\theta_E > 0$ , disclosure is strictly increasing in congenial alignment:*

$$\frac{\partial \Pr(d = 1)}{\partial \alpha} > 0.$$

Proposition 1a captures expressive disclosure. Holding factual belief correction and expected receiver movement fixed, senders with an expressive motive are more likely to disclose the state of the world when it affirms their own political position.

**Proposition 1b** (Persuasive selective disclosure). *If  $c(A) > 0$ , disclosure is strictly increasing in  $w(\Delta\pi)$ :*

$$\frac{\partial \Pr(d = 1)}{\partial w(\Delta\pi)} > 0.$$

*If, in addition,  $w$  is differentiable and strictly increasing, disclosure is strictly increasing in  $\Delta\pi$ :*

$$\frac{\partial \Pr(d = 1)}{\partial \Delta\pi} > 0.$$

Proposition 1b captures projective selective disclosure. Holding factual belief correction and expressive alignment fixed, senders are more likely to disclose the state of the world when they perceive it as moving the receiver toward their own position, and less likely to disclose it when they perceive it as moving the receiver away.

When  $\text{sign}(\alpha) = \text{sign}(\Delta\pi)$ —the typical case for a politically aligned fact—Propositions 1a and 1b predict the same direction. When the signs differ, they diverge: the expressive motive favors disclosure of facts the sender finds politically congenial ( $\alpha > 0$ ) regardless of persuasive direction, while the projective persuasion motive favors disclosure only when the fact is expected to move the receiver toward the sender’s position ( $\Delta\pi > 0$ ).

**Proposition 2** (Political action amplification). *If  $\theta_I > 0$ , political action strictly increases the marginal utility of persuasion:*

$$\frac{\partial^2 U(d=1)}{\partial w(\Delta\pi) \partial A} = \theta_I > 0.$$

Proposition 2 is stated at the utility-index level. This is the object directly disciplined by the utility function. Without additional restrictions on the distribution  $F$ , the model does not require a uniformly positive interaction effect in disclosure probabilities for all parameter values. The behavioral prediction is that the latent value of moving the receiver toward the sender’s position is larger when the receiver’s beliefs or preferences may translate into political action; this increment in the value of persuasion is the instrumental motive, a sub-motive distinct from the always-present projective persuasion motive.

The framework also yields a magnitude prediction: when  $w$  is strictly increasing, stronger expected persuasion raises disclosure for facts that move the receiver toward the sender’s position and lowers it for facts that move the receiver away. I develop this prediction formally in Appendix C, where the boundary case  $w(\Delta\pi) = \text{sign}(\Delta\pi)$  marks the threshold at which the sender responds only to the direction of persuasion and not to its magnitude.

## 4 Experiment 2

Experiment 1 establishes three findings: provision is high, there is a baseline alignment gap consistent with an expressive motive ( $\theta_E$ ) or a projective persuasion motive ( $\theta_P$ ).

Additionally, political action approximately doubles the gap consistent with an instrumental persuasion motive ( $\theta_I$ ) or expression increasing under action. Experiment 1's design cannot decompose these and therefore, three questions remain open: (i) which component of the baseline gap is expressive and which is projective persuasion, (ii) is the increase in selectivity explained by expression or instrumental persuasion, and (iii) what the H3 null implies about  $w(\Delta\pi)$ . We adjust the experimental design in Experiment 2 to answer these questions.

## 4.1 Design, Treatment, and Hypotheses

I run a pre-registered online experiment with  $N = 1,548$  German adults recruited on Bilendi using quota-sampling on age, gender, and state (Bundesland).<sup>13</sup> I study two policy domains assigned between subjects: immigration policy and climate policy, extending the topic space beyond Experiment 1.

**Experimental Timeline** The experiment proceeds in three parts, followed by a post-questionnaire (Figure 3). Part 1 elicits senders' demographics, stated policy preferences and revealed policy preferences. Part 2 is the information provision task: senders decide whether to share each of five assigned facts, one per decision, directed at receivers whose policy stance is visible. An attention check follows Part 2. Part 3 elicits incentivized beliefs about how each fact corrects receivers' factual beliefs and shifts their policy views. In the post-questionnaire subjects rate two facts from the other domain in terms of their congeniality and trustworthiness. Two timing design choices differ from Experiment 1. First, provision decisions precede belief elicitations: placing provision before elicitation ensures that elicitation does not influence provision behavior, and the Part 3 elicitations provide a timing-independent check on the behavioral results. Second, comprehension checks gate the Part 2 bonus: senders must answer all three questions correctly on the first attempt; the questions loop until all answers are

---

<sup>13</sup>Two cumulative exports of the final panel batch overlapped; eight completers appearing in both exports were deduplicated. The speeder cutoff (one quarter of the median completion time), recomputed on the deduplicated sample, excluded one additional participant.

correct.




Figure 3: Experimental timeline for Experiment 2 ( $N = 1,548$  quota-sampled German adults, Bilendi, 2026; quota by age, gender, and Bundesland). Shaded boxes indicate between-subject arm assignment. Each sender makes five provision decisions (one per assigned fact). Post-provision belief elicitations ( $\Delta\pi$ ,  $\Delta b$ ) are incentivized via Quadratic Scoring Rule.

**Facts** I select six verified-true facts per topic, balanced for credibility and relevance in a pre-registered pre-validation survey ( $N = 300$ , AsPredicted #290989), see Appendix B.3. Each fact is presented as a binary, invertible comparison with a documented correct answer as in Experiment 1, sourced from German statistical offices and research institutes. Table 3 lists the twelve statements. All participants are informed at the outset that every fact is verified true and sourced, and receivers know the same.

**Policy Preferences** I measure senders' policy preferences in two ways. First, senders indicate their policy position on a five-point Likert scale. Second, senders split €250 between the pro-policy and anti-policy advocacy organization for their topic in a real donation decision: one decision per session is implemented with actual money. For immigration, the organizations are ProAsyl (advocates open migration policy) and Republik21 (advocates more restrictive policy); for climate, BUND (supports ambitious

Table 3: Facts Used in Experiments 2 and 3

| Direction                 | Fact                                                                                                                                                       |
|---------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <i>Immigration policy</i> |                                                                                                                                                            |
| Pro                       | The share of people with a migration background in shortage-occupation fields is <i>higher</i> than their population share.                                |
| Pro                       | The absolute increase in employed foreign nationals between 2013 and 2023 was <i>larger</i> than for German nationals.                                     |
| Pro                       | <i>More</i> than half of the persons attending a general integration course for the first time complete it without interruption through the final module.  |
| Anti                      | Among persons suspected of violent crimes, the share of non-German nationals is <i>higher</i> than their population share.                                 |
| Anti                      | The share of persons receiving basic income support ( <i>Bürgergeld</i> ) is <i>higher</i> among foreign nationals than among German nationals.            |
| Anti                      | Germany records <i>more</i> persons without a valid residence permit per capita than the EU average.                                                       |
| <i>Climate policy</i>     |                                                                                                                                                            |
| Pro                       | In German electricity generation in 2023 and 2024, the share of renewable energies was <i>higher</i> than the share of fossil fuels.                       |
| Pro                       | The cost of generating one kilowatt-hour of electricity from new, ground-mounted solar installations is <i>lower</i> than from new gas-fired power plants. |
| Pro                       | The percentage increase in renewable energy jobs in Germany between 2022 and 2023 was <i>higher</i> than the global average.                               |
| Anti                      | Grid fees for household customers are expected to be <i>more</i> than double their current level by 2045 due to grid expansion.                            |
| Anti                      | There are <i>fewer</i> companies that view the effects of the energy transition positively than companies that view them negatively.                       |
| Anti                      | At higher shares of solar power, <i>more</i> price fluctuations are observed in European electricity markets.                                              |

Notes: Verbatim claims as shown to senders in English translation (correct answer in *italics*). Direction (Pro) = supports more-permissive immigration or more-ambitious climate policy; (Anti) = supports more-restrictive or more-cautious policy. Pool of six facts per topic, all verified true from German statistical offices and research institutes. Each sender draws five facts: two pro-policy, two anti-policy, and one random tiebreaker. Topics assigned between subjects; all senders see one topic. Climate policy extends the topic space beyond Experiment 1. Pre-validation details in Appendix B.3.

climate targets) and Stiftung Marktwirtschaft (emphasizes economic costs of the transition). Senders are shown each organizations position in their own words before deciding.

**Receiver Policy Stance** Each sender faces five receivers sequentially, one per position on the five-point Likert scale policy stance. Before each provision decision, the sender observes the receiver's stance as a full sentence and knows the receiver in turn is told their policy position. This mutually visible stance structure generates within-sender

variation in sender-receiver ideological distance. Senders know that receivers are told a fact may or may not have been shared with them — with 50% probability the match occurs and the receiver receives the fact (if sent) — so the receiver cannot attribute a missing fact to the sender’s deliberate withholding; silence remains only weakly informative.

**Beliefs about Belief Correction** For each of the five receiver-fact pairs, senders predict the share of receivers with the same stance who will identify the fact correctly after disclosure versus without disclosure. Predictions are scored against matched receiver outcomes via a quadratic scoring rule.

**Beliefs about Persuasiveness** Senders also predict how each fact shifts the assigned receiver’s policy position on the five-point scale. They report the Likert position the receiver will adopt after seeing the fact and the position the receiver would take without seeing it; both predictions are scored via a quadratic scoring rule against the receiver’s realized outcome.

**Beliefs about Congeniality** I measure expressive congeniality—how much a fact validates the sender’s own position—without researcher-imposed directional classification, using a cross-domain rating task. After the provision task, senders rate two facts from the topic they were *not* assigned to—one hypothesized pro-policy and one hypothesized anti-policy—answering: “*Does this information show that your opinion is correct?*” The instructions specify that senders should not assess whether their abstract opinion is correct, nor whether the fact would convince others. The cross-domain design limits contamination by the provision decisions. Senders also rate whether each cross-domain fact is correct.

**Treatments and Hypotheses** Senders are assigned between subjects to CONTROL or ACTION. In both arms, receivers re-state their policy opinion after potentially receiving

a fact.<sup>14</sup> In the ACTION arm, receivers additionally make a real donation decision (€250 split between the same two organizations as Part 1): one decision per session is implemented with actual money, so receivers' choices carry real financial consequences for these organizations. Additionally, I again exploit the within-sender variation in alignment to the facts.

I pre-registered three hypotheses with respect to the treatments, mirroring Experiment 1:

**Hypothesis 4** (Baseline provision). *Provision rates for unaligned facts exceed the indifference benchmark of 50% in each treatment-topic cell.*

**Hypothesis 5** (Selective provision). *Senders are more likely to provide a fact aligned with their own political position than one that is not, holding fact and topic fixed.*

**Hypothesis 6** (Action amplification). *The alignment gap in provision is larger when the receiver makes a real political donation (ACTION) than when the receiver only restates their opinion (CONTROL).*

## 4.2 Sample, Incentives, and Protocol

Data collection ran from May 27 to June 4, 2026 on Bilendi with quota-sampling for age, gender, and state (Bundesland) to approximate the German adult population. (See Table ?? for details.) This study was approved by the Ethics Council of the Max Planck Society under the Generalized Approval of Experiments Following the Protocol that is Standard in Experimental Economics (approval no. XXX Renewal XXX) and was pre-registered on AsPredicted (No. 293252, an update of No. 292898). Senders took a median completion time of 17 minutes and earned €XX on average.

## 4.3 Policy Preferences and Beliefs

**Policy preferences.** Figure 4 reports stated and revealed preference distributions for each topic. For immigration, 66.1% of senders state policy preference in favor of more

---

<sup>14</sup>The ACTION treatment may also shift task salience or decision construal beyond political consequentiality.

restricted immigration and 26.3% donate exclusively to the organization advocating more restrictions. The correlation between the two immigration measures is strong ( $r = 0.63$ ,  $p < 0.001$ ). For climate, a majority states preferences for increased national climate action (52.8%), while the plurality of donations leans toward the organization advocating fewer and more market-based restrictions. The within-person ranking is nonetheless strongly correlated ( $r = 0.50$ ,  $p < 0.001$ ). As pre-registered, I combine the two measures, each normalized to  $[-1, +1]$  with  $+1$  the most pro-policy position: revealed preference  $\text{Pref}_i^{\text{rev}} = \frac{\text{Donation}_i^{\text{pro}} - 125}{125}$  from the €250 split, stated preference  $\text{Pref}_i^{\text{st}} = \frac{3 - \text{Stance}_i}{2}$  from the five-point Likert stance, and their average  $\text{Pref}_i = (\text{Pref}_i^{\text{rev}} + \text{Pref}_i^{\text{st}})/2$ . Preferences are distributed across the full scale in both samples within each topic allowing for identification of selectivity.

**Senders’ fact congeniality rating.** The cross-domain rating task asks senders whether each other-topic fact “shows that your view is correct.” I average the yes-responses within cells of sender stated stances and individual fact direction and transform  $\hat{\alpha}_{ij} = \frac{\text{share\_yes}_{ij} - 0.5}{0.5} \in [-1, +1]$ , where  $+1$  denotes universal validation and  $-1$  no validation by senders with the stance. Panel (a) of Figure 5 plots the resulting pattern. The empirical range across the ten stance  $\times$  direction cells is  $[-0.04, +0.78]$  (Appendix Tables A.32 and A.33); only two of the ten cells are negative, and both are negligible ( $-0.01$  and  $-0.04$ ).<sup>15</sup> Senders almost never regard a fact as purely invalidating their position. Neutral senders (stance 3) also record large variation in perceived congeniality for the different facts,  $\hat{\alpha}_{ij} = +0.15$  on anti-direction and  $+0.43$  on pro-direction facts on average. Notably, positional alignment assigns zero to all facts for senders with  $\text{Pref}_i \approx 0$ , erasing their within-sender variation. The measure correlates only moderately with positional alignment ( $r = 0.184$ ): both carry the same directional signal, but congeniality adds gradation that alignment cannot express. Senders also rated factual credibility;

<sup>15</sup>The two negative cells belong exclusively to lean senders (stances 2 and 4). Extreme senders (stances 1 and 5) rate counter-directional facts positively:  $+0.32$  (stance 1  $\times$  anti) and  $+0.27$  (stance 5  $\times$  pro). For a stance-5 sender the within-sender contrast is  $+0.78$  congenial versus  $+0.27$  counter-directional; this 0.51-unit spread drives identification. Within individual topics lean-sender negative values are more pronounced (stance 2 reaches  $-0.31$  for immigration anti-facts; stance 4 reaches  $-0.22$  for climate pro-facts), confirming the pattern is not a pooling artifact.




Figure 4: Distribution of policy preferences among Experiment 2 senders ( $N = 1,557$ ). *Top row (Panels A–B):* stated preferences from the 5-point Likert item; tick labels complete the stance statement shown in the panel title (e.g., *Immigration to Germany should be ... greatly eased*). *Bottom row (Panels C–D):* revealed preferences from the real donation task, grouped into five buckets from All pro (€250 to the pro-policy organisation) to All anti (€0), displayed pro-to-anti left to right. *Left column:* immigration senders ( $N = 810$ ). *Right column:* climate senders ( $N = 747$ ). Immigration senders concentrate at the anti-policy end (stated modal stance 5: 40.5%; 64.6% at stances 4–5). Climate senders concentrate at the pro-policy end (stated modal stance 1: 27.9%; 52.8% at stances 1–2). Within-topic correlations between stated and revealed preferences: immigration  $r = 0.63$ , climate  $r = 0.50$ , pooled  $r = 0.57$  (all  $p < 0.001$ ).

trust ratings range 52%–85% and follow the same directional gradient ( $\beta = +0.119$ ,  $p < 0.001$ ). Since all twelve facts are verified true, this alignment-predicted variation in truthfulness judgments is consistent with motivated belief formation.

**Senders' beliefs about factual belief correction.** Without disclosure, senders expect receivers to answer correctly roughly half the time: the baseline prior that the receiver answers correctly averages 49.4% across fact–receiver pairs (Appendix Figure A.8). Disclosure is expected to raise this to 66.5%, a 17-percentage-point improvement. Panel (b)




Figure 5: Elicited returns to disclosure, Experiment 2 (Citizens). Panel (a): perceived congeniality  $\hat{\alpha}_{ij}$  (Expression) of each fact by the sender’s five-point stance on the rated topic; one marker per fact  $\times$  stance cell with at least 10 ratings (6 anti-direction facts, red circles; 6 pro-direction facts, green squares), each a cell mean from the cross-domain rating task (2,308 ratings by 1,147 respondents; all completed rating-task responses across recruitment batches, the sample defining the cell-level Expression measure); lines show the observation-weighted mean per fact direction at each stance. Panels (b) and (c): distributions of the two incentivized elicitations for positionally aligned facts ( $\text{Align}_{ij} > 0$ ; bars) and unaligned facts ( $\text{Align}_{ij} < 0$ ; black outline), as shares within each group; observations of senders with  $\text{Pref}_i = 0$  (260 of 7,740) are omitted from the split. Panel (b): expected factual belief correction  $\hat{\Delta}b_{ijk}$  (0.05-unit bins; group means 0.173 aligned, 0.175 unaligned). Panel (c): expected persuasion  $\hat{\Delta}\pi_{ijk}$  (0.125-unit bins; the bin at  $[-0.125, 0)$  contains exact-zero predictions; group means 0.046 aligned, 0.047 unaligned). Shaded regions mark negative values (disclosure expected to move the receiver’s factual belief away from the truth, respectively the receiver away from the sender’s position). Elicitations incentivized via Quadratic Scoring Rule. Appendix Figure A.9 reports both returns by sender–receiver distance.

of Figure 5 shows the distribution of the expected factual belief correction gain  $\widehat{\Delta b}_{ijk}$  for aligned and unaligned facts, where  $k$  indexes the receiver assigned to sender  $i$ 's fact  $j$ ; each sender–fact pair faces exactly one receiver, so  $k$  is implied by  $(i, j)$  and estimation remains at the sender–fact level. The mean expected factual belief correction is positive for all twelve facts (range: 0.139–0.199). Climate facts yield higher expected gains (0.188 vs. 0.155 for immigration). Senders expect partial updating: only 11.9% of fact–receiver pairs receive a post-disclosure prediction of 100%, and 23.7% remain below 50% even after disclosure.

About 26% of fact–receiver pairs carry negative expected factual belief correction, i.e., the sender expects disclosure to leave the receiver slightly further from the truth. This rate is flat across receiver stances (23.4%–26.9%), receiver–fact alignment (counter-attitudinal vs. same-side), and sender–receiver distance (all regression coefficients  $p > 0.35$ ). The rate is not politically patterned. The mean expected factual belief correction is similarly unrelated to receiver position and flat in sender–receiver distance (pooled slope  $\beta = -0.003$ ,  $p = 0.38$ , fact fixed effects, sender-clustered standard errors; Appendix Figure A.9, Panel (a)). Therefore, the factual belief correction return does not vary with the receiver's position relative to the sender.

The incentivized elicitation shows no evidence of politically motivated over-rating: the aligned and unaligned distributions in Figure 5, Panels (b) and (c), lie nearly on top of one another. Senders assign statistically indistinguishable expected factual belief correction gains to aligned and unaligned facts (0.173 vs. 0.175,  $p = 0.62$ , fact FE), ruling out the concern that senders shade their predictions toward politically congenial facts. The same null holds for the Persuasion elicitation.<sup>16</sup>

**Senders' beliefs about the persuasive effect of disclosure.** The persuasion return  $\widehat{\Delta \pi}_{ijk} = |\text{Stance}_i - \hat{\pi}_{ijk}^0| - |\text{Stance}_i - \hat{\pi}_{ijk}^1|$  captures expected movement of receiver  $k$ 's stance toward sender  $i$ 's own position from (non-)disclosure of fact  $j$ ;  $\hat{\pi}_{ijk}^0$  and  $\hat{\pi}_{ijk}^1$  are the sender's incentivized expectations of receiver  $k$ 's stance without and with disclo-

<sup>16</sup>Using Expression ( $\alpha_{ij} > 0$ ) as the conditioning variable yields  $\Delta \pi_{ijk} = 0.055$  (congenial) versus 0.022 (uncongenial), with  $\beta = +0.033$ ,  $p = 0.25$ , fact FE—directional but not statistically significant.

sure. Panel (c) of Figure 5 shows its distribution for positionally aligned and unaligned facts. I find that in 41% of pairs, senders predict the fact will shift the receiver's stance counter to the fact's political direction—a distinct question from whether the receiver moves toward the sender's own position. Computing the rate relative to sender congeniality rather than fact direction leaves it unchanged (40.8%), so it does not reflect a mechanical confound.<sup>17</sup> Receiver extremity does not predict this rate (range: 36.9%–43.9% across the five receiver stances); sender–receiver distance is not statistically distinguishable from zero ( $\beta = +0.002$ ,  $p = 0.62$ ).

Senders' persuasion expectations differ widely across fact–receiver pairs: the standard deviation of  $\widehat{\Delta\pi}_{ijk}$  is roughly 0.9 in both alignment groups (0.888 aligned, 0.906 unaligned) against group means of 0.046 and 0.047; about 17% of pairs are exact zeros. In 59% of fact–receiver pairs, senders expect disclosure to bring the receiver's stance weakly closer to their own relative to non-disclosure ( $\widehat{\Delta\pi}_{ijk} \geq 0$ ; 41.7% strictly closer, 17.3% unchanged). Receiver position enters mechanically: a receiver already at the sender's position leaves essentially no gap to close, and the average return at distance zero is essentially zero ( $-0.067$ ). From distance 2 onward, senders expect disclosure to close a roughly constant share of the sender–receiver gap, about six percent (5.6%, 5.8%, and 5.9% at distances 2, 3, and 4; equality Wald  $p = 0.98$ , sender-clustered), so the level of the return rises with distance (Appendix Figure A.9, Panel (b)). The aligned and unaligned mean profiles are statistically indistinguishable (distance  $\times$  alignment interaction  $p = 0.12$ ). Unlike the factual belief correction return, the persuasion return is positive only where a gap exists to close. What characterizes senders' persuasion beliefs is thus not their level, which is near zero, but their dispersion.

**Variation in the elicited returns.** The systematic structure just described accounts for little of the variation in the elicited returns: fact identity, receiver position, and sender–receiver distance together explain at most 1.6% of the variance of either return.

<sup>17</sup>Directional backfire does not predict provision choices ( $\beta = +0.002$ ,  $p = 0.84$ ). Facts expected to move the receiver *away from the sender's own position* ( $\widehat{\Delta\pi}_{ijk} < 0$ ) do predict lower provision ( $\beta = -0.044$ ,  $p < 0.001$ ), but this is the Persuasion motive's negative tail operating as a selection criterion, not a separate directional effect.

About half of the factual belief correction variance (53%) and a third of the persuasion-return variance (30%) lies between senders, who differ persistently in how much they expect facts to correct and to persuade. The remainder (roughly half of the factual belief correction return, two-thirds of the persuasion return) is idiosyncratic to the fact-receiver pair within sender. This pair-level component is the within-sender variation on which the estimation relies. Observable sender characteristics—political extremity, left-right position, gender, age, education, party affiliation, and topic—capture little of the persistent sender component: regressing the sender mean returns on the full set yields an  $R^2$  of 5.1% for the factual belief correction return and 3.9% for the persuasion return ( $N = 1,543$  senders).<sup>18</sup>

**Uncorrelatedness of motive proxies.** Expression and Persuasion are analytically distinct: Expression captures whether a fact validates the sender’s view, a property of the fact relative to the sender alone, while Persuasion captures whether a specific receiver will move toward the sender’s position—a property of the fact-receiver pair. They diverge whenever a congenial fact faces a receiver who already agrees. A stance-5 sender rates anti-policy facts highly congenial ( $\hat{a}_{ij} = +0.78$ ), but a receiver also at stance 5 is already at the sender’s position; the sender expects near-zero persuasive gain ( $\widehat{\Delta\pi}_{ijk} \approx 0$ ). Positional alignment carries no persuasion-expectation signal either. Conditional on a fact being positionally aligned rather than unaligned, the share of pairs in which the sender expects movement toward her own position is essentially identical (44.2% versus 43.2%), and the mean difference in  $\widehat{\Delta\pi}_{ijk}$  is  $-0.001$ . The factual belief correction proxy is uncorrelated with both. I find all pairwise correlations below  $|r| = 0.04$ : Expression and Factual Belief Correction ( $r = -0.005$ ,  $p = 0.69$ ), Expression and Persuasion ( $r = +0.016$ ,  $p = 0.15$ ), Factual Belief Correction and Persuasion ( $r = +0.033$ ,  $p = 0.004$ ); VIF  $\approx 1.001$ . The design generates this separation across the three predictors. The ACTION randomization does not sort senders differentially on

<sup>18</sup>A descriptive battery interacting each characteristic with the congeniality rating and with sender-receiver distance yields one interaction of twenty-four below a Benjamini-Hochberg  $q$  of 0.10: senders with more extreme issue positions expect larger persuasion returns overall ( $\beta = +0.114$ ,  $p < 0.001$ ), with a flatter distance gradient ( $\beta = -0.035$ ,  $p = 0.004$ ).




Figure 6: Mean provision rates by arm and fact alignment, Experiment 2 (Citizens;  $N = 1,557$  quota-sampled German adults, Bilendi, 2026). Panel (A): mean sender-level provision rates for aligned (dark green) and unaligned (dark red) facts, by experimental arm. Panel (B): alignment bias (aligned minus unaligned provision rate) by arm. Sample: senders with a directional preference (revealed donation  $\neq$  €125). Alignment:  $\text{Pref}_i \times \text{Dir}_j \in [-1, 1]$ . Error bars are 95% confidence intervals (sender-level means).

any of the three proxies or on policy preferences: all four balance tests yield  $p > 0.20$ .

#### 4.4 Treatment Effects

Figure 6 shows provision rates by arm and alignment. I find that provision rates for unaligned facts exceed the 50% indifference benchmark in all four treatment–topic cells (range: 60.9%–63.8%, one-sample  $t$ -tests, all  $p < 0.001$ ), confirming Hypothesis 4.

**Result 4** (Baseline provision). *Senders provide unaligned facts at rates significantly above the 50% indifference benchmark in all four treatment–topic cells ( $p < 0.001$ ).*

Provision rates are between 61.1% and 65.7%, much lower than in Experiment 1, where unaligned rates ranged from 81.9% to 86.5%. First, sharing is implemented with 50% probability under the Dye-type structure, halving the expected persuasive return per sharing act relative to Experiment 1’s deterministic implementation. Second, receivers know the sender’s political position; in Experiment 1 the receiver was unaware a disclosure decision was occurring at all. Third, senders have no information about receivers’ prior beliefs, removing the ability to target disclosure toward those most in need of correction.

Table 4 reports the pre-registered tests. I estimate

$$y_{i,j} = \beta_0 + \beta_1 \text{Align}_{i,j} + \beta_2 \text{ACTION}_i + \beta_3 (\text{Align}_{i,j} \times \text{ACTION}_i) + \mathbf{S}_j + \gamma_3 \mathbf{X}_i + \varepsilon_{i,j}, \quad (5)$$

with  $\text{ACTION}_i \in \{0, 1\}$  an indicator equal to one if sender  $i$  is assigned to the ACTION arm. The unit of observation is a sender–fact pair; standard errors are clustered at the sender level throughout.  $\beta_1$  identifies baseline alignment-based selectivity (Hypothesis 5) and  $\beta_3$  whether political action amplifies it (Hypothesis 6). Column (1) includes no controls, column (2) adds fact fixed effects  $\mathbf{S}_j$ , and column (3)—the pre-registered preferred specification—adds demographic and political controls  $\mathbf{X}_i$ .<sup>19</sup> The confirmatory tests use positional alignment  $\text{Align}_{i,j}$  (not the motive proxies), with H1 ( $\beta_1 > 0$ , one-sided) and H2 ( $\beta_3 > 0$ , one-sided) as the pre-registered directional tests.

Senders share aligned facts at higher rates ( $p = 0.015$ , one-sided, column (3)), confirming Hypothesis 5.

**Result 5** (Selective provision). *Senders provide aligned facts at higher rates than unaligned facts ( $p = 0.015$ , one-sided, column (3)).*

The pre-registered action amplification of the alignment gap is not statistically distinguishable from zero ( $\beta = -0.010$ , one-sided  $p = 0.73$ , column (3)); I do not confirm Hypothesis 6. This null is informative about the proxy rather than the treatment: positional alignment is uncorrelated with the Persuasion proxy within facts ( $r = 0.000$ ), so the alignment interaction can inherit almost none of an action-conditional persuasion effect. Under the estimated instrumental coefficient, the implied  $\text{Align} \times \text{ACTION}$  coefficient is 0.1% of the test’s minimum detectable effect (0.046), and even perfect correlation between the measures would leave it below that threshold ( $0.024 \times 1.53 = 0.037 < 0.046$ ). As pre-registered, I examine the motive channels directly in the following subsection.

---

<sup>19</sup>The pre-registration specifies a perceived believability control jointly with fact fixed effects in column (2). This control was not collected: the cross-domain rating module was shortened after pre-registration to reduce survey length and dropout, a deviation documented in the post-pre-registration decision log. Column (2) therefore reports fact fixed effects without a believability control.

Table 4: Alignment-Based Selective Provision — Experiment 2 (Citizens)

|                       | Probability of Fact Provision |                    |                    |
|-----------------------|-------------------------------|--------------------|--------------------|
|                       | (1)                           | (2)                | (3)                |
| Align                 | 0.024**<br>(0.011)            | 0.024**<br>(0.012) | 0.026**<br>(0.012) |
| Action                | 0.025**<br>(0.013)            | 0.026**<br>(0.013) | 0.023*<br>(0.012)  |
| Align $\times$ Action | −0.008<br>(0.016)             | −0.007<br>(0.016)  | −0.010<br>(0.017)  |
| Clustered SE          | Sender                        | Sender             | Sender             |
| Fact FE               | No                            | Yes                | Yes                |
| Controls              | No                            | No                 | Yes                |
| Num. obs.             | 7740                          | 7740               | 7715               |
| R <sup>2</sup>        | 0.001                         | 0.002              | 0.016              |
| Adj. R <sup>2</sup>   | 0.001                         | 0.001              | 0.009              |

*Sample:*  $N = 1,548$  Experiment 2 senders (Bilendi panel, Germany, 2026), 7,740 sender-fact observations. Each sender makes five provision decisions, one per fact. Column (3) restricts to non-missing demographic controls (7,715 observations). *Outcome:* binary provision indicator ( $y_{ij} \in \{0,1\}$ ). *Alignment:*  $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1,1]$  (pre-registered positional alignment measure). *Fact FE:* indicator for each of the 12 facts. *Controls:* age, gender, education, Bundesland, voting intention, left-right self-placement. Standard errors clustered at sender level. \*\*\*  $p < 0.01$ ; \*\*  $p < 0.05$ ; \*  $p < 0.1$ . *Pre-registered tests on column (3):* H1 (Alignment  $> 0$ , one-sided  $p = 0.015$ ); H2 (Alignment  $\times$  Action  $> 0$ , coefficient negative, one-sided  $p = 0.73$ , not confirmed).

## 4.5 Separating the Motive Channels

Unlike Experiment 1, the model’s three belief arguments—expected factual belief correction, expected expressive congeniality, and expected receiver shift toward the sender’s position—are directly elicited rather than inferred from revealed preference. I estimate a linear probability model at the sender-fact level, with standard errors clustered at the sender level throughout:

$$y_{ij} = \beta_0 + \beta_E \alpha_{ij} + \beta_P \Delta \pi_{ijk} + \beta_2 \text{ACTION}_i + \beta_A (\Delta \pi_{ijk} \times \text{ACTION}_i) + \beta_B \Delta b_{ijk} + \mathbf{S}_j + \varepsilon_{ij}, \quad (6)$$

where  $y_{ij} \in \{0,1\}$  is sender  $i$ ’s provision decision for fact  $j$  and  $\mathbf{S}_j$  are fact fixed effects. The fact fixed effects absorb all time-invariant fact-level characteristics, including political direction, topic-level variation, and fact-specific shareability. Within-sender varia-

Table 5: Action Decomposition: Expression and Persuasion — Experiment 2 (Citizens)

|                                | Probability of Fact Provision |                    |                     |                     |
|--------------------------------|-------------------------------|--------------------|---------------------|---------------------|
|                                | (1)                           | (2)                | (3)                 | (4)                 |
| Expression                     | 0.083***<br>(0.031)           |                    | 0.083***<br>(0.031) | 0.069***<br>(0.023) |
| Expression $\times$ Action     | -0.027<br>(0.042)             |                    | -0.027<br>(0.042)   |                     |
| Persuasion                     |                               | 0.003<br>(0.009)   | 0.003<br>(0.009)    | 0.003<br>(0.009)    |
| Persuasion $\times$ Action     |                               | 0.024*<br>(0.013)  | 0.024*<br>(0.013)   | 0.024*<br>(0.013)   |
| Action                         | 0.035*<br>(0.020)             | 0.025*<br>(0.013)  | 0.034*<br>(0.020)   | 0.024*<br>(0.013)   |
| Factual belief correction gain | 0.038**<br>(0.017)            | 0.037**<br>(0.017) | 0.037**<br>(0.017)  | 0.037**<br>(0.017)  |
| Clustered SE                   | Sender                        | Sender             | Sender              | Sender              |
| Fact FE                        | Yes                           | Yes                | Yes                 | Yes                 |
| Spec                           | (1)                           | (2)                | (3)                 | (4)                 |
| Num. obs.                      | 7740                          | 7740               | 7740                | 7740                |
| R <sup>2</sup>                 | 0.004                         | 0.004              | 0.005               | 0.005               |
| Adj. R <sup>2</sup>            | 0.002                         | 0.002              | 0.003               | 0.003               |

Sample:  $N = 1,548$  Experiment 2 senders (Bilendi panel, Germany, 2026), 7,740 sender-fact observations; requires non-missing expressive identity, projective persuasion, and factual belief correction gain measures. *Outcome*: binary provision indicator. *Expression*: cross-domain behavioral rating  $a_{ij}$ , elicited post-provision on a different topic. *Persuasion*: incentivized expected preference shift  $\Delta\pi_{ijk}$  toward the sender's position (positional, preference-free reading). *Factual belief correction gain*: incentivized expected correction of the receiver's factual belief  $\Delta b_{ijk}$ . *Receiver distance (dist)*: reported separately in Table A.18 (pre-registered heterogeneity). Column (1): expressive motive and its action interaction. Column (2): the persuasion channel—projective main effect and instrumental action interaction. Column (3): both interactions jointly. Column (4): expressive main effect plus the instrumental (persuasion  $\times$  action) interaction. All specifications include fact fixed effects and sender-clustered standard errors. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

tion across the five facts assigned to each sender provides identification.<sup>20</sup> Columns (1) and (2) isolate the expressive and persuasion channels in turn; column (3) includes both jointly. Column (4) is the specification implied by the utility function in the conceptual framework: Expression main effect plus the instrumental Persuasion  $\times$  ACTION interaction, omitting Expression  $\times$  ACTION whose null is established in column (1).

All three motive proxies predict provision and survive fact fixed effects, sender-clustered standard errors, and logit estimation (Table 5). Presentation order changes nothing either: fact order is randomized, and although provision declines mildly over the five positions ( $-0.8$  pp per position,  $p = 0.037$ ), adding presentation-position fixed effects leaves every coefficient in the table unchanged to the third decimal. Positional

<sup>20</sup>Because identification rests on within-sender variation across facts rather than across experimental conditions, experimenter-demand effects are limited (de Quidt et al. 2018).

alignment ( $\text{Align}_{ij}$ ), the pre-registered proxy, is distinct from all three: it correlates with Expression at  $r = 0.184$  but is uncorrelated with Persuasion ( $r = 0.000$ ); it therefore proxies Expression alone. This uncorrelatedness is the quantitative basis of the Hypothesis 2 null documented in the previous subsection.

Expression predicts provision at +13.8 pp across its full range (column (4)) and displaces positional alignment in a joint specification (Appendix Table A.23): Expression holds ( $\beta = +0.061$ ,  $p = 0.013$ ) while Alignment loses significance ( $\beta = +0.012$ ,  $p = 0.195$ ). The expressive motive responds to a fact's political direction alone.

**Result 6** (Expression predicts baseline provision). *Senders are more likely to share facts with higher expressive congeniality ( $\beta = +0.069$ , column (4)), adding +13.8 pp across the full range of Expression.*

The Expression  $\times$  ACTION interaction is null ( $\beta = -0.027$ ,  $p = 0.52$ , column (3)), consistent with an action-invariant expressive motive.

Persuasion ( $\Delta\pi_{ijk}$ ) does not predict provision in the CONTROL arm ( $\beta = +0.003$ ,  $p = 0.76$ ; column (4)). It activates under political action: the ACTION interaction is marginally positive ( $\beta = +0.024$ ,  $p = 0.064$ ; exact wild-cluster bootstrap  $p = 0.047$ ; column (4)), and a randomization test that reassigns the ACTION arm across senders (5,000 permutations) yields one-sided  $p = 0.030$  in the pre-registered direction.<sup>21</sup> Appendix Figure A.5 shows the pattern by arm.

**Result 7** (Persuasion is instrumental: it activates under political action). *Facts expected to move the receiver toward the sender's own position are shared at higher rates when the receiver takes a downstream political action ( $p = 0.064$ ; exact wild-cluster bootstrap  $p = 0.047$ ; randomization test, one-sided  $p = 0.030$ ; column (4)). Little evidence of selection on Persuasion in the CONTROL arm ( $p = 0.76$ ).*

The factual belief correction gain  $\Delta b_{ijk}$  predicts provision ( $\beta = +0.037$  on  $[0, 1]$ , +3.7 pp across its full range,  $p = 0.033$ ), with a positive sign across all 12 facts.<sup>22</sup> Factual belief correction plays a positive but secondary role in provision decisions.

<sup>21</sup>The treatment shifts how senders weight beliefs, not the beliefs themselves: the distribution of  $\widehat{\Delta\pi}_{ijk}$  is statistically indistinguishable between arms.

<sup>22</sup>The pooled factual belief correction coefficient ( $\beta = +0.037$ ,  $p = 0.033$ ) is directionally consistent

**Result 8** (Belief correction predicts provision). *Senders share facts with higher expected factual belief correction at greater rates ( $\beta = +0.037$ ,  $p = 0.033$ , column (4)); the sign is positive for all 12 facts.*

Expressive selection operates in both arms; persuasion-based selection appears only when the receiver’s beliefs carry consequences.

A direction-only analog  $\delta_{ijk}^{\text{dir}} \equiv -\Delta\text{stance}_{ijk} \cdot \text{sign}(\text{pref}_i)$  scales expected receiver movement by the sign rather than the magnitude of the sender’s preference. Appendix A.8 compares the two proxies and their binary codings (Appendix Tables A.30 and A.31): baseline provision responds to the direction of expected persuasion, whereas the instrumental interaction survives only under codings that carry magnitude.

## 4.6 Heterogeneity

I examine three heterogeneity dimensions: sender-receiver ideological distance, sender political extremity (both pre-registered exploratory), and political topic. Appendix Table A.18 covers the treatment-effects contrast; Appendix Tables A.19 and A.20 cover the motive channels. The analysis estimates twenty interaction terms across these specifications; Benjamini–Hochberg-adjusted  $q$ -values within this family are reported in the table notes. Nine of the twenty terms fall below  $p = 0.10$ , against two expected under the null; no individual term survives false-discovery control at conventional levels (all  $q \approx 0.20$ ), so I read the patterns below as suggestive. Jointly, the ACTION interactions of the audience margins and the persuasion channel (six terms) reject at  $p = 0.059$  (randomization  $p = 0.061$ ); this joint test uses the same variation as the main estimate and is corroborative rather than independent evidence.

---

across all twelve facts. The immigration subsample reaches conventional significance ( $\beta = +0.052$ ,  $p = 0.042$ ); the climate coefficient is directionally consistent but loses significance at reduced power ( $\beta = +0.023$ ,  $p = 0.338$ ). Within immigration, five of six facts maintain a positive sign. The sixth is the leveraged observation; without it,  $\beta = +0.025$  ( $p = 0.160$ ), though the direction is preserved. Within-arm estimates are directionally consistent throughout ( $\beta = +0.056$  in CONTROL,  $\beta = +0.020$  in ACTION), and the formal factual belief correction  $\times$  ACTION interaction does not reach significance ( $\beta = -0.034$ ,  $p = 0.32$ ).

**Treatment effects.** *Distance.* Sender–receiver ideological distance is  $\text{dist}_{ij} \equiv |\text{Stance}_i - \text{Stance}_{ij}^R|$ , the absolute gap on the five-point Likert scale between sender  $i$ 's self-placement and the revealed position  $\text{Stance}_{ij}^R \in \{1, \dots, 5\}$  of the receiver paired with that decision. Provision falls  $-9.5$  pp per Likert unit of distance ( $p < 0.001$ ; column (1) of Appendix Table A.18). The ACTION frame partially offsets this penalty ( $\text{dist} \times \text{ACTION}$ ,  $\beta = +0.019$ ,  $p = 0.049$ ). The alignment slope is invariant to distance.<sup>23</sup>

*Political extremity.* More polarised senders provide more overall ( $\beta = +0.010$ ,  $p < 0.10$ , column (2)), but neither the alignment slope nor the ACTION treatment effect varies with continuous polarisation. Categorically, extreme-party supporters (AfD, Linke, BSW) provide  $-4.6$  pp less than center voters ( $p = 0.019$ ) and non-affiliates  $-9.3$  pp less ( $p < 0.001$ ; column (3)). No group exhibits a differential alignment effect. ACTION lifts provision for non-affiliates ( $\beta = +0.066$ ,  $p = 0.075$ ).

*Topic.* Alignment is stronger in climate ( $\beta = +0.034$ ,  $p < 0.10$ , column (5)) than immigration ( $\beta = +0.013$ , column (4)). Action and the alignment-action interaction do not differ by topic.

**Result 9** (Provision falls with sender-receiver ideological distance). *Provision declines  $-9.5$  pp per Likert unit of sender-receiver ideological distance ( $p < 0.001$ ). The ACTION frame partially offsets the penalty ( $\beta = +0.019$ ,  $p = 0.049$ ); alignment and action do not differ by political group or topic.*

**Motive decomposition.** *Distance.* The distance penalty persists in the motive specification ( $-10.5$  pp per Likert unit; column (1) of Appendix Table A.19). Each motive slope is invariant to distance. The triple Persuasion  $\times$  ACTION  $\times$  dist is positive ( $\beta = +0.018$ ,  $p < 0.05$ ): when a distant receiver will act politically, senders select more strongly on expected persuasion direction, consistent with strategic targeting of persuasive facts at distant receivers under high-stakes conditions.

*Political extremity.* Polarisation is associated with stronger responses on both chan-

---

<sup>23</sup>Senders with extreme preferences face more high-distance receiver cells by construction, potentially conflating behavioral effects with a mechanical alignment-distance correlation. Within-sender variation in distance across the five receivers is the primary identifying variation and is independent of sender extremity.

nels: Expression  $\times$  pol  $\beta = +0.027$  ( $p < 0.10$ ) and Persuasion  $\times$  pol  $\beta = +0.011$  ( $p < 0.10$ ; column (2)). Under ACTION, the instrumental interaction is concentrated among center voters (Persuasion  $\times$  ACTION  $\beta = +0.047$ ,  $p = 0.011$ , center reference, column (3)); it is not detected for extreme-party supporters and is muted for non-affiliates (triple  $\beta = -0.065$ ,  $p < 0.10$ ). Center senders exhibit expressive bias at baseline; ACTION adds the instrumental channel on top.

*Topic.* Appendix Table A.20 points to domain-specific patterns: Expression predicts provision in climate ( $\beta = +0.086$ ,  $p = 0.007$ ) but not immigration in this specification, while Factual Belief Correction ( $\beta = +0.053$ ,  $p = 0.036$ ) and the instrumental interaction (Persuasion  $\times$  ACTION  $\beta = +0.037$ ,  $p = 0.035$ ) are detected in immigration only.

**Result 10** (Partisan and topic structure of the motives). *Polarisation is associated with stronger expressive and persuasion responses. The instrumental interaction is concentrated among center voters ( $\beta = +0.047$ ,  $p = 0.011$ ) and not detected for extreme-party supporters. Selection on expected persuasion under ACTION is stronger for ideologically distant receivers (Persuasion  $\times$  ACTION  $\times$  dist,  $\beta = +0.018$ ), consistent with strategic targeting. Expression is concentrated in climate; Factual Belief Correction and the instrumental interaction in immigration.*

## 5 Experiment 3 (Journalists)

Results below are based on  $N = 307$  journalists in the regression sample (1,535 sender-fact observations;  $N = 308$  collected), reaching the pre-registered target of 300.

Models of media bias attribute slanted coverage to extrinsic incentives: advertiser revenue, audience demand for confirmation, or owner preferences (Gentzkow and Shapiro 2006, Mullainathan and Shleifer 2005, Besley and Prat 2006). Experiment 3 supplements this literature by isolating the intrinsic component: *some* journalist selectivity persists when those incentives are absent from the task environment, and I ask whether the same motive structure that drives citizen fact-selection operates

among professional journalists. Journalists differ from citizen senders along two dimensions that bear on the motive structure. First, they work under explicit professional norms—codified in Germany’s *Pressekodex* and enforced by the Presserat [CITE: journalist norms research]—which impose veracity and balance obligations absent in lay information sharing. Second, journalists supply information to mass audiences: a single editorial decision reaches many receivers, not the single known receiver of Experiment 2, which may alter how strongly expressive and persuasion motives enter the provision decision relative to factual belief correction.

## 5.1 Design and Sample

I apply the ACTION arm of Experiment 2 to professional journalists: the fact set, incentivized elicitations, and receiver matrix are identical, and receivers are matched with 100 % probability. Two features differ: the attention check is omitted,<sup>24</sup> and the post-questionnaire replaces citizen items with journalist-specific measures of medium type, geographic reach, and editorial autonomy. Because all journalists are in the ACTION arm,  $\theta_I$  is not separately identified; I recover  $\theta_P + \theta_I$ .

Journalists were recruited via the Deutscher Journalisten-Verband (DJV) and Deutsche Journalistinnen- und Journalisten-Union (DJU) mailing lists (pre-registered target  $N = 300$ ; AsPredicted No. 292902).<sup>25</sup> Appendix Table A.34 compares the regression sample to Loosen et al. (2023): role composition and sector affiliation are close to the benchmark, but online journalists are substantially overrepresented (45.4 % vs. 3.1 % native-online), reflecting the digital recruitment channel.<sup>26</sup> The sample is politically left-leaning (58.5 % Grüne or Linke; mean left-right self-placement  $\approx 3.2$ ). Because journalists completed no cross-domain rating task, I merge the Expression measure  $\alpha_{ij}$  from citizen cross-domain ratings using (sender perceived stance, fact direction), extending

<sup>24</sup>The professional journalist population is small enough that exclusions would be costly and uninformative

<sup>25</sup>One journalist completed the study seven times, identified through duplicate payout codes; only the first participation is retained.

<sup>26</sup>The sample is a convenience sample of German journalists recruited via online channels, not a probability sample. Perfect random sampling of German journalists is infeasible; Loosen et al. required eighteen months of multi-source population tracing to assemble their sampling frame.

Table 6: Pre-Registered Treatment Effects: Journalists

|                     | Probability of Fact Provision |                  |                  |
|---------------------|-------------------------------|------------------|------------------|
|                     | (1)                           | (2)              | (3)              |
| Alignment           | 0.017<br>(0.014)              | 0.006<br>(0.019) | 0.007<br>(0.019) |
| Clustered SE        | Journalist                    | Journalist       | Journalist       |
| Fact FE             | No                            | Yes              | Yes              |
| Controls            | No                            | No               | Yes              |
| Num. obs.           | 1535                          | 1535             | 1535             |
| R <sup>2</sup>      | 0.001                         | 0.005            | 0.014            |
| Adj. R <sup>2</sup> | 0.000                         | −0.003           | 0.003            |

Sample:  $N = 307$  Experiment 3 journalists, 1,535 journalist-fact observations. All journalists in the ACTION arm; no between-subject Action treatment variation. Alignment:  $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, 1]$  (pre-registered expressive proxy). Controls (col. 3): age. Standard errors clustered at journalist level. Pre-registered H1 test, col. (3): Alignment > 0, one-sided  $p$  reported in text. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

the stance-equivalence assumption used for positional congeniality.

## 5.2 Results

Among journalists, the positional alignment gap has the sign H1 predicts: aligned facts are provided more often than unaligned ones ( $\beta = +0.007$ , one-sided  $p = 0.360$ , column (3) of Table 6), matching the direction of the citizen result in Experiment 2, but not statistically distinguishable from zero.<sup>27</sup> Because all journalists are in the ACTION arm, the between-subject amplification test (H2) is not identified from Experiment 3 alone.

**Result 11** (Directionally consistent alignment effect among journalists). *Journalists provide aligned facts more often than unaligned ones, consistent with H1 in direction, though not statistically distinguishable from zero.*

The three elicited proxies remain empirically uncorrelated in the journalist sample (pairwise  $|r| \leq 0.05$ , within-fact VIF  $\approx 1.00$ ), so the channels are separately estimable, as in Experiment 2. I find Persuasion positive and significant ( $\Delta\pi_{ijk}$ ,  $\beta = +0.047$ ,  $p = 0.001$ ; the journalist–citizen gap widens under the directional specification,  $\beta =$

<sup>27</sup>The coefficient is positive in all three specifications; column (1) omits fact fixed effects and controls, column (2) adds fact fixed effects, and column (3) adds controls.

Table 7: Motive Decomposition: Citizens versus Journalists — Experiments 2 and 3

|                                | Probability of Fact Provision |                     |                     |
|--------------------------------|-------------------------------|---------------------|---------------------|
|                                | (1)                           | (2)                 | (3)                 |
| Expression                     | 0.069***<br>(0.023)           | 0.058*<br>(0.032)   | 0.064<br>(0.058)    |
| Factual belief correction gain | 0.037**<br>(0.017)            | 0.020<br>(0.024)    | 0.102***<br>(0.039) |
| Persuasion                     | 0.015**<br>(0.007)            | 0.027***<br>(0.009) | 0.047***<br>(0.014) |
| Clustered SE                   | Sender                        | Sender              | Journalist          |
| Fact FE                        | Yes                           | Yes                 | Yes                 |
| Arm                            | Both                          | Action only         | Action only         |
| Population                     | Citizens                      | Citizens            | Journalists         |
| Num. obs.                      | 7740                          | 3965                | 1535                |
| R <sup>2</sup>                 | 0.004                         | 0.006               | 0.020               |
| Adj. R <sup>2</sup>            | 0.002                         | 0.002               | 0.011               |

*Sample:* Experiment 2 citizens (columns 1–2;  $N = 1,548$ , Bilendi panel, Germany, 2026) and Experiment 3 journalists (column 3;  $N = 307$  in the regression sample, recruited via DJV/DJU mailing lists). *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating  $\alpha_{ij}$ ; for journalists, merged from the citizen cross-domain rating cell lookup under the stance-equivalence assumption. *Factual belief correction gain:* incentivized expected correction of the receiver’s factual belief  $\Delta b_{ijk}$ . *Persuasion:* incentivized expected preference shift  $\Delta \pi_{ijk}$  toward the sender’s position (positional reading, `delta_gain_2`). All columns include fact fixed effects; standard errors clustered at the sender (journalist) level. All journalists are in the ACTION arm. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

+0.058\*\*\* vs +0.026\*\* in the citizen ACTION arm; specifications as in Appendix Table A.30, estimated on the journalist sample). The factual belief correction coefficient is  $5.1\times$  the citizen ACTION-arm estimate ( $\beta = +0.102$ ,  $p = 0.009$  vs.  $\beta = +0.020$ )—consistent with professional norms that preserve veracity as a selection criterion even when the receiver acts politically, though this cross-sample comparison is not within-study randomization.

Journalists also report roughly twice the expected factual belief correction per fact ( $\bar{\Delta b} = 0.322$  vs. 0.172 for citizens; Appendix Figure A.11): the climate-over-immigration hierarchy and the absence of motivated beliefs appear in both populations. Journalists thus both perceive larger factual belief correction gains per fact and respond more

steeply to them—neither comparison identifies a difference in underlying motives.<sup>28</sup> Appendix Figure A.12 shows the distributions of both elicited returns for aligned and unaligned facts. As for citizens, the two distributions nearly coincide (factual belief correction means 0.324 vs. 0.319; persuasion means 0.030 vs. 0.003), the visual counterpart of the uncorrelatedness between congeniality and the elicited returns.

Expression is positive but imprecisely estimated ( $\alpha_{ij}$ ,  $\beta = +0.064$ ,  $p = 0.266$ ); the point estimate is close to the citizen ACTION-arm estimate ( $\beta = +0.058$ ). Positional alignment is also not significant in the secondary specification ( $\beta = +0.006$ ,  $p = 0.75$ ).<sup>29</sup> The three estimates line up with the citizen ACTION arm throughout.<sup>30</sup>

**Result 12** (The motive structure appears among journalists). *Among professional journalists facing the same task without extrinsic, financial incentives, all three motive proxies are directionally consistent with the citizen ACTION arm. Persuasion and the factual belief correction gain are statistically significant ( $p = 0.001$  and  $p = 0.009$ ); Expression is positive and directionally consistent but not significant ( $p = 0.266$ ).*

I close with the pre-registered exploratory heterogeneity battery. Heterogeneity of the alignment gap by sender–receiver distance, polarisation, issue extremity, and the journalist background variables (medium type, broadcaster type, reach, role, and editorial autonomy) is reported in Appendix Table A.35; no interaction survives false-discovery control at the 5 % level (minimum  $q = 0.11$  in the 8-test alignment family). Extending the battery to the three motive channels (24 tests), one pattern survives at the 10 % level: journalists at online outlets select less on Expression ( $\beta = -0.29$ ,  $p = 0.003$ ,  $q = 0.074$ ), directionally echoed in the alignment battery ( $\beta = -0.070$ ,  $p = 0.017$ ,

<sup>28</sup>Approximately 21 % of journalist participants provided their names and outlets, enabling a planned match to their published articles over the past twelve months to test whether experimental selectivity predicts real-world reporting direction.

<sup>29</sup>For journalists, positional alignment correlates with the merged Expression measure at  $r = 0.48$  (citizens:  $r = 0.18$ ): because  $\alpha_{ij}$  is merged from the citizen cross-domain rating cells, it shares the sender–distance input with positional alignment. Positional alignment is not part of the channel specification, so this correlation does not affect the separability of the three elicited proxies.

<sup>30</sup>Ideological distance is a heterogeneity variable, not a motive in the utility function; I report it for the journalist sample for cross-sample comparison with Experiment 2. Appendix Figure A.13 plots both elicited returns by sender–receiver distance for journalists, the counterpart to Appendix Figure A.9. The  $\sim 2\times$  level gap in expected factual belief correction holds at every distance, and the by-distance shape—flat for  $\Delta b_{ijk}$ , rising for  $\Delta \pi_{ijk}$ —matches the citizen profile.

$q = 0.11$ ). I read this pattern as suggestive.

### 5.3 Pooled Comparison: Citizens and Journalists

I pool the Experiment 2 citizens with the Experiment 3 journalists, adding a journalist indicator and demographic controls; because journalists are observed only in the ACTION arm, the journalist indicator captures their provision level relative to citizens in that arm while alignment and action effects are identified from citizens. Table 8 shows the pre-registered treatment effects. The journalist indicator is large and precisely estimated (+9.5 pp, conditional on alignment and demographics), and the alignment slope does not differ by population (Align  $\times$  Journalist,  $\beta = +0.006$ ,  $p = 0.74$ ): journalists are more forthcoming in level but select along the same dimension and to the same degree.

All three motives are positive and significant in the pooled sample (Table 9, column (2)). Adding the action decomposition (column (3)), the Persuasion  $\times$  ACTION interaction is positive and significant ( $\beta = +0.030$ ,  $p < 0.01$ ), identified from the citizens. Column (4) allows each motive to differ for journalists: the factual belief correction slope is larger (Factual belief correction gain  $\times$  Journalist,  $p = 0.043$ ) and the expressive slope does not differ ( $p = 0.30$ ), a conservative null given the sample's left-political composition. The implied persuasion slope rises from the citizen CONTROL arm (0.019) to the citizen ACTION arm (0.039) to journalists (0.076); the journalist-specific increment is exploratory and the confidence interval is wide.

### 5.4 Stated Norms and Selection

The post-questionnaire of both experiments elicited agreement (1–5) with three normative statements: that whoever passes on political information should also pass on factually correct information contradicting their own position; that social media should be more transparent about why users see some political information and not others; and that platforms and media should be obligated to surface relevant information from different perspectives. These self-reports are exploratory and were not pre-registered.

Table 8: Provision Decisions: Citizens and Journalists Pooled

|                           | Probability of Fact Provision |                     |                     |
|---------------------------|-------------------------------|---------------------|---------------------|
|                           | (1)                           | (2)                 | (3)                 |
| Align                     | 0.025**<br>(0.012)            | 0.020*<br>(0.012)   | 0.020*<br>(0.012)   |
| Action                    | 0.026**<br>(0.013)            | 0.022*<br>(0.013)   | 0.022*<br>(0.013)   |
| Align $\times$ Action     | -0.009<br>(0.015)             | -0.006<br>(0.015)   | -0.008<br>(0.016)   |
| Journalist                | 0.116***<br>(0.016)           | 0.095***<br>(0.019) | 0.095***<br>(0.019) |
| Align $\times$ Journalist |                               |                     | 0.006<br>(0.019)    |
| Clustered SE              | Sender                        | Sender              | Sender              |
| Fact FE                   | Yes                           | Yes                 | Yes                 |
| Controls                  | No                            | Yes                 | Yes                 |
| Sample                    | Pooled                        | Pooled              | Pooled              |
| Num. obs.                 | 9275                          | 9245                | 9245                |
| R <sup>2</sup>            | 0.012                         | 0.066               | 0.066               |
| Adj. R <sup>2</sup>       | 0.010                         | 0.060               | 0.060               |

*Sample:* pooled Experiment 2 citizens ( $N = 1,548$ ; both arms; 7,740 observations) and Experiment 3 journalists ( $N = 307$ ; ACTION arm only; 1,535 observations). *Outcome:* binary provision indicator. *Align:* positional alignment  $\text{Pref}_i \times \text{Dir}_j \in [-1, +1]$ . *Journalist:* indicator for Experiment 3 participants (all in the ACTION arm). *Controls* (columns (2)–(3)): sender-receiver ideological distance, age, gender, education, Bundesland, left-right self-placement, and closest-party affiliation—the demographic variables measured in both samples. All specifications include fact fixed effects. Standard errors clustered at sender level. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Journalists endorse the two principle statements more strongly than citizens (means 4.60 vs. 4.13 and 4.61 vs. 4.05) but not the obligation statement (4.22 vs. 4.17). Appendix Table A.36 reports the distributions.

I find that stated norms predict how much senders share, not how selectively. All three items positively predict overall provision among citizens ( $\beta = +0.015$  to  $+0.023$  per scale point,  $p \leq 0.042$ ), as do the two principle items among journalists ( $+0.096$  and  $+0.044$ ). I interact each item with the three selection measures—positional alignment,

Table 9: Motive Decomposition: Citizens and Journalists Pooled

|                                                    | Probability of Fact Provision |                     |                     |                     |
|----------------------------------------------------|-------------------------------|---------------------|---------------------|---------------------|
|                                                    | (1)                           | (2)                 | (3)                 | (4)                 |
| Expression                                         | 0.070***<br>(0.021)           | 0.074***<br>(0.020) | 0.073***<br>(0.020) | 0.066***<br>(0.022) |
| Factual belief correction gain                     | 0.046***<br>(0.016)           | 0.052***<br>(0.016) | 0.052***<br>(0.016) | 0.041**<br>(0.017)  |
| Persuasion                                         | 0.019***<br>(0.006)           | 0.036***<br>(0.006) | 0.018**<br>(0.009)  | 0.019**<br>(0.009)  |
| Persuasion $\times$ Action                         |                               |                     | 0.030***<br>(0.012) | 0.020<br>(0.012)    |
| Action                                             |                               |                     | 0.021*<br>(0.013)   | 0.021*<br>(0.013)   |
| Journalist                                         | 0.123***<br>(0.015)           | 0.100***<br>(0.019) | 0.090***<br>(0.019) | 0.046<br>(0.030)    |
| Expression $\times$ Journalist                     |                               |                     |                     | 0.054<br>(0.052)    |
| Factual belief correction gain $\times$ Journalist |                               |                     |                     | 0.082**<br>(0.040)  |
| Persuasion $\times$ Action $\times$ Journalist     |                               |                     |                     | 0.037**<br>(0.017)  |
| Clustered SE                                       | Sender                        | Sender              | Sender              | Sender              |
| Fact FE                                            | Yes                           | Yes                 | Yes                 | Yes                 |
| Controls                                           | No                            | Yes                 | Yes                 | Yes                 |
| Sample                                             | Pooled                        | Pooled              | Pooled              | Pooled              |
| Num. obs.                                          | 9275                          | 9245                | 9245                | 9245                |
| R <sup>2</sup>                                     | 0.014                         | 0.072               | 0.073               | 0.074               |
| Adj. R <sup>2</sup>                                | 0.013                         | 0.066               | 0.067               | 0.068               |

*Sample:* pooled Experiment 2 citizens ( $N = 1,548$ ; both arms; 7,740 observations) and Experiment 3 journalists ( $N = 307$ ; ACTION arm only; 1,535 observations). *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating  $a_{ij} \in [-1, +1]$  from the citizen cross-domain rating cell lookup (merged for journalists under the stance-equivalence assumption). *Factual belief correction gain:* incentivized expected correction of the receiver's factual belief  $\Delta b_{ijk} \in [0, 1]$ . *Persuasion:* incentivized expected movement of the receiver toward the sender's position  $\Delta \pi_{ijk}$  (positional reading). *Journalist:* indicator for Experiment 3 participants. Column (3) adds the ACTION main effect and the Persuasion  $\times$  ACTION interaction, identified from the citizens (who vary across arms). Column (4) adds journalist heterogeneity: the Expression and Factual belief correction interactions with the journalist indicator, and the persuasion channel's heterogeneity as the Persuasion  $\times$  ACTION  $\times$  Journalist term; because journalists are observed only in the ACTION arm, the Persuasion  $\times$  Journalist double is collinear with this triple and omitted. *Controls* (columns (2)–(4): sender-receiver ideological distance, age, gender, education, Bundesland, left-right self-placement, and closest-party affiliation—the demographic variables measured in both samples. All specifications include fact fixed effects. Standard errors clustered at sender level. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Expression, and Persuasion—and apply a Benjamini–Hochberg correction across the resulting family of 18 tests spanning both samples. No citizen interaction is distinguishable from zero, and a pooled index of the three items is also null (all  $p > 0.13$ ). Among journalists, the transparency item is the exception: journalists who endorse platform transparency select less on positional alignment ( $\beta = -0.075$ ,  $p = 0.003$ ,  $q = 0.060$ ) and less on Expression ( $\beta = -0.194$ ,  $p = 0.007$ ,  $q = 0.060$ ). The correlation is consistent with the professional-norms reading of the factual-belief-correction result above, though journalist answers are compressed at the top of the scale (95 % agreement), so I read the moderation as suggestive. The full battery is in Appendix Table A.37.

## 6 Experiment 4 (Receiver)

### 6.1 Conceptual Framework: Receiver Updating

I adopt an epistemic welfare benchmark: the receiver values making the decision he would make under accurate beliefs. Under this criterion, any information structure that brings receiver beliefs closer to the truth weakly dominates one that does not—an application of the Blackwell ordering to political decision-making.<sup>31</sup> Since disclosure of any true fact weakly raises receiver accuracy ( $\Delta b \geq 0$ ), any fact withheld represents a missed accuracy improvement. The epistemic welfare cost of curated truth is therefore the set of facts that would have improved receiver beliefs but were selectively omitted. Because the distortion operates through unaligned omissions rather than aligned commissions, the residual cost concentrates in information the receiver never sees.

### 6.2 Receiver Updating from Silence

The provision environment of Experiments 2 and 3 leaves the receiver with an inference problem. A disclosed fact reaches the receiver with probability  $\tau = 0.5$ ; a withheld fact never arrives; the receiver knows  $\tau$  and the sender's position. Silence therefore pools two events the receiver cannot distinguish—a fact withheld and a fact that failed to transmit. The pooling is analogous to [Dye \(1985\)](#), where silence instead pools withholding with a sender who has nothing to disclose; in both environments a missing fact cannot be attributed to the sender's choice, which is what keeps disclosure from unraveling. Each receiver faces a single candidate fact, whose content she does not observe unless it arrives; because she knows the sender's position, the fact's realization either supports that position or cuts against it. Let  $\sigma_A$  and  $\sigma_U$  denote the sender's disclosure probabilities conditional on the two realizations, and let  $p$  denote the receiver's prior that the realization supports the sender's position. A rational receiver applies Bayes'

---

<sup>31</sup>See [Blackwell \(1953\)](#). I remain agnostic about which policy is socially optimal; the criterion asks only whether receiver beliefs reflect available evidence.

rule to the pooled silence event (Jung and Kwon 1988),

$$b_{\text{Bayes}} = \frac{p(1 - \tau\sigma_A)}{p(1 - \tau\sigma_A) + (1 - p)(1 - \tau\sigma_U)},$$

and because  $\sigma_A > \sigma_U$ —the selectivity documented in Sections 4 and 5—silence moves a rational receiver toward unaligned content: withholding is weakly informative, while  $\tau < 1$  keeps the inference partial. The inference conditions only on the realization’s direction relative to the sender’s position; it is agnostic about whether expressive or persuasive motives generate the direction-dependence in disclosure. Disclosure, by contrast, is degenerate: facts are verified, so the disclosure posterior is  $b_{\text{disc}} = 1$ . Shin (2003) characterizes the polar case in which the sender never discloses unfavorable outcomes and the rational silence discount is maximal; my senders disclose stochastically, so the rational discount is interior.

Receivers need not perform this inference. I parameterize departures from it with a single weight  $\chi \in [0, 1]$ :

$$b_{\text{silence}}(\chi) = \chi b_{\text{prior}} + (1 - \chi) b_{\text{Bayes}},$$

where  $\chi = 0$  is the rational benchmark and  $\chi = 1$  treats silence as carrying no information. The linear form is not a convenience: it is exactly the posterior of a  $\chi$ -cursed receiver (Eyster and Rabin 2005), and Fong et al. (2025) derive it in closed form from the  $\chi$ -cursed Bayes’ rule; because a single silence observation is a one-stage inference, the static and sequential formulations coincide there, so  $\chi$  below refers to cursedness generically. Bowen et al. (2023) reach the same two endpoints through a nonlinear likelihood-ratio form; I adopt the linear form and treat the choice between them as a robustness question.

The behavioral literature divides on whether  $\chi$  may depend on what silence implies. In cursed equilibrium the weight is a single parameter, common across opponents and inferences (Eyster and Rabin 2005, Fong et al. 2025); read into my environment, it supplies the symmetric benchmark. A separate literature documents updating that weak-

ens precisely where information is unwelcome: subjects underweight unfavorable signals about themselves (Eil and Rao 2011, Möbius et al. 2021), the asymmetry disappears when the ego stake is removed (Möbius et al. 2021, Barron 2021), ego-threatening feedback fades selectively over time (Zimmermann 2020), and in the political domain trust in a received message tracks the congeniality of its content (Thaler 2024). In my environment the uncongenial inference sits at co-partisan silence: a same-side sender withholds facts that cut against the shared position, so rational updating on her silence moves the receiver away from it. Writing  $\chi_{\text{co}}$  and  $\chi_{\text{opp}}$  for the weights applied to co-partisan and opposing-sender silence, the symmetric benchmark is  $\chi_{\text{co}} = \chi_{\text{opp}}$ ; the motivated hypothesis I take to the data is  $\chi_{\text{co}} > \chi_{\text{opp}}$ —discounting concentrated where the rational inference is unwelcome.

Estimating  $\chi$  requires a rational benchmark, and there are two. The equilibrium benchmark computes  $b_{\text{Bayes}}$  from the disclosure rates estimated in Experiments 2 and 3. The subjective benchmark computes it from the receiver’s own conjectured rates  $\hat{\sigma}_A, \hat{\sigma}_U$ , elicited directly, following Jin et al. (2021), who elicit receivers’ beliefs about state-contingent reporting rates and find them systematically too low. The gap between the two benchmarks measures calibration—whether receivers know how selective senders are; the gap between the elicited posterior and the subjective benchmark measures inference—whether receivers use what they believe. The layers carry different theory: a cursed receiver’s conjectures about average behavior are correct by construction (Eyster and Rabin 2005), so miscalibration is evidence for strategic naivety of the Jin et al. (2021) type rather than cursedness, and  $\chi$  measured against the subjective benchmark isolates the layer where cursed and motivated accounts live.

The distinction matters beyond model selection. Bowen et al. (2023) show that receivers who misperceive silence, embedded in unbalanced sharing networks, can polarize permanently even when every shared item is true—and their mechanism requires only symmetric misperception, so the epistemic stakes of  $\chi > 0$  do not hinge on the asymmetry test. The receiver experiment is designed to elicit incentivized posteriors after disclosure and after silence from senders of known stance, alongside the

conjectures  $\hat{\sigma}$ , and will estimate  $\chi_{\text{co}}$  and  $\chi_{\text{opp}}$  against the subjective benchmark and test their equality.

### 6.3 Design and Hypotheses

[PLACEHOLDER — September 2026. Receivers matched to Exp2/3 senders under Dye-type structure; curated-information condition vs. randomly-assigned-information condition; Action-arm receivers make incentive-compatible donation; QSR targets close sender elicitation loop.]

### 6.4 Sample, Incentives, and Protocol

[PLACEHOLDER —  $N = 1,000$  Bilendi, pre-registered separately after Exp2/3 sender collection completes, target July 2026.]

### 6.5 Results

...

## 7 Discussion

[TODO]

## 8 Concluding Remarks

[TODO]

## References

Abramitzky, R., Ager, P., Boustan, L., Cohen, E., and Hansen, C. W. (2024). The effect of immigration on the economy: Lessons from the 1920s border closure. *American Economic Review: Insights*.

- Alesina, A., Miano, A., and Stantcheva, S. (2023). Immigration and redistribution. *The Review of Economic Studies*, 90(1):1–39.
- Allcott, H. and Gentzkow, M. (2017). Social media and fake news in the 2016 election. *Journal of Economic Perspectives*, 31(2):211–236.
- Allen, J., Howland, B., Mobius, M., Rothschild, D., and Watts, D. J. (2020). Evaluating the fake news problem at the scale of the information ecosystem. *Science Advances*, 6(14):eaay3539.
- Allen, J., Mobius, M., Rothschild, D. M., and Watts, D. J. (2021). Research note: Examining potential bias in large-scale censored data. *Harvard Kennedy School Misinformation Review*, 2(4).
- Ambuehl, S., Bernheim, B. D., and Ockenfels, A. (2021). What motivates paternalism? an experimental study. *American economic review*, 111(3):787–830.
- Andreottola, G. and Louis-Sidois, C. (2025). Why us and not them? a theory of political fact-checking. SSRN Working Paper No. 5122378.
- Angelucci, C. and Prat, A. (2024). Is journalistic truth dead? measuring how informed voters are about political news. *American Economic Review*, 114(4):887–925.
- Aridor, G., Jiménez-Durán, R., Levy, R., and Song, L. (2024). The economics of social media. *Journal of Economic Literature*, 62(4):1422–1474.
- Barrera, O., Guriev, S., Henry, E., and Zhuravskaya, E. (2020). Facts, alternative facts, and fact checking in times of post-truth politics. *Journal of public economics*, 182:104123.
- Barron, K. (2021). Belief updating: Does the ‘good-news, bad-news’ asymmetry extend to purely financial domains? *Experimental Economics*, 24(1):31–58.
- Bartling, B., Cappelen, A. W., Hermes, H., Tungodden, B., and Skivenes, M. (2023). Free to fail? paternalistic preferences in the united states. *NHH Dept. of Economics Discussion Paper*, (09).
- Bénabou, R. and Tirole, J. (2006). Incentives and prosocial behavior. *American economic review*, 96(5):1652–1678.
- Berger, L. M. (2026). How digital media markets amplify news sentiment. Working paper, University of Cologne.
- Berger, L. M., Kerkhof, A., Mindl, F., and Münster, J. (2025). Debunking “fake news” on social media: Immediate and short-term effects of fact-checking and media literacy interven-

- tions. *Journal of Public Economics*, 245:105345. Open access, CC BY. Available online 5 April 2025. Pre-registered at AEA Registry AEARCTR-0008199.
- Besley, T. and Prat, A. (2006). Handcuffs for the grabbing hand? media capture and government accountability. *American Economic Review*, 96(3):720–736.
- Blackwell, D. (1953). Equivalent comparisons of experiments. *The annals of mathematical statistics*, pages 265–272.
- Bowen, T. R., Dmitriev, D., and Galperti, S. (2023). Learning from shared news: When abundant information leads to belief polarization. *The Quarterly Journal of Economics*, 138(2):955–1000.
- Boxell, L. and Conway, J. (2022). Journalist ideology and the production of news: Evidence from movers. Working paper.
- Braghieri, L. (2024). Political correctness, social image, and information transmission. *American Economic Review*, 114(12):3877–3904.
- Braghieri, L., Eichmeyer, S., Levy, R., Mobius, M. M., Steinhardt, J., and Zhong, R. (2025). Level slant and polarization of news consumption on social media. *SSRN Working Paper*.
- Budak, C., Nyhan, B., Rothschild, D. M., Thorson, E., and Watts, D. J. (2024). Misunderstanding the harms of online misinformation. *Nature*, 630:45–53.
- Bullock, J. G. and Lenz, G. (2019). Partisan bias in surveys. *Annual Review of Political Science*, 22(1):325–342.
- Burdea, V., Montero, M., and Sefton, M. (2023). Communication with partially verifiable information: An experiment. *Games and Economic Behavior*, 142:113–149.
- Bursztyn, L., Egorov, G., Haaland, I., Rao, A., and Roth, C. (2023a). Justifying dissent. *The Quarterly Journal of Economics*, 138(3):1403–1451.
- Bursztyn, L., Rao, A., Roth, C., and Yanagizawa-Drott, D. (2023b). Opinions as facts. *The Review of Economic Studies*, 90(4):1832–1864.
- Cagé, J., Gallo, N., Hengel, M., Henry, E., and Huang, Y. (2025). Fact-checking and misinformation: Evidence from the market leader. CESifo Working Paper 12319, CESifo.
- Cagé, J., Hengel, E., Hervé, N., and Urvoy, C. (2022). Hosting media bias: Evidence from the universe of French broadcasts 2002–2020. HAL preprint hal-03878119.

- Chopra, F., Haaland, I., and Roth, C. (2022). Do people demand fact-checked news? evidence from us democrats. *Journal of Public Economics*, 205:104549.
- Chopra, F., Haaland, I., and Roth, C. (2024). The demand for news: Accuracy concerns versus belief confirmation motives. *The Economic Journal*, 134(661):1806–1834.
- Crawford, V. P. and Sobel, J. (1982). Strategic information transmission. *Econometrica*, 50(6):1431–1451.
- de Quidt, J., Haushofer, J., and Roth, C. (2018). Measuring and bounding experimenter demand. *American Economic Review*, 108(11):3266–3302.
- Degan, A., Li, M., and Xie, H. (2023). An experimental investigation of persuasion through selective disclosure of evidence. *Canadian Journal of Economics/Revue canadienne d'économie*, 56(4):1490–1516.
- Durante, R. and Knight, B. (2012). Partisan control, media bias, and viewer responses: Evidence from berlusconis italy. *Journal of the European Economic Association*, 10(3):451–481.
- Dye, R. A. (1985). Disclosure of nonproprietary information. *Journal of accounting research*, pages 123–145.
- Eil, D. and Rao, J. M. (2011). The good news-bad news effect: Asymmetric processing of objective information about yourself. *American Economic Journal: Microeconomics*, 3(2):114–138.
- Ershov, D. and Morales, J. S. (2024). Sharing news left and right: Frictions and misinformation on twitter. *The Economic Journal*, 134:2391–2417.
- Eyster, E. and Rabin, M. (2005). Cursed equilibrium. *Econometrica*, 73(5):1623–1672.
- Farina, A., Fréchette, G. R., Ispano, A., Lizzeri, A., and Perego, J. (2026). The selective disclosure of evidence: An experiment. *The Review of Economic Studies*. Forthcoming.
- Farina, A. and Leccese, M. (2025). Hiding a flaw? experimental evidence on multi-dimensional information disclosure. Working paper, September 2025. Princeton University and Boston University.
- Feess, E., Jost, P.-J., and Ressi, A. (2024). Fake news and the asymmetry between type I and type II errors: Dismissing true information is more prevalent than following false information. Working paper. AEA RCT Registry AEARCTR-0010984.
- Fong, M.-J., Lin, P.-H., and Palfrey, T. R. (2025). Cursed sequential equilibrium. *American Economic Review*, 115(8):2616–2658.

- Freitag, J., Kerkhof, A., and Münster, J. (2021). Selective sharing of news items and the political position of news outlets. *Information Economics and Policy*, 56:100926.
- Gentzkow, M. and Shapiro, J. M. (2006). Media bias and reputation. *Journal of Political Economy*, 114(2):280–316.
- Gentzkow, M. and Shapiro, J. M. (2010). What drives media slant? evidence from us daily newspapers. *Econometrica*, 78(1):35–71.
- Golman, R., Loewenstein, G., Molnar, A., and Saccardo, S. (2022). The demand for, and avoidance of, information. *Management Science*, 68(9):6454–6476.
- González-Bailón, S., Lazer, D., Barberá, P., Zhang, M., Allcott, H., Brown, T., Crespo-Tenorio, A., Freelon, D., Gentzkow, M., Guess, A. M., et al. (2023). Asymmetric ideological segregation in exposure to political news on facebook. *Science*, 381(6656):392–398.
- Grigorieff, A., Roth, C., and Ubfal, D. (2020). Does information change attitudes toward immigrants? *Demography*, 57:1117–1143.
- Grossman, S. J. (1981). The informational role of warranties and private disclosure about product quality. *The Journal of law and Economics*, 24(3):461–483.
- Grossmann, M. R. P. (2024). Knowledge and freedom: Evidence on the relationship between information and paternalism.
- Guess, A., Nagler, J., and Tucker, J. (2019). Less than you think: Prevalence and predictors of fake news dissemination on facebook. *Science advances*, 5(1):eaau4586.
- Guriev, S., Henry, E., Marquis, T., and Zhuravskaya, E. (2026). Curtailing false news, amplifying truth. *Econometrica*. Forthcoming.
- Haaland, I. and Roth, C. (2020). Labor market concerns and support for immigration. *Journal of Public Economics*, 191:104256.
- Haaland, I., Roth, C., and Wohlfart, J. (2023). Designing information provision experiments. *Journal of Economic Literature*, 61(1):3–40.
- Henry, E., Zhuravskaya, E., and Guriev, S. (2022). Checking and sharing alt-facts. *American Economic Journal: Economic Policy*, 14(3):55–86.
- Hillman, A. L. (2010). Expressive behavior in economics and politics. *European Journal of political economy*, 26(4):403–418.
- Jin, G. Z., Luca, M., and Martin, D. (2021). Is no news (perceived as) bad news? an experiment-

- tal investigation of information disclosure. *American Economic Journal: Microeconomics*, 13(2):141–173.
- Jung, W.-O. and Kwon, Y. K. (1988). Disclosure when the market is unsure of information endowment of managers. *Journal of Accounting research*, pages 146–153.
- Kamenica, E. and Gentzkow, M. (2011). Bayesian persuasion. *American Economic Review*, 101(6):2590–2615.
- Kuziemko, I., Norton, M. I., Saez, E., and Stantcheva, S. (2015). How elastic are preferences for redistribution? evidence from randomized survey experiments. *American Economic Review*, 105(4):1478–1508.
- Lazer, D. M., Baum, M. A., Benkler, Y., Berinsky, A. J., Greenhill, K. M., Menczer, F., Metzger, M. J., Nyhan, B., Pennycook, G., Rothschild, D., et al. (2018). The science of fake news. *Science*, 359(6380):1094–1096.
- Light, M. T., He, J., and Robey, J. P. (2020). Comparing crime rates between undocumented immigrants, legal immigrants, and native-born us citizens in texas. *Proceedings of the National Academy of Sciences*, 117(51):32340–32347.
- Loosen, W., von Garmissen, A., Bartelt, E., and van Olphen, T. (2023). Journalismus in Deutschland 2023: Aktuelle Befunde zu Situation und Wandel. Arbeitspapiere des Hans-Bredow-Instituts 68, Hans-Bredow-Institut, Hamburg.
- Louis-Sidois, C. (2025). Both judge and party? investigating the political unbiasedness of fact-checkers. *Journal of the European Economic Association*, 23(6):2137–2164.
- Manili, J. (2024). Disclosure to an heterogeneous audience: The role of motivated beliefs. Working paper, Northwestern University. SSRN 4696951.
- Marie, A., Altay, S., and Strickland, B. (2023). Moralization and extremism robustly amplify myside sharing. *PNAS Nexus*, 2(4):pgad078.
- Marie, A. and Petersen, M. B. (2025). Motivations to connect with like-minded audiences increase partisan sharing on social media. *PNAS nexus*, 4(7):pgaf197.
- Martin, G. J. and McCrain, J. (2019). Local news and national politics. *American Political Science Review*, 113(2):372–384.
- Milgrom, P. R. (1981). Good news and bad news: Representation theorems and applications. *The Bell Journal of Economics*, pages 380–391.

- Möbius, M. M., Niederle, M., Niehaus, P., and Rosenblat, T. S. (2021). Managing self-confidence: Theory and experimental evidence. Working paper, October 2021.
- Mullainathan, S. and Shleifer, A. (2005). The market for news. *American Economic Review*, 95(4):1031–1053.
- National Academies of Sciences, Engineering, and Medicine (2017). *The Economic and Fiscal Consequences of Immigration*. The National Academies Press, Washington, DC.
- Osmundsen, M., Bor, A., Vahlstrup, P. B., Bechmann, A., and Petersen, M. B. (2021). Partisan polarization is the primary psychological motivation behind political fake news sharing on twitter. *American political science review*, 115(3):999–1015.
- Pennycook, G., Epstein, Z., Mosleh, M., Arechar, A. A., Eckles, D., and Rand, D. G. (2021). Shifting attention to accuracy can reduce misinformation online. *Nature*, 592:590–595.
- Peterson, E. and Iyengar, S. (2021). Partisan gaps in political information and information-seeking behavior: Motivated reasoning or cheerleading? *American Journal of Political Science*, 65(1):133–147.
- Pew Research Center (2024). How americans get news on TikTok, X, Facebook and Instagram. Report, June 12, 2024. <https://www.pewresearch.org/journalism/2024/06/12/>.
- Pew Research Center (2025). 42% of social media users say the sites are important for them getting involved with political or social issues. Report, September 16, 2025; survey of 5,097 U.S. adults, Feb. 10–17, 2025. <https://www.pewresearch.org/short-reads/2025/09/16/>.
- Prat, A. and Strömberg, D. (2013). The political economy of mass media. *Advances in economics and econometrics*, 2:135.
- Shin, H. S. (2003). Disclosures and asset returns. *Econometrica*, 71(1):105–133.
- Shin, J. and Thorson, K. (2017). Partisan selective sharing: The biased diffusion of fact-checking messages on social media. *Journal of Communication*, 67(2):233–255.
- Thaler, M. (2021). The supply of motivated beliefs. *arXiv preprint arXiv:2111.06062*.
- Thaler, M. (2024). The fake news effect: Experimentally identifying motivated reasoning using trust in news. *American Economic Journal: Microeconomics*, 16(2):1–38.
- Verrecchia, R. E. (1983). Discretionary disclosure. *Journal of Accounting and Economics*, 5:179–194.

- Vosoughi, S., Roy, D., and Aral, S. (2018). The spread of true and false news online. *Science*, 359(6380):1146–1151.
- Zhu, Y. and Dukes, A. (2015). Selective reporting of factual content by commercial media. *Journal of Marketing Research*, 52(1):56–76.
- Zimmermann, F. (2020). The dynamics of motivated beliefs. *American Economic Review*, 110(2):337–361.

Pending

## A Supplementary Tables and Figures

### A.1 Sample Characteristics

Table A.1: Sample characteristics and treatment balance

|                              | Population<br>(1) | Sample<br>(2) | BASE<br>(3) | A<br>(4) | I<br>(5) | A × I<br>(6) |
|------------------------------|-------------------|---------------|-------------|----------|----------|--------------|
| <i>Gender</i>                |                   |               |             |          |          |              |
| Female                       | 0.509             | 0.502         | 0.507       | 0.490    | 0.503    | 0.510        |
| Male                         | 0.491             | 0.478         | 0.470       | 0.477    | 0.477    | 0.487        |
| <i>Ethnicity</i>             |                   |               |             |          |          |              |
| White                        | 0.578             | 0.587         | 0.577       | 0.590    | 0.587    | 0.593        |
| Black                        | 0.121             | 0.129         | 0.137       | 0.120    | 0.140    | 0.120        |
| Hispanic                     | 0.187             | 0.182         | 0.180       | 0.183    | 0.177    | 0.187        |
| Asian                        | 0.059             | 0.058         | 0.063       | 0.063    | 0.057    | 0.050        |
| Other                        | 0.055             | 0.044         | 0.043       | 0.043    | 0.040    | 0.050        |
| <i>Political Affiliation</i> |                   |               |             |          |          |              |
| Democrat                     | 0.386             | 0.388         | 0.377       | 0.393    | 0.393    | 0.387        |
| Independent                  | 0.310             | 0.308         | 0.323       | 0.303    | 0.303    | 0.300        |
| Republican                   | 0.304             | 0.298         | 0.297       | 0.297    | 0.300    | 0.300        |

*Sample:*  $N = 1,200$  U.S. adults recruited on Prolific (quota-sampled), Experiment 1. Column (1): 2020 U.S. Census population proportions. Column (2): full sample proportions. Columns (3)–(6): proportions by treatment cell (BASE = no-Action, no-Info; A = Action only; I = Info only; A × I = Action and Info). Population benchmarks from the 2020 U.S. Census.

### A.2 Additional Summary Statistics

Table A.2: Prior knowledge of the four immigration statistics

|        | Share of Responses |       |            |
|--------|--------------------|-------|------------|
|        | Correct            | False | Don't know |
| Crime  | 0.700              | 0.223 | 0.076      |
| Labor  | 0.794              | 0.141 | 0.065      |
| Border | 0.612              | 0.317 | 0.071      |
| Fiscal | 0.304              | 0.553 | 0.143      |

*Sample:*  $N = 1,200$  senders, Experiment 1. Elicited before senders were informed they would make provision decisions. *Correct:* stated belief matches the verified fact direction; *False:* stated belief contradicts the fact. *Don't know:* explicit non-answer (senders who select "Don't know" receive a 25% probability of being scored as correct, incentivizing truthful reporting of ignorance).




Figure A.2: Distribution of senders donation allocations to the pro-immigration advocacy group ( $N = 1,200$ , Experiment 1). Horizontal axis: allocation to the pro-immigration group collapsed into five categories (“All pro = \$500, “Pro-leaning = \$251–499, “Equal-split = \$250, “Anti-leaning = \$1–249, “All anti = \$0) from a total \$500 stake. Vertical axis: proportion of senders. The incentive-compatible preference measure is  $\text{Pref}_i = (AIC_i - 250)/250 \in [-1, 1]$ ; senders with  $AIC_i = 250$  (“Equal-split) are excluded from alignment-based analyses ( $N = 86, 7.2\%$ ).




Figure A.1: Senders’ beliefs about the persuasive effect of information (pre-treatment). *Notes:* The green line marks the true effect of 69% from [Haaland and Roth \(2020\)](#).

### A.3 Robustness of Main Results

**Logistic Regression.** Table A.7 re-estimates specification (2) using logistic regression, which is the correctly specified model if the latent propensity to disclose follows a logistic distribution—as implied by the cumulative logistic  $F$  assumed in Section 3.2. All results replicate: the  $\text{Align} \times \text{ACTION}$  interaction is positive and significant across all

Table A.3: Perceived Credibility of Statistics by Political Affiliation

|        | Percentage Perceiving Statistic as Credible |           |              |             |
|--------|---------------------------------------------|-----------|--------------|-------------|
|        | All                                         | Democrats | Independents | Republicans |
| Crime  | 0.962                                       | 0.996     | 0.949        | 0.933       |
| Labor  | 0.978                                       | 0.998     | 0.965        | 0.964       |
| Border | 0.978                                       | 0.985     | 0.976        | 0.972       |
| Fiscal | 0.968                                       | 0.981     | 0.962        | 0.958       |

*Sample:*  $N = 1,200$  senders, Experiment 1. Credibility assessed post-provision (potentially endogenous to the provision decision). Values shown are the share of senders rating each statistic as credible (1 = credible, 0 = not credible).

Table A.4: Expected effects of information on political preferences

| Statistic             | Expect Effect of Information on Political Preferences |                     |                      |                      |
|-----------------------|-------------------------------------------------------|---------------------|----------------------|----------------------|
|                       | Crime<br>(1)                                          | Labor<br>(2)        | Border<br>(3)        | Fiscal<br>(4)        |
| Constant              | 0.248***<br>(0.016)                                   | 0.168***<br>(0.016) | -0.338***<br>(0.018) | -0.322***<br>(0.020) |
| INFO                  | 0.058***<br>(0.019)                                   | 0.038**<br>(0.019)  | -0.033<br>(0.021)    | 0.031<br>(0.023)     |
| Estimated Effect Size | 0.000<br>(0.000)                                      | 0.001***<br>(0.000) | 0.001***<br>(0.000)  | 0.002***<br>(0.000)  |
| Clustered SE          | No                                                    | No                  | No                   | No                   |
| Controls              | No                                                    | No                  | No                   | No                   |
| R <sup>2</sup>        | 0.009                                                 | 0.018               | 0.017                | 0.022                |
| Adj. R <sup>2</sup>   | 0.007                                                 | 0.016               | 0.015                | 0.020                |
| Num. obs.             | 1200                                                  | 1200                | 1200                 | 1200                 |

*Sample:*  $N = 1,200$  senders, Experiment 1. *Outcome:* expected effect of each statistic on receivers' immigration policy preferences,  $[-1, 1]$  (positive = pro-immigration shift). *Info:* binary indicator for the between-subject feedback treatment. *Estimated Effect Size:* sender's incentivized estimate of the treatment effect magnitude in [Haaland and Roth \(2020\)](#), included as a covariate. Standard errors are not clustered. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

specifications ( $p < 0.05$ ), and the Align  $\times$  INFO interaction remains null. I report the linear probability model in the main text for interpretability of coefficients as percentage-point effects.

Table A.5: Expected effects of information on political preferences with controls

| Statistic                               | Expect Effect of Information on Political Preferences |                     |                      |                      |
|-----------------------------------------|-------------------------------------------------------|---------------------|----------------------|----------------------|
|                                         | Crime<br>(1)                                          | Labor<br>(2)        | Border<br>(3)        | Fiscal<br>(4)        |
| <i>Panel (A): INFO</i>                  |                                                       |                     |                      |                      |
| Constant                                | 0.236***<br>(0.067)                                   | 0.033<br>(0.071)    | -0.336***<br>(0.073) | -0.259***<br>(0.080) |
| INFO                                    | 0.055***<br>(0.019)                                   | 0.035*<br>(0.019)   | -0.034<br>(0.021)    | 0.028<br>(0.023)     |
| Estimated Effect Size                   | 0.000<br>(0.000)                                      | 0.001***<br>(0.000) | 0.001***<br>(0.000)  | 0.001***<br>(0.000)  |
| Belief Accuracy                         | 0.116***<br>(0.023)                                   | 0.133***<br>(0.027) | 0.003<br>(0.024)     | 0.129***<br>(0.026)  |
| <i>Panel (B): Receiver Belief State</i> |                                                       |                     |                      |                      |
| Constant                                | 0.249***<br>(0.068)                                   | 0.043<br>(0.071)    | -0.347***<br>(0.074) | -0.255***<br>(0.080) |
| Receiver Belief False                   | 0.014<br>(0.019)                                      | 0.006<br>(0.019)    | -0.000<br>(0.021)    | 0.018<br>(0.023)     |
| Estimated Effect Size                   | 0.000<br>(0.000)                                      | 0.001***<br>(0.000) | 0.001***<br>(0.000)  | 0.001***<br>(0.000)  |
| Belief Accuracy                         | 0.116***<br>(0.023)                                   | 0.134***<br>(0.027) | 0.001<br>(0.024)     | 0.129***<br>(0.026)  |
| Clustered SE                            | No                                                    | No                  | No                   | No                   |
| Controls                                | Yes                                                   | Yes                 | Yes                  | Yes                  |
| <i>Panel (A): R<sup>2</sup></i>         | 0.049                                                 | 0.054               | 0.047                | 0.076                |
| <i>Panel (B): R<sup>2</sup></i>         | 0.042                                                 | 0.051               | 0.044                | 0.075                |
| <i>Panel (A) Adj. R<sup>2</sup></i>     | 0.031                                                 | 0.036               | 0.029                | 0.058                |
| <i>Panel (B) Adj. R<sup>2</sup></i>     | 0.025                                                 | 0.033               | 0.027                | 0.058                |
| Num. obs.                               | 1200                                                  | 1200                | 1200                 | 1200                 |

*Sample:*  $N = 1,200$  senders, Experiment 1. *Outcome:* expected effect of each statistic on receivers' immigration policy preferences,  $[-1, 1]$  (positive = pro-immigration shift). *Panel (A):* regressor is the INFO indicator. *Panel (B):* regressor is Receiver Belief False (= 1 if the matched receiver holds an incorrect prior; = 0 if the receiver reports not knowing). *Estimated Effect Size:* sender's incentivized estimate of the Haaland & Roth (2020) treatment effect. *Belief Accuracy:* incentivized prior knowledge indicator. Controls: demographic covariates. Standard errors are not clustered. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

**Alternative Information Provision Specification.** Table A.8 allows asymmetric INFO effects by fact direction to accommodate the asymmetric belief-shifting effect documented in Table A.4. The additional terms  $\text{INFO}_i \times \text{PROSTAT}_j$  and  $\text{Align}_{i,j} \times \text{INFO}_i \times \text{PROSTAT}_j$  are not statistically significant in the specifications with fact fixed effects.

Table A.6: Treatment Effects on Alignment Bias

|                         | Alignment Bias      |                    |
|-------------------------|---------------------|--------------------|
|                         | (1)                 | (2)                |
| Intercept               | 0.067***<br>(0.021) | −0.026<br>(0.071)  |
| <i>Treatment: A</i>     | 0.065**<br>(0.029)  | 0.066**<br>(0.029) |
| <i>Treatment: I</i>     | 0.008<br>(0.029)    | −0.000<br>(0.029)  |
| <i>Treatment: A × I</i> | 0.029<br>(0.029)    | 0.024<br>(0.029)   |
| Controls                | No                  | Yes                |
| R <sup>2</sup>          | 0.005               | 0.048              |
| Adj. R <sup>2</sup>     | 0.003               | 0.029              |
| Num. obs.               | 1114                | 1114               |

*Sample:*  $N = 1,114$  senders with revealed directional policy preference ( $AIC_i \neq 250$ ), Experiment 1. *Outcome:* sender-level alignment bias (mean aligned minus mean unaligned provision rate). Reference: BASE (no-Action, no-Info). *Treatment A:* Action only. *Treatment I:* Info only. *Treatment A × I:* Action and Info. Column (2) adds age, gender, education, ethnicity, and party identification. Standard errors are not clustered. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

The core result—Align × ACTION—is unchanged.

**Leave-One-Statistic-Out Robustness.** Table A.9 re-estimates the preferred specification (Table 2, column 3) four times, each time dropping one of the four statistics. Both coefficients are stable in magnitude across all four runs—Align stays within 0.031–0.043 (full sample: 0.037) and remains significant at  $p < 0.05$  in every run; Align × ACTION stays within 0.039–0.044 (full sample: 0.041) and is significant at  $p < 0.05$  in three of four runs and at  $p < 0.10$  in the fourth. No single statistic drives either result.

Table A.7: Treatment Effects on Information Provision Logistic

|                                     | Probability of Statistic Provision |                     |                     |
|-------------------------------------|------------------------------------|---------------------|---------------------|
|                                     | (1)                                | (2)                 | (3)                 |
| Intercept                           | 1.967***<br>(0.148)                | 1.800***<br>(0.163) | 2.120***<br>(0.502) |
| Align                               | 0.460***<br>(0.133)                | 0.416***<br>(0.134) | 0.430***<br>(0.137) |
| INFO                                | -0.126<br>(0.183)                  | -0.126<br>(0.184)   | -0.126<br>(0.188)   |
| ACTION                              | 0.173<br>(0.195)                   | 0.172<br>(0.196)    | 0.167<br>(0.200)    |
| Belief Accuracy                     | 0.399***<br>(0.103)                | 0.308***<br>(0.114) | 0.317***<br>(0.113) |
| Align $\times$ INFO                 | 0.005<br>(0.177)                   | -0.009<br>(0.177)   | -0.010<br>(0.183)   |
| Align $\times$ ACTION               | 0.535**<br>(0.210)                 | 0.538**<br>(0.209)  | 0.553***<br>(0.214) |
| INFO $\times$ ACTION                | -0.234<br>(0.259)                  | -0.233<br>(0.260)   | -0.241<br>(0.268)   |
| Align $\times$ INFO $\times$ ACTION | -0.434<br>(0.270)                  | -0.425<br>(0.270)   | -0.433<br>(0.277)   |
| Clustered SE                        | Subject                            | Subject             | Subject             |
| Statistics FE                       | No                                 | Yes                 | Yes                 |
| Controls                            | No                                 | No                  | Yes                 |
| No. of Senders                      | 1200                               | 1200                | 1200                |
| AIC                                 | 3258.054                           | 3247.312            | 3248.924            |
| BIC                                 | 3316.341                           | 3325.029            | 3449.691            |
| Log Likelihood                      | -1620.027                          | -1611.656           | -1593.462           |
| Deviance                            | 3240.054                           | 3223.312            | 3186.924            |
| Num. obs.                           | 4800                               | 4800                | 4800                |

Sample:  $N = 1,200$  senders, 4,800 sender-statistic observations, Experiment 1. Logistic regression; coefficients are log-odds (not percentage-point effects). Column structure matches Table 2: column (1) no fixed effects or controls; column (2) adds statistic fixed effects; column (3) adds statistic fixed effects and demographic controls. Standard errors clustered at sender level. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.8: Treatment Effects on Information Provision

|                                      | Probability of Statistic Provision |                     |                     |
|--------------------------------------|------------------------------------|---------------------|---------------------|
|                                      | (1)                                | (2)                 | (3)                 |
| Intercept                            | 0.871***<br>(0.015)                |                     |                     |
| Align                                | 0.041***<br>(0.012)                | 0.037***<br>(0.012) | 0.037***<br>(0.013) |
| INFO                                 | -0.034*<br>(0.019)                 | -0.018<br>(0.020)   | -0.018<br>(0.020)   |
| ACTION                               | -0.002<br>(0.016)                  | -0.002<br>(0.016)   | -0.003<br>(0.016)   |
| Belief Accuracy                      | 0.038***<br>(0.012)                | 0.035***<br>(0.012) | 0.034***<br>(0.012) |
| Align $\times$ INFO                  | -0.015<br>(0.021)                  | -0.009<br>(0.021)   | -0.011<br>(0.022)   |
| Align $\times$ ACTION                | 0.041**<br>(0.018)                 | 0.041**<br>(0.018)  | 0.041**<br>(0.018)  |
| INFO $\times$ ACTION                 | -0.007<br>(0.024)                  | -0.007<br>(0.024)   | -0.008<br>(0.024)   |
| Align $\times$ INFO $\times$ ACTION  | -0.027<br>(0.026)                  | -0.027<br>(0.026)   | -0.027<br>(0.026)   |
| INFO $\times$ PROSTAT                | 0.038***<br>(0.015)                | 0.005<br>(0.020)    | 0.005<br>(0.020)    |
| Align $\times$ INFO $\times$ PROSTAT | 0.025<br>(0.023)                   | 0.025<br>(0.023)    | 0.029<br>(0.025)    |
| Clustered SE                         | Subject                            | Subject             | Subject             |
| Statistics FE                        | No                                 | Yes                 | Yes                 |
| Controls                             | No                                 | No                  | Yes                 |
| No. of Senders                       | 1200                               | 1200                | 1200                |
| Num. obs.                            | 4800                               | 4800                | 4800                |
| R <sup>2</sup>                       | 0.032                              | 0.034               | 0.041               |
| Adj. R <sup>2</sup>                  | 0.030                              | 0.031               | 0.034               |

Sample:  $N = 1,200$  senders, 4,800 sender-statistic observations, Experiment 1. Extends Table 2 with asymmetric INFO effects by statistic direction. *ProStat*: binary indicator for pro-immigration statistics (Crime and Labor).  $\text{INFO} \times \text{ProStat}$  tests whether the INFO treatment shifts provision differentially by statistic direction;  $\text{Align} \times \text{INFO} \times \text{ProStat}$  tests whether the alignment gradient shifts differentially. Column structure matches Table 2. Standard errors clustered at sender level. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

## A.4 Heterogeneous Treatment Effects

Table A.11: Sender-level heterogeneous treatment effects, pooling across INFO conditions

| Heterogeneity | Alignment Bias    |                   |                     |                  |                    |                   |
|---------------|-------------------|-------------------|---------------------|------------------|--------------------|-------------------|
|               | Partisanship      |                   | Political Party     |                  | Issue Polarization |                   |
|               | (1)               | (2)               | (3)                 | (4)              | (5)                | (6)               |
| Intercept     | -0.003<br>(0.027) | -0.054<br>(0.068) | 0.114***<br>(0.023) | 0.062<br>(0.069) | 0.008<br>(0.033)   | -0.054<br>(0.070) |
| ACTION        | 0.073*<br>(0.038) | 0.070*<br>(0.038) | 0.027<br>(0.038)    | 0.030<br>(0.038) | -0.017<br>(0.038)  | -0.018<br>(0.038) |

Table A.9: Leave-One-Statistic-Out Robustness of the Alignment and Action-Amplification Effects — Experiment 1

| Coefficient           | Sign-positive | Min $\beta$ | Max $\beta$ | Max $p$ |
|-----------------------|---------------|-------------|-------------|---------|
| Align                 | 4/4           | 0.031       | 0.043       | 0.028   |
| Align $\times$ Action | 4/4           | 0.039       | 0.044       | 0.068   |
| Statistics FE         |               | Yes         |             |         |
| Controls              |               | Yes         |             |         |
| Clustered SE          |               | Sender      |             |         |

*Sample:* Experiment 1 senders. Each coefficient is re-estimated four times in the preferred specification (Table 2, column 3; statistic fixed effects, full demographic and political controls, sender-clustered standard errors), each time dropping one of the four statistics (crime, labor, border, fiscal). Columns report the number of leave-one-out estimates with a positive coefficient, the minimum and maximum coefficient across the four runs, and the largest  $p$ -value observed. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.10: Sender-level Heterogeneous Treatment Effects

| Heterogeneity                  | Alignment Bias      |                    |                        |                      |                           |                   |
|--------------------------------|---------------------|--------------------|------------------------|----------------------|---------------------------|-------------------|
|                                | <i>Partisanship</i> |                    | <i>Political Party</i> |                      | <i>Issue Polarization</i> |                   |
|                                | (1)                 | (2)                | (3)                    | (4)                  | (5)                       | (6)               |
| Intercept                      | −0.005<br>(0.036)   | −0.105<br>(0.096)  | 0.132***<br>(0.032)    | 0.042<br>(0.095)     | 0.008<br>(0.045)          | −0.102<br>(0.100) |
| A                              | 0.118**<br>(0.052)  | 0.120**<br>(0.052) | 0.030<br>(0.046)       | 0.028<br>(0.046)     | −0.006<br>(0.067)         | 0.003<br>(0.068)  |
| Partisanship                   | 0.140**<br>(0.058)  | 0.142**<br>(0.059) |                        |                      |                           |                   |
| A $\times$ Partisanship        | −0.102<br>(0.082)   | −0.101<br>(0.083)  |                        |                      |                           |                   |
| Independent                    |                     |                    | −0.132***<br>(0.048)   | −0.130***<br>(0.048) |                           |                   |
| Republican                     |                     |                    | −0.075<br>(0.050)      | −0.069<br>(0.050)    |                           |                   |
| A $\times$ Independent         |                     |                    | 0.146**<br>(0.068)     | 0.150**<br>(0.069)   |                           |                   |
| A $\times$ Republican          |                     |                    | −0.039<br>(0.069)      | −0.025<br>(0.071)    |                           |                   |
| Alignment Intensity            |                     |                    |                        |                      | 0.088<br>(0.061)          | 0.104*<br>(0.061) |
| A $\times$ Alignment Intensity |                     |                    |                        |                      | 0.088<br>(0.087)          | 0.079<br>(0.088)  |
| Controls                       | No                  | Yes                | No                     | Yes                  | No                        | Yes               |
| R <sup>2</sup>                 | 0.020               | 0.058              | 0.036                  | 0.069                | 0.027                     | 0.067             |
| Adj. R <sup>2</sup>            | 0.015               | 0.026              | 0.028                  | 0.035                | 0.022                     | 0.035             |
| Num. obs.                      | 554                 | 554                | 554                    | 554                  | 554                       | 554               |

*Sample:*  $N = 554$  senders with revealed directional policy preference ( $AIC_i \neq 250$ ), restricted to BASE and A treatment cells, Experiment 1. *Outcome:* alignment bias = sender's mean aligned provision rate minus mean unaligned provision rate. *Partisanship:* absolute deviation of self-reported political leaning from center (0 = pure Independent, 1 = strong partisan). *Political Party:* Democrat is the reference category; Independent and Republican are indicator contrasts. *Alignment Intensity:*  $|AIC_i - 250|/250$ , immigration-specific preference intensity. Controls (even columns): age, gender, education, ethnicity, party identification. Standard errors are not clustered. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.12: Receiver belief state: perceived effect of providing information

| Statistic                               | Expected Effect of Information on Political Preferences |                     |                      |                      |
|-----------------------------------------|---------------------------------------------------------|---------------------|----------------------|----------------------|
|                                         | Crime<br>(1)                                            | Labor<br>(2)        | Border<br>(3)        | Fiscal<br>(4)        |
| <i>Panel (A): Receiver Belief State</i> |                                                         |                     |                      |                      |
| Constant                                | 0.265***<br>(0.016)                                     | 0.185***<br>(0.017) | -0.352***<br>(0.018) | -0.313***<br>(0.020) |
| Receiver Belief False                   | 0.020<br>(0.019)                                        | 0.004<br>(0.019)    | -0.003<br>(0.021)    | 0.011<br>(0.023)     |
| Estimated Effect Size                   | 0.000<br>(0.000)                                        | 0.001***<br>(0.000) | 0.001***<br>(0.000)  | 0.002***<br>(0.000)  |
| R <sup>2</sup>                          | 0.002                                                   | 0.015               | 0.015                | 0.021                |
| Adj. R <sup>2</sup>                     | 0.000                                                   | 0.013               | 0.013                | 0.019                |
| Num. obs.                               | 1200                                                    | 1200                | 1200                 | 1200                 |

\*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ . Panel (A): Dependent variable is the sender's expected effect of each statistic on receivers' immigration preferences (range  $[-1, 1]$ ). Receiver Belief False equals one if the matched receiver holds a false belief, zero if the receiver reports not knowing. The null effect confirms that the  $\Delta b$  variation (receiver belief state) does not contaminate senders' perceived  $\delta$ : the two channels are empirically separable.




Figure A.3: Provision rates by expected effect on receiver's immigration preferences, Experiment 1 ( $N = 1,200$ ). Horizontal axis: *AlignPersuasion* = expected preference shift toward receiver's policy position, signed by sender direction,  $\in [-1, 1]$ . Vertical axis: share of statistics disclosed. Red dots: sender expects a preference-unaligned shift ( $\delta < 0$ ); gray squares: no expected shift ( $\delta = 0$ ); blue dots: sender expects a preference-aligned shift ( $\delta > 0$ ). Dashed horizontal lines indicate the within-category mean provision rate (red for unaligned, blue for aligned category). Error bars are 95% confidence intervals.




Figure A.4: Distribution of primary stated motivations across all senders ( $N = 1,200$ , Experiment 1, Prolific).

## A.5 Mechanism-Consistent Evidence

## A.6 Experiment 2 Supplementary Tables and Figures

I omit demographic controls from the main-text provision tables because the alignment effect is identified within sender and the ACTION effect is randomized. Adding the full set of sender demographics to the pre-registered tests and to the action decomposition leaves the coefficients stable (see Tables A.21 and A.22).

## A.7 Robustness of the Motive-Channel Estimates — Experiment 2

The action decomposition of Table 5 is robust to alternative inference, ordering, functional-form, and sample-restriction choices; all checks use the canonical positional persuasion measure  $\Delta\pi_{ijk}$ . Table A.24 re-estimates the preferred specification under two-way clustering by sender and fact and under an exact wild-cluster bootstrap over all  $2^{12} = 4,096$  fact-level sign patterns. Table A.25 adds the within-sender presentation order and a lagged provision indicator. Table A.26 probes the shape of  $w(\Delta\pi)$  by replacing  $\Delta\pi_{ijk}$  with its sign and with a version winsorized at the ninetieth percentile of  $|\Delta\pi|$ : the sign coding carries the Persuasion main effect ( $\beta = +0.018$ ,  $p < 0.05$ ) but not its ACTION interaction ( $\beta = -0.005$ ). Table A.27 drops low- $|\Delta\pi|$  senders, perfect sharers, and low-




Figure A.5: Mean provision rates by expected preference-aligned belief shift and arm, Experiment 2 (Citizens). Dark green: facts for which the sender expects the receiver’s policy position to move toward the sender’s own position ( $\Delta\pi_{ijk} > 0$ ). Dark red: facts for which no such shift is expected ( $\Delta\pi_{ijk} \leq 0$ ).  $x$ -axis: experimental arm (CONTROL or ACTION). Error bars are 95% confidence intervals (sender-level means). Sample: Experiment 2 senders with non-missing  $\Delta\pi_{ijk}$  elicitation;  $\Delta\pi_{ijk}$  is incentivized via Quadratic Scoring Rule.




Figure A.6: Mean provision rates by Expression congeniality and arm, Experiment 2 (Citizens;  $N = 1,557$  quota-sampled German adults). Dark green: facts for which the sender’s cross-domain congeniality ( $\alpha_{ij, \text{align\_control\_cell}}$ ) exceeds the median for that sender’s stance group ( $\tilde{\alpha}_s$ , computed within each sender\_perc\_stance value 1–5). Dark red: facts at or below that median.  $x$ -axis: experimental arm (CONTROL or ACTION). Error bars are 95% confidence intervals (sender-level means). The median split is preference-free: it identifies the more vs. less congenial fact direction for each sender position without requiring a non-zero preference score.

$\alpha$  facts in turn; the instrumental interaction stays positive throughout, though it is not separately significant in these restricted subsamples under the positional measure. Table A.28 reports leave-one-fact-out estimates for each motive, and Table A.29 reports the near-unity variance inflation factors. Figures A.7 and A.8 show the distributions of the two incentivized elicitations.

## A.8 Directional vs. Positional Persuasion Proxy

The canonical persuasion proxy  $\Delta\pi_{ijk}$  measures reduction in Likert distance between the receiver's predicted post-disclosure stance and the sender's own position—a positional measure. A direction-only analog,  $\delta_{ijk}^{\text{dir}}$  (defined in Section 4.5), strips preference magnitude from the signal: it is positive whenever the receiver moves toward the sender's preferred side, regardless of how far the receiver is from the sender's exact position or how strong the sender's preference is. Because  $\delta_{ijk}^{\text{dir}}$  is undefined for neutral senders, 260 observations from 52 neutral senders (3.4 % of observations) are excluded from the directional columns.

Table A.30 re-estimates the specification replacing  $\Delta\pi_{ijk}$  with  $\delta_{ijk}^{\text{dir}}$ , pooled (columns 1–3) and with the ACTION interaction (columns 4–6). Both proxies survive standalone specifications with similar coefficients (positional  $\beta = +0.015$ ,  $p = 0.025$ ; directional  $\beta = +0.016$ ,  $p < 0.01$ ; columns 1–2). When they enter jointly, the directional proxy absorbs the shared variance ( $\beta_{\text{dir}} = +0.018$ ,  $p < 0.05$ ;  $\beta_{\text{pos}} = -0.004$ , not statistically distinguishable from zero; column 3). The two measures correlate at  $r = 0.74$  after fact fixed effects, so the horse race partitions shared variance; it is descriptive and does not adjudicate between the proxies. A common-sample check rules out sample composition as the driver of the joint-specification pattern: positional estimates on the non-neutral subsample remain close to their full-sample values (pooled  $\beta = +0.012$ ,  $p = 0.079$ ;  $\times$  ACTION  $\beta = +0.022$ ,  $p = 0.095$ ). The  $\theta_I$  interaction is directionally consistent under both codings ( $\beta = +0.024$ ,  $p = 0.064$  positional;  $\beta = +0.019$ ,  $p = 0.103$  directional; columns 4–5).

Table A.31 isolates the functional form: it repeats the ACTION-interaction specifi-




Figure A.7: Distribution of the elicited expected preference shift  $|\delta_{ijk}|$ , Experiment 2 (Citizens).  $N = 7,785$  incentivized elicitations ( $1,557$  senders  $\times 5$  facts).  $\delta_{ijk}$  is the sender’s predicted change in the matched receiver’s five-point policy position from disclosure versus non-disclosure, elicited post-provision and incentivized via Quadratic Scoring Rule; the figure plots its absolute value. The mass near zero reflects senders who expect little movement; identification of the persuasion motive comes from the right tail.

cation under binary codings of both proxies, replacing each continuous measure with an indicator for a strictly positive value. The two codings separate the persuasion response by arm. In the CONTROL arm, senders respond to the *direction* of expected persuasion: a binary indicator for movement toward the sender’s preferred side predicts provision on its own ( $\beta = +0.035$ ,  $p < 0.05$ ), consistent with a sign-like  $w$  when  $A = 0$ . The ACTION amplification—the instrumental motive  $\theta_I$ —requires *magnitude*: replacing either continuous proxy with its binary version eliminates the interaction entirely (positional binary  $\times$  ACTION  $\beta = +0.001$ ,  $p = 0.967$ ; directional binary  $\times$  ACTION  $\beta = +0.013$ ,  $p = 0.586$ ), while both continuous codings retain it (positional  $p = 0.064$ ; directional  $p = 0.103$ ).

## A.9 Experiment 3 Sample Composition

Table A.34 compares the demographic and occupational characteristics of the Experiment 3 journalist sample to Loosen et al. (2023), the most recent representative survey of the German journalist population. The post-questionnaire collected gender, age, education, medium type, conditional broadcaster affiliation, geographic reach, role (multiple selection), and editorial autonomy. Fields not elicited in Experiment 3—years of




Figure A.8: Expected factual belief correction  $\Delta b_{ijk}$ , Experiment 2 (Citizens,  $N = 7,740$  elicitations). Panel (A): mean  $\Delta b_{ijk}$  by fact, ranked low to high; error bars are 95% CIs. Dark bars = immigration facts; light bars = climate facts. All twelve facts yield positive expected factual belief correction. Climate facts ( $\bar{\Delta b} = 0.189$ ) generate higher expected factual belief correction than immigration facts ( $\bar{\Delta b} = 0.156$ ), consistent with lower baseline knowledge of climate statistics in the quota sample. Panel (B): mean  $\Delta b_{ijk}$  by fact alignment (left pair) and topic (right pair). Alignment does not predict expected factual belief correction—aligned (0.173) and unaligned (0.175) facts are statistically indistinguishable ( $p = 0.73$ )—ruling out motivated optimism about belief-correcting ability as a driver of provision decisions.

experience, full-time versus freelance status, and topics covered—cannot be compared to the benchmark.

## A.10 Experiment 3 Heterogeneity

Table A.35 reports the pre-registered exploratory heterogeneity of the alignment gap among journalists—by sender–receiver distance, polarisation, issue extremity, and the journalist background variables (medium type, broadcaster type, reach, role, editorial autonomy)—together with the corresponding interactions for the three motive channels.




Figure A.9: Elicited returns to disclosure by sender-receiver distance, Experiment 2 (Citizens). One marker per fact  $\times$  sender-receiver distance  $\times$  alignment (cell mean over  $\geq 20$  elicitations), positionally aligned facts ( $\text{Align}_{ij} > 0$ ) as filled black circles, unaligned ( $\text{Align}_{ij} < 0$ ) as open gray squares, horizontally offset for legibility; facts of senders with  $\text{Pref}_i = 0$  are omitted from the split. Solid (dashed) lines connect the observation-weighted means of aligned (unaligned) elicitations at each distance. Panel (a): expected factual belief correction  $\widehat{\Delta b}_{ijk}$ —flat in distance (pooled slope  $\beta = -0.003$ ,  $p = 0.38$ , fact fixed effects, sender-clustered standard errors). Panel (b): expected persuasion  $\widehat{\Delta \pi}_{ijk}$ —rising in distance ( $\beta = +0.072$ ,  $p < 0.001$ ); the aligned and unaligned mean profiles are statistically indistinguishable (distance  $\times$  alignment interaction  $p = 0.12$ ). Dashed lines mark zero. Elicitations incentivized via Quadratic Scoring Rule.

## A.11 Stated Norms and Perceived Design Balance

Both experiments closed with the question whether respondents perceived a political imbalance in the study. A large majority perceived none (85.5% of citizens, 82.4% of journalists), and the shares are flat across treatment arms and topics. Among the minority perceiving an imbalance, a left-favoring reading outnumbers a right-favoring one roughly two to one in both samples (9.5% vs. 5.0% among citizens; 11.7% vs. 5.9% among journalists), against a fact menu balanced by construction (two pro and two anti facts, plus one randomly drawn). Table A.36, Panel A, reports the shares.

Table A.36, Panel B, reports the distributions of the three normative items, and Table A.37 the item  $\times$  selection interaction battery discussed in Section 5.4.




Figure A.10: Distribution of the absolute expected preference shift  $|\delta_{ijk}|$ , citizens (Experiment 2) vs. journalists (Experiment 3).  $\delta_{ijk}$  is the sender's predicted change in the matched receiver's five-point policy position from disclosure versus non-disclosure, elicited post-provision and incentivized via Quadratic Scoring Rule; the figure plots the absolute value. Both panels share identical axes. The mass-near-zero pattern is consistent across populations: 43.0 % of citizen elicitations and 39.7 % of journalist elicitations fall in the  $[0, 0.25)$  bin; medians equal zero in both groups. The slight rightward shift for journalists (heavier tails above 0.5 Likert units) indicates marginally more decisive directional predictions but does not alter the qualitative sign-function interpretation of  $w(\Delta\pi)$ .

Table A.13: Heterogeneous Treatment Effects of Receiver Beliefs

|                                               | Probability of Statistic Provision |                     |
|-----------------------------------------------|------------------------------------|---------------------|
|                                               | (1)                                | (2)                 |
| Intercept                                     | 0.863***<br>(0.018)                | 0.894***<br>(0.045) |
| Align                                         | 0.026*<br>(0.014)                  | 0.025*<br>(0.014)   |
| INFO                                          | -0.034<br>(0.021)                  | -0.034<br>(0.022)   |
| ACTION                                        | -0.021<br>(0.020)                  | -0.022<br>(0.021)   |
| Receiver Belief False                         | -0.025<br>(0.018)                  | -0.024<br>(0.018)   |
| Accuracy Prior Belief                         | 0.035***<br>(0.012)                | 0.034***<br>(0.012) |
| Align × INFO                                  | 0.025<br>(0.023)                   | 0.025<br>(0.023)    |
| Align × ACTION                                | 0.053**<br>(0.024)                 | 0.054**<br>(0.024)  |
| INFO × ACTION                                 | 0.014<br>(0.030)                   | 0.014<br>(0.031)    |
| Align × Receiver Belief False                 | 0.019<br>(0.022)                   | 0.020<br>(0.022)    |
| INFO × Receiver Belief False                  | 0.043<br>(0.027)                   | 0.043<br>(0.027)    |
| ACTION × Receiver Belief False                | 0.034<br>(0.026)                   | 0.034<br>(0.026)    |
| Align × INFO × ACTION                         | -0.050<br>(0.035)                  | -0.051<br>(0.035)   |
| Align × INFO × Receiver Belief False          | -0.042<br>(0.032)                  | -0.042<br>(0.032)   |
| Align × ACTION × Receiver Belief False        | -0.020<br>(0.033)                  | -0.022<br>(0.033)   |
| INFO × ACTION × Receiver Belief False         | -0.041<br>(0.038)                  | -0.043<br>(0.038)   |
| Align × INFO × ACTION × Receiver Belief False | 0.047<br>(0.045)                   | 0.048<br>(0.045)    |
| Clustered SE                                  | Subject                            | Subject             |
| Statistics FE                                 | Yes                                | Yes                 |
| Controls                                      | No                                 | Yes                 |
| No. of Senders                                | 1200                               | 1200                |
| R <sup>2</sup>                                | 0.034                              | 0.041               |
| Adj. R <sup>2</sup>                           | 0.030                              | 0.033               |
| Num. obs.                                     | 4800                               | 4800                |

Sample:  $N = 1,200$  senders, 4,800 sender-statistic observations, Experiment 1. Outcome: binary provision indicator. Full four-way interaction: Align, ACTION, INFO, and Receiver Belief False (= 1 if the matched receiver holds a factually incorrect prior belief; = 0 if the receiver reports not knowing). Column (2) adds demographic controls. Summary statistics from this specification are reported in Appendix Table A.12, Panel (B). Statistic fixed effects and sender-clustered standard errors throughout. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.14: Individual-level Heterogeneous Treatment Effects

| Heterogeneity                 | Probability of Statistic Provision |                     |                        |                     |
|-------------------------------|------------------------------------|---------------------|------------------------|---------------------|
|                               | <i>Partisanship</i>                |                     | <i>Political Party</i> |                     |
|                               | (1)                                | (2)                 | (3)                    | (4)                 |
| Align                         | 0.001<br>(0.016)                   | 0.001<br>(0.016)    | 0.057***<br>(0.019)    | 0.057***<br>(0.019) |
| ACTION                        | 0.026<br>(0.024)                   | 0.026<br>(0.025)    | 0.018<br>(0.023)       | 0.014<br>(0.023)    |
| Partisanship                  | -0.003<br>(0.026)                  | -0.001<br>(0.027)   |                        |                     |
| Accuracy Prior Belief         | 0.031**<br>(0.012)                 | 0.031**<br>(0.012)  | 0.033***<br>(0.012)    | 0.034***<br>(0.012) |
| Align × ACTION                | 0.057**<br>(0.027)                 | 0.057**<br>(0.027)  | 0.019<br>(0.025)       | 0.019<br>(0.025)    |
| Align × Partisanship          | 0.065***<br>(0.023)                | 0.065***<br>(0.023) |                        |                     |
| ACTION × Partisanship         | -0.053<br>(0.036)                  | -0.054<br>(0.037)   |                        |                     |
| Align × ACTION × Partisanship | -0.029<br>(0.037)                  | -0.029<br>(0.037)   |                        |                     |
| Independent                   |                                    |                     | 0.009<br>(0.021)       | 0.005<br>(0.021)    |
| Republican                    |                                    |                     | 0.005<br>(0.021)       | 0.001<br>(0.021)    |
| Align × Independent           |                                    |                     | -0.043**<br>(0.020)    | -0.043**<br>(0.020) |
| Align × Republican            |                                    |                     | -0.022<br>(0.023)      | -0.022<br>(0.023)   |
| ACTION × Independent          |                                    |                     | -0.034<br>(0.029)      | -0.031<br>(0.029)   |
| ACTION × Republican           |                                    |                     | -0.032<br>(0.029)      | -0.027<br>(0.029)   |
| Align × ACTION × Independent  |                                    |                     | 0.061**<br>(0.030)     | 0.061**<br>(0.030)  |
| Align × ACTION × Republican   |                                    |                     | 0.009<br>(0.032)       | 0.009<br>(0.032)    |
| Clustered SE                  | Subject                            | Subject             | Subject                | Subject             |
| Statistics FE                 | Yes                                | Yes                 | Yes                    | Yes                 |
| Controls                      | No                                 | Yes                 | No                     | Yes                 |
| R <sup>2</sup>                | 0.037                              | 0.042               | 0.035                  | 0.040               |
| Adj. R <sup>2</sup>           | 0.034                              | 0.036               | 0.032                  | 0.033               |
| Num. obs.                     | 4800                               | 4800                | 4800                   | 4800                |

Sample:  $N = 1,200$  senders, 4,800 sender-statistic observations, Experiment 1. Outcome: binary provision indicator. Columns (1)–(2): three-way interaction Align×ACTION×Partisanship (Democrat reference for ACTION interaction). Columns (3)–(4): three-way interaction Align×ACTION×Political Party (Democrat is reference). Partisanship: absolute deviation of self-reported leaning from center. Statistic fixed effects and sender-clustered standard errors throughout. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.15: Mechanisms

| Mechanism Evidence          | Probability of Statistic Provision |                     |                       |                     |
|-----------------------------|------------------------------------|---------------------|-----------------------|---------------------|
|                             | <i>Belief correction</i>           |                     | <i>Persuasiveness</i> |                     |
|                             | (1)                                | (2)                 | (3)                   | (4)                 |
| Intercept                   | 0.600***<br>(0.062)                | 0.592***<br>(0.091) | 0.581***<br>(0.062)   | 0.566***<br>(0.065) |
| Align                       | 0.030**<br>(0.013)                 | 0.031**<br>(0.012)  |                       |                     |
| INFO                        | -0.014<br>(0.017)                  | -0.014<br>(0.017)   | -0.014<br>(0.017)     | -0.014<br>(0.017)   |
| ACTION                      | 0.010<br>(0.023)                   | 0.023<br>(0.122)    | 0.001<br>(0.016)      | 0.012<br>(0.033)    |
| Belief Accuracy             | 0.032*<br>(0.017)                  | 0.025**<br>(0.012)  | 0.039***<br>(0.012)   | 0.038***<br>(0.012) |
| Credibility                 | 0.256***<br>(0.060)                | 0.269***<br>(0.090) | 0.264***<br>(0.060)   | 0.260***<br>(0.060) |
| Align × INFO                | 0.007<br>(0.017)                   | 0.007<br>(0.017)    |                       |                     |
| Align × ACTION              | 0.047**<br>(0.019)                 | 0.045**<br>(0.018)  |                       |                     |
| INFO × ACTION               | -0.008<br>(0.023)                  | -0.008<br>(0.023)   | -0.008<br>(0.023)     | -0.008<br>(0.024)   |
| ACTION × Belief Accuracy    | -0.014<br>(0.022)                  |                     |                       |                     |
| Align × INFO × ACTION       | -0.032<br>(0.025)                  | -0.032<br>(0.026)   |                       |                     |
| ACTION × Credibility        |                                    | -0.022<br>(0.122)   |                       |                     |
| AlignPersuasion             |                                    |                     | 0.021<br>(0.016)      |                     |
| AlignPersuasion × ACTION    |                                    |                     | 0.042*<br>(0.024)     |                     |
| AlignPersuasion(-)          |                                    |                     |                       | 0.006<br>(0.025)    |
| AlignPersuasion(+)          |                                    |                     |                       | 0.040*<br>(0.024)   |
| AlignPersuasion(-) × ACTION |                                    |                     |                       | -0.022<br>(0.036)   |
| AlignPersuasion(+) × ACTION |                                    |                     |                       | -0.004<br>(0.034)   |
| Clustered SE                | Subject                            | Subject             | Subject               | Subject             |
| Statistics FE               | Yes                                | Yes                 | Yes                   | Yes                 |
| Controls                    | No                                 | No                  | No                    | No                  |
| No. of Senders              | 1200                               | 1200                | 1200                  | 1200                |
| R <sup>2</sup>              | 0.051                              | 0.051               | 0.041                 | 0.041               |
| Adj. R <sup>2</sup>         | 0.048                              | 0.048               | 0.039                 | 0.039               |
| Num. obs.                   | 4800                               | 4800                | 4800                  | 4800                |

*Sample:*  $N = 1,200$  senders, 4,800 sender-statistic observations, Experiment 1. *Outcome:* binary provision indicator. Columns (1)–(2) test belief-correction channels: column (1) adds Action×Belief Accuracy to the main specification; column (2) replaces it with Action×Credibility. *Credibility:* sender's unincentivized posterior assessment of statistic credibility (binary), elicited post-provision. *Belief Accuracy:* incentivized prior knowledge indicator, elicited pre-treatment. Columns (3)–(4) replace Align with persuasion-based measures. *AlignPersuasion:* expected preference shift signed by sender direction,  $[-1, 1]$ . Column (4) decomposes into AlignPersuasion(+) and AlignPersuasion(-) binary indicators (zero = reference). All specifications include statistic fixed effects and sender-clustered standard errors. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.16: Self-Reported Motivations Effects on Alignment Bias

|                                     | Alignment Bias       |                      |                      |
|-------------------------------------|----------------------|----------------------|----------------------|
|                                     | (1)                  | (2)                  | (3)                  |
| Intercept                           | 0.107***<br>(0.014)  | 0.089***<br>(0.017)  | 0.090***<br>(0.017)  |
| <i>Belief correction</i>            | -0.094***<br>(0.022) | -0.095***<br>(0.022) | -0.092***<br>(0.022) |
| <i>Projection</i>                   | 0.216***<br>(0.040)  | 0.216***<br>(0.040)  | 0.219***<br>(0.040)  |
| <i>Instrumental</i>                 | 0.249***<br>(0.068)  | 0.236***<br>(0.068)  | -0.076<br>(0.168)    |
| <i>Miscellaneous</i>                | -0.041<br>(0.035)    | -0.042<br>(0.035)    | -0.041<br>(0.035)    |
| ACTION                              |                      | 0.037*<br>(0.020)    | 0.032<br>(0.020)     |
| <i>Instrumental</i> $\times$ ACTION |                      |                      | 0.372**<br>(0.183)   |
| Receiver Action                     | No                   | Yes                  | Yes                  |
| R <sup>2</sup>                      | 0.059                | 0.062                | 0.065                |
| Adj. R <sup>2</sup>                 | 0.056                | 0.058                | 0.060                |
| Num. obs.                           | 1113                 | 1113                 | 1113                 |

*Sample:*  $N = 1,113$  senders with revealed directional policy preference ( $AIC_i \neq 250$ ), Experiment 1. *Outcome:* sender-level alignment bias (mean aligned minus mean unaligned provision rate). Motive categories are primary classifications from LLM-based coding of open-ended responses (Appendix ??). *Reference category:* Did Not Refer to Disclosure. *Receiver Action:* binary indicator for the ACTION treatment. Column (3) adds the Instrumental  $\times$  ACTION interaction. Standard errors are not clustered. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.17: Treatment Effects on Information Provision

| Statement               | Probability of Statistic Provision                   |                                                     |                                           |                                                    |
|-------------------------|------------------------------------------------------|-----------------------------------------------------|-------------------------------------------|----------------------------------------------------|
|                         | <i>False factual beliefs<br/>should be corrected</i> | <i>Facts are important<br/>for forming opinions</i> | <i>Social media should<br/>fact check</i> | <i>Disagreement stems<br/>from facts not views</i> |
|                         | (1)                                                  | (2)                                                 | (3)                                       | (4)                                                |
| Intercept               | 0.538***<br>(0.056)                                  | 0.415***<br>(0.083)                                 | 0.734***<br>(0.033)                       | 0.898***<br>(0.027)                                |
| Align                   | 0.197***<br>(0.052)                                  | 0.213**<br>(0.088)                                  | 0.109***<br>(0.028)                       | 0.007<br>(0.021)                                   |
| INFO                    | -0.009<br>(0.016)                                    | -0.007<br>(0.016)                                   | -0.011<br>(0.017)                         | -0.014<br>(0.017)                                  |
| ACTION                  | -0.003<br>(0.016)                                    | -0.005<br>(0.016)                                   | -0.004<br>(0.016)                         | -0.003<br>(0.016)                                  |
| Accuracy Belief         | 0.034***<br>(0.012)                                  | 0.032***<br>(0.012)                                 | 0.035***<br>(0.012)                       | 0.034***<br>(0.012)                                |
| Statement_Agree         | 0.069***<br>(0.011)                                  | 0.091***<br>(0.016)                                 | 0.026***<br>(0.006)                       | -0.013**<br>(0.006)                                |
| Align × INFO            | 0.002<br>(0.017)                                     | 0.001<br>(0.017)                                    | 0.002<br>(0.017)                          | 0.004<br>(0.017)                                   |
| Align × ACTION          | 0.039**<br>(0.018)                                   | 0.039**<br>(0.018)                                  | 0.043**<br>(0.018)                        | 0.041**<br>(0.018)                                 |
| INFO × ACTION           | -0.011<br>(0.023)                                    | -0.010<br>(0.023)                                   | -0.010<br>(0.024)                         | -0.006<br>(0.024)                                  |
| Align × Statement_Agree | -0.036***<br>(0.011)                                 | -0.036**<br>(0.018)                                 | -0.018***<br>(0.006)                      | 0.009<br>(0.006)                                   |
| Align × INFO × ACTION   | -0.024<br>(0.026)                                    | -0.023<br>(0.026)                                   | -0.024<br>(0.026)                         | -0.027<br>(0.026)                                  |
| Clustered SE            | Subject                                              | Subject                                             | Subject                                   | Subject                                            |
| Statistics FE           | Yes                                                  | Yes                                                 | Yes                                       | Yes                                                |
| Controls                | Yes                                                  | Yes                                                 | Yes                                       | Yes                                                |
| No. of Senders          | 1200                                                 | 1200                                                | 1200                                      | 1200                                               |
| R <sup>2</sup>          | 0.065                                                | 0.066                                               | 0.044                                     | 0.036                                              |
| Adj. R <sup>2</sup>     | 0.063                                                | 0.063                                               | 0.042                                     | 0.033                                              |
| Num. obs.               | 4800                                                 | 4800                                                | 4800                                      | 4800                                               |

Sample:  $N = 1,200$  senders, 4,800 sender-statistic observations, Experiment 1. Outcome: binary provision indicator. Each column corresponds to a different pro-accuracy norm statement (column headers). *Statement\_Agree*: sender's agreement with the column-header statement, measured post-provision on a five-point scale (1 = strongly disagree; 5 = strongly agree). *Align* × *Statement\_Agree* tests whether norm endorsement moderates the alignment gradient. All specifications include statistic fixed effects, demographic controls, and sender-clustered standard errors. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.18: Heterogeneity by Sender Profile: Crude Alignment and Action — Experiment 2 (Citizens)

Table A.19: Heterogeneity by Sender Profile: Motive Decomposition — Experiment 2 (Citizens)

|                                    | Probability of Fact Provision |                    |                     |
|------------------------------------|-------------------------------|--------------------|---------------------|
|                                    | (1)                           | (2)                | (3)                 |
| Expression                         | 0.038<br>(0.033)              | 0.028<br>(0.031)   | 0.071*<br>(0.037)   |
| Persuasion                         | 0.021<br>(0.014)              | −0.014<br>(0.013)  | −0.011<br>(0.013)   |
| Action                             | −0.007<br>(0.018)             | 0.037**<br>(0.018) | 0.001<br>(0.019)    |
| Persuasion × Action                | −0.018<br>(0.020)             | 0.025<br>(0.018)   | 0.047**<br>(0.018)  |
| Receiver distance                  | −0.105***<br>(0.009)          |                    |                     |
| Expression × Distance              | 0.021<br>(0.017)              |                    |                     |
| Persuasion × Distance              | −0.001<br>(0.007)             |                    |                     |
| Action × Distance                  | 0.018*<br>(0.010)             |                    |                     |
| Persuasion × Action × Distance     | 0.018**<br>(0.009)            |                    |                     |
| Sender polarisation                |                               | −0.000<br>(0.008)  |                     |
| Expression × Polarisation          |                               | 0.027*<br>(0.014)  |                     |
| Persuasion × Polarisation          |                               | 0.011*<br>(0.006)  |                     |
| Action × Polarisation              |                               | −0.009<br>(0.009)  |                     |
| Persuasion × Action × Polarisation |                               | −0.001<br>(0.008)  |                     |
| Extremes (ref: center)             |                               |                    | −0.065**<br>(0.026) |
| Others (ref: center)               |                               |                    | −0.072**<br>(0.034) |
| Expression × Extremes              |                               |                    | 0.037<br>(0.048)    |
| Expression × Others                |                               |                    | −0.071<br>(0.061)   |
| Persuasion × Extremes              |                               |                    | 0.014<br>(0.019)    |
| Persuasion × Others                |                               |                    | 0.043*<br>(0.026)   |
| Action × Extremes                  |                               |                    | 0.025<br>(0.027)    |
| Action × Others                    |                               |                    | 0.070*<br>(0.037)   |
| Persuasion × Action × Extremes     |                               |                    | −0.028<br>(0.028)   |
| Persuasion × Action × Others       |                               |                    | −0.065*<br>(0.038)  |
| Clustered SE                       | Sender                        | Sender             | Sender              |
| Fact FE                            | Yes                           | Yes                | Yes                 |
| Moderator                          | Distance                      | Polarisation       | Political group     |
| Spec                               | (1)                           | (2)                | (3)                 |
| Num. obs.                          | 7740                          | 7740               | 7715                |
| R <sup>2</sup>                     | 0.056                         | 0.007              | 0.009               |
| Adj. R <sup>2</sup>                | 0.053                         | 0.004              | 0.005               |

Sample:  $N = 1,548$  Experiment 2 citizens (Bilendi panel, Germany, 2026); 7,740 sender-fact observations in columns (1)–(2), 7,715 in column (3) (non-missing political group). Outcome: binary provision indicator. Expression: cross-domain behavioral rating  $\alpha_{ij}$ , elicited post-provision on a different topic. Persuasion: incentive  $\delta_i$  expected preference shift  $\Delta\pi_{ijk}$  toward the sender's position. Receiver distance:  $|\text{sender\_likert} - \text{recv\_stance}| \in \{0, 1, 2, 3, 4\}$ . Sender polarisation: absolute left-right distance from centre. Political group: Extremes = AfD, Linke, BSW; Others = other/no vote; Center = CDU/CSU, SPD, Grüne, FDP (reference). Each column includes the moderator interacted with Expression, Persuasion, ACTION, and the Persuasion × ACTION × Moderator three-way term. All specifications include fact fixed effects and sender-clustered standard errors. Multiple testing: the exploratory heterogeneity family spans twenty interaction terms across this table and Table A.18; Benjamini–Hochberg-adjusted  $q$ -values within that family are approximately 0.20 for every interaction significant at raw  $\alpha < 0.10$ , \*\*\* $\alpha < 0.01$ , \*\* $\alpha < 0.05$ , \* $\alpha < 0.1$ .

Table A.20: Motive Decomposition by Topic — Experiment 2 (Citizens)

|                            | Probability of Fact Provision |                     |
|----------------------------|-------------------------------|---------------------|
|                            | (1)                           | (2)                 |
| Expression                 | 0.049<br>(0.033)              | 0.086***<br>(0.032) |
| Factual belief correction  | 0.053**<br>(0.025)            | 0.023<br>(0.024)    |
| Persuasion                 | 0.002<br>(0.013)              | 0.003<br>(0.013)    |
| Action                     | 0.025<br>(0.018)              | 0.023<br>(0.017)    |
| Persuasion $\times$ Action | 0.037**<br>(0.017)            | 0.010<br>(0.019)    |
| Clustered SE               | Sender                        | Sender              |
| Fact FE                    | Yes                           | Yes                 |
| Topic                      | Immigration                   | Climate             |
| Spec                       | (1)                           | (2)                 |
| Num. obs.                  | 4045                          | 3695                |
| R <sup>2</sup>             | 0.007                         | 0.004               |
| Adj. R <sup>2</sup>        | 0.004                         | 0.001               |

Sample:  $N = 1,548$  Experiment 2 citizens (Bilendi panel, Germany, 2026). Column (1): immigration subsample ( $N = 809$  senders). Column (2): climate subsample ( $N = 739$  senders). Outcome: binary provision indicator. Each column applies the preferred motive specification (col. (4) of Table 5) within topic subsample: Expression ( $\alpha_{ij}$ , cross-domain behavioral), Factual belief correction ( $\Delta b_{ijk}$ , incentivized), Persuasion ( $\Delta \pi_{ijk}$ , incentivized positional), and the Persuasion  $\times$  ACTION interaction. All specifications include fact fixed effects and sender-clustered standard errors. \*\*\*  $p < 0.01$ ; \*\*  $p < 0.05$ ; \*  $p < 0.1$ .

Table A.21: Provision Decisions with Demographic Controls — Experiment 2 (Citizens)

|                       | Probability of Fact Provision |                    |
|-----------------------|-------------------------------|--------------------|
|                       | (1)                           | (2)                |
| Align                 | 0.019*<br>(0.011)             | 0.026**<br>(0.012) |
| Action                | 0.021*<br>(0.012)             | 0.023*<br>(0.012)  |
| Align $\times$ Action | −0.008<br>(0.016)             | −0.010<br>(0.017)  |
| Clustered SE          | Sender                        | Sender             |
| Fact FE               | No                            | Yes                |
| Controls              | Yes                           | Yes                |
| Num. obs.             | 7715                          | 7715               |
| R <sup>2</sup>        | 0.061                         | 0.016              |
| Adj. R <sup>2</sup>   | 0.056                         | 0.009              |

*Sample:*  $N = 1,548$  Experiment 2 citizens, both arms; 7,715 observations with non-missing controls. *Outcome:* binary provision indicator. *Align:* positional alignment  $\text{Pref}_i \times \text{Dir}_i \in [-1, +1]$ . *Controls:* sender-receiver ideological distance, age, gender, education, Bundesland, closest-party affiliation, and left-right self-placement. Column (1) omits fact fixed effects; column (2) includes them and reproduces the pre-registered preferred specification. Standard errors clustered at sender level. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.22: Action Decomposition with Demographic Controls — Experiment 2 (Citizens)

|                                | Probability of Fact Provision |                     |
|--------------------------------|-------------------------------|---------------------|
|                                | (1)                           | (2)                 |
| Expression                     | 0.069***<br>(0.023)           | 0.069***<br>(0.023) |
| Persuasion                     | 0.003<br>(0.009)              | 0.004<br>(0.009)    |
| Persuasion $\times$ Action     | 0.024*<br>(0.013)             | 0.023*<br>(0.013)   |
| Action                         | 0.024*<br>(0.013)             | 0.021*<br>(0.012)   |
| Factual belief correction gain | 0.037**<br>(0.017)            | 0.046***<br>(0.017) |
| Clustered SE                   | Sender                        | Sender              |
| Fact FE                        | Yes                           | Yes                 |
| Controls                       | No                            | Yes                 |
| Num. obs.                      | 7740                          | 7715                |
| R <sup>2</sup>                 | 0.005                         | 0.018               |
| Adj. R <sup>2</sup>            | 0.003                         | 0.012               |

*Sample:*  $N = 1,548$  Experiment 2 citizens, both arms; column (1) 7,740 observations, column (2) 7,715 with non-missing controls. *Outcome:* binary provision indicator. *Expression:* cross-domain behavioral rating  $\alpha_{ij} \in [-1, +1]$ . *Persuasion:* incentivized expected movement of the receiver toward the sender's position  $\Delta\pi_{ijk}$  (positional reading). *Factual belief correction gain:* incentivized expected correction of the receiver's factual belief  $\Delta b_{ijk} \in [0, 1]$ . Column (1) reproduces the preferred action-decomposition specification; column (2) adds demographic controls (age, gender, education, Bundesland, closest-party affiliation, left-right self-placement). Sender-receiver ideological distance is excluded as a heterogeneity variable. All specifications include fact fixed effects. Standard errors clustered at sender level. \*\*\*  $p < 0.01$ ; \*\*  $p < 0.05$ ; \*  $p < 0.1$ .

Table A.23: Expression vs. Positional Alignment: Horse Race — Experiment 2 (Citizens)

|                     | Probability of Fact Provision |                     |                    |
|---------------------|-------------------------------|---------------------|--------------------|
|                     | (1)                           | (2)                 | (3)                |
| Alignment           | 0.020**<br>(0.009)            |                     | 0.012<br>(0.009)   |
| Expression          |                               | 0.070***<br>(0.023) | 0.061**<br>(0.025) |
| Clustered SE        | Sender                        | Sender              | Sender             |
| Fact FE             | Yes                           | Yes                 | Yes                |
| Expressive measure  | Positional                    | Behavioral          | Both               |
| Num. obs.           | 7740                          | 7740                | 7740               |
| R <sup>2</sup>      | 0.002                         | 0.002               | 0.002              |
| Adj. R <sup>2</sup> | 0.000                         | 0.001               | 0.001              |

*Sample:*  $N = 1,548$  Experiment 2 senders (Bilendi panel, Germany, 2026), 7,740 sender-fact observations. *Outcome:* binary provision indicator. *Alignment:* positional alignment  $\text{Align}_{ij} = \text{Pref}_i \times \text{Dir}_j \in [-1, +1]$ , the pre-registered confirmatory proxy. *Expression:* cross-domain behavioral rating  $\alpha_{ij}$ , elicited post-provision on a different topic. Column (1): Alignment only. Column (2): Expression only. Column (3): horse race including both. All specifications include fact fixed effects and sender-clustered standard errors. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.24: Inference Robustness: Clustering and Wild Bootstrap — Experiment 2 (Citizens)

|                                            | (1) Sender SE (preferred) | (2) Two-way SE (Sender + Fact) |
|--------------------------------------------|---------------------------|--------------------------------|
| Expression                                 | 0.069***<br>(0.023)       | 0.069**<br>(0.026)             |
| Factual belief correction gain             | 0.037**<br>(0.017)        | 0.037*<br>(0.018)              |
| Persuasion                                 | 0.003<br>(0.009)          | 0.003<br>(0.009)               |
| Persuasion $\times$ Action                 | 0.024*<br>(0.013)         | 0.024*<br>(0.011)              |
| Action                                     | 0.024*<br>(0.013)         | 0.024**<br>(0.010)             |
| Clustered SE                               | Sender                    | Sender + Fact                  |
| Fact FE                                    | Yes                       | Yes                            |
| Wild bootstrap $p$ (fact, exact $2^{12}$ ) | —                         | 0.047                          |
| Num. obs.                                  | 7,740                     | 7,740                          |

*Sample:* Experiment 2 citizens; 7,740 sender-fact observations. Both columns estimate the preferred action-decomposition specification (*Expression*, *Factual belief correction gain*, *Persuasion*, *Persuasion  $\times$  Action*, *Action*) with fact fixed effects. Column (1): standard errors clustered at the sender level (preferred). Column (2): two-way clustering by sender and fact (Cameron–Gelbach–Miller;  $K = 12$  fact clusters, asymptotic standard errors indicative only). *Wild bootstrap:* exact enumeration of all  $2^{12} = 4,096$  Rademacher sign patterns at the fact level (impose-null variant);  $p$ -value for *Persuasion  $\times$  Action*. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.25: Order and Serial-Autocorrelation Robustness — Experiment 2 (Citizens)

|                                | (1) + Order         | (2) + Lag           | (3) + Both          |
|--------------------------------|---------------------|---------------------|---------------------|
| Expression                     | 0.068***<br>(0.023) | 0.089***<br>(0.025) | 0.089***<br>(0.025) |
| Factual belief correction gain | 0.036**<br>(0.017)  | 0.042**<br>(0.019)  | 0.043**<br>(0.019)  |
| Persuasion                     | 0.003<br>(0.009)    | 0.009<br>(0.011)    | 0.010<br>(0.011)    |
| Persuasion $\times$ Action     | 0.024*<br>(0.013)   | 0.029*<br>(0.015)   | 0.029*<br>(0.015)   |
| Action                         | 0.024*<br>(0.013)   | 0.024*<br>(0.013)   | 0.024*<br>(0.013)   |
| Presentation order             | -0.007*<br>(0.004)  | —                   | 0.006<br>(0.005)    |
| Lag(provision)                 | —                   | 0.075***<br>(0.014) | 0.076***<br>(0.014) |
| Clustered SE                   | Sender              | Sender              | Sender              |
| Fact FE                        | Yes                 | Yes                 | Yes                 |
| Num. obs.                      | 7,740               | 6,192               | 6,192               |

*Sample:* Experiment 2 citizens. Column (1): 7,740 observations (all five slots). Columns (2)–(3): slots 2–5 only (slot 1 has no lag). All columns add the listed control(s) to the preferred action-decomposition specification with fact fixed effects and sender-clustered standard errors. *Presentation order:* linear slot index  $\in \{1, \dots, 5\}$ . *Lag(provision):* the sender's provision decision for the immediately preceding fact. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.26: Functional Form of  $w(\delta)$ : Direction versus Magnitude — Experiment 2 (Citizens)

|                                        | (1) $\text{sign}(\delta)$ | (2) Winsorized $\delta$ |
|----------------------------------------|---------------------------|-------------------------|
| Expression                             | 0.069***<br>(0.023)       | 0.069***<br>(0.023)     |
| Factual belief correction gain         | 0.037**<br>(0.017)        | 0.037**<br>(0.017)      |
| Persuasion (transform)                 | 0.018**<br>(0.009)        | 0.010<br>(0.011)        |
| Persuasion (transform) $\times$ Action | -0.004<br>(0.012)         | 0.020<br>(0.016)        |
| Action                                 | 0.025**<br>(0.013)        | 0.025*<br>(0.013)       |
| Persuasion transform                   | $\text{sign}(\delta)$     | Winsorized $\delta$     |
| Fact FE                                | Yes                       | Yes                     |
| Num. obs.                              | 7,740                     | 7,740                   |

*Sample:* Experiment 2 citizens; 7,740 observations. Both columns replace *Persuasion* ( $\Delta\pi_{ijk}$ , positional reading) in the preferred action-decomposition specification with a transformed version, keeping fact fixed effects and sender-clustered standard errors. Column (1):  $\text{sign}(\delta) \in \{-1, 0, +1\}$  (direction only). Column (2):  $\delta$  winsorized at the 90th percentile of  $|\delta|$  (magnitude, extremes capped;  $N$  unchanged). A significant direction main effect with a null direction  $\times$  Action interaction, alongside a surviving winsorized magnitude  $\times$  Action interaction, indicates  $w$  is strictly increasing rather than a step function. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.27: Sample-Exclusion Robustness for the Instrumental Motive — Experiment 2 (Citizens)

|                                | (1) Drop low- $ \delta $ | (2) Drop perfect sharers | (3) Drop low- $\alpha$ facts |
|--------------------------------|--------------------------|--------------------------|------------------------------|
| Expression                     | 0.067**<br>(0.026)       | 0.067***<br>(0.025)      | 0.087***<br>(0.027)          |
| Factual belief correction gain | 0.027<br>(0.019)         | 0.047***<br>(0.018)      | 0.043**<br>(0.021)           |
| Persuasion                     | 0.003<br>(0.009)         | 0.006<br>(0.010)         | 0.008<br>(0.011)             |
| Persuasion $\times$ Action     | 0.023*<br>(0.013)        | 0.021<br>(0.014)         | 0.022<br>(0.015)             |
| Action                         | 0.029**<br>(0.014)       | 0.015<br>(0.011)         | 0.022<br>(0.015)             |
| Fact FE                        | Yes                      | Yes                      | Yes                          |
| Num. obs.                      | 5,805                    | 6,355                    | 5,233                        |

*Sample:* Experiment 2 citizens; preferred action-decomposition specification, fact fixed effects, sender-clustered standard errors. Column (1) drops senders in the bottom quartile of mean  $|\delta|$  (expect near-zero persuasion across facts). Column (2) drops perfect sharers (provision rate = 1.0, all five facts shared). Column (3) drops the bottom tertile of facts by mean expressive magnitude  $|\alpha|$ . The instrumental motive (*Persuasion*  $\times$  *Action*) is positive in every restriction. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.28: Leave-One-Fact-Out Robustness of the Motive Decomposition — Experiment 2 (Citizens)

| Motive                         | Sign-positive | Min $\beta$ | Max $\beta$ | Max $p$ |
|--------------------------------|---------------|-------------|-------------|---------|
| Expression                     | 12/12         | 0.053       | 0.082       | 0.031   |
| Factual belief correction gain | 12/12         | 0.025       | 0.043       | 0.160   |
| Persuasion                     | 12/12         | 0.011       | 0.018       | 0.102   |
| Fact FE                        | Yes           |             |             |         |
| Clustered SE                   | Sender        |             |             |         |

*Sample:* Experiment 2 citizens. Each motive is re-estimated twelve times in the three-motive specification (*Expression*, *Factual belief correction gain*, *Persuasion*), each time dropping one of the twelve facts; fact fixed effects, sender-clustered standard errors. Columns report the number of leave-one-out estimates with a positive coefficient, the minimum and maximum coefficient across the twelve runs, and the largest  $p$ -value observed.

Table A.29: Near-Orthogonality of the Motive Proxies — Experiment 2 (Citizens)

|                                          | Expression ( $\alpha$ ) | Persuasion ( $\delta$ ) | Factual belief correction ( $\Delta b$ ) | Distance    |
|------------------------------------------|-------------------------|-------------------------|------------------------------------------|-------------|
| Expression ( $\alpha$ )                  | 1.00                    |                         |                                          |             |
| Persuasion ( $\delta$ )                  | $\approx 0.00$          | 1.00                    |                                          |             |
| Factual belief correction ( $\Delta b$ ) | $\approx 0.00$          | $\approx 0.00$          | 1.00                                     |             |
| Distance                                 | $\approx 0.00$          | $\approx 0.00$          | $\approx 0.00$                           | 1.00        |
| <b>VIF</b>                               | $\approx 1$             | $\approx 1$             | $\approx 1$                              | $\approx 1$ |

*Sample:* Experiment 2 citizens; 7,740 sender-fact observations. Entries are within-fact residual correlations: each predictor is first residualized on the other three motives and fact fixed effects, then pairwise correlations are computed. The four motive proxies are near-uncorrelated after fact fixed effects, with variance inflation factors near 1. This corroborates the design logic—each proxy was constructed to be independent of the others within a sender’s five-fact portfolio—but the identification of separate motives rests on the design, not on the variance inflation factors alone.

|                                                       | Pooled              |                     |                     | With Action Interaction |                     |                     |
|-------------------------------------------------------|---------------------|---------------------|---------------------|-------------------------|---------------------|---------------------|
|                                                       | Positional          | Directional         | Horse race          | Positional              | Directional         | Horse race          |
| Expression $\alpha_{ij}$                              | 0.069***<br>(0.023) | 0.065***<br>(0.023) | 0.065***<br>(0.023) | 0.069***<br>(0.023)     | 0.064***<br>(0.023) | 0.064***<br>(0.023) |
| Factual belief correction $\Delta b_{ijk}$            | 0.037**<br>(0.017)  | 0.036**<br>(0.018)  | 0.036**<br>(0.018)  | 0.037**<br>(0.017)      | 0.036**<br>(0.018)  | 0.036**<br>(0.018)  |
| Persuasion-positional $\Delta\pi_{ijk}^{\text{pos}}$  | 0.015**<br>(0.007)  |                     | -0.004<br>(0.009)   | 0.003<br>(0.009)        |                     | -0.010<br>(0.013)   |
| Persuasion-directional $\Delta\pi_{ijk}^{\text{dir}}$ |                     | 0.016***<br>(0.006) | 0.018**<br>(0.008)  |                         | 0.006<br>(0.008)    | 0.013<br>(0.012)    |
| ACTION                                                |                     |                     |                     | 0.024*<br>(0.013)       | 0.024*<br>(0.013)   | 0.024*<br>(0.013)   |
| Persuasion-pos. $\times$ ACTION                       |                     |                     |                     | 0.024*<br>(0.013)       |                     | 0.013<br>(0.019)    |
| Persuasion-dir. $\times$ ACTION                       |                     |                     |                     |                         | 0.019<br>(0.011)    | 0.010<br>(0.017)    |
| Clustered SE                                          | Sender              | Sender              | Sender              | Sender                  | Sender              | Sender              |
| Fact FE                                               | Yes                 | Yes                 | Yes                 | Yes                     | Yes                 | Yes                 |
| Action level + interaction                            | No                  | No                  | No                  | Yes                     | Yes                 | Yes                 |
| Sample                                                | All senders         | Non-neutral         | Non-neutral         | All senders             | Non-neutral         | Non-neutral         |
| Num. obs.                                             | 7740                | 7480                | 7480                | 7740                    | 7480                | 7480                |
| R <sup>2</sup>                                        | 0.004               | 0.004               | 0.004               | 0.005                   | 0.005               | 0.005               |
| Adj. R <sup>2</sup>                                   | 0.002               | 0.002               | 0.002               | 0.003                   | 0.003               | 0.003               |

Dependent variable: binary provision decision ( $y_{ij} \in \{0,1\}$ ). Sample: Experiment 2,  $N = 1,548$  senders. Directional columns exclude 260 observations from senders with  $\text{pref}_i = 0$  (neutral donors;  $N = 7,480$ ); positional columns retain the full sample ( $N = 7,740$ ). Persuasion-positional ( $\Delta\pi_{ij}^{\text{pos}}$ ) is `delta_gain_2`: reduction in Likert distance to the sender's own position. Persuasion-directional ( $\Delta\pi_{ij}^{\text{dir}}$ ) is  $\delta_{ij}^{\text{dir}} \equiv -\Delta\text{stance}_{ij} \cdot \text{sign}(\text{pref}_i)$ : whether the receiver moves toward the sender's preferred side, preference-magnitude-free. The two proxies correlate at  $r = 0.74$  after fact fixed effects; horse-race columns partition shared variance and are descriptive, not an adjudication between the proxies. Standard errors clustered at the sender level. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.30: Positional vs. directional persuasion proxies, Experiment 2.

|                                            | Positional          |                     | Directional         |                     |
|--------------------------------------------|---------------------|---------------------|---------------------|---------------------|
|                                            | Continuous          | Binary              | Continuous          | Binary              |
| Expression $\alpha_{ij}$                   | 0.069***<br>(0.023) | 0.069***<br>(0.023) | 0.064***<br>(0.023) | 0.065***<br>(0.023) |
| Factual belief correction $\Delta b_{ijk}$ | 0.037**<br>(0.017)  | 0.038**<br>(0.017)  | 0.036**<br>(0.018)  | 0.036**<br>(0.018)  |
| Persuasion-pos. (continuous)               | 0.003<br>(0.009)    |                     |                     |                     |
| Persuasion-pos. (binary)                   |                     | 0.023<br>(0.016)    |                     |                     |
| Persuasion-dir. (continuous)               |                     |                     | 0.006<br>(0.008)    |                     |
| Persuasion-dir. (binary)                   |                     |                     |                     | 0.035**<br>(0.017)  |
| ACTION                                     | 0.024*<br>(0.013)   | 0.025<br>(0.016)    | 0.024*<br>(0.013)   | 0.020<br>(0.017)    |
| Persuasion-pos. (cont.) $\times$ ACTION    | 0.024*<br>(0.013)   |                     |                     |                     |
| Persuasion-pos. (bin.) $\times$ ACTION     |                     | 0.001<br>(0.023)    |                     |                     |
| Persuasion-dir. (cont.) $\times$ ACTION    |                     |                     | 0.019<br>(0.011)    |                     |
| Persuasion-dir. (bin.) $\times$ ACTION     |                     |                     |                     | 0.013<br>(0.024)    |
| Clustered SE                               | Sender              | Sender              | Sender              | Sender              |
| Fact FE                                    | Yes                 | Yes                 | Yes                 | Yes                 |
| Sample                                     | All senders         | All senders         | Non-neutral         | Non-neutral         |
| Num. obs.                                  | 7740                | 7740                | 7480                | 7480                |
| R <sup>2</sup>                             | 0.005               | 0.004               | 0.005               | 0.005               |
| Adj. R <sup>2</sup>                        | 0.003               | 0.002               | 0.003               | 0.003               |

Dependent variable: binary provision decision ( $y_{ij} \in \{0,1\}$ ). Sample:  $N = 1,548$  Experiment 2 senders (7,740 observations); directional columns exclude 260 observations from senders with  $\text{pref}_i = 0$  (neutral donors;  $N = 7,480$ ). Each column re-estimates the ACTION-interaction specification with one coding of the persuasion proxy. Continuous codings: positional  $\Delta\pi_{ijk}^{\text{pos}}$  ( $\text{delta\_gain\_2}$ ) and directional  $\Delta\pi_{ijk}^{\text{dir}}$ . Binary codings: an indicator equal to one when the corresponding continuous measure is strictly positive. All specifications include Expression, Factual belief correction, fact fixed effects, and sender-clustered standard errors. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.31: Persuasion proxy functional form: continuous vs. binary codings, Experiment 2.

Table A.32: Expression Measure ( $\alpha_{ij}$ ) by Sender Stance and Fact Direction — Experiment 2*Panel A: Pooled (all topics)*

| Sender stance        | Anti-policy fact | Pro-policy fact |
|----------------------|------------------|-----------------|
| 1 (most pro-policy)  | +0.32            | +0.56           |
| 2                    | −0.01            | +0.51           |
| 3 (neutral)          | +0.15            | +0.43           |
| 4                    | +0.65            | −0.04           |
| 5 (most anti-policy) | +0.78            | +0.27           |

*Panel B: By topic*

| Sender stance | Immigration |       | Climate |       |
|---------------|-------------|-------|---------|-------|
|               | Anti        | Pro   | Anti    | Pro   |
| 1 (most pro)  | +0.12       | +0.41 | +0.37   | +0.59 |
| 2             | −0.31       | +0.51 | +0.11   | +0.51 |
| 3 (neutral)   | +0.09       | +0.53 | +0.19   | +0.37 |
| 4             | +0.67       | +0.05 | +0.61   | −0.22 |
| 5 (most anti) | +0.81       | +0.31 | +0.71   | +0.15 |

Notes: Expression ( $\alpha_{ij}$ ) = mean cross-domain rating for each (sender stance, fact direction) cell, scaled to  $[-1, +1]$ :  $\alpha_{ij} = 2(\bar{g}_{ij} - 0.5)$ , where  $g = 1$  if the sender rates the cross-domain fact as “supporting my opinion.” Sender stance is the stated Likert position on the topic elicited after the provision task (1 = most pro-policy, 5 = most anti-policy). Cross-domain ratings are elicited on the *other* topic; fact direction refers to the rated fact, not the provision topic. Panel A pools immigration and climate cross-domain ratings; cell sizes range from 195 to 294 observations (per stance, pooled). Panel B splits by topic; per-cell sizes range from approximately 90 to 160 observations. See Appendix Table A.33 for per-fact breakdown.

Table A.33: Expression Measure ( $\alpha_{ij}$ ) by Sender Stance and Individual Fact — Experiment 2

## Panel A: Immigration facts

| Sender stance | Anti-policy direction (–) |            |              | Pro-policy direction (+) |            |             |
|---------------|---------------------------|------------|--------------|--------------------------|------------|-------------|
|               | Violent crime             | Bürgergeld | Res. permits | Shortage fields          | Employment | Integration |
| 1 (most pro)  | +0.00                     | +0.29      | +0.00        | +0.43                    | +0.00      | +0.67       |
| 2             | –0.33                     | –0.16      | –0.43        | +0.44                    | +0.37      | +0.77       |
| 3 (neutral)   | +0.29                     | –0.07      | +0.00        | +0.64                    | +0.33      | +0.59       |
| 4             | +0.66                     | +0.83      | +0.54        | +0.00                    | +0.28      | –0.16       |
| 5 (most anti) | +0.82                     | +0.82      | +0.79        | +0.38                    | +0.34      | +0.24       |

## Panel B: Climate facts

| Sender stance | Anti-policy direction (–) |                |              | Pro-policy direction (+) |            |             |
|---------------|---------------------------|----------------|--------------|--------------------------|------------|-------------|
|               | Grid fees                 | Business surv. | Price fluct. | Renewable share          | Solar cost | Renew. jobs |
| 1 (most pro)  | +0.23                     | +0.45          | +0.38        | +0.49                    | +0.71      | +0.55       |
| 2             | –0.15                     | +0.23          | +0.26        | +0.46                    | +0.59      | +0.46       |
| 3 (neutral)   | +0.14                     | +0.10          | +0.30        | +0.31                    | +0.50      | +0.30       |
| 4             | +0.64                     | +0.58          | +0.61        | +0.00                    | –0.40      | –0.33       |
| 5 (most anti) | +0.78                     | +0.79          | +0.54        | +0.26                    | +0.20      | –0.10       |

Notes: Expression  $(\alpha_{ij}) = 2 (\bar{g}_{ij} - 0.5)$ , where  $g = 1$  if the sender rates the cross-domain fact as “supporting my opinion.” Per-cell sizes are approximately 30–100 observations (individual fact  $\times$  stance); values in small cells (fewer than 40 observations) should be interpreted with caution. Fact identifiers in brackets. Sender stance = stated Likert position on the provision topic (1 = most pro-policy; 5 = most anti-policy). Cross-domain ratings are elicited on the *other* topic; all senders in the cell have the same stated stance. Pooled table in Appendix Table A.32.




Figure A.11: Expected factual belief correction  $\Delta b_{ijk}$ : citizens (Experiment 2) vs. journalists (Experiment 3). *Panels (A)–(B)*: mean  $\Delta b_{ijk}$  by fact, each population ranked independently from lowest (rank 1) to highest (rank 12); error bars are  $\pm 1$  SE. Facts are ranked independently within each population; the x-axis scale is identical across panels, making the  $\approx 1.9\times$  level difference directly visible. *Panel (C)*: mean  $\Delta b_{ijk}$  by alignment (Aligned vs. Unaligned, left pair) and topic (Climate vs. Immigration, right pair) for each population. The climate-over-immigration hierarchy replicates across populations (gap  $+0.033$  pp citizens,  $+0.053$  pp journalists). Aligned and unaligned facts are statistically indistinguishable within each population ( $p = 0.729$  citizens;  $p = 0.732$  journalists), ruling out motivated beliefs about belief-correcting ability as a driver of provision differences between groups.




Figure A.12: Distributions of the elicited returns to disclosure, Experiment 3 (Journalists). Distributions of the two incentivized elicitations for positionally aligned facts ( $\text{Align}_{ij} > 0$ ; bars) and unaligned facts ( $\text{Align}_{ij} < 0$ ; black outline), as shares within each group; observations of senders with  $\text{Pref}_i = 0$  (50 of 1,540) are omitted from the split, and 5 observations of one journalist with missing elicitations are excluded. Panel (a): expected factual belief correction  $\hat{\Delta}b_{ijk}$  (0.05-unit bins; group means 0.324 aligned, 0.319 unaligned). Panel (b): expected persuasion  $\hat{\Delta}\pi_{ijk}$  (0.125-unit bins; the bin at  $[-0.125, 0)$  contains exact-zero predictions; group means 0.030 aligned, 0.003 unaligned). Shaded regions mark negative values (disclosure expected to move the receiver's factual belief away from the truth, respectively the receiver away from the sender's position). Journalists are always in the ACTION arm. Elicitations incentivized via Quadratic Scoring Rule. This is the Experiment 3 counterpart of Figure 5, panels (b) and (c).




Figure A.13: Expected returns to disclosure by sender-receiver distance, Experiment 3 (Journalists). One marker per fact  $\times$  sender-receiver distance  $\times$  alignment (cell mean over  $\geq 3$  elicitations), positionally aligned facts ( $Align_{ij} > 0$ ) as filled black circles, unaligned ( $Align_{ij} < 0$ ) as open gray squares, horizontally offset for legibility; facts of senders with  $Pref_i = 0$  are omitted from the split. Panel (a): expected factual belief correction  $\widehat{\Delta b}_{ijk}$ . Panel (b): expected persuasion  $\widehat{\Delta \pi}_{ijk}$  (predicted reduction in Likert distance between the receiver's stance and the sender's own position). Journalists are always in the ACTION arm. Dashed lines mark zero. Elicitations incentivized via Quadratic Scoring Rule. This is the Experiment 3 counterpart to Appendix Figure A.9.

Table A.34: Experiment 3 Sample Characteristics: Comparison to Population Benchmark

|                                                        | Experiment 3<br>( <i>N</i> = 307) | Loosen et al. (2023)<br>( <i>N</i> = 1,221) |
|--------------------------------------------------------|-----------------------------------|---------------------------------------------|
| <i>Gender</i>                                          |                                   |                                             |
| Female                                                 | 53.4                              | 44.0                                        |
| Male                                                   | 44.6                              | 55.9                                        |
| <i>Age</i>                                             |                                   |                                             |
| Under 50                                               | 77.3                              | 62.4                                        |
| <i>Education</i>                                       |                                   |                                             |
| University degree (B.A.+)                              | 84.6                              | 67.1                                        |
| <i>Medium type</i>                                     |                                   |                                             |
| Print                                                  | 28.4                              | 57.0                                        |
| Online                                                 | 45.4                              | 3.1                                         |
| Television                                             | 7.2                               | 17.3                                        |
| Radio/podcast                                          | 10.8                              | 17.2                                        |
| News agency                                            | 3.6                               | 3.5                                         |
| <i>Broadcaster type</i> (TV/radio respondents only)    |                                   |                                             |
| Public broadcaster                                     | 80.0                              | 61.4                                        |
| <i>Role</i> (multiple selection; shares exceed 100 %)  |                                   |                                             |
| Editor                                                 | 57.2                              | ≈ 61                                        |
| <i>Editorial autonomy</i>                              |                                   |                                             |
| High ( $\geq 4/5$ )                                    | 75.8                              | 62.0                                        |
| <i>Political affiliation</i> (no population benchmark) |                                   |                                             |
| Grüne or Linke (closest party)                         | 58.5                              | —                                           |
| Mean left-right self-placement (0–10)                  | 3.2                               | —                                           |

*Sample:* *N* = 307 professional German journalists recruited via DJV and DJU mailing lists, Experiment 3 (convenience sample, not a probability sample). *Benchmark:* Loosen et al. (2023), stratified random sample of active German journalists (*N* = 1,221, September 2022–February 2023, HBI Hamburg; Loosen et al.). All figures are percentages. *Medium type:* Experiment 3 uses a single online category for all digital-native outlets; the Loosen et al. benchmark reports native-online journalists (3.1%) as a distinct population segment; print aggregates newspapers and magazines (57.0%). *Broadcaster type* is conditional on TV or radio respondents only. *Role* allows multiple selection; the Loosen et al. editor figure is approximate. *Age:* Loosen et al. report a population mean of 45.3 years; Experiment 3 elicited age as a categorical variable. *Political affiliation* and left-right self-placement were not elicited in Loosen et al. and appear without a benchmark.

| Moderator                                           | Interaction of moderator with: |                      |                    |                    | Obs. |
|-----------------------------------------------------|--------------------------------|----------------------|--------------------|--------------------|------|
|                                                     | Align <sub>ij</sub>            | $\alpha_{ij}$        | $\Delta b_{ijk}$   | $\Delta \pi_{ijk}$ |      |
| Distance  Stance <sub>i</sub> – Stance <sub>k</sub> | –0.009<br>(0.011)              | –0.082**<br>(0.038)  | 0.021<br>(0.029)   | 0.016<br>(0.012)   | 1535 |
| Polarisation                                        | –0.010<br>(0.010)              | –0.029<br>(0.036)    | –0.028<br>(0.027)  | 0.012<br>(0.012)   | 1530 |
| Issue extremity                                     | –0.023<br>(0.030)              | –0.126*<br>(0.072)   | 0.006<br>(0.057)   | 0.020<br>(0.023)   | 1535 |
| Online medium                                       | –0.070**<br>(0.029)            | –0.288***<br>(0.096) | 0.013<br>(0.080)   | 0.006<br>(0.029)   | 1530 |
| Public broadcaster                                  | –0.013<br>(0.077)              | –0.302<br>(0.236)    | 0.075<br>(0.243)   | –0.069<br>(0.071)  | 275  |
| National/international reach                        | 0.068**<br>(0.031)             | 0.082<br>(0.102)     | 0.035<br>(0.080)   | 0.027<br>(0.032)   | 1535 |
| Editor role                                         | –0.027<br>(0.030)              | –0.063<br>(0.099)    | –0.068<br>(0.077)  | 0.012<br>(0.029)   | 1535 |
| Editorial autonomy                                  | 0.005<br>(0.015)               | –0.009<br>(0.055)    | 0.095**<br>(0.040) | 0.017<br>(0.012)   | 1530 |

Dependent variable: binary provision decision ( $y_{ij} \in \{0, 1\}$ ), Experiment 3 journalists. Each row reports interaction coefficients between one moderator and the selection measures: column (1) is from a regression of provision on positional alignment, the moderator, and their interaction; columns (2)–(4) are from a single regression interacting the moderator with Expression ( $\alpha_{ij}$ ), the factual belief correction gain ( $\Delta b_{ijk}$ ), and Persuasion ( $\Delta \pi_{ijk}$ ) jointly. All specifications include fact fixed effects and journalist-clustered standard errors (in parentheses); moderators are centered at the estimation-sample mean. Moderators: sender–receiver distance, polarisation (|left–right – 5|), issue extremity (|stance – 3|), and the pre-registered journalist background variables, coarsened as online medium (vs. all other medium types), public broadcaster (defined only for television and radio senders), national or international reach (vs. local or regional), any editor role in the multi-select role item, and editorial autonomy (Likert 1–5). Benjamini–Hochberg correction is applied within two families: the pre-registered alignment battery (column 1; 8 tests, minimum  $q = 0.11$ ) and the channel battery (columns 2–4; 24 tests), in which only Expression  $\times$  Online medium survives at the 10 % level ( $q = 0.074$ ). Stars mark unadjusted  $p$ -values: \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.35: Heterogeneity of alignment- and channel-based selection among journalists, Experiment 3.

|                                                                    | Citizens (Exp. 2) |      |         | Journalists (Exp. 3) |      |         |
|--------------------------------------------------------------------|-------------------|------|---------|----------------------|------|---------|
|                                                                    | Mean              | SD   | % Agree | Mean                 | SD   | % Agree |
| <i>Panel A: Perceived political imbalance (% of valid answers)</i> |                   |      |         |                      |      |         |
| No imbalance perceived                                             |                   | 85.5 |         |                      | 82.4 |         |
| Imbalance favoring the left                                        |                   | 9.5  |         |                      | 11.7 |         |
| Imbalance favoring the right                                       |                   | 5.0  |         |                      | 5.9  |         |
| <i>Panel B: Normative and policy items (1–5 agreement)</i>         |                   |      |         |                      |      |         |
| Share opposing information                                         | 4.13              | 0.90 | 80.0    | 4.60                 | 0.68 | 95.4    |
| Social-media transparency                                          | 4.05              | 0.96 | 75.7    | 4.61                 | 0.65 | 95.1    |
| Viewpoint-diversity obligation                                     | 4.17              | 0.90 | 80.4    | 4.22                 | 0.91 | 82.7    |

Post-questionnaire self-reports, one per sender. Panel A: answers to “Did you perceive a political imbalance in the study?” (options: no imbalance, imbalance favoring the left, imbalance favoring the right), in percent of valid answers; Experiment 2  $N = 1,548$  (no nonresponse), Experiment 3  $N = 307$  (one nonresponse). Panel B: agreement on a 1–5 scale (1 = fully disagree, 5 = fully agree) with three items: whoever passes on political information should also pass on factually correct information that contradicts their own position (worded for journalists in Experiment 3); social media should be more transparent about why users see some political information and not others; platforms and media should be obligated to surface relevant information from different perspectives. % Agree is the share answering 4 or 5. Experiment 2  $N = 1,548$ ; Experiment 3  $N = 306$  (two nonresponses).

Table A.36: Perceived design balance and normative items, Experiments 2 and 3.

| Item                                       | Interaction of item with: |                      |                   | Obs. |
|--------------------------------------------|---------------------------|----------------------|-------------------|------|
|                                            | Align <sub>ij</sub>       | $\alpha_{ij}$        | $\Delta\pi_{ijk}$ |      |
| <i>Panel A: Citizens (Experiment 2)</i>    |                           |                      |                   |      |
| Share opposing information                 | −0.008<br>(0.009)         | −0.014<br>(0.022)    | +0.006<br>(0.007) | 7740 |
| Social-media transparency                  | −0.007<br>(0.008)         | +0.029<br>(0.022)    | +0.007<br>(0.006) | 7740 |
| Viewpoint-diversity obligation             | −0.004<br>(0.009)         | −0.008<br>(0.023)    | +0.011<br>(0.007) | 7740 |
| <i>Panel B: Journalists (Experiment 3)</i> |                           |                      |                   |      |
| Share opposing information                 | −0.041<br>(0.026)         | −0.095<br>(0.078)    | −0.038<br>(0.029) | 1530 |
| Social-media transparency                  | −0.075***<br>(0.025)      | −0.194***<br>(0.071) | −0.041<br>(0.026) | 1530 |
| Viewpoint-diversity obligation             | +0.002<br>(0.019)         | −0.008<br>(0.060)    | −0.014<br>(0.017) | 1530 |

Dependent variable: binary provision decision ( $y_{ij} \in \{0, 1\}$ ). Each cell reports the interaction coefficient between the row item (centered at 3, the scale midpoint) and one selection measure, from a separate regression of provision on the measure, the item, and their interaction, with fact fixed effects and sender-clustered standard errors (in parentheses). Selection measures: positional alignment (Align<sub>ij</sub>), Expression ( $\alpha_{ij}$ ), and Persuasion ( $\Delta\pi_{ijk}$ ). Items as in Table A.36, Panel B. Benjamini–Hochberg correction is applied across the 18-test family (3 items  $\times$  3 measures  $\times$  2 samples): the two Experiment-3 transparency interactions in columns (1)–(2) survive at the 10% level ( $q = 0.060$ ); no other test has  $q < 0.40$ . Stars mark unadjusted  $p$ -values: \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

Table A.37: Normative items and selection: channel  $\times$  item interactions, Experiments 2 and 3.

## B Experimental Design Details

### B.1 Statistic Selection and Validation Survey

I selected statistics for the main experiment from a larger pool of candidates using a pre-registered validation survey ( $N = 100$ , AsPredicted No. 232064) conducted on Prolific between June 10 and June 21, 2025. The survey evaluated eight candidate statistics—four hypothesized as pro-immigration and four as anti-immigration—against three survey-measured criteria (perceived credibility, policy relevance, and perceived political direction) and one researcher-verified criterion (veracity). The selection rule was pre-specified: within each direction, select the two statistics with the highest policy relevance among those with adequate perceived credibility, subject to the constraint that the final set contains two pro-immigration and two anti-immigration statistics.

**Candidate Statistics and Selection Criteria.** Table B.38 reports the eight candidate statistics with their survey-measured scores. The four criteria were operationalized as follows:

*Veracity* was established by the researcher through source verification. Each selected statistic is supported by multiple independent, authoritative sources: peer-reviewed academic publications, federal statistical agencies, or nonpartisan research institutions. Source documentation is provided below.

*Policy relevance* was measured by asking respondents: “How relevant do others consider the statistic when forming immigration policy positions?” (1 = completely irrelevant, 5 = extremely relevant).

*Perceived credibility* was measured by asking respondents: “How credible is the statistic?” (1 = not at all, 5 = entirely).

*Perceived political direction* was measured through two items: *usability* (“For whom would the statistic be most useful in a policy debate?”; 1 = exclusively supporting expansive immigration, 5 = exclusively supporting restrictive immigration) and *opinionability* (“How does learning about the statistic change someone’s immigration policy

views?"; 1 = convince to support expansion, 5 = convince to support limitation). These directional measures serve the classification of statistics as pro- or anti-immigration and the balance criterion.

Table B.38: Candidate Statistics: Survey-Measured Indicators

| Statistic                     | Policy      | Perceived   | Perceived Direction |                |
|-------------------------------|-------------|-------------|---------------------|----------------|
|                               | Relevance   | Credibility | Usability           | Opinionability |
| <i>Pro-immigration</i>        |             |             |                     |                |
| Crime Rate <sup>†</sup>       | 3.74 (1.18) | 3.50 (0.97) | 2.23 (1.06)         | 2.36 (1.11)    |
| Labor <sup>†</sup>            | 3.64 (1.04) | 3.81 (1.02) | 2.40 (1.09)         | 2.57 (1.23)    |
| Entrepreneurship              | 3.12 (1.08) | 3.54 (1.00) | 2.09 (1.05)         | 2.18 (1.15)    |
| Healthcare                    | 3.28 (1.03) | 3.59 (1.07) | 2.03 (1.05)         | 2.19 (1.15)    |
| <i>Anti-immigration</i>       |             |             |                     |                |
| Border Crossings <sup>†</sup> | 3.85 (1.14) | 3.97 (0.90) | 3.94 (1.14)         | 4.08 (1.11)    |
| Fiscal Impact <sup>†</sup>    | 3.79 (1.01) | 3.55 (0.93) | 3.68 (1.15)         | 3.73 (1.01)    |
| Poverty Rate                  | 3.41 (1.01) | 3.93 (0.93) | 3.40 (1.13)         | 3.52 (1.11)    |
| Wage Inequality               | 3.35 (0.95) | 3.53 (0.90) | 3.41 (1.06)         | 3.43 (0.99)    |

Notes:  $N = 100$ . Cells report means with standard deviations in parentheses. Selected statistics with <sup>†</sup>. Policy relevance: 1 = completely irrelevant, 5 = extremely relevant. Perceived credibility: 1 = not at all credible, 5 = entirely credible. Usability and opinionability: 1 = supports expansive immigration, 5 = supports restrictive immigration. One-sample  $t$ -tests against the neutral midpoint of 3 reject the null for all eight statistics on both directional measures ( $p < 0.001$  for all selected statistics; see text).

**Selection.** The selection rule applied within each direction independently. Among the four pro-immigration candidates, Crime Rate (3.74) and Labor Force Participation (3.64) had the highest policy relevance scores, both with adequate perceived credibility ( $\geq 3.50$ ). Among the four anti-immigration candidates, Border Crossings (3.85) and Fiscal Impact (3.79) had the highest policy relevance, again with adequate credibility. This rule uniquely determines the four selected statistics reported in the main text (Section 2.1). The one rejected statistic with notably high credibility—Poverty Rate (3.93)—scored lower on policy relevance (3.41) than both selected anti-immigration statistics,

confirming that the selection was not driven by credibility alone.

Perceived directional scores did not vary significantly by respondents' political affiliation (Democrat, Republican, Independent), indicating that the directional classification reflects consensus perception rather than partisan projection.

**Direction Validation.** I classify Crime Rate and Labor Force Participation as pro-immigration and Border Crossings and Fiscal Impact as anti-immigration. This classification is validated by the survey's directional measures. For usability and opinionability respectively, the pro-immigration statistics score below the neutral midpoint of 3 (Crime: 2.23, 2.36; Labor: 2.40, 2.57) and the anti-immigration statistics score above it (Border: 3.94, 4.08; Fiscal: 3.68, 3.73). One-sample  $t$ -tests against the midpoint of 3 reject the null of no directional content for all four selected statistics on both measures: Crime ( $t = -7.25$  and  $t = -5.74$ ), Labor ( $t = -5.50$  and  $t = -3.49$ ), Border Crossings ( $t = 8.21$  and  $t = 9.75$ ), and Fiscal Impact ( $t = 5.89$  and  $t = 7.20$ ), all  $p < 0.001$ .

This classification is independently confirmed by the main experiment. In the expected attitude shift elicitation (Section 2.3.1, Table A.4), senders in the main experiment ( $N = 1,200$ ) expect Crime and Labor to shift receiver preferences toward pro-immigration positions (25 and 17 pp, both  $p < 0.001$ ) and Border and Fiscal to shift preferences toward anti-immigration positions ( $-34$  and  $-32$  pp, both  $p < 0.001$ ). The convergence of two independent samples using two independent methods—survey-based directional perception in the validation study and incentivized attitude-shift predictions in the main experiment—makes the directional classification robust.

**Source Documentation.** Veracity was established through source verification rather than survey measurement. Each selected statistic is supported by multiple independent, authoritative sources. I report the specific values shown to participants, the primary source for each, and corroborating evidence. All sources were available at the time of the experiment (July–August 2025).

*Crime Rate* (immigrants have lower crime rates than U.S.-born Americans). Participants were shown violent crime *arrest* rates per 100,000 residents in Texas (2012–2018):

213.0 for U.S.-born citizens, 185.3 for legal immigrants (13% lower), and 96.2 for undocumented immigrants (55% lower). These figures are from [Light et al. \(2020\)](#), who analyze the universe of felony arrests recorded in the Texas Computerized Criminal History database cross-referenced with Department of Homeland Security immigration status checks.<sup>32</sup> The finding is independently corroborated by [?](#), who uses Texas Department of Public Safety conviction data over a longer period (2013–2022) and reports homicide conviction rates of 3.0 per 100,000 for native-born Americans, 2.2 for unauthorized immigrants (26% lower), and 1.2 for legal immigrants (61% lower).<sup>33</sup> At the national level, [Abramitzky et al. \(2024\)](#) analyze census-linked incarceration records from 1870–2020 and find that immigrants have been consistently less likely to be incarcerated than the U.S.-born population throughout this period.<sup>34</sup> Texas is the only U.S. state that systematically records the immigration status of all arrestees, making it the sole source of population-level crime-rate comparisons by immigration status. The national-level incarceration evidence in [Abramitzky et al. \(2024\)](#) suggests the Texas pattern is not anomalous.

*Labor Force Participation* (immigrants have higher labor force participation than U.S.-born Americans). Participants were shown 2023 annual average labor force participation rates: 66.6% for the foreign-born population and 61.8% for the native-born population (4.8 percentage points higher). These figures are from Table 1 of the Bureau of Labor Statistics annual report *Foreign-Born Workers: Labor Force Characteristics—2023*.<sup>35</sup> The same data are available through the Federal Reserve Bank of St. Louis FRED database, Current Population Survey Table A-7.<sup>36</sup> The foreign-born participation rate has exceeded the native-born rate in every year of the BLS series.

*Border Crossings* (unauthorized border crossings have been higher in the past five

---

<sup>32</sup>Published in *Proceedings of the National Academy of Sciences*; data and replication materials available at <https://www.openicpsr.org/openicpsr/project/124923>.

<sup>33</sup>Cato Institute Policy Analysis no. 977; available at <https://www.cato.org/policy-analysis/illegal-immigrant-murderers-texas-2013-2022>.

<sup>34</sup>Published in *American Economic Review: Insights*.

<sup>35</sup>The annual report is available at <https://www.bls.gov/news.release/pdf/forbrn.pdf>. At the time of writing, the 2024 edition (released May 2025) reports 66.5% versus 61.7%, confirming the same pattern.

<sup>36</sup><https://fred.stlouisfed.org/release/tables?rid=50&eid=1736>

years relative to historical averages). Participants were shown two figures computed from U.S. Customs and Border Protection (CBP) enforcement statistics:<sup>37</sup> a historical average of 890,435 apprehensions per year over fiscal years 1970–2019, and a recent five-year average of 1,584,573 encounters per year over fiscal years 2020–2024 (78% higher). The comparison requires combining two CBP reporting series: “apprehensions” (reported through FY2019) and “encounters” (reported from FY2020 onward, which combine apprehensions and expulsions under Title 42 and Title 8 processing).<sup>38</sup> The statistic reflects data available at experiment time (July–August 2025). Subsequent enforcement changes under the second Trump administration reduced encounters to record lows in FY2025, as documented by Pew Research Center analysis of CBP data.<sup>39</sup> The FY2025 decline does not affect the validity of the statistic as presented to participants, which concerned the five-year average through FY2024.

*Fiscal Impact* (net fiscal costs of first-generation immigrants are higher than U.S.-born Americans at the state and local level). Participants were shown that first-generation immigrant households generate \$57.4 billion more in annual state and local government costs than they contribute in taxes, based on 2011–2013 data. This figure is from the National Academies of Sciences, Engineering, and Medicine comprehensive report on the economic and fiscal consequences of immigration, which remains the most authoritative and comprehensive fiscal analysis of immigration in the United States.<sup>40</sup> The report notes that this fiscal gap reverses at the federal level and across generations: the children and grandchildren of immigrants are among the strongest net fiscal contributors. I present only the state-and-local first-generation finding because it represents the anti-immigration direction in my balanced design. The Cato Institute provides complementary analysis of immigrants’ fiscal effects over a longer period (1994–

---

<sup>37</sup>CBP historical data available at <https://www.cbp.gov/newsroom/stats/cbp-enforcement-statistics>.

<sup>38</sup>Beginning in March 2020, CBP shifted from reporting apprehensions to reporting “encounters,” a broader category. The historical series (1970–2019) uses the apprehension metric; the recent series (2020–2024) uses the encounter metric. Since encounters  $\geq$  apprehensions by construction, the 78% comparison is an upper bound—i.e., it may slightly overstate the increase by comparing a narrower historical measure to a broader recent measure.

<sup>39</sup><https://www.pewresearch.org/short-reads/2026/02/02/migrant-encounters-at-the-us-mexico-border-ar>

<sup>40</sup>National Academies of Sciences, Engineering, and Medicine (2017); the state and local fiscal impact analysis is in Chapter 8. Available at <https://www.nationalacademies.org/read/23550/chapter/14>.

2023), confirming that the fiscal picture varies substantially by level of government and generational horizon.<sup>41</sup>

## B.2 Belief Elicitation and Donation Task (Experiment 1)

This section documents two implementation details of the Experiment 1 sender protocol referenced in Section 2.1: the incentivization of the prior-knowledge elicitation and the revealed-preference donation task.

**Prior-Knowledge Elicitation.** Senders complete each comparative statement by choosing “higher,” “lower,” or “don’t know.” Correct completions receive a bonus, incorrect completions do not, and “don’t know” responses receive the bonus with 25% probability. The 25% payoff for “don’t know” is strictly dominated under expected utility maximization with binary alternatives, where uninformed guessing yields a 50% expected bonus. The option may nonetheless be selected by highly ambiguity-averse participants who prefer an objective 25% lottery over an ambiguous 50% guess. This implies that reported “don’t know” rates likely understate genuine ignorance among expected-utility maximizers.

**Donation Task.** Senders allocate \$500 between the *American Immigration Council* (AIC), which advocates for increased immigration, and *NumbersUSA*, which advocates for decreased immigration—both well-established organizations representing opposing positions in the immigration debate. I briefly describe the stated goals of each organization and present their missions in their own words, with the display order randomized. Participants are informed that one donation decision from the experiment will be randomly selected and implemented, and they receive a verification link to confirm the donation was placed.

---

<sup>41</sup><https://www.cato.org/white-paper/immigrants-recent-effects-government-budgets-1994-2023>

### B.3 Fact Selection and Pre-Validation Survey — Experiments 2 and 3

I selected facts for Experiments 2 and 3 from a larger candidate pool using a pre-registered pre-validation survey (AsPredicted #290989, quota target  $N = 330$ , final analysis sample  $N \approx 300$  after speeder and consent exclusions) conducted on the Bilendi panel in Germany in May 2026. The survey evaluated 36 candidate facts—six hypothesized as pro-policy and six as anti-policy in each of three domains (immigration, climate, redistribution)—against four criteria, and additionally validated the domain-specific advocacy organizations used in the ACTION treatment.

**Criteria and Selection Rule.** Four criteria were applied to each candidate fact. *Prior knowledge* was elicited via an incentivized fill-in-the-blank completion of the fact statement with a confidence slider [50, 100]; the probabilistic prior knowledge score is  $PK = \text{confidence}/100$  for correct completions and  $1 - \text{confidence}/100$  for incorrect completions. *Directional consistency* was measured by a 5-point rating (1 = strongly pro-policy, 5 = strongly anti-policy); a fact passes if the share of respondents rating it in its assigned direction exceeds the share rating it in the opposite direction. *Policy relevance* and *perceived credibility* were each measured on a 1–5 scale.

The pre-specified selection rule selected the top three facts by policy relevance within each topic  $\times$  direction cell among those passing directional consistency, subject to balance constraints across the pro and anti sets:  $|\bar{c}_{\text{pro}} - \bar{c}_{\text{anti}}| \leq 1.0$  for credibility,  $|\overline{PK}_{\text{pro}} - \overline{PK}_{\text{anti}}| \leq 0.25$ , and  $|\bar{r}_{\text{pro}} - \bar{r}_{\text{anti}}| \leq 1.0$  for policy relevance.

The ACTION treatment organizations were additionally validated against two criteria: (i) at least 50% of respondents correctly recognize each organization as representing the pro-policy or anti-policy position respectively, and (ii) a positive Spearman correlation between topic-specific attitude and hypothetical donation to the pro-organization, restricted to respondents with directional attitudes.

**Domain Selection.** The redistribution domain failed on both grounds and is excluded. All six redistribution anti-direction candidate facts failed the directional consistency criterion, and the redistribution anti-organization had a recognition rate of 33.3%, well

below the 50% threshold. Immigration passed all fact and organization criteria (anti-organization recognition: 93.6%; Spearman  $r = 0.57$ ). Climate's facts all passed directionality and the organization correlation criterion ( $r = 0.57$ ); the climate anti-organization fell marginally below the recognition threshold (42.6% vs. 50%). Since climate's failure was marginal and isolated to one metric while redistribution's failures were compounding, climate was retained and this deviation is noted here.

**Deviations from Algorithmic Fact Selection.** *Immigration.* The algorithm selected *Bildungsabschluss* (PK = 0.81), which produced a pro-anti PK gap of 0.29, exceeding the 0.25 balance threshold. It was replaced by *Illegale Aufenthalte* (PK = 0.51), closing the gap to 0.15 while satisfying all other balance constraints. The label of one anti-fact was updated from *Arbeitslosigkeit* to *Bürgergeld* to reflect the 2026 German policy context; the underlying statistic is unchanged.

*Climate.* The pro set was revised to *Erneuerbare Energien*, *Stromgestehungskosten*, and *Grüne Beschäftigung* for higher directional consistency. *Netzausbaukosten* (anti) was reformulated between the pre-validation survey and the main experiment; directional ratings in Table B.39 reflect the pre-reformulation version, which failed the consistency criterion (17.5% consistent vs. 42.1% inconsistent). The reformulated version is what was fielded in the main experiment.

**Selected Facts and Pre-Validation Ratings.** Table B.39 reports the 12 selected facts with their pre-validation survey scores. Table B.40 reports the balance check.

Table B.39: Selected Facts: Pre-Validation Ratings (Experiments 2 and 3)

| Fact                          | Dir. | Relevance | Credibility | PK   | % Consistent |
|-------------------------------|------|-----------|-------------|------|--------------|
| <i>Immigration</i>            |      |           |             |      |              |
| Fachkräftemangel              | Pro  | 3.79      | 3.65        | 0.52 | 45.8         |
| Beschäftigungswachstum        | Pro  | 3.43      | 3.45        | 0.55 | 38.8         |
| Integrationskurs              | Pro  | 3.61      | 3.29        | 0.37 | 32.1         |
| Kriminalität                  | Anti | 4.04      | 3.76        | 0.70 | 90.0         |
| Bürgergeld <sup>†</sup>       | Anti | 3.98      | 3.74        | 0.67 | 78.7         |
| Illegale Aufenthalte          | Anti | 3.69      | 3.45        | 0.51 | 76.5         |
| <i>Climate policy</i>         |      |           |             |      |              |
| Erneuerbare Energien          | Pro  | 3.67      | 3.61        | 0.52 | 51.0         |
| Stromgestehungskosten         | Pro  | 3.46      | 3.51        | 0.56 | 61.4         |
| Grüne Beschäftigung           | Pro  | 3.28      | 3.46        | 0.57 | 47.8         |
| Netzausbaukosten <sup>†</sup> | Anti | 3.84      | 3.53        | 0.72 | 17.5         |
| Energiekosten Industrie       | Anti | 4.09      | 3.60        | 0.81 | 35.8         |
| Preisvolatilität              | Anti | 3.73      | 3.36        | 0.66 | 33.9         |

Notes:  $N \approx 300$ . Policy relevance and perceived credibility: 1 = lowest, 5 = highest. PK is the confidence-weighted prior knowledge probability (see text). % Consistent is the share of respondents whose directional rating matches the assigned fact direction. <sup>†</sup> Reformulated after pre-validation; ratings are pre-reformulation and failed the consistency criterion. <sup>‡</sup> Relabeled from *Arbeitslosigkeit*; underlying statistic unchanged.

Table B.40: Balance Check: Selected Facts by Domain (Experiments 2 and 3)

| Domain      | $\Delta$ Relevance | $\Delta$ Credibility | $\Delta$ PK | Passes |
|-------------|--------------------|----------------------|-------------|--------|
| Immigration | 0.29               | 0.19                 | 0.15        | Yes    |
| Climate     | 0.42               | 0.03                 | 0.18        | Yes    |

Notes:  $\Delta$  = absolute difference between mean pro-direction and mean anti-direction facts. Pre-specified thresholds:  $\Delta$ Relevance  $\leq 1.0$ ,  $\Delta$ Credibility  $\leq 1.0$ ,  $\Delta$ PK  $\leq 0.25$ . Climate  $\Delta$ PK uses pre-reformulation statistics for *Netzausbaukosten*.

## B.4 Information Condition Implementation

Part 2 of the experiment proceeds in three stages: study presentation and elicitation (all senders), information feedback (INFO = 1 only), and expected attitude change elicitation (all senders). I describe each stage and reproduce the text shown to participants.

**Stage 1: Study Presentation and Elicitation (All Senders).** All senders read the following description of the [Haaland and Roth \(2020\)](#) study:

A recent study published in the prestigious academic journal, *Journal of Public Economics*, examined how people form political opinions.

The researchers surveyed more than 3,000 Americans that are representative of the US population.

People were told about a 1980 event when 55,000 Cuban immigrants unexpectedly arrived in Miami, increasing the city's workforce by 8% almost overnight. They were asked how this affected the wages and employment of U.S.-born workers.

Then, half of the participants learned that research found this wave of immigration had virtually no effect on the wages or unemployment for U.S.-born workers, because the new immigrants increased demand for goods and services, which created more jobs.

The other half of the participants received no information.

Senders then make two incentivized predictions:

1. *Behavioral effect (continuous)*: "By what percentage did learning about the research findings change people's petition signing rates?" Senders enter a percentage (positive or negative).
2. *Differential updating (categorical)*: "After learning about the research findings, who do you think changed their attitudes toward immigration more?" Senders choose among: (a) people who initially believed immigration harmed U.S.-born workers,

- (b) people who initially believed immigration did not harm U.S.-born workers, or
- (c) both groups changed about the same amount.

Each question is equally likely to be selected for the bonus payment. For the continuous estimate, the probability of receiving the bonus is 100% if the estimate equals the true value (69%) and decreases linearly to 0% at a deviation of 20 percentage points or more.<sup>42</sup> For the categorical estimate, the probability is 100% if correct and 0% otherwise. Senders are instructed not to search for the answers and tab-switching is monitored to enforce compliance; senders who switched tabs during this stage were excluded as pre-registered.

**Stage 2: Information Feedback (INFO = 1 Only).** Senders in the INFO condition receive personalized feedback comparing their predictions to the actual results. Senders in the no-INFO condition proceed directly to Stage 3. The feedback screen displays the following:

**Petition Signatures**

The percentage of people who signed the White House petition actually increased by **69%**.

You estimated that the percentage of people who signed the White House petition changed by *[participant's estimate]%*.

That means *you* *[underestimated/overestimated]* the change by *[participant-specific gap] percentage points*.

**Effect on Attitudes towards Immigration**

Actual Result: People who initially believed immigration harmed U.S.-born workers changed their attitudes towards immigration **significantly more** than those who already held correct beliefs.

Your Answer: *[participant's answer]*.

The feedback is accompanied by a bar chart that displays the sender's estimate and

<sup>42</sup>This implements a piecewise-linear proper scoring rule that is incentive-compatible under risk neutrality and approximately so under moderate risk aversion.

the true value (69%) side by side, with bar heights scaled proportionally. The sender's categorical prediction is marked as correct or incorrect with a color-coded indicator. A summary reiterates the two key findings: (i) receiving information strongly changed behavior and (ii) the effect was strongest for participants who initially held false beliefs.<sup>43</sup>

**Stage 3: Expected Attitude Change Elicitation (All Senders).** All senders then predict how receiving each of the four immigration statistics would shift the attitudes of hypothetical receivers. For each statistic  $j$  and each receiver belief state (false belief or “don’t know”), senders allocate 10 hypothetical participants among three categories: those who become *more supportive* of immigration, those who become *less supportive* of immigration, and those whose opinions *do not change*. The allocations must sum to 10. This procedure generates the expected attitude change measure  $\Delta\text{Attitude}_{i,j}$  defined in Section 2.3.1, which provides within-sender variation in perceived persuasiveness across statistics and receiver types.

You will estimate how ten participants who hold **false beliefs** or **do not know the answer** would change their attitudes about immigration if they **are provided with research findings**.

In particular, you will estimate how many of them...

- become less supportive of immigration.
- don't change their opinion.
- become more supportive of immigration.

The interface displays 10 person icons to make the allocation task concrete. Senders complete this task eight times: once for each of the four statistics crossed with each of the two receiver belief states. The order of statistics is randomized across senders.

<sup>43</sup>The gap between the participant's estimate and the true value is computed individually and displayed dynamically in the experimental interface. Senders who estimated the effect exactly see a confirmation message rather than a deviation.

## **B.5 Experimental Instructions and Interface**

Available upon request.

Pending

## B.6 Classification Procedure

I classify the 1,200 open-ended responses to the question “What factors or considerations influenced your decisions during the study?” into theory-derived motive categories using large language models. The primary classification uses OpenAI’s GPT-4o (accessed via the OpenAI API, temperature = 0). To assess robustness, I independently replicate the classification using Anthropic’s Claude Sonnet 4.6 (claude-sonnet-4-6, temperature = 0) with an identical prompt. Both models receive the same structured prompt containing category definitions, strict decision rules, and explicit boundary conditions.

The classification scheme comprises six categories, summarized in Table B.41. Categories are not mutually exclusive: multiple labels are assigned when a response contains distinct and separable justifications for each. No motives are inferred beyond the text. In addition to binary category indicators, I assign a primary motive corresponding to the most central justification.

The classification proceeds in five steps, enforced through the prompt structure. First, each response is checked for any reference to information disclosure (sharing, withholding, or a receiver). Responses without such reference are classified as *did not refer to disclosure* and receive no further coding. Second, only explicitly stated reasons are extracted. Third, responses are evaluated against category-specific criteria, including explicit boundary rules designed to distinguish closely related motivations. Fourth, multiple categories are assigned only when each is supported by a distinct reason. Fifth, a primary motive is selected.

The theoretically most important boundary is between *projection* and *instrumental*. Projection captures the sender’s preference for the receiver to hold aligned beliefs or attitudes—wanting the receiver to *think* a certain way. The instrumental category captures the sender’s concern for the receiver’s downstream behavior—wanting the receiver to *do* something specific (e.g., donate, vote, support an organization). This distinction maps directly onto the separation of  $\theta_P$  from  $\theta_I$  in the theoretical framework. The prompt enforces this boundary through explicit decision rules: references to influ-

encing what others *think* or *believe* are classified as projection, whereas references to influencing what others *do* are classified as instrumental.

A second important boundary separates *projection* from *did not refer to disclosure*. Projection requires that the respondent’s political beliefs are explicitly linked to the sharing or withholding decision—not merely stated. Responses that express political opinions without connecting them to the disclosure decision (e.g., “My personal beliefs”) are classified as *did not refer to disclosure*. This boundary is enforced through a strict necessary condition in the prompt: projection requires (i) explicit reference to own beliefs, (ii) an explicit link between those beliefs and the sharing decision, and (iii) that the direction of the information relative to those beliefs matters.

Table B.41: Coding Scheme and Decision Rules

| Category                    | Definition                                                                                                    | Key Decision Rule / Boundary                                                                                                                                                                                       |
|-----------------------------|---------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Belief correction           | Truth-telling, honesty, fairness, or the receiver’s right to accurate information; non-directional motivation | Must not depend on political direction; sender shares all facts, even those contradicting own beliefs                                                                                                              |
| Projection                  | Own beliefs or ideology explicitly linked to disclosure decision, including alignment of information          | Requires explicit link between beliefs and sharing decision; mere opinions without disclosure reference → <i>did not refer to disclosure</i> ; if the sender references influencing behavior → <i>instrumental</i> |
| Instrumental                | Intent to influence receiver’s actions (e.g., donating, voting, supporting an organization)                   | Requires explicit behavioral outcome; influencing beliefs only → <i>projection</i>                                                                                                                                 |
| Disbelief                   | Doubts about accuracy or credibility used to justify withholding                                              | Must explicitly refer to credibility concerns; politically motivated skepticism → <i>projection</i>                                                                                                                |
| Did not refer to disclosure | No reference to sharing, withholding, or receiver                                                             | Triggered in first decision step; excludes further classification                                                                                                                                                  |
| Miscellaneous               | Refers to decision but lacks interpretable reasoning                                                          | Residual category for non-classifiable responses                                                                                                                                                                   |

Applied to open-ended responses from Experiment 1 senders. GPT-4o (primary classifier) and Claude Sonnet 4.6 (validation classifier) each receive the identical prompt independently. Multiple labels may be assigned per response when distinct justifications are present; the primary motive identifies the most central justification.

The complete classification prompt, including all decision rules, boundary condi-

tions, and worked examples, is available upon request.

## B.7 Validation

To assess the robustness of the classification, I replicate the coding using a second LLM (Claude Sonnet 4.6) with the identical prompt and compare the two independent classifications. This inter-model agreement test checks whether the category definitions and decision rules are sufficiently precise to produce consistent classifications across different model architectures.

**Overall Agreement.** Table B.42 reports Cohen’s  $\kappa$  and disagreement rates for all categories. Inter-model agreement is substantial to near-perfect across the categories central to the analysis.  $\kappa$  exceeds 0.75 for four of the six binary categories: belief correction ( $\kappa = 0.887$ ), instrumental ( $\kappa = 0.788$ ), did not refer to disclosure ( $\kappa = 0.771$ ), and disbelief ( $\kappa = 0.708$ ). Agreement is somewhat lower for projection ( $\kappa = 0.670$ ) and miscellaneous ( $\kappa = 0.524$ ). The lower agreement on projection reflects the inherent difficulty of the projection–did-not-refer boundary: responses that state political opinions without clearly linking them to the sharing decision occupy a gray zone that different models resolve differently. The low agreement on miscellaneous is expected given its residual nature; this category enters no regression. For the composite primary motive classification,  $\kappa = 0.755$  with a disagreement rate of 16.3%.

Table B.42: Inter-Model Reliability: Cohen’s  $\kappa$  and Disagreement Rates

| Category                    | Cohen’s $\kappa$ | Disagreement Rate |
|-----------------------------|------------------|-------------------|
| Belief correction           | 0.887            | 5.0%              |
| Projection                  | 0.670            | 3.8%              |
| Instrumental                | 0.788            | 0.9%              |
| Disbelief                   | 0.708            | 1.6%              |
| Did Not Refer to Disclosure | 0.771            | 11.4%             |
| Miscellaneous               | 0.524            | 10.2%             |
| Primary Motive (overall)    | 0.755            | 16.3%             |

GPT-4o is the primary classifier; Claude Sonnet 4.6 is the independent validation classifier. Cohen’s  $\kappa$  and disagreement rates computed over Experiment 1 open-ended responses. Disagreement Rate = share of responses where the two classifiers assign different labels for that category.

**Precision, Recall, and F1.** Table B.43 reports category-level precision, recall, and F1-score, treating the GPT-4o classification as the positive reference for the purpose of computing directional metrics. For the three categories central to the main analysis—belief correction, projection, and instrumental—performance is strong. Belief correction achieves F1-scores of 0.96 (negative class) and 0.92 (positive class). Instrumental yields near-perfect agreement for the negative class ( $F1 = 1.00$ ) and strong agreement for the positive class ( $F1 = 0.79$ ), consistent with its low base rate (2.3% of responses). Projection shows somewhat lower recall for positive cases (0.63), reflecting the boundary ambiguity noted above. Miscellaneous shows the weakest positive-class F1 (0.58), expected given its residual nature. This does not affect the main findings, which do not rely on this category.

Table B.43: Category-Level Precision, Recall, and F1-Score

| Category                    | Class | Precision | Recall | F1   | Support |
|-----------------------------|-------|-----------|--------|------|---------|
| Belief correction           | 0     | 0.97      | 0.95   | 0.96 | 809     |
|                             | 1     | 0.91      | 0.94   | 0.92 | 390     |
| Projection                  | 0     | 0.97      | 0.99   | 0.98 | 1120    |
|                             | 1     | 0.76      | 0.63   | 0.69 | 79      |
| Instrumental                | 0     | 0.99      | 1.00   | 1.00 | 1171    |
|                             | 1     | 0.84      | 0.75   | 0.79 | 28      |
| Disbelief                   | 0     | 1.00      | 0.99   | 0.99 | 1171    |
|                             | 1     | 0.62      | 0.86   | 0.72 | 28      |
| Did Not Refer to Disclosure | 0     | 0.86      | 0.93   | 0.89 | 615     |
|                             | 1     | 0.92      | 0.84   | 0.88 | 584     |
| Miscellaneous               | 0     | 0.97      | 0.92   | 0.94 | 1083    |
|                             | 1     | 0.48      | 0.72   | 0.58 | 116     |

GPT-4o treated as the reference classifier (positive label = GPT-4o assigns the category). Claude Sonnet 4.6 is the comparison classifier. Support counts are from the GPT-4o classification. Class 0 = GPT-4o does not assign the category; Class 1 = GPT-4o assigns the category.

**Confusion Matrices.** Table B.44 reports confusion matrices for each binary category. Disagreements are rare and largely symmetric for the categories central to the analysis. For belief correction, 38 cases coded 0 by GPT-4o receive a 1 from Claude, and 22 cases show the reverse pattern—a small and roughly balanced discrepancy. Instrumental and disbelief show very few disagreements in absolute terms (11 and 19 cases respectively), consistent with their high  $\kappa$  values. The largest absolute disagreements occur

in did not refer to disclosure (137 cases) and miscellaneous (122 cases), both of which exhibit higher base rates and broader definitional scope.

Table B.44: Confusion Matrices: Binary Categories (rows = GPT-4o, columns = Claude Sonnet)

| Category                    |         | Claude = 0 | Claude = 1 |
|-----------------------------|---------|------------|------------|
| Belief correction           | GPT = 0 | 771        | 38         |
|                             | GPT = 1 | 22         | 368        |
| Projection                  | GPT = 0 | 1104       | 16         |
|                             | GPT = 1 | 29         | 50         |
| Instrumental                | GPT = 0 | 1167       | 4          |
|                             | GPT = 1 | 7          | 21         |
| Disbelief                   | GPT = 0 | 1156       | 15         |
|                             | GPT = 1 | 4          | 24         |
| Did Not Refer to Disclosure | GPT = 0 | 570        | 45         |
|                             | GPT = 1 | 92         | 492        |
| Miscellaneous               | GPT = 0 | 993        | 90         |
|                             | GPT = 1 | 32         | 84         |

Rows = GPT-4o (primary) classification; columns = Claude Sonnet 4.6 (validation) classification. Off-diagonal cells are inter-model disagreements. Experiment 1 open-ended responses.

Table B.45 reports the confusion matrix for the primary motive classification. The diagonal entries confirm strong agreement for the most common categories: belief correction (351 of 383 GPT-coded cases, 91.6%), did not refer to disclosure (492 of 584 cases, 84.2%), and miscellaneous (84 of 116 cases, 72.4%). The most frequent off-diagonal pattern is the confusion between did not refer to disclosure and miscellaneous (72 cases), which is theoretically expected: both categories capture responses that do not engage substantively with the disclosure decision, and their boundary is inherently ambiguous. Confusion between projection and other motivational categories remains limited, and instrumental is rarely misclassified, consistent with the high  $\kappa$  reported above.

Table B.45: Confusion Matrix: Primary Motive (rows = GPT-4o, columns = Claude Sonnet)

| GPT-4o                      | Claude Sonnet |           |           |            |           |                   |
|-----------------------------|---------------|-----------|-----------|------------|-----------|-------------------|
|                             | DNRD          | Disbelief | Misc.     | Projection | Instr.    | Belief correction |
| Did Not Refer to Disclosure | <b>492</b>    | 3         | 72        | 5          | 0         | 12                |
| Disbelief                   | 2             | <b>20</b> | 2         | 0          | 0         | 2                 |
| Miscellaneous               | 17            | 2         | <b>84</b> | 2          | 0         | 11                |
| Projection                  | 8             | 1         | 11        | <b>40</b>  | 3         | 9                 |
| Instrumental                | 1             | 0         | 0         | 2          | <b>17</b> | 8                 |
| Belief correction           | 17            | 1         | 2         | 2          | 0         | <b>351</b>        |

DNRD = Did Not Refer to Disclosure; Instr. = Instrumental.

Rows = GPT-4o primary motive; columns = Claude Sonnet 4.6 primary motive. Bold diagonal entries are agreements. Experiment 1 open-ended responses.

## B.8 Coding Examples

Table B.46 presents representative examples of open-ended responses classified into each motive category, illustrating the classification boundaries.

Pending

Table B.46: Coding Examples

| Motive                      | Example                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
|-----------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Belief correction           | <p><i>"For the questions, I went off my knowledge and my gut feelings. For whether or not to tell people the correct answer to their missed question, I always picked for them to learn the correct answer as I want our society to be as educated as it can be whether I like the stats or not."</i></p> <p><i>"I prefer facts even if they go against what I thought."</i></p> <p><i>"People should always have the right information."</i></p> <p><i>"I think it is important that people have all the facts."</i></p> <p><i>"I believe that everyone should see their answer and the research. What influenced my decisions was the research (or the answers)."</i></p>                                                                                                                                                                                                                                                                                                                                                                                                                      |
| Projection                  | <p><i>"I think it's important in this political environment that people advocate for immigrants regardless of statistics."</i></p> <p><i>"Immigration is such a polarizing topic right now. People are literally being sent to concentration camps. I didn't want to give anyone any information that would negatively impact their opinion against immigrants so I tried to give them positives and not share negatives."</i></p> <p><i>"I know how harmless immigrants are, and I also know how much harder it's been to get into this country. I want everyone else to know how harmless they are as well."</i></p> <p><i>"Especially in this day and age where the general public strays away from fact-checking and relies heavily on AI generated search results and/or answers via apps which tends to spew out right-wing rhetoric, it is so important to fight against anti-intellectualism and continue to fight against fake news/misinformation/propaganda..."</i></p> <p><i>"Knowing how illegal immigrants change communities for the worse was my biggest consideration."</i></p> |
| Instrumental                | <p><i>"Wanting people to be more well informed, especially when making donation decisions towards a big cause."</i></p> <p><i>"Some answers factored more into the organization I wanted donated more to. Other than that just using the information provided."</i></p> <p><i>"I do not wish to have any funds going to the pro-immigration group. I only selected information to give to the 'other participant' that might discourage them from donating."</i></p> <p><i>"I believe with the information, the participants will be able to make a more informed decision on how to allocate the \$500."</i></p> <p><i>"Consideration of how that person's opinions will directly affect a donation."</i></p>                                                                                                                                                                                                                                                                                                                                                                                   |
| Miscellaneous               | <p><i>"The facts."</i></p> <p><i>"Money."</i></p> <p><i>"Try to use common sense."</i></p> <p><i>"Information provided and my own moral compass."</i></p> <p><i>"Whether correcting or giving information would be beneficial and if they would be open to hearing it."</i></p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| Did Not Refer to Disclosure | <p><i>"Research findings and the inhumane treatment that immigrants often receive."</i></p> <p><i>"Estimating based on my knowledge of human nature but mostly instinctive."</i></p> <p><i>"Mainly my own knowledge or opinions, partially the compensation and desire to get things correct."</i></p> <p><i>"Being an empathetic person."</i></p> <p><i>"Own beliefs and presented findings."</i></p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| Disbelief                   | <p><i>"In terms of what influenced my decisions, I would say strength of evidence, the credibility of it, and if it had an intent to educate and not manipulate us."</i></p> <p><i>"I think the premise of some of the questions is wrong. For instance, illegally entering the country is a crime; therefore 100% of illegal immigrants are criminals even if the government doesn't universally prosecute them."</i></p> <p><i>"Doesn't really matter as significant portions of said study used effectively weaponized study propaganda."</i></p> <p><i>"I don't think people are going to change their mind on immigration based on outdated information from 15 years ago."</i></p> <p><i>"How believable is the fact."</i></p>                                                                                                                                                                                                                                                                                                                                                             |

Representative examples of open-ended responses from Experiment 1 senders, illustrating classification boundaries between motive categories. Responses are lightly excerpted for length; full text available upon request. Primary motive classification by GPT-4o.

## C Proofs

This appendix collects the proofs of the propositions in Section 3.4 and develops the magnitude prediction in full.

**Proof of Proposition 0 (Belief correction).** From equation (4),

$$\frac{\partial \Pr(d = 1)}{\partial \Delta b} = f(U(d = 1)) \theta_B.$$

Since  $f > 0$  everywhere and  $\theta_B > 0$ , the derivative is strictly positive.  $\square$

**Proof of Proposition 1a (Expressive selective disclosure).** From equation (4),

$$\frac{\partial \Pr(d = 1)}{\partial \alpha} = f(U(d = 1)) \theta_E.$$

Since  $f > 0$  everywhere and  $\theta_E > 0$ , the derivative is strictly positive.  $\square$

**Proof of Proposition 1b (Projective selective disclosure).** From equation (4),

$$\frac{\partial \Pr(d = 1)}{\partial w(\Delta \pi)} = f(U(d = 1)) c(A).$$

Since  $f > 0$  everywhere and  $c(A) > 0$ , this derivative is strictly positive.

If  $w$  is differentiable and strictly increasing, then  $w'(\Delta \pi) > 0$  on the relevant domain. Therefore,

$$\frac{\partial \Pr(d = 1)}{\partial \Delta \pi} = f(U(d = 1)) c(A) w'(\Delta \pi) > 0.$$

$\square$

**Proof of Proposition 2 (Political action amplification).** Differentiating equation (3) with respect to  $w(\Delta\pi)$  gives

$$\frac{\partial U(d=1)}{\partial w(\Delta\pi)} = \theta_P + \theta_I A.$$

The difference in this marginal utility between the action and no-action states is

$$\left[ \frac{\partial U(d=1)}{\partial w(\Delta\pi)} \mid A=1 \right] - \left[ \frac{\partial U(d=1)}{\partial w(\Delta\pi)} \mid A=0 \right] = (\theta_P + \theta_I) - \theta_P = \theta_I.$$

Since  $\theta_I > 0$ , political action strictly increases the marginal utility of persuasion; this increment is the instrumental motive.  $\square$

**Persuasiveness amplification.** The framework also yields a magnitude prediction, stated here in full.

**Proposition 3** (Persuasiveness amplification). *Fix  $A \in \{0, 1\}$  and suppose  $c(A) > 0$ . If  $w$  is differentiable and strictly increasing, then stronger perceived persuasion increases disclosure when persuasion moves the receiver toward the sender's own position and decreases disclosure when persuasion moves the receiver away from it. Specifically, for  $\Delta\pi > 0$ ,*

$$\frac{\partial \Pr(d=1)}{\partial \Delta\pi} > 0,$$

*while for  $p < 0$ , disclosure is strictly decreasing in  $|\Delta\pi|$ :*

$$\frac{\partial \Pr(d=1)}{\partial |\Delta\pi|} < 0.$$

*Proof.* For  $\Delta\pi > 0$ , equation (4) implies

$$\frac{\partial \Pr(d=1)}{\partial \Delta\pi} = f(U(d=1)) c(A) w'(\Delta\pi).$$

Since  $f > 0$ ,  $c(A) > 0$ , and  $w'(\Delta\pi) > 0$ , this derivative is strictly positive.

For  $p < 0$ , by the chain rule,

$$\frac{\partial \Pr(d = 1)}{\partial |\Delta\pi|} = \frac{\partial \Pr(d = 1)}{\partial \Delta\pi} \cdot \frac{\partial \Delta\pi}{\partial |\Delta\pi|} = f(U(d = 1)) c(A) w'(\Delta\pi) \cdot (-1).$$

Since  $f > 0$ ,  $c(A) > 0$ , and  $w'(\Delta\pi) > 0$ , this derivative is strictly negative.  $\square$

Proposition 3 adds a magnitude prediction to the directional prediction in Proposition 1b. When the state of the world is perceived to move the receiver toward the sender's position, disclosure increases with the perceived strength of that movement. When it is perceived to move the receiver away, disclosure decreases with the perceived strength of that movement. The boundary case is a function such as  $w(\Delta\pi) = \text{sign}(\Delta\pi)$ , where the sender responds only to the direction of persuasion and not to its magnitude.

## D Pre-Registration

This study was pre-registered on AsPredicted (No. 238343). The pre-registration specified four directional hypotheses (Hypotheses 0–3), a target sample size of  $N = 1,200$  (300 per treatment cell), and the decision-level linear probability model (2) as the primary specification.

A deviation from the pre-analysis plan is noted: the alignment measure was reparameterized from  $[0, 1]$  to  $[-1, 1]$  by centering at the indifference point ( $AIC_i = 250$ ). This is a one-to-one transformation that leaves all hypothesis tests invariant.